

Empirical Software Engineering

Guidelines for Empirical Studies in Software Engineering involving Large Language Models

--Manuscript Draft--

Manuscript Number:	EMSE-D-25-00637R1
Full Title:	Guidelines for Empirical Studies in Software Engineering involving Large Language Models
Article Type:	Research Papers
Keywords:	Software Engineering; Large language models; empirical research; meta-science; guidelines
Corresponding Author:	Sebastian Baltes Heidelberg University Bayreuth, GERMANY
Corresponding Author Secondary Information:	
Corresponding Author's Institution:	Heidelberg University
Corresponding Author's Secondary Institution:	
First Author:	Sebastian Baltes
First Author Secondary Information:	
Order of Authors:	Sebastian Baltes
	Florian Angermeir
	Chetan Arora
	Marvin Muñoz Barón
	Chunyang Chen
	Lukas Böhme
	Fabio Calefato
	Neil Ernst
	Davide Falessi
	Brian Fitzgerald
	Davide Fucci
	Junda He
	Christoph Treude
	Marcos Kalinowski
	Stefano Lambiase
	Daniel Russo
	Mircea Lungu
	Cristina Martinez Montes
	Lutz Prechelt
	Paul Ralph
	Rijnard van Tonder
	Stefan Wagner

Order of Authors Secondary Information:		
Funding Information:	Carlsbergfondet (CF24-0693)	Prof. Dr. Daniel Russo
	Alfred P. Sloan Foundation (G-2024-22586)	Prof. Dr. Daniel Russo
Abstract:	Large Language Models (LLMs) are now ubiquitous in software engineering (SE) research and practice, yet their non-determinism, opaque training data, and rapidly evolving models threaten the reproducibility and replicability of empirical studies. We address this challenge through a collaborative effort of 22 researchers, presenting a taxonomy of seven study types that organizes the landscape of LLM involvement in SE research, together with eight guidelines for designing and reporting such studies. Each guideline distinguishes requirements (must) from recommended practices (should) and is contextualized by the study types it applies to. Our guidelines recommend that researchers: (1)~declare LLM usage and role; (2) report model versions, configurations, and customizations; (3) document the tool architecture beyond the model; (4) disclose prompts, their development, and interaction logs; (5) validate LLM outputs with humans; (6) include an open LLM as a baseline; (7) use suitable baselines, benchmarks, and metrics; and (8) articulate limitations and mitigations. We complement the guidelines with an applicability matrix mapping guidelines to study types and a reporting checklist for authors and reviewers. We maintain the study types and guidelines online as a living resource for the community to use and shape (https://llm-guidelines.org/).	

Response Letter

Guidelines for Empirical Studies in Software Engineering Involving Large Language Models

EMSE-D-25-00637

S. Baltes et al.

Dear editor and reviewers,

Thank you for giving us the opportunity to submit a major revision of our manuscript. We have thoroughly revised the manuscript based on the feedback provided by the reviewers. In the following, you can find detailed responses to each comment.

As part of this revision, we welcomed new co-authors (Cristina Martinez Montes, Rijnard van Tonder, and Junda He) who contributed additional expertise to the major rewrite of the guidelines. This project is a community effort, and we remain open to new contributors who can strengthen the breadth and depth of our recommendations.

We have considerably revised the paper based on the reviewers' feedback. Due to the extensive nature of the changes (structural reorganization, new sections, and cross-cutting edits), the provided `latexdiff` document may be difficult to navigate. We therefore provide detailed responses to each comment below, referencing the relevant sections of the revised manuscript.

Best regards,
Sebastian Baltes (on behalf of all co-authors)

RESPONSE TO EDITOR

Comment 0.1:

The reviewers all recommend a major revision of the paper. For the revision, please focus on:

- Structure & Presentation. Make the paper more concise, reduce repetition, provide a summary tables of the guidelines.
- Terminology #1. Reconsider the use of RFC terminology. As R2 makes the case, peer review is not a standardization method.
- Terminology #2. Be careful claiming "undisputable truths" (see comments by R2). Guidelines should allow flexibility given resource constraints.
- Methodological Improvements. Provide additional details on the methodology as requested in the reviews. For the revision, we do not expect any additional experiments,

evaluations, but please increase the rigor as much as possible.

Response:

Thanks for your guidance. We addressed your focus points in the following way:

- **Structure & Presentation:** We conducted a comprehensive revision of the existing content, focusing on deduplication. We moved advantages and challenges to a cross-cutting summary section (Section 4.3), added a “How to Navigate this Paper” subsection, created Table 3 (guideline–study type matrix), Table 4 (rationale–recommendations mapping), and a comprehensive checklist (Appendix A).
- **Terminology #1:** We dropped the reference to RFC 2119/8174. We retained *must* and *should* as plain indicators of essential requirements vs. desired practices, and converted all instances of *may* to prose or *should* to reduce complexity.
- **Terminology #2:** We revisited all *must* items and limited them to practices that our author group unanimously considered both essential and broadly feasible. We added flexibility language (Section 5) encouraging sensible trade-offs and justification for deviations.
- **Methodological Improvements:** We expanded Section 3 (Methodology) with details on the example selection process. We added a self-assessment against the SIGSOFT Empirical Standards quality dimensions (conclusion). Each guideline now includes dedicated “Rationale” subsections that explicitly motivate the guideline, as well as “Advice for Reviewers” and “See Also” subsections.

We provide more detailed responses to the individual reviewer comments below. Thank you for your valuable comments on our manuscript. We have done our best to incorporate changes to reflect the suggestions, which allowed us to improve the quality of our work. We provide a `latexdiff` document as part of the submission; however, given the extent of the changes, it may be difficult to navigate.

RESPONSE TO REVIEWERS

Reviewer 1

Comment 1.1:

This is worthwhile work that can benefit from a number of improvements mainly in terms of its structure.

Response:

Thank you for your feedback. We have carefully addressed all the issues item by item as follows.

Comment 1.2:

Regarding significance, the motivation behind this work is clearly set in the introduction. The presented contributions can clearly impact the field of software engineering theory and practice as described in the concluding sections.

Response:

We thank the reviewer for these supportive comments.

Comment 1.3:

Regarding novelty, the work’s contributions are original with respect to the state-of-the-art. Existing work in the area is coherently analyzed and synthesized, and the study’s contributions are clearly identified in its context.

Response:

We thank the reviewer for these supportive comments.

Comment 1.4:

Regarding soundness, the study’s contributions address its research questions, but in some cases the applied research methods lack the required rigor. In particular the methods described in Section 3 could be improved by having the process informed in a systematic way from other disciplines. The contribution’s evaluation can be improved perhaps with a provision (as an appendix) of a model study following the recommendations, in order to judge the recommendations’ practicality and value. Threats to validity and their effect on the results’ soundness are not explicitly presented.

Response:

We thank the reviewer for their feedback. However, we want to note that this is a meta-scientific paper. It doesn’t have “threats to validity” or “soundness” in the same sense as an empirical paper. The SIGSOFT Empirical Standards suggest the following quality criteria for meta-science:¹

- comprehensiveness of analysis or guidance provided
- usefulness to the research community
- quality of argumentation supporting analysis of guidance
- degree of integration with previous work, both in software engineering and in reference disciplines

We have added a paragraph to the conclusion that discusses these aspects from the authors’ point of view, but we are open to consider the reviewer’s feedback along these dimensions, especially regarding the quality of our argumentation.

Regarding the “model study”, we want to point the reviewer to the “Example(s)” sub-sections of the guidelines, which contain pointers to papers following (part of) our guidelines, highlighting the integration with previous work. Extending these sub-sections would significantly increase an already high page count.

¹<https://www2.sigsoft.org/EmpiricalStandards/docs/standards?standard=MetaScience>

However, the guidelines exhibit high face validity in that subject matter experts derived these guidelines. Also, the process for establishing the guidelines is similar to that of guidelines with similar goals in other disciplines (see, e.g., [Begg et al \(1996\)](#)).

Comment 1.5:

Regarding verifiability, independent verification of the paper’s claimed contributions is difficult, though this is a difficult to address issue.

Response:

Following our response to [Comment 1.4](#), we argue, that verifiability does not align with the quality criteria for meta-science established in the SIGSOFT Empirical Standards.

Comment 1.6:

Regarding presentation, the paper’s structure can be improved. Several challenges and advantages are repeated in many sections. Consider codifying them, analyzing them once, and then cross-referencing them from a table. Furthermore, the identified challenges are given without any proposed means to address them. Most are later addressed through the proposed guidelines. Consider codifying and linking the two, again in tabular form.

Response:

We agree with the reviewer and moved the advantages and the challenges to a separate section and linked, where relevant, to study types. We deliberately avoided bi-directional links between study types and guidelines for long-term maintainability reasons, but introduced with Table 3 a matrix that maps study types to guidelines, creating an implicit link.

Comment 1.7:

Along similar lines, consider splitting the provided recommendations into a hierarchically numbered set of recommendations and a separate rationale exposition. In some places the exposition is not as clear as it could be.

Response:

To ensure the recommendations get the necessary exposition, we summarized them into a dedicated checklist inspired by CONSORT [Schulz et al \(2010\)](#), structuring items into requirements (*must*) and recommendations (*should*).

Comment 1.8:

Furthermore, many paragraphs, such as those in Section 5.8.1, start with a long-winded rationale before arriving at the recommendation. Consider using recommendations as each paragraph’s thematic sentence. The text is free from typos and style errors. The formatting follows the provided instructions, but could be improved. Please fix the capitalization of referenced titles, see e.g. reference 96 (should be Python, ChatGPT).

Response:

We agree, that a thematic start improves the reading flow. Due to the number of requirements per recommendation section we refrained from using thematic sentences, but rather started each with a short rationale subsection.

Comment 1.9:

Consider converting the RFC 2119 MUST/SHALL/MAY verbs into lowercase. Their uppercase appearance is distracting, and the paper contains guidelines, not a technical standard.

Response:

We agree, that the paper is not intended as a standardization, but as a community-driven effort. Hence we dropped the technical standards background (RFC). MUST and SHOULD are kept as indicators towards essential and desirable requirements. To reduce complexity, we converted instances of MAY to prose or SHOULD.

Comment 1.10:

Consider reducing in-text cross-references to other sections—they are distracting. If needed, add them in a separate section, as is done in many Design Pattern catalogues. The abstract is self-standing and summarizes succinctly the paper’s essence. The paper could also benefit from the addition of a few figures to help navigating it.

Response:

We see the point about the length of cross-references. We intentionally did not reduce in-text cross-references to sections, as we feel this would require higher cognitive load of the reader. After consideration, we shortened the in-text cross-references significantly (e.g. from "Use Human Validation for LLM Outputs" to "Human Validation"). This should also facilitate navigation across the entire paper.

Regarding figures, the paper now contains additional tables that serve a navigational purpose (Tables 1–4). We considered adding figures but found that the tabular format better served the paper’s reference-style usage. The checklist and matrix table collectively address the navigation concern.

Comment 1.11:

Below is a list of more detailed comments.

Page 3, line 16: Consider providing some details or examples as motivation, perhaps specifically targetting LLMs.

Response:

We restructured the introduction to foreground concrete LLM-specific reproducibility threats—non-determinism, model evolution beyond version identifiers, and opaque training data—as motivation, rather than burying them later in the gap paragraph. Prompt formatting and temperature effects are now also cited early as examples of unreported settings that directly impact reproducibility. The revised introduction is shorter overall (three paragraphs instead of four) while providing the concrete LLM-specific examples the reviewer suggested.

Comment 1.12:

Page 3, line 29: Consider also mentioning the ACM SIGSOFT Empirical Standards for Software Engineering.

Response:

We added the ACM SIGSOFT Empirical Standards.

Comment 1.13:

Page 3, line 33: LLMs are also a new tool, which requires fresh methodological guidance.

Response:

We restructured the introduction to explicitly motivate why LLMs require fresh methodological guidance. The revised text now foregrounds three LLM-specific threats (non-determinism, evolving models, opaque training data) and concrete examples (prompt formatting, temperature) before identifying the gap in existing SE guidelines.

Comment 1.14:

Page 4, line 7: Consider changing "recent version" into "recent version of this document".

Response:

We are referring to continuously evolving guidelines as published online. We changed it as follows:

The most recent version of the guidelines is always available online.

Comment 1.15:

Page 4, line 9: Remove "xw"? Page 4, line 38: Consider changing "advise" into "advice".

Response:

Fixed.

Comment 1.16:

Section 2: The scope includes software code in addition to natural language, right?

Response:

Yes, absolutely. We changed it to “textual use cases.”

Comment 1.17:

Page 5, line 23: Also data generation, as in 4.1.4

Response:

We added ”generation“.

Comment 1.18:

Page 6, line 29: The numbering is wrong.

Response:

Fixed.

Comment 1.19:

Page 7, line 38: Consider changing "inter-rater agreement" into "accuracy".

Response:

We kept “inter-rater agreement” here because we are referring to a specific paper in which they use that term.

Comment 1.20:

Page 8, line 16: Please clarify.

Response:

We have now clarified it further and added an example to improve the sentence’s readability.

Studies have shown that LLMs can be unreliable for high-stakes labeling, i.e., labels where misclassification has substantial downstream consequences (e.g., safety-critical or compliance-related decisions). In addition, LLMs can exhibit uneven per-class performance, with accuracy differing markedly across label categories, often performing worse on rare or ambiguous labels and systematically over- or under-assigning specific annotation categories.

Comment 1.21:

Section 4.1: It might be worth investigating the following philosophical question. Is it sound to consider human judgment as the gold standard against which LLM performance is evaluated? For example, this is not the approach scientists employ when evaluating the accuracy of spectrometer measurements.

Response:

This is quite an interesting remark. We discuss this in Section 5.5.5 (challenges of human validation for LLM outputs).

Comment 1.22:

Section 4.1.2: Several of the advantages listed here also apply to 4.1.1. Consider some form of merging to avoid gratuitous differences or repetitions.

Response:

See our response to [Comment 1.6](#). We consolidated the advantages and challenges from individual study types into a cross-cutting Section 4.3, which directly addresses the overlap observed by the reviewer.

Comment 1.23:

Page 10, line 11: Please explain why this is an issue.

Response:

We added:

The underlying statistical framework of an LLM usually pushes outputs towards the most likely (majority) answer, which might not be the best answer. Questions about bias, accuracy, and trust remain

Comment 1.24:

Page 10, line 32: The word synthesis also applies to this use, but under a different meaning: "the production of a substance from simpler materials after a chemical reaction" rather than "the mixing of different ideas, influences, or things to make a whole that is different, or new" as used at the beginning of the paragraph. To avoid confusion, consider splitting this case into a separate section with a title like "synthetic data generation".

Response:

Thanks for this remark. We have revised the section to make the two different notions of “synthesis” clearer. Please note the examples we reference, showing that SE researchers use both meanings in research papers.

Although synthesis in the preceding notion refers to abstraction and interpretation across multiple data sources, the term is sometimes also used to refer to generating synthetic content (e.g., source code, bug-fix pairs, requirements, etc.) that are then used in downstream tasks to train, fine-tune, or evaluate existing models or tools. In this case, the synthesis is done primarily using the LLM and its training data; the input is limited to basic instructions and examples.

Comment 1.25:

Section 4.1.3: The examples provided here are the the wrong (meta) level compared to the other sections. If no direct examples can be found, consider mentioning this distinction.

Response:

We have now added concrete examples from software engineering, from scoping studies to specific proposed implementation frameworks.

Comment 1.26:

Section 4.1.3: Several of the challenges listed here also apply to other sections. Consider some form of merging to avoid gratuitous differences or repetitions.

Response:

See actions and response to [Comment 1.6](#).

Comment 1.27:

Page 12, line 32: Please clarify the meaning and implications of "should". Is this something the researchers should verify (if so, how) or is it a currently unattainable ideal?

Response:

We understand the point being made. However, these 4 criteria are drawn from reference [Argyle et al \(2022\)](#) so we do not feel we can/should qualify an individual criterion.

Comment 1.28:

Page 16, line 22: A key advantage is the automation of tasks previously executed by humans.

Response:

Added this as an overarching advantage.

Comment 1.29:

Page 16, line 44: Please ... and replicability.

Response:

We added this point on replicability to the list of challenges.

Comment 1.30:

Section 5: Consider also including guidelines for reviewing papers that use LLMs. Following the providing guidelines is a start, but there are also things the reviewers should avoid, such as excessive requirements regarding verification. Furthermore, consider the space needed for fulfilling all the requirements; could much of the data be supplied as part of a replication package? If so which?

Response:

We added a paragraph to Section 5. indicating that reviewers should use the guidelines to help write reviews, emphasizing a focus on doing what is practically possible rather than ideal in principle.

We also added detailed new “advice for reviewers” subsections to each guideline advising reviewers on what to focus on, how to deal with common problems, and how to differentiate minor vs. major problems. We use these sections to encourage good reviewing thinking rather than repeat each guideline’s individual must/should items.

Comment 1.31:

Section 5.2: Is the position regarding the requirement to disclose version too timid? As stated, many providers do not make it possible to use a past version. I think more epistemological thought, deliberation, and discussion is needed regarding the fact that research with SAAS models may not be deterministically reproducible. This is not necessarily a fatal problem. Many research methods, such as ethnography, grounded theory, and action research, share this feature. However, in such cases additional emphasis is placed in methodological elements such as triangulation, reflexivity, audit trails, and peer debriefing. These measures aim to ensure trustworthiness, encompassing credibility, transferability, dependability, and confirmability—the qualitative counterparts to validity, generalizability, reliability, and objectivity in deterministic quantitative research. Consider discussing these issues and adjusting the recommendations.

Response:

Thank you for this feedback. We now explicitly draw on qualitative research epistemology (Guba (1981)), framing SaaS-based LLM research as facing analogous conditions to ethnography, grounded theory, and action research. We added trustworthiness criteria (credibility, transferability, dependability, confirmability) and measures such as triangulation, reflexivity, and audit trails to Section 5.8.

Comment 1.32:

Section 5.2.4: Consider expanding and explaining the drawbacks, advantages, and limitations of a 0 temperature setting regarding reproducibility.

Response:

We added a corresponding discussion at the end of Section 5.2.5.

Comment 1.33:

Section 5.2.5: I recommend summarizing recommendations by study type also in a table. (Also applies to the other .5 subsections.)

Response:

See response to [Comment 1.6](#) about avoiding bi-directional links between study types and guidelines.

Comment 1.34:

Section 5.3: In general, for replicability we tend to distribute source code. So why here is the requirement only for abstract models (e.g. architecture diagrams, coordination logic), rather than concrete implementations? Consider requiring both.

Response:

We agree with the reviewer, that a replication package is to be expected (see discussions in Section 5.6 and Section 5.8).

Comment 1.35:

Page 29, line 1: Why is "keeping a record of the actual conversation" "crucial for accuracy"? There are several other such statements in the paragraphs' concluding phrases, which read like AI-generated slop. Please review them carefully, removing those that don't add anything substantial and fixing those, such as this one, that read nicely but are actually incorrect.

Response:

We thank the reviewer for this feedback and agree with their assessment of the corresponding sentences. We have reformulated them as follows:

The rationale for this is similar to the rationale for reporting interview transcripts in qualitative research. Just as a human participant might give different answers to the same question asked two months apart, the responses from tools such as ChatGPT can also vary over time.

Comment 1.36:

Section 5.5.2: The text in the quotes appears to be paraphrased, so remove them (or adjust the text). Clarify what are "the first three categories".

Response:

We removed the quote and summarized it, to also mitigating any confusion about "the first three categories".

Comment 1.37:

Page 32, line 25: Please clarify how this is relevant to human validation.

Response:

We agree, that the example was not perfectly aligned with the section. Hence we replaced it with a better fitting example on human validation of LLM generated content.

Comment 1.38:

Section 5.5.2: What is your recommendation given the cited low human-model and human-human agreements?

Response:

This example underlines our recommendation to use human validation for LLM outputs. LLMs, in their current state, are not suitable to judge each others outputs, as they share systematic biases. We rephrase this part to reason deeper about the low human-model and high model-model agreements.

Comment 1.39:

Section 5.5.4: Please clarify the relevance of the "struggle with adapting to AI-assisted work" to this section.

Response:

We thank the reviewer for their remark and have removed the corresponding article from the corresponding section.

Comment 1.40:

Section 5.6.1: Please clarify why including an open LLM baseline might not be possible if the study involves human participants.

Response:

In studies involving human participants, including an open LLM baseline may require participants to repeat the same tasks with multiple models. This can be infeasible, for example, due to time constraints and learning effects. Therefore, while open baselines are strongly encouraged, they may not always be practical in human-centered study designs. We are open to strengthening Section 5.6.1 based on additional feedback.

Comment 1.41:

Page 35, line 45: Consider changing "managed cloud APIs provided by proprietary vendors" into "vendor-provided managed proprietary cloud APIs".

Response:

We clarified our statement by removing "managed cloud". And added example of one vendor for clarity.

Comment 1.42:

Section 5.6: Consider tabulating the scientific advantages of moving from proprietary, to open-weight, to full-open models.

Response:

Thank you for the suggestion. We considered adding such a table but concluded that a meaningful comparison would require multiple dimensions (e.g., reproducibility, cost, customizability, transparency of training data) that vary significantly across specific models and change rapidly. Instead, the rationale and recommendations in Section 5.6 discuss these trade-offs in prose, which we believe is more maintainable given the pace of change in this field.

Comment 1.43:

Page 37, line 15: Static analysis would be an even better example.

Response:

Thanks, we added static analysis as an example in addition to program repair.

Comment 1.44:

Page 37, line 4: Please clarify: what benchmarks is the example referring to?

Response:

Thanks, we added references (Chen et al (2021), Austin et al (2021)) for improved clarity.

Comment 1.45:

Section 5.7.2: Most of the information presented here is (useful) background, not concrete examples. Consider moving it to an appropriate section.

Response:

We moved the background explanations of metrics to the recommendation subsection.

Comment 1.46:

Page 38, line 37: Please use a display equation

Response:

Fixed.

Comment 1.47:

Page 38, line 49: The tasks mentioned here do not adhere to the problem types introduced at the beginning of Section 5.7.2

Response:

We restructured the example section according to the problem types. We hope this clarifies the relation of tasks to problem types.

Comment 1.48:

Section 5.7: This section can benefit from a structured classification of the benchmarks described in Section 5.7.2.

Response:

Thanks for the suggestion. To not increase the length of the paper further we refer the interested reader to the classifications in the initial references of 5.7.1.

Comment 1.49:

Page 40, line 18: This is one of the most important issues. Consider placing it more prominently and providing concrete

Response:

Thank you for raising this point. The topic of benchmark contamination/inter-dataset duplication was discussed mainly in Section 5.8 (Report Limitations and Mitigations) with a brief mentioning in Section 5.7 (Use Suitable Baselines, Benchmarks, and Metrics). We performed the following changes: First, a detailed discussion of benchmark contamination and related mitigations in 5.7. Second, we reduced the mentioning in 5.8 on the broader topic of inter-dataset duplication and the necessity of authors to discuss those threats to the validity. We hope this restructuring give the topic the appropriate attention and makes it more concrete for authors.

Comment 1.50:

Page 42, line 7: GPT-4o is a very coarse model identifier. Clarify whether the problem also applies to models sharing the same fingerprint.

Response:

Thank you for the feedback, we extensively rewrote the limitations and mitigations section to avoid repetitions with the other recommendation sections. This specific part is described on a more abstract level now. To our knowledge, there is no research on the stability of model performance based on fingerprinting over time.

Comment 1.51:

Page 42, line 35: Please start the paragraph with a thematic sentence.

Response:

Thank you for the feedback, we extensively rewrote the limitations and mitigations section to avoid repetitions with the other recommendation sections.

Comment 1.52:

Page 44, line 26: The paragraph starts by discussing generalizability and then switches to a discussion of code duplication. Please structure accordingly.

Response:

Thank you for the feedback, we extensively rewrote the limitations and mitigations section to avoid repetitions with the other recommendation sections. The topic of code duplication was mentioned as one example for dataset leakage affecting generalizability, however we also describe this matter on a more abstract level now.

Comment 1.53:

Section 5.8: To aid transparency and provide context, the cost should be reported by breaking down its components, e.g. number of input tokens, output tokens, and corresponding unit costs. This allows predicting future costs.

Response:

Thank you for the thoughtful suggestion. We included it as the preferred way of reporting costs.

Comment 1.54:

Section 5.8: Given that its benefits are very difficult to judge, I am not convinced regarding the value of an environmental impact assessment for scientific research. How would one evaluate CERN LHC, Mars space missions, ITER, NIF, and LLM model training in this context?

Response:

The reviewer raises an important point here. We agree, that SHOULD was too strong. Given the increased environmental impact of LLM-focused research compared to traditional research, we nevertheless deem it noteworthy to raise awareness. Hence we shifted the discussion from impact assessment to the tradeoff of performance improvement vs. increase resource consumption.

Reviewer 2

Comment 2.1:

The paper is about LLM usage in software engineering research. It first describes the main uses cases of LLMs in SE research and then proceeds with guidelines. These guidelines recommend that researchers declare LLM usage, report model versions and configurations, document tool architectures, disclose prompts, use human validation, employ open LLMs and suitable base-lines, and articulate study limitations. The ultimate goal is to enable open science and ensure the reliability of research findings.

Main strengths:

- **timeliness:** This is a timely and well-structured contribution to the software engineering methodology literature. There are already zillions of SE papers with LLMs, and there will be many more. The paper provides a clear and actionable framework for improving the rigor and transparency of empirical studies involving LLMs.
- **Relevance and Clear Motivation:** The paper addresses a critical, contemporary problem in software engineering research: the challenge of maintaining scientific rigor (specifically reproducibility and replicability) in studies involving rapidly evolving, non-deterministic, and often opaque LLMs. The motivation is clear, compelling, and immediately relevant to any researcher or practitioner in the field.
- **Credible, Community-Driven Methodology:** The guidelines are not the product of a single research group but are presented as a community effort that originated from discussions at a research and was refined through workshops and reviews. This collaborative approach lends the guidelines authority and increases the likelihood of their adoption by the community.
- **Actionable Guidelines:** The eight guidelines are thorough, practical, and well-articulated.

Response:

We thank the reviewer for their encouraging feedback.

Comment 2.2:

Too long: we want to encourage the community to read this, the longer the paper, the fewer the readers. Trim, cut. For example, the "study types" subsections in Sec 5 are heavy and confusing (it took me some time to understand they're referring to Sec 4), I would simply get rid of those subsections.

Response:

We agree with the reviewer, that the length of the paper could be counterproductive to the goal of encouraging the community to read it. Hence, we introduced a checklist summarizing all MUST/SHOULD recommendations as an initial starting point, state the key recommendations in Table 4, integrate a matrix (Table 3) to navigate between study types and guidelines, and introduced a "How to Navigate" subsection.

We hope those changes enable easy consumption of the guidelines while not losing any of the necessary context and rationales that contribute to the length of the paper.

Comment 2.3:

Lack of overview: we should be able to read the paper in one page. Consolidate all guidelines in a table at the beginning

Response:

We aggregated all recommendations in a new checklist (see appendix).

Comment 2.4:

Usage of RFC terminology MUST/SHOULD: a paper, even community driven, cannot speak for the whole community, the paper is not a standard and EMSE is not a standardization body. Peer-review is not a standardization method. I would recommend to NOT take the style a standard, this is unsound and confusing.

Response:

Also the first reviewer pointed out something similar. We agree, that this manuscript should not act as a standardization. We nevertheless deem it important to distinguish between different recommendations. Hence we explained the difference between MUST and SHOULD in terms of desired practices, and dropped the reference to the RFC. Further, we converted MAY to SHOULD or prose to reduce complexity.

Comment 2.5:

- Usage of MUST: this is too strong. even if I agree with the utmost importance of some guidelines, my general philosophy of science is to be very cautious with undisputable truth and hard rules. Science is about taking a critical standpoint at everything, and about being able to take one's own path, regardless of the others command. I would remove the MUST.

Response:

We understand the reviewer's concern; given the size of the research team, there are indeed differing perspectives within our group, including views aligned with the one expressed by the reviewer. That said, through internal discussion we converged on the conclusion that, in order to make this contribution concrete and genuinely usable, the adoption of such terminology (not only MUST, but also SHOULD, etc.) was necessary. We deliberately limited the use of MUST to guidelines that we, as a group, unanimously considered both essential and broadly acceptable. However, to make it truly usable in practice, we provide more flexibility now on a sensible selection of recommendations and critical discussions on the impact of those decisions in Section 5 (in line with the reviewers suggestion).

Comment 2.6:

If some guidelines are too strong, it decreases the credibility of the whole effort. for example the suggestion to test results on commercial models "over an extended period of time" is methodologically sound but may be infeasible given the time and cost constraints of typical academic projects.

Response:

We went over all recommendations again and ensured, that MUST is limited to instances where requirements are essential and feasible (within the boundaries of sensible decisions by the authors in the individual study context).

Comment 2.7:

Cherry-picking examples: there are many more example papers to cite, but it is not the goal of the paper to be a survey. Yet, how were the examples selected? bias towards the authors' paper? their collaborators

Response:

The paper selection process was conducted as follows:

1. During an initial meeting, at least two authors were assigned to each specific part of the guidelines (study types or individual guidelines), with the responsibility of drafting the corresponding section.
2. Each pair of authors independently started drafting their section by searching for relevant examples using search strings on academic databases.
3. After collecting an initial set of examples, the pool was further extended based on co-author's suggestions.
4. The identified papers were then individually assessed by one author and subsequently reviewed by the second author. Papers that did not add new insights beyond those already covered by the existing examples were not included.
5. Finally, multiple final review rounds were carried out on the guidelines as a whole, during which the selected papers were cross-checked by additional authors and, where necessary, further examples were added.
6. As part of this journal revision, we then checked the references again based on the reviewers' feedback.

We added this to the methodology section. As correctly noted by the reviewer, the purpose of the examples was not to be systematic, but illustrative. Since this work is not intended to be a systematic review, we believe that this approach is appropriate and sufficient. Moreover, the involvement of multiple authors at different stages of the process helped mitigate individual bias and ensured a balanced selection of examples. Finally, this paper is a community initiative that will extend beyond this paper. Therefore, we invite everything to suggest better examples than the ones we referred to.

Comment 2.8:

Focus on Natural Language Use Cases -> Focus on Textual Use Cases. Code is not natural language.

Response:
Fixed!

Reviewer 3

Comment 3.1:

This paper establishes guidelines and recommendations for software engineering researchers conducting empirical studies that use Large Language Models (LLMs). The guidelines are structured around a taxonomy of study types where LLMs serve various roles, such as tools for SE researchers (Annotators, Judges, Synthesis, and Subjects) or tools for software engineers (Studying LLM Usage, LLMs for New Tools, and Benchmarking). Furthermore, the paper presents concrete reporting guidelines covering essential aspects of reproducibility, including declaring LLM usage, documenting model versions and configurations, describing tool architectures, reporting prompts and interaction logs, advocating for human validation, using open LLMs as baselines, and meticulously reporting limitations and mitigations. Overall, the paper aims to promote validity and transparency in LLM-related SE research, recognizing the challenges posed by their non-deterministic nature, evolving versions, and potential biases.

Strengths:

This research is timely for the SE community.

Covering the SE cycle and benchmarks highlights the importance of benchmarking every step and the challenges faced.

The guidelines enhance the validity, transparency, reproducibility, and replicability of empirical studies in SE that involve LLMs.

The taxonomy organizing LLM usage into seven distinct study types is well-structured and covers the landscape systematically.

The eight guidelines use MUST, SHOULD, MAY terminology to distinguish between essential and desired practices.

The methodology demonstrates real community engagement.

Response:

Thank you for the summary and the acknowledged strengths!

Comment 3.2:

Insufficient Discussion of Trade-offs and Conflicts Between Guidelines. For example, using an open LLM baseline (Guideline 6) while also requiring multiple experimental repetitions (Guideline 7) and human validation (Guideline 5) could be prohibitively expensive for many researchers. How to prioritize the guidelines? Do you recommend any decision tree to define which study type is the most appropriate?

Response:

We agree that there exists an inherent trade-off between guidelines. Where possible, we mentioned those trade-offs. Additionally, we added a remark in Section 5 advocating for sensible trade-offs and related discussions. Due to the variety in different study setups, we cannot offer a prescriptive way of prioritizing guidelines beyond our classification into requirements and desired practices.

Comment 3.3:

An Insufficient Treatment of Data Contamination and Leakage. While Section 5.8.1 mentions data leakage, the discussion is buried in the limitations section rather than being a standalone guideline. Given that benchmark contamination is acknowledged as a "major challenge," this deserves more prominence. What recommendations or validation strategies do you recommend?

Response:

Thank you for highlighting this improvement point. We agree, that this topic is a major challenge. To ensure this issue receives the necessary attention, we expanded the discussion of benchmark contamination and gave concrete mitigations in more detail in Section 5.7. In the same line, we streamlined the discussion in 5.8 on the broader topic of inter-dataset duplication, also here discussing the necessity of authors to discuss those threats to the validity while providing initial guidance to mitigate those threats. We hope this restructuring gives the topic the appropriate attention and makes it more concrete for authors.

Comment 3.4:

Insufficient Guidance on Statistical Power and Sample Sizes. While the paper mentions power analysis in Guideline 5, there's no systematic guidance on determining appropriate sample sizes for LLM experiments, especially given non-determinism.

Response:

Providing systematic guidance on statistical power analysis or sample size estimation for LLM experiments is beyond the scope of these guidelines, as it depends heavily on the specific study design, task variability, and target precision. For general guidance on sampling, we refer the reader to the ACM SIGSOFT Empirical Standards. However, we have expanded Section 5.7 to discuss the number of experimental repetitions, recommending that researchers justify their chosen number through power analysis or convergence monitoring of descriptive statistics across incremental runs.

Comment 3.5:

DETAILED COMMENTS: This is a highly pertinent and timely paper for empirical SE community. While I think this paper should definitely be accepted, I have a few suggestions for the authors to improve before a possible publication. Hence, I recommend a Major Revision.

Response:

Thanks for this positive stance towards our manuscript and the chance to improve it!

Comment 3.6:

Several guidelines contain SHOULD recommendations that may be difficult to follow:

- * Guideline 5.4 (Report full interaction logs): For large-scale studies with thousands of LLM calls, storage and privacy concerns may prohibit full disclosure
- * Guideline 5.6 (Open LLM baseline): When state-of-the-art proprietary models significantly outperform open alternatives, including an underperforming baseline may provide limited scientific value
- * Guideline 5.8 (Report costs): Cloud API costs fluctuate; reported costs may not reflect what future researchers will pay

Response:

As per the new definition of SHOULD, such recommendations present a desired practice, hence deviation is acceptable, and likely will often happen. Nonetheless, we believe, that making desired practices explicit adds value to a study description.

Guideline 5.4: Guideline 5.4 already addresses exceptional situations by specifying that: "If full prompt disclosure is not feasible... summaries or examples should be provided." However, the broader principle is that, when a paper's analysis is based on interaction logs, finding ways to share that data is a challenge to embrace, not avoid. Shared data enables verification of claims and follow-up research, and increases the value of the work itself, as every published dataset becomes a resource for meta-analyses and replication studies.

For privacy concerns, sensitive content can be redacted while preserving methodologically relevant details, as is standard practice in human-subjects research.

Guideline 5.6: While the open LLM might be under-performing it provides a more stable baseline than proprietary models, which not only change explicitly (model versions), but also implicitly (e.g. internal prompt filtering/reformulation functionalities). We do not dare to predict the future of open LLMs in comparison to proprietary models.

Guideline 5.8: While the costs fluctuate, reporting the input and output tokens enables a preliminary estimation of costs, supporting the study design phase (e.g. for resource allocation).

Comment 3.7:

The paper does not adequately address situations where guidelines conflict or where following all guidelines simultaneously is infeasible due to resource, privacy, or technical constraints. For instance, guideline 5 (human validation) and guideline 6 (open baseline) imply validating both commercial and open models, which increases human effort and cost. Small research groups might not even be able to run open-source baseline models, and models might be outdated compared to the access to commercial models. For studies where commercial models are 2-3x better than open alternatives, the baseline may not provide meaningful scientific insight.

Response:

Please see [Comment 3.2](#) and [Comment 3.6](#). In summary, we have added a remark in Section 5 encouraging researchers to make sensible trade-offs between guidelines and justify deviations. We distinguish between requirements (*must*) and desired practices (*should*) precisely to acknowledge that not all guidelines can always be followed simultaneously. Where two guidelines conflict in practice, researchers should prioritize *must* items and justify omitting *should* items.

Comment 3.8:

Similarly, some guidelines depend on others in non-obvious ways; for example, guideline 5.3 assumes a documented model per guideline 5.2. Guideline 5.7 interacts with guideline 5.6 for the open LLM baseline. Please elaborate or add a dependency diagram or a checklist showing the logical flow of applying the guidelines.

Response:

We agree that visualization could help see dependencies clearer, but we decided against adding it to no further add complexity to the paper. However, we consolidated all cross-references in new “See also” subsections. We further added a table that summarizes rationale and key recommendations for each guideline (Table 4).

Comment 3.9:

While Guideline 5.7 mentions repeating experiments due to non-determinism, there is no systematic guidance on determining appropriate sample sizes or the number of repetitions for different study types. For a researcher working with LLM experiments, what would be a valid statistical sample and parameters to reproduce the experiment?

Response:

For sample sizes, please refer to [Comment 3.4](#). The number of required repetitions depends on factors such as the study type, the variability of the task, and the desired precision of estimates. As with sample sizes for human validation, researchers should justify their chosen number of repetitions, for example through a power analysis [Cohen \(1992\)](#); [Bjarnason et al \(2026\)](#) or by monitoring the convergence of descriptive statistics across incremental runs [Blackwell et al \(2024\)](#). A pilot study can help estimate the expected variability and inform this decision. We clarified this in the paper.

Comment 3.10:

What are the recommendations to compare models or outputs? How to select a metric to compare models or experiments as researchers mitigate data contamination, and metrics evolve?

Response:

The reviewer raises an important and largely open challenge. We expanded Section 5.7 to discuss concrete mitigation strategies, including selecting benchmarks released after model training cutoff dates, using private or newly created test sets, and cross-validating with complementary evaluation methods. We also emphasize the importance of transparent reporting on potential contamination risks so that readers can assess result validity. As this remains an evolving area, we invite community contributions through our living guidelines at <https://llm-guidelines.org>.

Comment 3.11:

While the RFC 2119 distinction is clear in principle, how should it be enforced? What constitutes an acceptable justification for not meeting “should” criteria? How should reviewers handle partial compliance? Provide reviewer guidance or a compliance checklist to operationalize the guidelines in peer review.

Response:

We added a paragraph to Section 5 indicating that reviewers should use the guidelines to help write reviews, emphasizing a focus on doing what is practically possible rather than ideal in principle.

We also added detailed new “advice for reviewers” subsections to each guideline advising reviewers on what to focus on, how to deal with common problems, and how to differentiate minor vs. major problems. We use these sections to encourage good reviewing thinking rather than repeat each guideline’s individual must/should items.

While some of the authors are big fans of checklists and very much agree with Reviewer 3’s instincts here, we think converting the current guidelines into a reviewer checklist is out of scope for this paper and possibly premature. If the reviewer compares our guidelines to one of the SIGSOFT empirical standards checklists, for example, three problems are apparent: (1) the sum of our ‘must’ and ‘should’ items would be much too long to compose a reviewer checklist; (2) our guidelines cover a wide variety of LLM uses in different study types, which all have to be disaggregated and organized into separate usable checklists; (3) guidelines phrases from the author perspective must be significantly reworded and reworked for a reviewer perspective.

This is probably why other widely cited guidelines papers do not also provide reviewer checklists, and why synthesizing reviewer checklists based on them is so much work. In sum, we agree with R3 that reviewer checklists are a natural next step, but we hope R3 can see why we think it’s out of scope for this revision.

Comment 3.12:

While including examples from other domains (healthcare, social sciences) provides context, some readers may question their relevance to SE-specific challenges. Ensure every cross-domain example is explicitly connected back to SE contexts.

Response:

Thanks for pointing this out! We removed examples that were too far away from SE. The remaining examples are directly from SE or are so general that the application to SE is clear.

Comment 3.13:

Please consider the following references:
<https://arxiv.org/html/2406.04244v1>
<https://arxiv.org/abs/2502.00678>

Response:

We have incorporated the reviewer-suggested references at the end of Section 5.7.5.

References

- Argyle LP, Busby EC, Fulda N, Gubler J, Rytting CM, Wingate D (2022) Out of one, many: Using language models to simulate human samples. CoRR abs/2209.06899, DOI 10.48550/ARXIV.2209.06899, URL <https://doi.org/10.48550/arXiv.2209.06899>, 2209.06899
- Austin J, Odena A, Nye MI, Bosma M, Michalewski H, Dohan D, Jiang E, Cai CJ, Terry M, Le QV, Sutton C (2021) Program synthesis with large language models. CoRR abs/2108.07732, URL <https://arxiv.org/abs/2108.07732>, 2108.07732
- Begg C, Cho M, Eastwood S, Horton R, Moher D, Olkin I, Pitkin R, Rennie D, Schulz KF, Simel D, Stroup DF (1996) Improving the quality of reporting of randomized controlled trials. The CONSORT statement. JAMA 276(8):637–639, DOI 10.1001/jama.276.8.637
- Bjarnason BH, Silva A, Monperrus M (2026) On randomness in agentic evals. URL <https://arxiv.org/abs/2602.07150>, 2602.07150
- Blackwell RE, Barry J, Cohn AG (2024) Towards reproducible llm evaluation: Quantifying uncertainty in llm benchmark scores. URL <https://arxiv.org/abs/2410.03492>, 2410.03492
- Chen M, Tworek J, Jun H, Yuan Q, de Oliveira Pinto HP, Kaplan J, Edwards H, Burda Y, Joseph N, Brockman G, Ray A, Puri R, Krueger G, Petrov M, Khlaaf H, Sastry G, Mishkin P, Chan B, Gray S, Ryder N, Pavlov M, Power A, Kaiser L, Bavarian M, Winter C, Tillet P, Such FP, Cummings D, Plappert M, Chantzis F, Barnes E, Herbert-Voss A, Guss WH, Nichol A, Paino A, Tezak N, Tang J, Babuschkin I, Balaji S, Jain S, Saunders W, Hesse C, Carr AN, Leike J, Achiam J, Misra V, Morikawa E, Radford A, Knight M, Brundage M, Murati M, Mayer K, Welinder P, McGrew B, Amodei D, McCandlish S, Sutskever I, Zaremba W (2021) Evaluating large language models trained on code. CoRR abs/2107.03374, URL <https://arxiv.org/abs/2107.03374>, 2107.03374
- Cohen J (1992) A power primer. Psychological Bulletin 112(1):155–159, DOI 10.1037/0033-2909.112.1.155
- Guba EG (1981) Criteria for assessing the trustworthiness of naturalistic inquiries. ECTJ 29(2):75–91, DOI 10.1007/BF02766777
- Schulz KF, Altman DG, Moher D (2010) CONSORT 2010 statement: updated guidelines for reporting parallel group randomised trials. BMJ 340:c332, DOI 10.1136/bmj.c332

Guidelines for Empirical Studies in Software Engineering involving Large Language Models

Sebastian Baltes, Florian Angermeir, Chetan Arora, Marvin Muñoz Barón,
Chunyang Chen, Lukas Böhme, Fabio Calefato, Neil Ernst, Davide Falessi,
Brian Fitzgerald, Davide Fucci, Junda He, Christoph Treude,
Marcos Kalinowski, Stefano Lambiase, Daniel Russo, Mircea Lungu,
Cristina Martinez Montes, Lutz Prechelt, Paul Ralph, Rijnard van Tonder,
and Stefan Wagner

Detailed author information can be found on the next page.

Guidelines for Empirical Studies in Software Engineering involving Large Language Models

Sebastian Baltes · Florian Angermeir ·
Chetan Arora · Marvin Muñoz Barón ·
Chunyang Chen · Lukas Böhme ·
Fabio Calefato · Neil Ernst ·
Davide Falessi · Brian Fitzgerald ·
Davide Fucci · Junda He ·
Christoph Treude ·
Marcos Kalinowski · Stefano Lambiase ·
Daniel Russo · Mircea Lungu ·
Cristina Martinez Montes ·
Lutz Prechelt · Paul Ralph ·
Rijnard van Tonder · Stefan Wagner

Submitted: 26/08/2025

Revised: 27/02/2026

Sebastian Baltes
Heidelberg University, Germany
E-mail: sebastian.baltes@uni-heidelberg.de
(corresponding author)

Florian Angermeir
fortiss, Germany
Blekinge Institute of Technology, Sweden
E-mail: angermeir@fortiss.org

Chetan Arora
Monash University, Australia
E-mail: chetan.arora@monash.edu

Marvin Muñoz Barón · Chunyang Chen
Technical University of Munich, Germany
E-mail: {marvin.munoz-baron, chun-yang.chen}@tum.de

Lukas Böhme
Hasso-Plattner-Institut, Germany
University of Potsdam, Germany
E-mail: lukas.boehme@hpi.de

Fabio Calefato
University of Bari, Italy
E-mail: fabio.calefato@uniba.it

Neil Ernst
University of Victoria, Canada
E-mail: nernst@uvic.ca

Abstract Large Language Models (LLMs) are now ubiquitous in software engineering (SE) research and practice, yet their non-determinism, opaque training data, and rapidly evolving models threaten the reproducibility and replicability of empirical studies. We address this challenge through a collaborative effort of 22 researchers, presenting a taxonomy of seven study types that organizes the landscape of LLM involvement in SE research, together with eight guidelines for designing and reporting such studies. Each guideline distinguishes requirements (MUST) from recommended practices (SHOULD) and is contextualized by the study types it applies to. Our guidelines recommend that researchers: (1) declare LLM usage and role; (2) report model versions,

Davide Falessi

University of Rome Tor Vergata, Italy

E-mail: falessi@ing.uniroma2.it

Brian Fitzgerald

Lero, Ireland

University of Limerick, Ireland

E-mail: brian.fitzgerald@ul.ie

Davide Fucci

Blekinge Institute of Technology, Sweden

E-mail: davide.fucci@bth.se

Junda He · Christoph Treude

Singapore Management University, Singapore

E-mail: {jundahe, ctreude}@smu.edu.sg

Marcos Kalinowski

PUC Rio de Janeiro, Brazil

E-mail: kalinowski@inf.puc-rio.br

Stefano Lambiase · Daniel Russo

Aalborg University in Copenhagen, Denmark

E-mail: {stla, daniel.russo}@cs.aau.dk

Mircea Lungu

IT University of Copenhagen, Denmark

E-mail: mlun@itu.dk

Cristina Martinez Montes

Chalmers University of Technology, Sweden

University of Gothenburg, Sweden

E-mail: montesc@chalmers.se

Lutz Prechelt

Freie Universität Berlin, Germany

E-mail: prechelt@inf.fu-berlin.de

Paul Ralph

Dalhousie University, Canada

E-mail: paulralph@dal.ca

Rijnard van Tonder

Independent Researcher, Antigua and Barbuda

E-mail: rvantonder@gmail.com

Stefan Wagner

Technical University of Munich, Germany

E-mail: stefan.wagner@tum.de

configurations, and customizations; (3) document the tool architecture beyond the model; (4) disclose prompts, their development, and interaction logs; (5) validate LLM outputs with humans; (6) include an open LLM as a baseline; (7) use suitable baselines, benchmarks, and metrics; and (8) articulate limitations and mitigations. We complement the guidelines with an applicability matrix mapping guidelines to study types and a reporting checklist for authors and reviewers. We maintain the study types and guidelines online as a living resource for the community to use and shape (llm-guidelines.org).

Keywords software engineering, large language models, empirical research, meta-science, guidelines

1 Introduction

Since the release of ChatGPT in November 2022, large language models (LLMs) have rapidly become a major focus of software engineering (SE) research [48], yet the reproducibility and replicability of empirical studies involving LLMs remains uncertain. Recent findings indicate a low reproducibility of SE studies involving LLMs [7]. Three characteristics of LLMs make this particularly challenging: their inherent non-determinism causes variability across runs, with slight changes leading to significantly different results [2, 18, 112, 120]; commercial models evolve beyond version identifiers, so reported performance can change over time [23]; and even for “open” models, training data and fine-tuning details often remain undisclosed [38]. Moreover, prompt formatting choices alone can shift accuracy by tens of percentage points [116], and configured parameters such as temperature affect output variability [101]. Hence, not reporting these settings directly affects reproducibility.

Although the SE research community has developed guidelines for conducting and reporting specific types of empirical studies such as controlled experiments (e.g., [117, 137]), their replications (e.g., [107]), or empirical studies in general (e.g., the *ACM SIGSOFT Empirical Standards* [96]), none of these address the LLM-specific aspects described above. Previously, a position paper highlighted these issues [128], but there was no comprehensive community-developed guidance for designing and reporting empirical studies involving LLMs in SE.

Therefore, we present community-developed guidelines for conducting and reporting studies involving LLMs in SE research, co-developed by 22 researchers. After outlining our **Scope**, we introduce a taxonomy of **Study Types**, then present eight **Guidelines**. We complement these with an applicability matrix mapping guidelines to study types and a reporting checklist for authors and reviewers. For each study type and guideline, we identify relevant examples, both within and outside of SE research. We maintain the guidelines online as a living resource for the community to use and shape.¹

¹ <https://llm-guidelines.org>

2 Scope of Guidelines

Software Engineering as our Target Discipline: We target SE research because existing cross-disciplinary guidelines do not address its needs. For instance, Gallifant et al.’s [36] guidelines for LLM use in healthcare research include discipline-specific items irrelevant to most SE studies (e.g., clinical-care usability) while lacking guidance on tool architectures, benchmarking, and code-specific evaluation that SE research requires. Moreover, SE research employs a wide variety of empirical methods [96, 137]. Our work builds on Sallou et al.’s vision paper on mitigating validity threats in LLM-based SE research [106] and our own position paper [128], providing more detailed advice for a wider variety of research methods organized into a taxonomy of *Study Types* (Section 4).

Focus on Text-Based Use Cases: While multi-modal foundation models that use or generate images, audio, or video may also support SE research and practice, we focus on textual use cases of LLMs (e.g., in natural language or programming languages). Many of our guidelines—particularly those concerning model reporting, prompt documentation, and reproducibility—are likely applicable to multi-modal settings as well, but we leave their explicit validation to future work.

Focus on Direct Development or Research Support: While researchers may use LLMs for many peripheral tasks (e.g., proof-reading, spell-checking, translation), our guidelines focus on their direct role in empirical research and engineering practice. For engineers, we focus on the use of LLMs to automate SE tasks, that is, artificial intelligence (AI) for software engineering (AI4SE) (see *LLMs for SE*). This includes agentic systems that autonomously plan and execute multi-step tasks using LLMs (see *Architecture*). For researchers, we focus on the use of LLMs to automate empirical research tasks such as data collection, processing, or analysis (see *LLMs for Research*).

Researchers as our Target Audience: Our guidelines are intended to help SE researchers design, plan, conduct, and report empirical studies involving LLMs, and to support scholarly peer review of such studies. Each guideline includes an *Advice for Reviewers* subsection with targeted assessment suggestions. Our guidelines focus on what to report and how; they complement but do not replace methodological guidance for designing specific types of empirical studies.

How to Navigate this Paper: This paper is structured to support different reading strategies depending on the reader’s goal. *Researchers planning a new study* may start with the taxonomy of study types (see *Study Types*) to identify which types apply to their planned work, then consult Table 3 to determine which guidelines are requirements (MUST) and which are recommendations (SHOULD) for those study types. Each guideline section opens with a

tl;dr summary in a shaded box, allowing readers to quickly assess relevance before reading the full text. *Researchers writing up results* may prefer to start with the checklist in Appendix A, which organizes actionable items by typical paper sections (Introduction, Research Design and Methods, Results, etc.). Table 2 provides a quick reference to all eight guidelines, and Table 1 maps each guideline’s rationale to its key recommendations. *Reviewers* can use the *Advice for Reviewers* subsection at the end of each guideline for targeted guidance on assessing manuscripts. Table 3 helps reviewers identify which guidelines apply to the study type under review.

3 Methodology

This project was initiated at the 2024 meeting of the *International Software Engineering Research Network* (ISERN) in Barcelona, Spain, where the first and last author organized a session on guidelines for empirical studies involving LLMs. These discussions led to a position paper at *WSESE 2025* [128], which was presented at the *2nd Copenhagen Symposium on Human-Centered Software Engineering AI* (CHASEAI 2024), forming a workstream to collaboratively refine and extend the preliminary guidelines. Our process follows an established tradition of developing reporting guidelines through expert collaboration, as exemplified by CONSORT for clinical trials [17] and, more recently, TRIPOD-LLM for LLM-based prediction studies in medicine [36].

We held bi-weekly meetings, in which we decided on the structure of the study type taxonomy and the guidelines. During these discussions, we added a new study type (*LLMs for Synthesis*), refined and extended the guideline sections (in particular *Version and Configuration*, *Architecture*, *Prompts and Logs*) and covered *Benchmarks and Metrics* as well as *Limitations and Mitigations*.

Then, at least two co-authors were assigned to each study type and guideline to refine and extend the content, followed by a peer review and refinement phase. Afterwards, the first author reviewed all sections and discussed potential changes with the topic leads. A requirement for co-authorship was that each author contributed or reviewed content and, most importantly, read the complete guidelines and confirmed that they stand behind all recommendations, not only their own contributions.

To select illustrative examples for each study type and guideline, co-author pairs independently searched for relevant papers using academic databases. This initial pool was extended based on suggestions from other co-authors. Identified papers were assessed by one author and subsequently reviewed by a second author; papers that did not add new insights beyond those already covered were not included. Multiple review rounds were then carried out on the guidelines as a whole, during which selected papers were cross-checked by additional authors and, where necessary, further examples were added or replaced. The examples are intended to be illustrative, not the result of a sys-

Table 1 Overview of study types.

ID	Section	Topic
	4.1	LLMs as Tools for Software Engineering Researchers
S1	4.1.1	LLMs as Annotators
S2	4.1.2	LLMs as Judges
S3	4.1.3	LLMs for Synthesis
S4	4.1.4	LLMs as Subjects
	4.2	LLMs as Tools for Software Engineers
S5	4.2.1	Studying LLM Usage in Software Engineering
S6	4.2.2	LLMs for New Software Engineering Tools
S7	4.2.3	Benchmarking LLMs for Software Engineering Tasks

tematic review; the involvement of multiple authors at different stages helped mitigate individual selection bias.

After completing the internal review, we invited external experts in empirical SE to review the study types and guidelines. During the review process of the *Empirical Software Engineering* journal, the structure and content were improved based on reviewers’ suggestions. We self-assessed our guidelines against the quality dimensions of the *SIGSOFT Empirical Standards* [96]; this assessment is reported in the conclusion.

4 Study Types

Empirical guidelines are needed for SE studies involving LLMs to ensure the validity and reproducibility of their results. However, such guidelines must be tailored to different study types that each pose unique challenges. Therefore, we created a taxonomy of study types to contextualize our recommendations. We use the term *study type* to refer to categories of LLM involvement in empirical research, rather than to research methodologies (e.g., experiments, case studies). A single empirical study may involve multiple such study types. Each study type section starts with a **description**, followed by **examples** from the SE research community and beyond. The **advantages** and **challenges** of using LLMs are discussed in a summary subsection at the end of this section, synthesizing cross-cutting themes across all study types. See Table 1 for an overview of study types.

4.1 LLMs as Tools for Software Engineering Researchers

LLMs can serve as tools to help researchers conduct empirical studies. They can automate various tasks such as data collection, pre-processing, and analysis. For example, LLMs can apply predefined coding guides to qualitative datasets (*LLMs as Annotators*), assess the quality of software artifacts (*LLMs as Judges*), generate summaries of research papers (*LLMs for Synthesis*), or simulate human behavior in empirical studies (*LLMs as Subjects*). If LLMs

could complete these tasks well enough, they would reduce the time and effort required to conduct a study. However, these applications each pose challenges to validity and reproducibility.

4.1.1 LLMs as Annotators (S1)

Description: In qualitative data analysis, manually annotating (“coding”) natural language text (e.g., requirements, interview transcripts, open-ended survey responses) is a time-consuming manual process [13]. LLMs can augment human coding, suggest new codes, and label artifacts based on a pre-defined coding guide much faster than humans can [46]. The extent to which, or under what conditions, LLMs can perform these tasks *effectively* remains an open research question [4]. Indeed, measuring their effectiveness is practically and philosophically challenging. From an interpretivist philosophical perspective, one cannot measure the quality of analysis by comparing one (human or machine) analyst’s work to another. From a realist perspective, triangulating across multiple human judges, LLM judges, and other data sources (e.g., whether a pull request is marked as having resolved an issue) improves confidence in the findings but does not prove that any one judge is valid.

Examples: Huang et al [51] utilized multiple LLMs for joint annotation of mobile application reviews. They used three models of comparable size with an absolute majority voting rule (i.e., a label is only accepted if it receives more than half of the total votes from the models). This approach slightly outperformed the best individual model tested. Meanwhile, Ahmed et al [4] examined LLMs as annotators in SE research across five datasets, six LLMs, and ten annotation tasks. They found that model-model agreement strongly correlates with human-model agreement; models performed poorly in tasks where humans also frequently disagreed. They proposed to use model confidence scores to identify specific samples that could be safely delegated to LLMs, potentially reducing human annotation effort without compromising inter-rater agreement.

4.1.2 LLMs as Judges (S2)

Description: LLMs can rate properties of software artifacts (e.g. assess code readability, adherence to coding standards, or comment quality) or sort multiple solutions by some attribute. These kinds of judgments are distinct from annotating unstructured text (see *LLMs as Annotators*).

Example(s): Lubos et al [72] leveraged Llama-2 to evaluate the quality of software requirements statements. They prompted the LLM with the text below, where the words in braces reflect the study parameters:

```
Your task is to evaluate the quality of a software requirement.
Evaluate whether the following requirement is
{quality_characteristic}.
```

```

1  {quality_characteristic} means:
2  {quality_characteristic_explanation}
3  The evaluation result must be: 'yes' or 'no'.
4  Request: Based on the following description of the project:
5  {project_description}
6  Evaluate the quality of the following requirement:
7  {requirement}.
8  Explain your decision and suggest an improved version.

```

They evaluated LLM output against expert human judges and found moderate agreement for simple requirements and poor agreement for more complex requirements. In contrast, Wang et al. [132] used an LLM to generate acceptance criteria for user stories, and provided a rubric to an LLM to judge the generated acceptance criteria on interpretable scales (0 to 4).

4.1.3 LLMs for Synthesis (S3)

Description: Unlike annotation (see [LLMs as Annotators](#)), which focuses on categorizing or labeling individual data points, synthesis refers to the process of integrating and interpreting information from multiple sources to generate higher-level insights, identify patterns across datasets, and develop conceptual frameworks or theories. LLMs may be able to support synthesis tasks in SE research by processing and distilling information from qualitative data sources. Although synthesis in the preceding notion refers to abstraction and interpretation across multiple data sources, the term is sometimes also used to refer to generating synthetic content (e.g., source code, bug-fix pairs, requirements, etc.) that are then used in downstream tasks to train, fine-tune, or evaluate existing models or tools. In this case, the synthesis is done primarily using the LLM and its training data; the input is limited to basic instructions and examples.

Example(s): Published examples of applying LLMs for synthesis in SE remain scarce; however, some recent work in other domains is instructive [13]. Barros et al. [15] conducted a systematic mapping study on using LLMs for qualitative research and found that most studies focused on the healthcare and social sciences. In these studies, LLMs supported qualitative methods, such as grounded theory and thematic analysis, by aiding in pattern identification. In SE, de Morais Leca et al. [82] explored how LLMs have been applied for qualitative data analysis (QDA) and proposed general strategies and guidelines for their application. Ornelas et al. [88] complemented this perspective by studying the opportunities and limitations of introducing LLM-based support into QDA, and by formulating recommendations for embedding humanAI collaboration across the thematic analysis phases. Building on these, subsequent work has proposed hybrid frameworks combining LLM support with human-led QDA. Rashid et al. [98] designed an LLM-driven multi-agent system that integrates AI with human decision-making to automate qualitative data analysis methods. Their system generated initial codes, developed themes, and summarized text. Similarly, Montes et al. [81] compared the performance of

humans and LLMs in coding, theme development, definition, and refinement, creating guidelines for a hybrid-LLM framework. Finally, [El-Hajjani and Salinesi](#)'s work is an example of using LLMs to create synthetic datasets. They present an approach to generate synthetic requirements, showing that they “*can match or surpass human-authored requirements for specific classification tasks*” [\[34\]](#).

4.1.4 LLMs as Subjects (S4)

Description: In empirical studies, data is collected from participants through methods such as surveys, interviews, or controlled experiments. LLMs can serve as virtual *subjects* by simulating human behavior and interactions. If LLMs can generate responses that approximate those of human participants, they could be valuable for research involving user interactions, collaborative coding environments, and software usability assessments [\[145\]](#). To achieve this, prompt engineering techniques are widely employed; for instance, the *Personas Pattern* [\[63\]](#) involves tailoring LLM responses to align with predefined profiles or roles that emulate specific user archetypes. To serve as virtual subjects, generated responses should be indistinguishable from human-produced texts, consistent with the attitudes and sociodemographic information of the conditioning context (e.g., junior vs. senior developers), naturally aligned with the form, tone, and content of the simulated scenario, and reflect patterns in relationships between ideas, demographics, and behavior observed in comparable human data [\[8\]](#).

Example(s): [Xu et al](#) [\[141\]](#) compiled a list of ways LLMs can support social science research, some of which transfer to empirical SE research. For example, LLMs can emulate human responses and behaviors in simulated interviews and focus groups [Gerosa et al](#) [\[37\]](#). Similarly, [Bano et al](#) [\[14\]](#), investigated biases in LLM-generated candidate profiles in SE recruitment processes. They found biases favoring male candidates, lighter skin tones, and slim physiques, particularly for senior roles. LLMs may be able to simulate end-user feedback and behavior in usability studies, identify usability issues and offering suggestions for improvement based on predefined user personas.

4.2 LLMs as Tools for Software Engineers

LLM-based assistants have become a widely adopted tool for software engineers [\[122\]](#), supporting them in various tasks such as code generation and debugging. Researchers have studied how software engineers use LLMs ([Studying LLM Usage](#)), developed new tools that integrate LLMs ([LLMs for Tools](#)), and benchmarked LLMs for software engineering tasks ([Benchmarking LLMs](#)).

4.2.1 Studying LLM Usage in Software Engineering (S5)

Description: Studying how software engineers use LLMs and LLM-based tools is important for understanding the current state of practice in SE. Researchers can observe software engineers’ usage of LLM-based tools in the field, or study if and how they adopt such tools, their usage patterns, as well as perceived benefits and challenges. Surveys, interviews, observational studies, or analysis of usage logs can provide insights into how LLMs are integrated into development processes, how they influence decision making, and what factors affect their acceptance and effectiveness. Such studies can inform improvements for existing LLM-based tools, motivate the design of novel tools, or derive best practices for LLM-assisted software engineering. They can also uncover risks or deficiencies of existing tools.

Example(s): Based on a convergent mixed-methods study, [Russo](#) has found that early adoption of generative AI by software engineers is primarily driven by compatibility with existing workflows [\[105\]](#). [Khojah et al](#) investigated the use of ChatGPT (GPT-3.5) by professional software engineers in a week-long observational study [\[60\]](#). They found that most developers do not use the code generated by ChatGPT directly but instead use the output as a guide to implement their own solutions. [Azanza et al](#)’s case study found that LLMs could “enhance personalized, instant onboarding support; however, relying on proprietary external LLMs poses significant data privacy risks” [\[11\]](#). [Jahic and Sami](#) found that most participants from 15 software companies they surveyed had already adopted AI (especially ChatGPT) for SE tasks; but cited copyright and privacy issues, as well as inconsistent or low-quality outputs, as barriers to adoption [\[65\]](#). Retrospective studies that analyze data generated while developers use LLMs can provide additional insights into human-LLM interactions. For example, researchers can employ data mining methods to build large-scale conversation datasets, such as the DevGPT dataset introduced by [Xiao et al](#) [\[139\]](#). Conversations can then be analyzed using quantitative [\[94\]](#) and qualitative [\[79\]](#) analysis methods.

4.2.2 LLMs for New Software Engineering Tools (S6)

Description: LLMs are being integrated into new tools that support software engineers in their daily tasks (e.g., to assist in code comprehension [\[142\]](#) and test case generation [\[108\]](#)). One way of integrating LLM-based tools into software engineers’ workflows is using GenAI agents. Unlike traditional LLM-based tools, these agents are capable of acting autonomously and proactively, are often tailored to meet specific user needs, and can interact with external environments [\[126, 136\]](#). From an architectural perspective, GenAI agents can be implemented in various ways [\[136\]](#). However, they generally share three key components: (1) a reasoning mechanism that guides the LLM (often enabled by advanced prompt engineering), (2) a set of tools to interact with external systems (e.g., APIs or databases), and (3) a user communication interface

that extends beyond traditional chat-based interactions [102, 125, 147]. Researchers can also test and compare different tool architectures to increase artifact quality and developer satisfaction.

Example(s): Yan et al. [142] proposed *IVIE*, a tool integrated into the VS Code graphical interface that generates and explains code using LLMs [142]. The authors focused more on the presentation, providing a user-friendly interface to interact with the LLM. Schäfer et al. [108] presented a large-scale empirical evaluation on the effectiveness of LLMs for automated unit test generation. They presented *TestPilot*, a tool that implements an approach in which the LLM is provided with prompts that include the signature and implementation of a function under test, along with usage examples extracted from the documentation. Richards and Wessel [102] introduced a preliminary GenAI agent designed to assist developers in understanding source code by incorporating a reasoning component grounded in the theory of mind [102].

4.2.3 Benchmarking LLMs for Software Engineering Tasks (S7)

Description: Benchmarking is the process of evaluating an LLM's performance on standardized tasks, using standardized metrics, on a standardized dataset. The LLM output is compared to a supposed ground truth from the benchmark dataset. Typical tasks include code generation, code summarization, code completion, and code repair, but also natural language processing tasks, such as anaphora resolution (i.e., the task of identifying the referring expression of a word or phrase occurring earlier in the text). Metrics may include general metrics for text generation, such as *ROUGE*, *BLEU*, or *METEOR* [48], or task-specific metrics, such as *CodeBLEU* for code generation. Benchmarking requires high-quality reference datasets.

Example(s): In SE, benchmarking may include evaluating an LLM's ability to produce accurate and reliable outputs for a given input (usually a task description, possibly accompanied by data obtained from curated real-world projects or from synthetic SE-specific datasets). *RepairBench* [118], for example, contains 574 buggy Java methods and their corresponding fixed versions, which can be used to evaluate the performance of LLMs in code repair tasks. It uses the *Plausible@1* metric (i.e., the probability that the first generated patch passes all test cases) and the *AST Match@1* metric (i.e., the probability that the abstract syntax tree of the first generated patch matches the ground truth patch). *SWE-Bench* [57] is a more generic benchmark that contains 2,294 SE Python tasks extracted from GitHub pull requests. To score the LLM's task performance, the benchmark validates whether the generated patch successfully compiles and calculates the percentage of passed test cases. Meanwhile, *HumanEval* [26] is often used to assess code generation.

4.3 Advantages and Challenges

Using LLMs in empirical SE research, whether as tools for researchers or for software engineers, offers several potential advantages but also raises fundamental challenges that cut across the study types described above.

Advantages: The primary advantage of using LLMs for research tasks is *speed, cost reduction, and scalability*. LLMs can annotate, judge, synthesize, and simulate faster and at lower cost than human researchers, with studies showing cost reductions of 50–96% on various natural language tasks [L33] (e.g., He et al found that GPT-4 annotation required only two days and 122.08 USD compared to several weeks and 4,508 USD for a comparable MTurk pipeline [46]). Similarly, augmenting or replacing human participants with LLM-generated virtual participants would reduce recruitment effort [73]. This efficiency unlocks scalability: qualitative research traditionally does not scale well to large samples, but LLMs allow researchers to analyze larger datasets [L3, L5, 82] and support larger judgment datasets.

LLMs can also *automate* tasks such as coding qualitative data, assessing artifact quality, and generating summaries, reducing cognitive demands and resources required for qualitative research. Some studies suggest that LLMs may improve *consistency*: ChatGPT’s accuracy exceeded crowd workers by approximately 25%, and LLMs can achieve higher inter-rater agreement than crowd workers and trained annotators [39]. However, no compelling evidence supports claims that LLMs are less biased than human annotators [39].

LLMs could potentially provide *access to otherwise-inaccessible research contexts*. If virtual participants’ behavior is sufficiently similar to human participants, LLMs could access underrepresented and hard-to-reach populations, strengthen generalizability and inclusiveness, impute missing data (see *LLMs for Synthesis*), and enable research that is ethically problematic with real humans (e.g., questions that would force a human to relive past trauma do not harm an LLM).

Studying real-world LLM-based tools allows researchers to *understand the state of practice*, uncovering usage patterns, adoption rates, and contextual factors, and *generate hypotheses* about how LLMs affect developer productivity, collaboration, and decision-making. From an engineering perspective, developing LLM-based tools *requires less task-specific engineering* than traditional SE approaches such as static analysis or symbolic execution, because a single model can handle diverse inputs without language-specific parsers or hand-crafted rules. Good benchmarks provide *standardized evaluation and model comparison*, reducing research effort and fostering open science by providing common ground for sharing data and results. Benchmarks built for specific SE tasks can help identify LLM weaknesses and support optimization and fine-tuning.

Challenges: The stochastic nature of LLM responses, where identical prompts may yield different outputs, *undermines replicability* across all study types,

complicating interpretation of experimental results and test-retest reliability. More broadly, *reliability* is a persistent concern: conceptualizing an LLM-as-judge or LLM-as-annotator as a measurement instrument, researchers should expect validity, test-retest reliability, inter-rater reliability, minimal error, and measurement invariance [97]. However, LLMs show significant variability depending on the dataset and task [90], are sensitive to prompt variations [99] and option order [93], can behave differently when reviewing their own outputs [91], and can be unreliable for high-stakes labeling [134]. Crucially, *reliability does not imply validity*: a reliable LLM might be reliably inaccurate [146], and for tasks with no single correct answer, the statistical framework pushes outputs toward the most likely answer, which may not be the best one [16]. See [Version and Configuration](#) and [Limitations and Mitigation](#) for reporting guidance on replicability and reliability.

The *rapid evolution* of LLMs and LLM-based tools, combined with professionals rapidly learning how to employ them, complicates longitudinal comparisons and may quickly render study findings obsolete (e.g., GPT-4’s accuracy on a simple mathematical task dropped from 84% to 51% within three months [25]).

The prevalence of *proprietary tools and opaque training data* limits researchers’ ability to assess and mitigate biases, and enables benchmark contamination [6]. LLMs exhibit *bias* in multiple forms: tendencies to overestimate certain labels [148], well-documented fairness issues [35], and the potential to reinforce prejudices. When simulating human participants, LLMs are “*likely to both misportray and flatten the representations of demographic groups*” [130]. See [Open LLMs](#) and [Limitations and Mitigations](#) for mitigation strategies.

Evaluation difficulty is inherent in studying LLM-based tools and benchmarks. Benchmarks may lack construct validity [97], usually do not capture the full complexity of software engineering work [23], and may lead to *overconfidence* and overfitting [129]. LLMs are also susceptible to adversarial manipulation where semantics-preserving changes may deceive them into accepting flawed artifacts [45]. See [Benchmarks and Metrics](#) for detailed guidance on benchmark design, including Cao et al.’s [22] guidelines for coding task benchmarks.

A fundamental challenge is *philosophical and methodological incongruence* with qualitative research. In a recent open letter, 419 experienced qualitative researchers argued “*that analytic approaches such as reflexive thematic analysis are human research practices requiring a subjective, positioned, and reflexive researcher and therefore the use of GenAI in such approaches is not methodologically congruent*” [58]. LLMs lack the capacity for genuinely reflexive qualitative analysis because they operate on statistical prediction without understanding the meaning of the data being analyzed [12, 58, 81]. Effective use requires structured prompts and careful human oversight [15]; see [Human Validation](#) and [Limitations and Mitigations](#) for guidance.

There is *insufficient evidence* for the effectiveness of LLMs in most research roles. No compelling evidence exists that LLMs can accurately simulate human participants [111] or reliably judge most relevant properties of

SE artifacts; gathering such evidence for each specific usage may be quite difficult [44]. LLMs may be better suited for augmenting rather than replacing human researchers [46, 133], but even then, limited evidence exists that augmentation increases *effectiveness* as well as *efficiency* [111]. When LLM outputs are incorrect, they can negatively *affect human judgment* [50, 89]; see *Human Validation* for mitigation strategies.

Pooling multiple outputs for majority voting improves reliability but increases *cost and environmental impact* [99, 133]. While open models are available, the most capable ones require substantial hardware; relying on *cloud-based APIs* introduces concerns related to data privacy, security, and replicability. See *Limitations and Mitigations* for sustainability considerations.

Field study findings face *generalizability* challenges because outcomes may be highly sensitive to individual differences, usage patterns, goals, and contexts. Field studies must be “dependable” [124] beyond traditional validity criteria, which complicates methodology; see *Limitations and Mitigations* for a detailed threat taxonomy.

5 Guidelines

A primary goal of our guidelines is to *enable reproducibility and replicability* of empirical SE studies involving LLMs. As repeating LLM-focused research to verify results lies somewhere between the ACM’s definitions of reproducibility (different team, same research artifacts) and repeatability (different team, different research artifacts) [11] due to potential model changes and imperfect research artifacts, we follow Angermeir et al. [7] and use the terms interchangeably. While previous guidelines regarding open science and empirical studies still apply, LLM-specific characteristics (e.g. inherent non-determinism [120, 144], opaque and proprietary models) present additional replicability challenges, which, in turn, demand new guidance.

Each guideline below begins with a brief *tl;dr* summary, followed by its **rationale**, **recommendations**, **examples**, **benefits**, **challenges**, links to the relevant **study types**. The *rationale* articulates the underlying principle, i.e., why the guideline matters, while the *recommendations* provide concrete, actionable practices. Table 1 provides a compact overview of this mapping.

The guidelines further contain **advice for reviewers**. Broadly, reviewers should use our guidelines to help them interrogate the extent to which a manuscript’s authors have done what was practically possible to improve reproducibility. We must neither accept research absent reasonable efforts to improve reproducibility, nor reject research for failing to obtain an impossible goal. We borrow this principle of “reasonable efforts” from the SIGSOFT Empirical Standards [96], where it applies to methodological rigor more generally. Of course, papers should acknowledge their limitations, but to determine whether these limitations are reasonable, reviewers should ask “have the authors done what they could to minimize limitations?” Our guidelines attempt to capture what “reasonable efforts” practically means for LLM-based studies.

Table 2 Overview of guidelines.

ID	Section	Guideline
G1	5.1	<u>Declare LLM Usage and Role</u>
G2	5.2	<u>Report Model Version, Configuration, and Customizations</u>
G3	5.3	<u>Report Tool Architecture Beyond Models</u>
G4	5.4	<u>Report Prompts, their Development, and Interaction Logs</u>
G5	5.5	<u>Use Human Validation for LLM Outputs</u>
G6	5.6	<u>Use an Open LLM as a Baseline</u>
G7	5.7	<u>Use Suitable Baselines, Benchmarks, and Metrics</u>
G8	5.8	<u>Report Limitations and Mitigations</u>

To distinguish essential criteria from recommendations, our guidelines use two tiers. A MUST criterion is a *requirement*. Studies that intend to follow these guidelines are expected to meet all MUST criteria. A SHOULD criterion represents a desired practice that strengthens a study’s rigor or transparency. However, there may be valid reasons to deviate in particular circumstances (e.g., resource constraints, inapplicability to a specific study context or type). Nonetheless, studies that deviate from a SHOULD criterion should briefly justify the deviation and discuss its potential impact (e.g., on validity or reproducibility). Table 2 lists our eight guidelines. Table 3 provides an overview of how each guideline applies to each study type. Cells marked ● indicate that the guideline is a MUST requirement for the study type and cells marked ● indicate that the guideline SHOULD be followed. Cells marked – indicate that the guideline is typically not applicable or not feasible for the study type.

The following sections indicate which information we expect researchers to report, and whether it should be in the *paper* or *supplementary material*. Where a publication venue’s page limits hinder reporting all expected elements in the *paper*, it is better to report essential information in the *supplementary material* than not at all. The *supplementary material* should be published according to the *ACM SIGSOFT Open Science Policies* [41].

5.1 Declare LLM Usage and Role (G1)

tl;dr: *Researchers MUST disclose any use of LLMs to support empirical studies in their paper, specifying which LLM was used, how it was used, and where in the research process it was employed. They SHOULD report the exact purpose, the tasks that were automated, and the expected benefits in the paper. When the LLM plays a central role, the declaration SHOULD be prominent and detailed in the methodology section; for tangential uses, a brief statement in the methodology or acknowledgments suffices.*

Table 3 Applicability of guidelines to study types. ● = the guideline’s core recommendations MUST be followed for this study type, ● = SHOULD be followed, – = are not directly applicable. Study type IDs: S1 = [LLMs as Annotators](#), S2 = [LLMs as Judges](#), S3 = [LLMs for Synthesis](#), S4 = [LLMs as Subjects](#), S5 = [Studying LLM Usage](#), S6 = [LLMs for Tools](#), S7 = [Benchmarking LLMs](#). Each guideline’s study-type-specific guidance is detailed in the corresponding subsection.

	S1	S2	S3	S4	S5	S6	S7
G1: Declare Usage	●	●	●	●	●	●	●
G2: Model Version	●	●	●	●	●	●	●
G3: Architecture	●	●	●	●	●	●	●
G4: Prompts	●	●	●	●	●	●	●
G5: Human Valid	●	●	●	●	●	●	●
G6: Open LLM	●	●	●	–	–	●	●
G7: Benchmarks	●	●	●	●	●	●	●
G8: Limitations	●	●	●	●	●	●	●

5.1.1 Rationale:

Transparency about LLM involvement is a prerequisite for informed assessment of a study’s scope, limitations, and potential biases. Without explicit disclosure, readers cannot evaluate how the LLM’s characteristics may have influenced the research process or its outcomes.

5.1.2 Recommendations:

When conducting any kind of empirical study involving LLMs, researchers MUST clearly declare that an LLM was used (see [Scope](#) for what we consider relevant research support). This SHOULD be done in a suitable section of the *paper*, for example, in the introduction or research methods section. For authoring scientific articles, this transparency is, for example, required by the ACM Policy on Authorship: “*The use of generative AI tools and technologies to create content is permitted but must be fully disclosed in the Work*” [9].

Beyond generic declarations, researchers SHOULD report the exact purpose of using an LLM in a study, the tasks it was used to automate, and the expected benefits in the *paper*. A sufficient declaration specifies not only *that* an LLM was used, but also *which* LLM (name and version), *how* it was used (e.g., as an annotator, code generator, or judge), and *where* in the research process it was employed (e.g., data collection, analysis, or synthesis).

When the LLM plays a central role in the study (e.g., as the main tool being evaluated or as a core component of the research method), the declaration SHOULD be prominent and detailed, appearing in the methodology section with cross-references to the specific guidelines that apply (e.g., Sections [Version and Configuration](#), [Architecture](#), and [Prompts and Logs](#)). When the LLM’s role is more tangential (e.g., used for a single preprocessing step), a brief but explicit statement in the methodology or acknowledgments section is sufficient. In

either case, the disclosure MUST be specific enough for readers to assess how the LLM's involvement may affect the study's validity and reproducibility.

5.1.3 *Example(s):*

The *ACM Policy on Authorship* [9] suggests disclosing GenAI usage in the acknowledgments section of the *paper*, advising to “err on the side of caution, and include a disclosure in the acknowledgments section of the Work” when uncertain about the need [9]. For double-blind review, researchers can add a temporary “AI Disclosure” section where the acknowledgments would appear. An example of an LLM disclosure beyond writing support can be found in a recent paper by [Lubos et al](#) [72], in which they write in the methodology section:

“We conducted an LLM-based evaluation of requirements utilizing the Llama 2 language model with 70 billion parameters, fine-tuned to complete chat responses...”

5.1.4 *Benefits:*

Transparency in the use of LLMs helps other researchers understand the context and scope of the study, facilitating better interpretation and comparison of the results. Beyond this declaration, we recommend researchers to be explicit about the LLM version they used (see [Version and Configuration](#)) and the LLM's exact role.

5.1.5 *Challenges:*

Declaring LLM usage requires only a brief statement and no additional experiments, making compliance straightforward. One challenge might be authors' reluctance to disclose LLM usage for valid use cases, because they fear that AI-generated content makes reviewers think that the authors' work is less original. In fact, there is evidence suggesting that AI disclosure can negatively affect trust in authors [109]. However, the *ACM Policy on Authorship* is very clear in that any use of GenAI tools to create content MUST be disclosed. Our guidelines focus on research support beyond proof-reading and writing support (see [Scope](#)), but the threshold of what must be declared continues to evolve as organizations such as the ACM update their authorship policies.

5.1.6 *Study Types:*

Researchers MUST follow this guideline for all study types. The specific focus of the declaration varies by study type. For [LLMs as Annotators](#), [LLMs as Judges](#), [LLMs for Synthesis](#), and [LLMs as Subjects](#), researchers MUST declare the specific role assigned to the LLM (e.g., annotator, judge, synthesizer, or simulated participant). For [Studying LLM Usage](#), researchers MUST clarify

which LLM(s) the observed participants used and under which conditions. For [LLMs for Tools](#), researchers MUST declare the LLM’s role within the tool architecture and its contribution to the tool’s functionality. For [Benchmarking LLMs](#), researchers MUST declare which LLMs were benchmarked and for which tasks.

5.1.7 Advice for Reviewers:

The most common problem with disclosure is incompleteness or vagueness about how the LLM was used. If the paper says “we used LLM X to help with task Y” without specifying how, reviewers should request clarification. Such requests are typically minor revisions unless the missing details may reveal methodological problems.

Using an LLM as a copyeditor becomes problematic when authors do not or cannot take responsibility for the resulting text. If a reviewer finds clear evidence that LLM-generated text was not carefully reviewed (e.g., text like “As a Large Language Model, I...”), this may warrant rejection. If a reviewer suspects an author *cannot* take responsibility for AI-generated text, they should raise the concern with their editor or program chair. Due process requires that authors not be accused of misconduct without clear evidence.

If undisclosed LLM use is suspected, the reviewer should similarly consult their editor or program chair. When the evidence is conclusive, the key question is the degree to which undisclosed use affects the study’s contribution, ranging from negligible (e.g., word choice in a single sentence) to severe (e.g., generating ostensibly empirical data or statistical analyses).

5.1.8 See Also:

- Section [Version and Configuration](#): model version and configuration details.
- Section [Architecture](#): system-level role documentation.
- Section [Prompts and Logs](#): prompt documentation and interaction logs.

5.2 Report Model Version, Configuration, and Customizations (G2)

tl;dr: Researchers MUST report the exact LLM model or tool version, configuration, and experiment date in the paper. When using quantized models, researchers SHOULD report the quantization level and method. For fine-tuned models, they MUST describe the fine-tuning goal, dataset, and procedure in the paper. Researchers SHOULD include default parameters, explain model choices, compare base- and fine-tuned model using suitable metrics and benchmarks, and share fine-tuning data and weights as supplementary material (or alternatively justify in the paper why they cannot share them).

5.2.1 *Rationale:*

LLMs and LLM-based tools are frequently updated, and configuration parameters such as temperature or seed values affect content generation. This guideline focuses on documenting the *model-specific* aspects of empirical studies involving LLMs, concentrating on the models themselves, their version, configuration parameters, and customizations (e.g., fine-tuning). While the *Architecture* section addresses system-level integration, the information outlined here is always essential for reproducibility whenever an LLM is involved.

5.2.2 *Recommendations:*

Researchers MUST document in the *paper* which model or tool version they used in their study, along with the date when the experiments were carried out and the configured parameters that affect output generation. Since default values might change over time, researchers SHOULD always report all configuration values, even if they used the defaults. Checksums and fingerprints SHOULD be reported since they identify specific versions and configurations. Depending on the study context, other properties such as the context window size (number of tokens) SHOULD be reported. When using quantized models, researchers SHOULD report the quantization level (e.g., 4-bit, 8-bit) and method (e.g., GPTQ or AWQ), as different quantization approaches produce different outputs, affecting both output quality and reproducibility. Researchers SHOULD motivate in the *paper* why they selected certain models, versions, and configurations. Reasons may be monetary, technical, or methodological (e.g., planned comparison to previous work). Depending on the specific study context, additional information regarding the experiment or tool architecture SHOULD be reported.

A common customization approach for existing LLMs is fine-tuning. If a model was fine-tuned, researchers MUST describe the fine-tuning goal (e.g., improving the performance for a specific task), the fine-tuning procedure (e.g., full fine-tuning vs. Low-Rank Adaptation (LoRA), selected hyperparameters, loss function, learning rate, batch size, etc.), and the fine-tuning dataset (e.g., data sources, the preprocessing pipeline, dataset size) in the *paper*. Researchers SHOULD either share the fine-tuning dataset as part of the *supplementary material* or explain in the *paper* why the data cannot be shared (e.g., because it contains confidential or personal data that could not be anonymized). The same applies to the fine-tuned model weights. Suitable benchmarks and metrics SHOULD be used to compare the base model with the fine-tuned model.

In summary, our recommendation is to report: (1) Model/tool name and version (MUST in *paper*); (2) All relevant configured parameters that affect output generation (MUST in *paper*); (3) Default values of all available parameters (SHOULD); (4) Checksum/fingerprint of used model version and configuration (SHOULD); (5) Additional properties such as context window size (SHOULD); (6) Model quantization level and method, if applicable (SHOULD).

For fine-tuned models, additional recommendations apply: (1) Fine-tuning goal (MUST in *paper*); (2) Fine-tuning dataset creation and characterization (MUST in *paper*); (3) Fine-tuning parameters and procedure (MUST in *paper*); (4) Fine-tuning dataset and fine-tuned model weights (SHOULD); (5) Validation metrics and benchmarks (SHOULD).

Commercial models (e.g., GPT-5) or LLM-based tools (e.g., ChatGPT) might not give researchers access to all required information. Our suggestion is to report what is available and openly acknowledge limitations that hinder reproducibility.

5.2.3 *Example(s)*:

Based on the documentation that OpenAI and Azure provide [77, 86], researchers might, for example, report:

*“We integrated a **gpt-4** model in version 0125-Preview via the Azure OpenAI Service, and configured it with a temperature of 0.7, top_p set to 0.8, a maximum token length of 512, and the seed value 23487. We ran our experiment on 10th January 2025. The system fingerprint was fp_6b68a8204b.*

[Kang et al.] provide a similar statement in their paper on exploring LLM-based bug reproduction [59]:

“We access OpenAI Codex via its closed beta API, using the code-davinci-002 model. For Codex, we set the temperature to 0.7, and the maximum number of tokens to 256.”

Our guidelines additionally suggest to report a checksum/fingerprint and exact dates, but otherwise this example is close to our recommendations.

[Dhar et al.] [30] assessed whether LLMs can generate architectural design decisions, detailing the system architecture and the LLM’s role within it. They provide information on the fine-tuning approach and datasets, including the source of architectural decision records, preprocessing methods, and data selection criteria.

For self-hosted models, the *supplementary material* can become a true replication package. For example, for models provisioned using **ollama**, one can report the specific tag and checksum, e.g., “*llama3.3, tag 70b-instruct-q8_0, checksum d5b5e1b84868.*” Given suitable hardware, running the model is then as easy as executing the following command:

```
ollama run llama3.3:70b-instruct-q8_0
```

5.2.4 *Benefits*:

Reporting this information is a prerequisite for the verification, reproduction, and replication of LLM-based studies. While LLMs are inherently non-deterministic, this cannot excuse dismissing reproducibility. Although exact reproducibility is hard to achieve, reporting the information outlined in this guideline helps researchers come as close as possible to that standard.

5.2.5 *Challenges:*

Different model providers and modes of operating the models allow for varying degrees of information. For example, OpenAI provides a model version and a system fingerprint describing the backend configuration, which can also influence the output. However, in fact, the fingerprint is intended only to detect changes in the model or its configuration; one cannot go back to a certain fingerprint. As a beta feature, OpenAI lets users set a seed parameter to receive “(mostly) consistent output” [85]. However, the seed value does not allow for full reproducibility and the fingerprint changes frequently. Although, as motivated above, open models significantly simplify re-running experiments, they also come with challenges in terms of reproducibility, as generated outputs can be inconsistent despite setting the temperature to 0 and using a seed value (see [GitHub issue for Llama3](#)). Setting the temperature to 0 configures greedy decoding (always selecting the most probable next token), which minimizes output variability but can degrade quality by producing repetitive text and missing higher-quality responses [47].

Even with a temperature of 0, full determinism is rarely guaranteed: floating-point arithmetic on GPUs causes slight numerical differences that cascade into divergent token selections [144], Sparse Mixture-of-Experts routing amplifies this effect [24], silent backend changes in commercial APIs produce different outputs over time [25], and even self-hosted open models with identical settings do not always yield consistent outputs [4, 120]. Researchers should therefore not treat a temperature of 0 as a guarantee of reproducibility, but as one measure among several, including fixed seed values [85], system fingerprints, and archiving of raw outputs. When a temperature of 0 is chosen primarily for reproducibility, this motivation SHOULD be stated explicitly, along with an acknowledgment of its potential impact on output quality.

5.2.6 *Study Types:*

This guideline MUST be followed for all study types for which the researcher has access to (parts of) the model’s configuration. They MUST always report the configuration that is visible to them, acknowledging the reproducibility challenges of commercial tools and models that are offered as-a-service. Depending on the specific study type, researchers SHOULD provide additional information on the architecture of a tool they built (see [Architecture](#)), prompts and interactions logs (see [Prompts and Logs](#)), and specific limitations and mitigations (see [Limitations and Mitigations](#)).

For example, when [Studying LLM Usage](#) by focusing on commercial tools such as ChatGPT or GitHub Copilot, researchers MUST be as specific as possible in describing their study setup. The configured model name, version, and the date when the experiment was conducted MUST always be reported. In those cases, reporting other aspects, such as prompts and interaction logs, is essential.

For LLMs as Annotators, LLMs as Judges, and LLMs for Synthesis, researchers MUST report the model configuration used for the respective annotation, judging, or synthesis tasks, including temperature and other sampling parameters that affect output variability. For LLMs as Subjects, researchers MUST report any persona-related configuration settings and parameters that shape the simulated behavior. For LLMs for Tools, researchers MUST report the configuration for each model integrated in the tool’s architecture, including any model-specific parameter choices. For Benchmarking LLMs, researchers MUST report the configuration for all benchmarked models to enable fair cross-model comparisons.

5.2.7 Advice for Reviewers:

Missing version, configuration, or parameter information is typically a minor revision request. Before concluding that information is absent, reviewers should check appendices and supplementary materials, as details are sometimes reported there rather than in the main text. Rejection over missing details is rarely warranted unless the omissions obscure deeper methodological problems.

5.2.8 See Also:

- Section Architecture: system-level architecture beyond the model.
- Section Prompts and Logs: prompt reporting and development.
- Section Benchmarks and Metrics: benchmarks for comparing base and fine-tuned models.
- Section Open LLMs: open models for self-hosted reproducibility.
- Section Limitations and Mitigations: reporting reproducibility limitations.

5.3 Report Tool Architecture Beyond Models (G3)

tl;dr: Researchers MUST describe the full architecture of LLM-based tools that they develop in their paper. This includes the role of the LLM, interactions with other components, and the overall system behavior. If autonomous agents are used, researchers MUST specify agent roles, reasoning frameworks, and communication flows. Hosting, hardware setup, and latency implications MUST be reported. For tools using retrieval or augmentation methods, data sources, integration mechanisms, and update and versioning strategies MUST be described. For ensemble architectures, the coordination logic between models MUST be explained in the paper. Researchers SHOULD include architectural diagrams and justify design decisions. Non-disclosed confidential or proprietary components MUST be acknowledged as reproducibility limitations.

5.3.1 *Rationale:*

LLM-based tools often have *complex software layers* around the model(s) that pre-process data, prepare prompts, filter user requests, or post-process responses. For example, ChatGPT and GitHub Copilot use the same underlying models, but their outputs differ significantly because Copilot automatically adds project context; researchers can also build tools using models directly via APIs. The infrastructure and business logic around the bare model can significantly contribute to the performance of a tool for a given task, for instance, RAG pipelines, output parsers, and retry logic all shape the final output independently of the model itself.

This section addresses the *system-level* aspects of LLM-based tools that researchers develop, complementing [Version and Configuration](#), which focuses on model-specific details. While standalone LLMs require limited architectural documentation, a more detailed description is required for study setups and tool architectures that integrate LLMs with other components to create more complex systems. This section provides guidelines for documenting these broader architectures.

5.3.2 *Recommendations:*

Researchers **MUST** clearly describe the tool architecture and what exactly the LLM (or ensemble of LLMs) contributes to the tool or method presented in a research paper. Researchers **SHOULD** provide a high-level architectural diagram to improve transparency. To improve clarity, researchers **SHOULD** explain design decisions, particularly regarding how the models were hosted and accessed (API-based, self-hosted, etc.) and which retrieval mechanisms were implemented (keyword search, semantic similarity matching, rule-based extraction, etc.). Researchers **MUST NOT** omit critical architectural details that could affect reproducibility, such as dependencies on proprietary tools that influence tool behavior. Especially for *time-sensitive measurements*, the description of the hosting environment is central, as it can significantly impact the results. Researchers **MUST** clarify whether local infrastructure or cloud services were used, including detailed infrastructure specifications and latency considerations.

Agent-Based and Complex Systems: If an LLM is used as a *standalone system*, for example, by sending prompts directly to a GPT-4o model via the OpenAI API without pre-processing the prompts or post-processing the responses, a brief explanation of this approach is usually sufficient. However, if LLMs are integrated into more *complex systems* with pre-processing, retrieval mechanisms, or autonomous agents, researchers **MUST** provide a detailed description of the system architecture in the *paper*. Aspects to consider are how the LLM interacts with other components such as databases, external APIs, and frameworks. If the LLM is part of an *agent-based system* that autonomously plans or executes tasks, researchers **MUST** describe its exact architecture, including the

agents' roles (e.g., planner, executor, coordinator), whether it is a single-agent or multi-agent system, how it interacts with external tools and users, and the reasoning framework used (e.g., chain-of-thought, self-reflection, multi-turn dialogue). Any agent configuration via context files (e.g., `AGENTS.md`) SHOULD also be documented as part of the system architecture (see [Prompts and Logs](#) for detailed reporting requirements). For agent-based systems that use external tools (e.g., Claude Code and its [subagents](#)), researchers MUST discuss *agent behavior traceability* by clearly delineating three distinct components: (1) *LLM Input/Output*: Internal deliberation, planning, and interpretation performed by the model. (2) *External tool calls*: Specific invocations of APIs, databases, file systems, or other external services that the agent explicitly triggers (e.g., via the [Model Context Protocol \(MCP\)](#)). (3) *User or system interactions*: Human-in-the-loop feedback, environment responses, or multi-agent communication. Researchers SHOULD report complete execution traces that show the sequence and causality between these components, including inputs and outputs of all external tool calls. This traceability is essential for determining whether task success is attributable to the LLM's output and tool-calling capabilities, the external tools' functionality, or their interaction patterns. When full tool call logging is not feasible due to privacy or proprietary constraints, researchers SHOULD provide representative examples or anonymized traces that demonstrate the agent's decision-making process.

Retrieval, Augmentation, and Ensembles: If *retrieval or augmentation methods* were used (e.g., retrieval-augmented generation (RAG), rule-based retrieval, structured query generation, or hybrid approaches), researchers MUST describe how external data is retrieved, stored, and integrated into the LLM's responses. This includes specifying the type of storage or database used (e.g., vector databases, relational databases, knowledge graphs) and how the retrieved information is selected and used. Stored data used for context augmentation MUST be reported, including details on data preprocessing, versioning, and update frequency. If this data is not confidential, an anonymized snapshot of the data used for context augmentation SHOULD be made available as *supplementary material*. For *ensemble models*, in addition to following the [Version and Configuration](#) guideline for each model, the researchers MUST describe the architecture that connects the models. The *paper* MUST at least contain a high-level description, and details can be reported in the *supplementary material*. Aspects to consider include documenting the logic that determines which model handles which input, the interaction between models, and the architecture for combining outputs (e.g., majority voting, weighted averaging, sequential processing).

5.3.3 Example(s):

[Schäfer et al.](#) [108] evaluated LLMs for automated unit test generation, providing a comprehensive description of the system architecture including code parsing, prompt formulation, LLM interaction, and test suite integration. They

also detail the datasets used, including sources, selection criteria, and preprocessing steps.

A second example is [Van et al.](#)'s *IVIE* tool [142], which integrates LLMs into the VS Code interface. The authors document the tool architecture, detailing the IDE integration, context extraction from code editors, and the formatting pipeline for LLM-generated explanations. This documentation illustrates how architectural components beyond the core LLM affect the overall tool performance and user experience.

5.3.4 *Benefits:*

The software layers around bare LLMs significantly impact tool performance and hence need detailed documentation. Documenting the architecture and supplemental data of LLM-based systems enhances reproducibility and transparency [7], enabling experiment replication, result validation, and cross-study comparison.

5.3.5 *Challenges:*

Researchers face challenges in documenting LLM-based architectures, including proprietary APIs and dependencies that restrict disclosure, managing large-scale retrieval databases, and ensuring efficient query execution. They must also balance transparency with data privacy concerns, adapt to the evolving nature of LLM integrations, and, depending on the context, handle the complexity of multi-agent interactions and decision-making logic, all of which can impact reproducibility and system clarity.

5.3.6 *Study Types:*

This guideline MUST be followed for all studies that involve tools with system-level components beyond bare LLMs, from lightweight wrappers that preprocess user input or post-process model outputs, to systems employing retrieval-augmented methods or complex agent-based architectures.

For [LLMs for Tools](#), this guideline is of primary importance: researchers MUST describe the tool's full architecture, including the role of each LLM, interactions with other components, and the overall system behavior. For [Benchmarking LLMs](#), researchers SHOULD describe the evaluation harness and infrastructure if it goes beyond bare model API calls (e.g., custom sandboxing, orchestration layers, or post-processing pipelines). For [Studying LLM Usage](#), researchers SHOULD describe the tool architecture of the studied tool to the extent that it is accessible, as architectural details may influence observed usage patterns. For [LLMs as Annotators](#), [LLMs as Judges](#), [LLMs for Synthesis](#), and [LLMs as Subjects](#), if the research setup involves a custom pipeline (e.g., retrieval-augmented generation for annotation or chained prompts for synthesis), the architecture SHOULD also be reported.

5.3.7 Advice for Reviewers:

As with other guidelines, missing architectural information is typically a minor revision request unless so much is missing that methodological rigor cannot be assessed. Reviewers may ask authors to move key details from supplements into the paper body to ensure the main text is self-contained.

Regarding proprietary or confidential components, three principles apply: (1) empirically evaluating an opaque artifact is a valid scientific contribution; (2) the less available and inspectable the artifact, the weaker the contribution; (3) *scientific instruments* including tools, metrics, scales, and experimental materials, must be fully disclosed, even when the object under study cannot be.

5.4 Report Prompts, their Development, and Interaction Logs (G4)

tl;dr: Researchers MUST publish all prompts or, when using templates, prompt templates with representative instances, including their structure, content, formatting, and dynamic components, as supplementary material. If full prompt disclosure is not feasible, for example, due to privacy or confidentiality concerns, summaries or examples SHOULD be provided. Prompt development strategies (e.g., zero-shot, few-shot), rationale, and selection process MUST be described in the paper. When prompts are long or complex, input handling and token optimization strategies MUST be documented. For dynamically generated or user-authored prompts, generation and collection processes MUST be reported. Prompt reuse across models and configurations MUST be specified. Researchers SHOULD report prompt revisions and pilot testing insights. For agentic systems, developed plans SHOULD be reported as supplementary material and interaction logs SHOULD be generalized to include human-agent interaction traces as supplementary material. All context files used to configure AI agent behavior SHOULD be reported as supplementary material. To address model non-determinism and ensure reproducibility, especially when targeting SaaS-based commercial tools, full interaction logs (prompts and responses) SHOULD be included as supplementary material if privacy and confidentiality can be ensured.

5.4.1 Rationale:

Prompts are critical for any study involving LLMs [116]. A *prompt* is a concrete input to an LLM that guides its output [113]. Depending on the task, a prompt may include instructions (e.g., “*classify the following bug report*”), context (e.g., a taxonomy of bug categories, related source code), input data (e.g., the bug report text), and output format specifications (e.g., “*respond as JSON with fields ‘category’ and ‘justification’*”), with outputs ranging from unstruc-

1 tured text to structured formats such as JSON.² A *prompt template* is a param-
2 eterized structure containing static elements (e.g., instructions, output format
3 specifications) and placeholders for dynamic content (e.g., source code under
4 analysis) that are filled in at runtime to construct concrete prompts [113]. In
5 automated studies, researchers typically design prompt templates from which
6 individual prompts are then instantiated. Prompts significantly influence the
7 quality of the output of a model, and understanding how exactly they were
8 formatted and integrated into an LLM-based study is essential to ensure trans-
9 parency, verifiability, and reproducibility. Sclar et al. question the methodolog-
10 ical validity of comparing models with “an arbitrarily chosen, fixed prompt
11 format”, because their research has shown that the performance of different
12 prompt formats only weakly correlates between different models [116]. We re-
13 mind the reader that these guidelines do not apply when LLMs are used solely
14 for language polishing, paraphrasing, translation, tone or style adaptation, or
15 layout edits (see *Scope*). This implies that our guidelines do not recommend
16 that researcher disclose the prompts they have used for such tasks.

19 5.4.2 Recommendations:

20 *Prompt Reporting:* Researchers MUST report all prompts that were used for
21 an empirical study, including all instructions, context, input data, and output
22 indicators. The only exception to this is if transcripts might identify anony-
23 mous participants or reveal personal or confidential information, for example,
24 when working with industry partners. When prompt templates are used, re-
25 searchers MUST report them alongside representative prompt instances, spec-
26 ifying which parts are static and which are dynamically filled. Moreover, re-
27 searchers SHOULD specify the exact formatting of prompts, including how code
28 snippets were enclosed (e.g., markdown-style code blocks, triple backticks),
29 whether whitespace was preserved, and how other artifacts such as error mes-
30 sages and stack traces were presented.

31 *Prompt Development:* Researchers MUST also explain in the *paper* how they
32 developed the prompts and why they decided to follow certain prompting
33 strategies. If prompts from the early phases of a research project are unavail-
34 able, they MUST at least summarize the prompt evolution. However, given that
35 prompts can be stored as plain text, there are established SE techniques such
36 as version control systems that can be used to collect the required provenance
37 information.

38 Prompt development is often iterative, involving collaboration between hu-
39 man researchers and AI tools. Researchers SHOULD report any instances in
40 which LLMs were used to suggest prompt refinements and how these sug-
41 gestions were incorporated. Prompts may need revision in response to fail-
42 ure cases, and pilot testing is vital to ensure reliable results. If such testing
43 was conducted, researchers SHOULD summarize key insights, including how

44 ² <https://www.promptingguide.ai/introduction/elements>

different prompt variations affected output quality and which criteria were used to finalize the prompt design. A prompt changelog can track prompt evolution, including key revisions and reasons for changes (e.g., v1.0: initial prompt; v2.0: incorporated examples of ideal responses). Since prompt effectiveness varies between models and model versions, researchers MUST make clear which prompts were used for which models in which versions and with which configuration.

Prompting Strategy and Input Handling: Researchers MUST specify whether zero-shot, one-shot, or few-shot prompting was used. For few-shot prompts, the examples provided to the model SHOULD be clearly outlined, along with the rationale for selecting them. If multiple versions of a prompt were tested, researchers SHOULD describe how these variations were evaluated and how the final design was chosen.

When dealing with extensive or complex prompt context, researchers MUST describe the strategies they used to handle input length constraints. Approaches might include truncating, summarizing, or splitting prompts into multiple parts. Token optimization measures, such as simplifying code formatting or removing unnecessary comments, MUST also be documented if applied.

Dynamic and User-Authored Prompts: In cases where prompts are generated dynamically, such as through preprocessing, template structures, or retrieval-augmented generation (RAG), the process MUST be thoroughly documented. This includes explaining any automated algorithms or rules that influenced prompt generation. For studies involving human participants who create or modify prompts, researchers MUST describe how these prompts were collected and analyzed.

Supplementary Material and Anonymization: To ensure complete reproducibility, researchers MUST make all prompts and prompt variations publicly available as part of their *supplementary material*. If the complete set of prompts cannot be included in the *paper*, researchers SHOULD provide summaries and representative examples. This also applies if complete disclosure is not possible due to privacy or confidentiality concerns. For prompts containing sensitive information, researchers MUST: (1) anonymize personal identifiers, (ii) replace proprietary code with placeholders, and (iii) clearly highlight modified sections.

Interaction Logs: When trying to verify results, even with the exact same prompts, decoding strategies, and parameters, LLMs can still behave non-deterministically. Non-determinism can arise from batching, input preprocessing, and floating point arithmetic on GPUs [24]. Thus, in order to enable other researchers to verify the conclusions that researchers have drawn from LLM interactions, researchers SHOULD report the full interaction logs (prompts and responses) as part of their *supplementary material*. Reporting this is especially

important for studies targeting commercial software-as-a-service (SaaS) solutions such as ChatGPT. The rationale for this is similar to the rationale for reporting interview transcripts in qualitative research. Just as a human participant might give different answers to the same question asked two months apart, the responses from tools such as ChatGPT can also vary over time.

Agentic Systems: For agentic systems that autonomously plan and execute tasks, the developed plans SHOULD be reported as *supplementary material* if available, as they document the sequence of actions the system chose to pursue and the intermediate goals it set. Interaction logs SHOULD be generalized to include full interaction traces between humans and agentic systems, capturing human-in-the-loop feedback, approval or rejection decisions, and iterative refinements. These traces SHOULD also be reported as *supplementary material* so that readers can reconstruct which actions the agent took, verify human oversight decisions, and assess the quality of tool-generated outputs.

Agentic coding tools can be configured using version-controlled Markdown files such as `AGENTS.md`, which contain persistent, project-specific instructions [80]. These *context files* are automatically injected into prompts and can significantly influence agent behavior by specifying coding conventions, architectural constraints, tool preferences, and task-specific instructions. Researchers SHOULD report all context files used to configure AI agent behavior as part of their *supplementary material*, as these files are a form of prompt artifact that shapes agent behavior similarly to system prompts.

5.4.3 Example(s):

A paper by Anandayuvraj et al [6] is a good example of making prompts available online. In that paper, the authors analyze software failures reported in news articles and use prompting to automate tasks such as filtering relevant articles, merging reports, and extracting detailed failure information. Their online appendix contains all the prompts used in the study, providing transparency and supporting reproducibility.

Liang et al's paper is another good example of comprehensive prompt reporting. The authors make the exact prompts available in their *supplementary material* on Figshare, including details such as code blocks being enclosed in triple backticks. The *paper* thoroughly explains the rationale behind the prompt design and the data output format. It also includes an overview figure and two concrete examples, enabling transparency and reproducibility while keeping the main text concise.

An example of reporting full interaction logs is the study by Ronanki et al [103], for which they reported the full answers of ChatGPT and uploaded them to Zenodo.

5.4.4 *Benefits:*

Detailed prompt documentation improves verifiability, reproducibility, and comparability of LLM-based studies, allowing other researchers to replicate studies, refine prompts, and evaluate how different content types and formatting choices influence LLM behavior.

Unlike human participant conversations, which often cannot be reported due to confidentiality, LLM interaction logs can be shared. This enables reproduction studies, tracking of response changes over time or across model versions, and secondary research on LLM consistency for specific SE tasks.

5.4.5 *Challenges:*

One challenge is the complexity of prompts that combine multiple components, such as code, error messages, and explanatory text. Formatting differences (e.g., Markdown vs. plain text) can affect how LLMs interpret input. Additionally, prompt length constraints may require careful context management, particularly for tasks involving extensive artifacts such as large codebases. Privacy and confidentiality concerns can hinder prompt sharing, especially when sensitive data is involved.

Not all systems allow the reporting of complete (system) prompts and interaction logs with ease. This hinders transparency and verifiability. We also recommend exploring open-source tools such as *OpenCode* [87]. For commercial tools, researchers MUST report all available information and acknowledge unknown aspects as limitations. Understanding suggestions of commercial tools such as *GitHub Copilot* might require recreating the exact state of the codebase at the time the suggestion was made, a challenging context to report. One solution could be to use version control to capture the exact state of the codebase when a recommendation was made, keeping track of the files that were automatically added as context.

5.4.6 *Study Types:*

Researchers MUST report the complete prompts that were used, including all instructions, context, and input data. This applies across all study types, but reporting requirements and specific focus areas vary.

For *LLMs for Tools*, researchers MUST explain how prompts were generated and structured within the tool. For *Studying LLM Usage*, especially for controlled experiments, exact prompts MUST be reported for all conditions. For observational studies, especially the ones targeting commercial tools, researchers MUST report the full interaction logs except for when transcripts might identify anonymous participants or reveal personal or confidential information. As mentioned before, if the complete interaction logs cannot be shared, e.g., because they contain confidential information, the prompts and responses MUST at least be summarized and described in the *paper*. For *LLMs as Annotators*, researchers MUST document any predefined coding guides or

instructions included in prompts, as these influence how the model labels artifacts. For *LLMs as Judge*, researchers MUST report the evaluation criteria, scales, and examples embedded in the prompt to ensure a consistent interpretation. For *LLMs for Synthesis* tasks (e.g., summarization, aggregation), researchers MUST document the sequence of prompts used to generate and refine outputs, including follow-ups and clarification queries. For *LLMs as Subject* (e.g., simulating human participants), researchers MUST report any role-playing instructions, constraints, or personas used to guide LLM behavior. For *Benchmarking LLMs* studies using pre-defined prompts (e.g., *HumanEval*, *SWE-Bench*), researchers MUST specify the benchmark version and any modifications made to the prompts or evaluation setup. If prompt tuning, retrieval-augmented generation (RAG), or other methods were used to adapt prompts, researchers MUST disclose and justify those changes, and SHOULD make the relevant code publicly available.

5.4.7 Advice for Reviewers:

As with other guidelines, missing prompt information is typically a minor revision request unless so much is missing that methodological rigor cannot be assessed. The most challenging aspect is the description of prompt development, which is inherently iterative and creative. Reviewers should *not* expect justification of each word choice or post-hoc rationalized accounts of prompt generation. Instead, reviewers should focus on whether (1) a new research team could *use* (not reproduce) exactly the same prompts in exactly the same way; and (2) potential biases or validity issues are transparent.

5.4.8 See Also:

- Section *Version and Configuration*: model versions and configuration for prompt-model mapping.
- Section *Architecture*: agent behavior traceability requirements.
- Section *Human Validation*: evaluation of agentic tool outputs.
- Section *Open LLMs*: open-source tools for transparent prompt logging.
- Section *Limitations and Mitigations*: reporting limitations of commercial tool transparency.

5.5 Use Human Validation for LLM Outputs (G5)

tl;dr: *If assessing the quality of generated artifacts is important and no reference datasets or suitable comparison metrics exist, researchers SHOULD use human validation for LLM outputs. If they do, they MUST define the measured construct (e.g., usability, maintainability) and describe the measurement instrument in the paper. Researchers SHOULD consider human validation early in the study design (not as an afterthought) and build on*

established reference models for human-LLM comparison. When developing or adapting measurement instruments, researchers MUST share them. When aggregating LLM judgments, methods and rationale SHOULD be reported and inter-rater agreement SHOULD be assessed. Confounding factors SHOULD be controlled for, and power analysis SHOULD be conducted to ensure statistical robustness.

5.5.1 *Rationale:*

Automated metrics alone cannot fully capture the validity of subjective or complex constructs. When LLMs are used to automate tasks that involve human judgment, their outputs must be validated against human assessments to ensure that the results are not only reliable but also valid.

5.5.2 *Recommendations:*

When LLMs are used to automate research or software development tasks previously performed by humans, it is essential to assess the LLM's performance. When existing reference datasets or comparison metrics cannot fully capture the target qualities relevant in the study context, researchers SHOULD rely on human judgment to validate the LLM outputs.

Study Design Considerations: Integrating human participants in the study design will require additional considerations, including a recruitment strategy, annotation guidelines, training sessions, or ethical approvals. Therefore, researchers SHOULD consider human validation early in the study design, not as an afterthought. Authors MUST clearly define in the *paper* the constructs that the human and LLM annotators evaluate [95]. When designing custom instruments to assess LLM output (e.g., questionnaires, scales), researchers MUST share their instruments.

Subjective Judgment and Agreement: When the judgment is subjective (i.e., depends on the judge's values or theories), the LLM output SHOULD be evaluated against judgments aggregated from multiple human judges, and researchers SHOULD clearly describe their aggregation method and reasoning. The preferred method is to randomly order the objects requiring judgment and put them in groups small enough to judge in one to three hours. Two or three human experts review the first group together, simultaneously judging, discussing, and creating a set of decision rules documenting their reasoning. Then, the experts iterate among rounds of:

1. independently rating one group of objects,
2. calculating inter-rater agreement (IRA) or reliability (IRR),
3. meeting to discuss disagreements and reach consensus,
4. updating the decision rules so the disagreement will not recur.

The goal is to reach a target IRA or IRR (e.g., Krippendorff's $\alpha > 0.8$), indicating sufficient decision rules. Once this target is reached and sustained for two to three groups, it is permissible to continue with a single rater. Researchers SHOULD report measures of IRA or IRR, preferably broken down by round.

Confounding factors SHOULD be discussed and, where feasible, controlled for (e.g., by categorizing participants according to their level of experience or expertise). Where applicable, researchers may perform a power analysis [29, 32] to estimate the required sample size, ensuring sufficient statistical power in their experimental design. However, we also note that to our knowledge, there is limited established guidance specific to determining sample sizes for LLM-human comparisons. In other fields, 100 comparisons are often made without further explanation [127].

Researchers SHOULD use established reference models to compare humans with LLMs. For example, Schneider et al [110] outline design considerations for studies comparing LLMs with humans. If studies aim to automate annotating software artifacts using an LLM, researchers SHOULD follow systematic approaches to decide whether and how human annotators can be replaced (e.g., showing that a jury of three LLMs exhibit model-to-model agreement comparable to trained human annotators, such as Krippendorff's $\alpha > 0.8$). Model-to-model agreement should not be held to a lower standard than human-to-human agreement, as low thresholds may simply reflect shared model biases rather than annotation quality.

Agentic Tools Besides generating and modifying content, agentic software development tools such as *Claude Code* can autonomously call command-line tools or pull in additional information from MCP servers. For such tools, it makes sense to separate textual output (e.g., file changes) from output related to tool execution. When evaluating agentic tools, researchers SHOULD assess the feedback that users provided for proposed changes, report statistics on how frequently they accepted the content, and how they modified it. This agentic human-in-the-loop interaction approach is essentially a built-in human validation, even though the degrees of freedom are larger than in more traditional experiments that use LLMs directly.

5.5.3 *Example(s)*:

Ahmed et al [4] proposed a systematic method for deciding whether LLMs can replace human annotators on a given task, using model-to-model agreement as an initial screening criterion (with a threshold of $\alpha > 0.5$) and model confidence for sample-level decisions. However, their threshold is well below the levels generally considered acceptable for inter-rater agreement. Krippendorff recommends discarding data with $\alpha < 0.667$, considers $0.667 \leq \alpha < 0.8$ sufficient only for tentative conclusions, and requires $\alpha \geq 0.8$ for reliable data [64]. Researchers adopting similar screening approaches should apply thresholds

consistent with these established standards. Notably, while three LLMs exhibited higher inter-model agreement than human annotators on some tasks, human-model agreement remained low on others. This illustrates that high model-model reliability does not guarantee alignment with human judgment, reinforcing the need for human validation before assuming that LLM annotations are valid. Moreover, a high model-model agreement could merely indicate that the models share systematic biases and hence reliably agree on the wrong answer.

Hymel and Johnson [54] [54] evaluated ChatGPT’s capability to generate requirements documents by comparing an LLM-generated and a human-generated document based on the same business use case. Domain experts reviewed both documents and attempted to distinguish their origin.

5.5.4 *Benefits:*

Validating LLMs against human judgments builds confidence in the LLM’s accuracy and validity, increasing the trustworthiness of findings, especially when IRA or IRR metrics are reported [61]. Furthermore, incorporating feedback from human judges may help to improve the LLM or LLM-based tool based on the reported experiences.

5.5.5 *Challenges:*

Assessing a construct like LLM performance is challenging because ensuring that a construct is defined well and operationalized using an appropriate measurement model requires a deep understanding of (1) the construct, (2) construct validity in general, and (3) instrumentation [97, 119]. Comparing an LLM to human judges is typically slower and more expensive than machine-generated measures. More fundamentally, neither human judgment nor machine-generated measures provides an objective ground truth against which LLM accuracy can be firmly determined.

Human judgments exhibit variability due to differences in experience, expertise, interpretations, and personal biases [74]. When diverse humans rate items reliably given clear decision rules, we assume that reliability implies validity, but it does not. Measuring reliability is *much* easier than measuring validity, and often the best researchers can do is argue conceptually for why their judges, decision rules, and constructs *should* produce valid ratings.

Researchers must weigh the cost and time need for human validation against the expected corresponding improvement in validity over machine-generated measures.

5.5.6 *Study Types:*

This guideline applies to all study types, although the need for human validation varies. For LLMs as Annotators, researchers SHOULD validate LLM-generated annotations against human annotators to assess labeling quality and

identify systematic biases. When using *LLMs as Judges*, researchers SHOULD co-create initial rating criteria with humans and validate a sample of LLM judgments against human expert assessments. For *LLMs for Synthesis*, researchers SHOULD employ human oversight to verify that qualitative interpretations and synthesized outputs faithfully represent the underlying data. For *LLMs as Subjects*, researchers SHOULD validate simulated responses against real human data to assess the fidelity of the simulation. For *Studying LLM Usage*, researchers SHOULD carefully reflect on the validity of their evaluation criteria and validate subjective assessments with human experts. In developing *LLMs for Tools*, human validation of the LLM output is critical when the output is supposed to match human expectations. For *Benchmarking LLMs*, there is less need for human validation when using extensively validated and widely-used benchmarks, but researchers SHOULD employ human validation when creating or adapting new benchmarks.

5.5.7 Advice for Reviewers:

Human validation may be the most challenging of our guidelines to assess because it often requires evaluating conceptual arguments. If LLM output is validated only by comparison with other LLMs, reviewers should look for *quantitative empirical evidence* that such comparison is reliable *and* valid. High inter-model agreement alone is insufficient, as reliability does not imply validity. Similarly, reviewers should expect evidence that any employed benchmarks are reliable and valid. Absent such evidence, human validation is warranted. A single human judge is appropriate only when judgments depend on widely accepted theories and involve limited value conflict (e.g., tagging method names containing abbreviations). For multiple judges, reviewers should expect IRA/IRR improvement techniques as described in the recommendations above (experienced raters, organized rounds, consensus meetings, updated decision rules). Low IRA or IRR (e.g., Krippendorff's $\alpha < 0.8$) without these techniques is a concern. Conversely, if authors have followed best practices and still obtained mediocre results (e.g., $0.66 < \alpha < 0.8$), this should be noted as a limitation. Beyond reliability, reviewers should expect authors to explain conceptually why their human judgments should be valid, considering construct definitions, decision rules, and judge expertise. As with other guidelines, missing information is typically a revision request. Absent judge instructions, instruments, decision rules, or construct definitions may prevent assessment of rigor and validity, while missing recruitment details are less critical. Clarification requests about construct definitions are routine and should not alone warrant rejection.

5.6 Use an Open LLM as a Baseline (G6)

tl;dr: Researchers *SHOULD* include an open LLM as a baseline when using commercial models and report inter-model agreement. By open LLM, we mean a model that everyone with the required hardware can deploy and operate. Such models are usually “open weight.” A full replication package with step-by-step instructions *SHOULD* be provided as part of the supplementary material.

5.6.1 *Rationale:*

Reproducibility depends on access to the model under study. When research relies exclusively on proprietary models, other researchers cannot independently verify or build upon the findings. Including an open LLM as a baseline ensures that at least part of the study can be fully replicated.

5.6.2 *Recommendations:*

Empirical studies using LLMs in SE, especially those that target commercial tools or models, *SHOULD* incorporate an open LLM as a baseline and report established metrics for inter-model agreement. We acknowledge that including an open LLM baseline might not always be possible, for example, if the study involves human participants, and letting them work on the tasks using two different models might not be feasible. Using an open model as a baseline is also not necessary if the use of the LLM is tangential to the study goal.

Open models allow other researchers to verify research results and build upon them, even without access to commercial models. A comparison of commercial and open models also allows researchers to contextualize model performance. Researchers *SHOULD* provide a complete replication package as part of their *supplementary material*, including clear step-by-step instructions on how to verify and reproduce the results reported in the paper.

Open LLMs are available on platforms such as [Hugging Face](#) and can be hosted locally using frameworks such as [Ollama](#) or [LM Studio](#), or on cloud services such as [Together AI](#), AWS, Azure, and Google Cloud.

The term “open” can have different meanings in the context of LLMs. [Widder et al.](#) discuss three types of openness (transparency, reusability, and extensibility) and what openness can and cannot provide [135]. The *Open Source Initiative* (OSI) [84] defines open-source AI as having access to everything needed to understand, modify, share, retrain, and recreate the model.

Researchers can also explore open-source tools such as Continue, Cline, and opencode as alternatives to commercial tools like GitHub Copilot and Claude Code, enabling instrumentation, architectural transparency, and detailed telemetry collection.

5.6.3 *Example(s):*

An increasing number of studies have adopted open LLMs as baseline models. For example, [Wang et al.](#) evaluated seven advanced LLMs, six of which were

open-source, testing 145 API mappings drawn from eight popular Python libraries across 28,125 completion prompts aimed at detecting deprecated API usage in code completion [131]. Moumoula et al. compared four LLMs on a cross-language code clone detection task [83]. Three evaluated models were open-source. Gonçalves et al. fine-tuned the open LLM *LLaMA* 3.2 on a refined version of the *DiverseVul* dataset to benchmark vulnerability detection performance [40]. *CodeBERT*, a bimodal transformer pre-trained by Microsoft Research, has been widely used as an open baseline [21, 121, 138, 143], with model weights, source code, and data-processing scripts published on GitHub [76].

5.6.4 *Benefits:*

Using a true open LLM as a baseline improves the reproducibility of scientific research by providing access to model architectures and parameter settings, and ideally training data, thereby allowing independent reconstruction and verification of experimental results. Moreover, by adopting an open-source baseline, researchers can directly compare novel methods against established performance metrics without the variability introduced by proprietary systems. The transparent nature of these models allows for detailed inspection of data processing pipelines and decision-making routines, which is essential for identifying potential sources of bias and delineating model limitations. Furthermore, unlike closed-source alternatives, which can be withdrawn or altered without notice, open LLMs ensure long-term accessibility and stability, preserving critical resources for future studies. Finally, the licensing requirements associated with open-source implementations lower financial barriers, making advanced language models attainable for research groups operating under constrained budgets.

5.6.5 *Challenges:*

Open-source LLMs face several notable challenges. First, they often lag behind the most advanced proprietary “frontier” models in common benchmarks, making it difficult for researchers to demonstrate clear improvements when evaluating new methods using open LLMs alone. Additionally, deploying and experimenting with these models typically requires substantial hardware resources, in particular high-performance GPUs, which may be beyond reach for many academic groups. The notion of “openness” remains in flux: many models release only trained weights without training data or methodological details (“open weight” openness) [38], which is why we reference the OSI definition in our recommendations [84]. Finally, unlike APIs provided by proprietary vendors (e.g., the OpenAI API), installing, configuring, and fine-tuning open-source models can be technically demanding. Documentation is often sparse or fragmented, placing a high barrier to entry for researchers without specialized engineering support.

5.6.6 Study Types:

This guideline applies primarily to study types in which the researcher controls which LLM is used. In formal benchmarking studies and controlled experiments (see [Benchmarking LLMs](#)), an open LLM MUST be one of the models under evaluation. When evaluating [LLMs for Tools](#), researchers SHOULD use an open LLM as a baseline whenever it is technically feasible; if integration proves too complex, they SHOULD report the initial benchmarking results of open models. For [LLMs as Annotators](#) and [LLMs as Judges](#), researchers SHOULD compare annotation or judgment quality from open vs. commercial models to assess the extent to which results depend on a specific proprietary model. For [LLMs for Synthesis](#), researchers SHOULD compare synthesis results from open and commercial models to evaluate robustness of the findings. For [Studying LLM Usage](#), using an open LLM as a baseline is often not feasible when the study observes participants using specific commercial tools; in such cases, investigators SHOULD explicitly acknowledge its absence and discuss how this limitation might affect their conclusions. For [LLMs as Subjects](#), using an open LLM as a baseline may similarly be impractical if the study design requires the capabilities of a specific model; researchers SHOULD acknowledge this limitation when applicable.

5.6.7 Advice for Reviewers:

Reviewers should distinguish between LLM use that is central to the research (e.g., building LLM-driven SE tools, using LLMs for synthesis) and use that is tangential (e.g., generating recruiting materials). An open LLM baseline is expected only when LLM use is central; otherwise, its absence need not be justified. When an open baseline is expected, reviewers should look for either the use of an open LLM or a convincing argument for why it is impractical. If authors claim openness, some justification for that characterization is appropriate. Where practical, reviewers should examine whether the replication package contains sufficiently detailed instructions. Reviewers should *not* penalize studies for performance differences between open and proprietary models, as this is beyond the authors' control.

5.6.8 See Also:

- Section [Version and Configuration](#): model version, configuration, and parameter settings.
- Section [Architecture](#): hardware and hosting requirements for open models.
- Section [Benchmarks and Metrics](#): metrics for inter-model agreement and benchmarking.
- Section [Limitations and Mitigations](#): reporting limitations when open baselines are absent.

5.7 Use Suitable Baselines, Benchmarks, and Metrics (G7)

tl;dr: *Researchers MUST justify all benchmark and metric choices in the paper and SHOULD summarize benchmark structure, task types, and limitations. Where possible, traditional (non-LLM) baselines SHOULD be used for comparison. Researchers MUST explain in the paper why the selected metrics are suitable for the specific study; prior adoption in related work alone does not constitute sufficient justification. They SHOULD report established metrics to make study results comparable, but can report additional metrics that they consider appropriate. Due to the inherent non-determinism of LLMs, experiments SHOULD be repeated; the result distribution SHOULD then be reported using descriptive statistics. If comparing models or tools, researchers MUST use appropriate inferential statistics rather than relying solely on summary statistics. Researchers SHOULD justify the number of experiment repetitions, for example through a power analysis or by monitoring convergence of descriptive statistics.*

5.7.1 *Rationale:*

Meaningful evaluation requires well-understood, valid measurement instruments. Without justified benchmarks and metrics, claims about LLM performance lack the rigor needed for scientific comparison and cumulative progress. Benchmarks, baselines, and metrics play an important role in assessing the effectiveness of LLMs and LLM-based tools. A *benchmark* is a “standard tool for the competitive evaluation and comparison of competing systems or components according to specific characteristics, such as performance, dependability, or security” [62]. Meanwhile, “a *metric* is a method, algorithm, or procedure for assigning one or more numbers to a phenomenon” [97]. A baseline is a reference point, enabling comparison of LLMs against traditional algorithms with lower computational costs.

5.7.2 *Recommendations:*

When selecting benchmarks and metrics, it is important to fully understand the benchmark tasks, what exactly is being measured, and how it relates to the (often latent) variables researchers actually care about. If one or more metrics or benchmarks are used, researchers:

- MUST briefly justify (in the *paper*) why selected benchmarks and metrics are suitable for the given task or study;
- MUST discuss the reliability and validity (especially construct validity) of selected metrics and benchmarks;
- SHOULD summarize the structure and tasks of the selected benchmark(s), including the programming language(s) and descriptive statistics such as the number of contained tasks and test cases;

- SHOULD discuss the limitations of the selected benchmark(s) (e.g. widely used benchmarks such as *HumanEval* [26] and *MBPP* [10] Python functions assess a very specific part of software development, which is not representative of the full breadth of SE work [23]).
- SHOULD include an example of a task and the corresponding test case(s) to illustrate the structure of the benchmark.

If multiple benchmarks exist for the same task, researchers SHOULD compare performance between benchmarks. When selecting only a subset of all available benchmarks, researchers SHOULD use the most specific benchmarks given the context.

Furthermore, researchers SHOULD check whether a less resource-intensive approach (e.g., for static analysis tasks or program repair) can serve as a baseline. If so, the LLM or LLM-based tool should be compared with such baselines using suitable metrics. Even if LLM-based tools outperform baselines, researchers SHOULD discuss whether the resources consumed justify the (often marginal) improvements [75].

To compare traditional and LLM-based approaches or different LLM-based tools, researchers SHOULD report established metrics whenever possible, as this allows secondary research. They can, of course, report additional metrics that they consider appropriate. We briefly discuss common metrics in the Example(s) subsection below. As mentioned, researchers MUST motivate why they chose a certain metric or variant thereof for their particular study. Prior adoption alone does not constitute sufficient justification; researchers MUST NOT justify metric choices solely by citing their use in prior work.

Due to LLM non-determinism, researchers SHOULD repeat experiments and report descriptive statistics of model or tool performance (e.g., arithmetic mean, median, confidence intervals, standard deviations) [2, 18]. If comparing models or tools, researchers MUST use appropriate inferential statistics (e.g., hypothesis tests such as the Mann-Whitney U test, McNemar’s test for binary outcomes [63], or bootstrap-based comparisons, along with effect sizes) rather than relying solely on differences in means or other summary statistics.

The number of required repetitions depends on factors such as the study type, the variability of the task, and the desired precision of estimates. As with sample sizes for human validation (*Human Validation*), researchers SHOULD justify their chosen number of repetitions, for example, through a power analysis [18, 29] or by monitoring the convergence of descriptive statistics across incremental runs [19]. A pilot study can help estimate the expected variability and inform this decision.

From a measurement perspective, researchers SHOULD reflect on the theories, values, and measurement models on which the benchmarks and metrics they have selected for their study are based. For example, the phenomena labeled “bugs” in a large open dataset of software bugs are according to a certain theory of what constitutes a bug, and from the values and perspectives of the people who labeled the dataset. Reflecting on the context in which these labels

were assigned and discussing whether and how the labels generalize to a new study context is crucial.

5.7.3 *Example(s)*:

Common Metrics: Two common metrics used for generation tasks are *BLEU-N* and *pass@k*. *BLEU-N* [92] was originally developed for evaluating machine translation quality by measuring n-gram precision between a candidate and reference text, ranging from 0 (dissimilar) to 1 (similar). It has been widely adopted in SE for code generation tasks, though its validity in this context is debatable: n-gram overlap does not capture functional correctness, and syntactically different code can be semantically equivalent (see Challenges below). Code-specific variations such as *CodeBLEU* [100] and *CrystalBLEU* [33] attempt to address these limitations by introducing additional heuristics such as AST matching.

The metric *pass@k* reports the likelihood that a model correctly completes a code snippet at least once within k attempts. Originally introduced as *success rate at B* by Kula et al [66] and popularized by Chen et al [26], it ranges from 0 (no correct solution in k tries) to 1 (at least one correct solution), with correctness typically defined by passing test cases.

The *pass@k* metric is defined as:

$$\text{pass@k} = 1 - \frac{\binom{n-c}{k}}{\binom{n}{k}}, \text{ where:}$$

- n is the total number of generated samples per prompt,
- c is the number of correct samples among n , and
- k is the number of samples drawn.

The choice of k depends on the downstream task: *pass@1* is critical for single-suggestion scenarios like code completion, while higher k values (e.g., 2, 5, 10) assess multi-attempt capability [43, 52, 69, 104]. However, *pass@k* is not a universal metric suitable for all generation tasks. It requires a binary notion of correctness, making the metric appropriate for code synthesis evaluated via unit tests, but unsuitable for open-ended generation tasks such as comment generation, where multiple valid outputs exist.

If a study evaluates an LLM-based tool for supporting humans, a relevant metric is the acceptance rate, meaning the ratio of all accepted artifacts (e.g., test cases, code snippets) in relation to all artifacts that were generated and presented to the user. Another way of evaluating LLM-based tools is calculating inter-model agreement. This allows researchers to assess how dependent a tool’s performance is on specific models and versions. Metrics used to measure inter-model agreements include general agreement (percentage), *Cohen’s kappa*, and *Krippendorff’s α* (see [Human Validation](#) for recommended thresholds and best practices for measuring agreement).

Common problem types for LLM-based studies are classification, recommendation, and generation, each requiring different metrics [48]. Hu et al [49]

categorized 191 LLM benchmarks by SE task, providing a valuable reference. Common metrics include *BLEU*, *pass@k*, *Accuracy@k*, and *Exact Match* for generation; *Mean Reciprocal Rank* for recommendation; and *Precision*, *Recall*, *F1-score*, and *Accuracy* for classification.

Benchmark Examples: Benchmarks used for code generation include *HumanEval* (available on [GitHub](#)) [26], *MBPP* (available on [Hugging Face](#)) [10], *ClassEval* (available on [GitHub](#)) [31], *LiveCodeBench* (available on [GitHub](#)) [56], and *SWE-bench* (available on [GitHub](#)) [57]. An example of a code translation benchmark is *TransCoder* [62] (available on [GitHub](#)).

Most benchmarks focus on generation tasks, but benchmarks for classification and recommendation also exist. For classification, *DiverseVul* (available on [GitHub](#)) [27] provides vulnerable and non-vulnerable functions evaluated using standard classification metrics. For recommendation, *CodeSearchNet* (available on [GitHub](#)) [53] contains code-documentation pairs evaluated using *Mean Reciprocal Rank*.

5.7.4 *Benefits:*

Benchmarks and metrics improve reproducibility and comparability across studies, leading to faster improvements in the cumulative body of knowledge. It is simply easier to assess a new system against one or a few benchmarks than hundreds of competing systems, and when most systems are assessed against the same benchmarks and metrics, we can see which approaches work best (at least, on those benchmarks). This comparison enables progress tracking; for example, when researchers iteratively improve a new LLM-based tool and test it against benchmarks after significant changes. For practitioners, leaderboards (published benchmark results of models) support selecting models for downstream tasks. Meanwhile, baselines help us understand just how much LLMs improve performance over non-LLM-enabled alternatives.

5.7.5 *Challenges:*

Computer scientists have a long history of using oversimplified, unvalidated metrics [95] (e.g., lines of code, CPU time) as proxies for complex, multi-dimensional latent variables (e.g., system size, environmental sustainability). In fields such as psychology, where measurement theory has a longer tradition [20], metrics are backed by foundational theories or extensive empirical validation of their psychometric properties. The fundamental challenge with AI metrics and benchmarks is that many have neither firmly understood theoretical underpinnings nor extensive empirical validation of their construct, measurement, and ecological validity.

For example, *pass@k* is intended to measure a model’s *functional correctness*, that is, its ability to generate code that produces the expected output for a given specification. However, functional correctness is only one dimension of code quality. *pass@k* does not capture maintainability, readability, security, or

1 efficiency of the generated code, all of which are critical for downstream use.
2 Furthermore, “correctness” is defined entirely by test suites, whose coverage
3 is itself unvalidated: a solution that passes all provided tests may still be in-
4 correct for untested inputs. Whether a test suite adequately operationalizes
5 correctness for a given task is a subjective judgment that is rarely examined.
6

7 Similarly, *HumanEval* is intended to measure a model’s ability to *synthe-*
8 *size short Python functions from docstring specifications*. This operationalizes
9 a narrow slice of “code generation ability”: it covers neither multi-file tasks, nor
10 debugging, nor the use of existing codebases—tasks that dominate real-world
11 software engineering [23]. Notably, neither *pass@k* nor *HumanEval* has under-
12 gone rigorous empirical validation of its construct validity; their widespread
13 adoption rests on face validity and convenience rather than on evidence that
14 they reliably measure what researchers intend them to measure [22].

15 Benchmarks and metrics also have generalizability problems. Prominent
16 LLM benchmarks such as *HumanEval* and *MBPP* use Python, so researchers
17 can optimize for Python’s idiosyncrasies, creating illusory improvements that
18 do not generalize to other languages. Similarly, the metric *BLEU-N* is a syn-
19 tactic metric, so code can score highly without being executable. The metric
20 *Exact Match*, meanwhile, does not account for functional equivalence of syn-
21 tactically different code. Both *BLEU-N* and *Exact Match* are influenced by
22 code formatting, which confounds their intended use.
23

24 Execution-based metrics such as *pass@k* directly evaluate correctness by
25 running test cases, but they require a setup with an execution environment.
26 When researchers observe unexpected values for certain metrics, the specific
27 results should be investigated in more detail to uncover potential problems.

28 Finally, benchmark data contamination—where the benchmark itself is
29 part of the training data—may lead to artificially high performance if the
30 model remembers the solution from the training data rather than actually
31 solving the task based on the input data [140]. Therefore, for proprietary
32 LLMs that do not release their training data, researchers SHOULD consider
33 using human validation, curating new data, or refactoring existing data [140].

34 To mitigate contamination, researchers can create new benchmark datasets
35 by collecting data after a specified cutoff date; researchers MUST disclose data
36 sources and collection dates for each release. However, this temporal approach
37 requires continuous updates as new models may include the benchmark in
38 future training data. Alternatively, keeping the benchmark private prevents
39 inclusion in training sets, but requires trust in the benchmark creator and a
40 system to execute it without leaking data. Researchers can also evaluate how
41 contaminated their benchmark is [28].
42

43 5.7.6 Study Types:

44 This guideline MUST be followed for all study types that automatically evaluate
45 the performance of LLMs or LLM-based tools. The design of a benchmark and
46 the selection of appropriate metrics are highly dependent on the specific study
47 type and research goal. Recommending specific metrics for specific study types
48
49
50
51
52
53
54
55
56
57
58
59
60
61
62
63
64
65

is beyond the scope of these guidelines, but [Hu et al.](#) provide a good overview of existing metrics for evaluating LLMs [49].

For [Benchmarking LLMs](#), this guideline is of primary importance: researchers MUST use established benchmarks or rigorously justify the creation of new ones, and MUST report standard metrics to enable cross-study comparison. For [LLMs as Annotators](#), the research goal might be to assess which model comes close to a ground truth dataset created by human annotators. Especially for open annotation tasks, selecting suitable metrics to compare LLM-generated and human-generated labels is important. In general, annotation tasks can vary significantly. Are multiple labels allowed for the same sequence? Are the available labels predefined, or should the LLM generate a set of labels independently? Due to this task dependence, researchers MUST justify their metric choice, explaining what aspects of the task it captures together with known limitations. For [LLMs for Tools](#), researchers SHOULD benchmark the tool against suitable baselines using appropriate metrics that capture the tool’s intended contribution. For [LLMs as Judges](#), researchers SHOULD report inter-rater agreement metrics and validity measures to demonstrate the reliability and quality of LLM judgments. For [LLMs for Synthesis](#), researchers SHOULD specify metrics for comparing synthesized outputs (e.g., coverage, faithfulness) and justify their appropriateness for the synthesis task. For [Studying LLM Usage](#), researchers SHOULD justify the measurement instruments and metrics used for studying LLM usage patterns, including any survey scales or behavioral measures. For [LLMs as Subjects](#), researchers SHOULD compare simulated and real human responses using appropriate metrics to assess simulation fidelity. If researchers assess a well-established task such as code generation, they SHOULD report standard metrics such as *pass@k* and compare the performance between models. If non-standard metrics are used, researchers MUST state their reasoning.

5.7.7 Advice for Reviewers:

Reviewers should expect manuscripts to: 1. clearly identify the constructs or variables the study aims to measure (e.g., LLM performance, quality of generated code), including independent, dependent, and control variables; 2. present their measurement model, i.e., which metrics, benchmarks, or baselines are used and how they relate to the target constructs; 3. justify the selection of any metrics, benchmarks, and baselines used; 4. discuss *in detail* the assumptions, reliability, and validity—*especially construct validity*—of each benchmark and metric; 5. articulate any limitations regarding construct and measurement validity. As with other guidelines, missing information about baselines or metrics is typically a revision request. However, vague descriptions that conflate broad concerns (e.g., effectiveness, quality) with specific counting methods should be questioned. Ubiquity of a benchmark or metric does not imply validity or appropriateness for a given context. Manuscripts should convey a solid understanding of construct and measurement validity by explaining and justifying their measurement models.

5.7.8 *See Also:*

- Section [Open LLMs](#): open models for inter-model comparison.
- Section [Human Validation](#): human evaluation to complement automated metrics.
- Section [Limitations and Mitigations](#): reporting construct and measurement validity limitations.

5.8 Report Limitations and Mitigations (G8)

tl;dr: Researchers MUST transparently report study limitations, including the impact of non-determinism and generalizability constraints. The paper MUST specify whether generalization across LLMs or across time was assessed, and discuss model and version differences. Authors MUST describe measurement constructs and methods, disclose any data leakage risks, and avoid leaking evaluation data into LLM improvement pipelines. They MUST provide model outputs and discuss sensitive data handling. They SHOULD justify LLM usage in light of its resource demands. Mitigation strategies such as replication packages, human validation, longitudinal re-runs, triangulation, and sensitivity analysis SHOULD be employed and reported where applicable. Where full data sharing is not possible, a subset of the validation data SHOULD be included to enable partial replication.

5.8.1 *Rationale:*

When using LLMs for empirical studies in SE, researchers face unique challenges and potential limitations that can influence the validity, reliability, and reproducibility of their findings [106]. It is important to openly discuss these limitations and explain how their impact was mitigated. These limitations are relative to current LLM capabilities and tool architectures; speculating about future improvements is beyond the scope of a paper's limitation section. Nevertheless, risk management and threat mitigation should be planned during study design, not as an afterthought.

5.8.2 *Recommendations:*

Researchers MUST make every effort to present the limitations of their work clearly without defensiveness or obfuscation. These limitations may concern diverse topics including generalizability, internal validity (e.g. data leakage), reliability (i.e. non-determinism), and reproducibility (e.g. resource requirements). When deterministic reproducibility is structurally unattainable (e.g., SaaS-based models with opaque versioning), researchers SHOULD adopt trustworthiness criteria from qualitative research to substantiate dependability and

confirmability of findings. LLM-based studies may involve many kinds of generalization including the following:

1. From tested LLMs to others; that is, do the performance characteristics of the LLM(s) studied generalize to LLM(s) not included in the study.
2. From tested configurations to others; that is, how sensitive are the results to the specific configuration(s) of the LLM under test.
3. From research to practice; for instance, just because researchers can get a certain level of code generation performance under ideal conditions does not mean software developers, given the same tools, will get similar results without having the researchers present to show them exactly what to do.
4. From a sample of people (e.g. human validators; human participants) to a larger population of people.
5. From one time period to another; i.e., does the same LLM, or other versions of the same LLM, produce similar results at a different point in time?

The following concerns present an overview that has to be tailored to the individual study context. This section does not repeat requirements from other recommendations.

Internal Validity: The primary threats to internal validity are:

- Data leakage and contamination, including inter-dataset duplication, potentially resulting in a training-evaluation overlaps, leading to overly optimistic evaluation results.
- Unintended inclusion of evaluation data in model-improvement pipelines, contributing to data leakage. This is especially relevant for longitudinal studies including LLMs.
- Incomplete architecture, prompt, or pipeline reporting, posing hidden confounding factors.

Inter-dataset duplication is prevalent in SE, particularly for code-related benchmarks. As transparency on training data is limited for LLMs, researchers MUST discuss potential data leakage effects and their impact on results.

Construct Validity: The primary threats to construct validity are:

- Metric-construct mismatch (e.g., BLEU/ROUGE vs. functional correctness). More traditional metrics might not capture all relevant SE-specific aspects.
- Over-reliance on benchmark-specific metrics. Optimizing for a single benchmark might result in dataset-specific heuristics rather than the intended construct, overstating real-world utility.
- Benchmark scope limitations. Benchmarks commonly ignore runtime behaviors, security implications, readability, testability, and maintainability, yielding results that may not transfer to realistic development settings.

If constructs are based on subjective interpretations, purely automated metrics are insufficient. Researchers MUST discuss how they ensured quality of subjective results, similarly to qualitative research.

External Validity: The primary threats to external validity are:

- Cross-model transfer limitations. Results obtained with one LLM (or family of LLMs) may not generalize to others due to differences in training data, architecture, and capabilities.
- Tool-architecture specificity. Tools built around a specific LLM’s API or capabilities may not transfer to other models without significant re-engineering.
- Limited domain coverage. Studies often focus on a narrow set of programming languages, task types, or application domains, limiting generalizability to other SE contexts.
- Limited participant diversity. Study participants (e.g., human validators, developers) may not represent the broader population in terms of expertise, geographic location, or cultural background.
- Cross-time instability (model evolution). Performance of proprietary models can change over time, leading to non-generalizable study outcomes [25, 68].

Generalizability is particularly critical for proprietary and non-deterministic systems whose behavior is subject to drift (i.e., silent changes in model output over time) [25]. The limitations and mitigations used for external validity MUST be discussed by the researchers.

Reliability & Reproducibility: The primary threats to reliability and reproducibility are:

- Non-deterministic outputs. Identical prompts and configurations can yield different outputs across runs due to factors such as floating-point arithmetic, batching, and stochastic decoding strategies.
- Infrastructure dependence. Results may vary depending on the hardware, software stack, and hosting environment used, making exact replication challenging across different infrastructure setups.
- Resource inequality preventing replication. LLM research is resource-intensive and hence remains predominantly in the domain of private companies or well-funded research institutions [3, 115], excluding researchers from under-resourced institutions.

Researchers MUST discuss the measures taken to increase reliability and reproducibility. However, non-deterministic reproducibility is not inherently disqualifying. Qualitative traditions such as ethnography, grounded theory, and action research ensure trustworthiness through *credibility*, *transferability*, *dependability*, and *confirmability* [42]. SaaS-based LLM research faces analogous conditions, as providers frequently deprecate model versions without guaranteeing stable behavior.

Ethical & Regulatory Boundaries: The primary concerns for ethical and regulatory matters are:

- Use of sensitive or proprietary data. Studies involving proprietary code, confidential business data, or personally identifiable information may face restrictions on data sharing that limit reproducibility.

- Jurisdictional obligations. Data protection regulations (e.g., GDPR, CCPA) and institutional policies may impose constraints on data collection, processing, and sharing in LLM-based studies.
- Implicit model bias. Especially for qualitative research LLM might “reinforce dominant paradigms and biases” and “identify, replicate and reinforce dominant language and patterns” [58].

Studies involving sensitive data MUST discuss data governance mechanisms tailored towards LLM environments, compliant with applicable juristical obligations. If applicable, researchers SHOULD discuss how model biases potentially impact the study outcomes and how those biases were evaluated.

Environmental & Sustainability Constraints:

- Energy consumption. With growing model size, the environmental impact of experiments with LLMs increases. Considering the environmental impact in the study design is important given the substantial energy costs of LLM experiments [23].
- Trade-off between repetition and sustainability. Increased reliability through repetition and the aim of cautious energy consumption pose an inherent tension field to navigate during the study design.

The study findings SHOULD be discussed in the context of the resource-intensiveness of LLMs to justify the increased resource consumption over traditional approaches.

Mitigation Strategies The following mitigation strategies can help address the threats described above, depending on the study scope and feasibility.

- *Replication Packages* covering prompt and architecture specifications, model outputs, and representative examples for partial replicability, accompanied by an implementation using an open model for long-term stability.
- *Human Validation* of subjective constructs, following quality criteria known from qualitative research.
- *Longitudinal Re-Runs* to repeat experiments with LLMs over time, complemented by statistical analyses.
- *Methodological Trustworthiness Measures*. Researchers SHOULD consider triangulation, reflexivity, audit trails, and peer debriefing as complementary measures when deterministic reproduction is structurally impossible.
- *Triangulation* via multiple models (e.g., proprietary and open models), multiple independent datasets, and multiple complementary metrics.
- *Cost Accounting* through reporting of input/output tokens, token and service costs, or hardware specifications.
- *Energy Preservation* by selecting smaller or newer less resource-intensive models and employing techniques such as input/output token reduction, model pruning, quantization, or knowledge distillation [48] where feasible. Carbon footprint estimation is desirable, but still difficult.

- *Ethical and Regulatory Considerations* through data governance mechanisms, ethical reviews, and bias assessment procedures.
- *Sensitivity Analysis* through variation of LLM configurations, prompts, architecture decisions, datasets, and if applicable human participant backgrounds.

5.8.3 Study Types:

Researchers MUST follow this guideline for all study types. Transparently reporting limitations and mitigations is a universal requirement, but specific concerns vary by study type. For *LLMs as Annotators*, researchers MUST discuss potential biases in label assignment, label reliability limitations, and sensitivity of annotations to prompt wording and model choice. For *LLMs as Judges*, researchers MUST address measurement validity concerns, known biases such as position bias or verbosity bias, and the extent to which LLM judgments align with human expert assessments. For *LLMs for Synthesis*, researchers MUST discuss the risk of contextual misinterpretation, potential loss of nuance in summarized or aggregated outputs, and reflexivity limitations inherent in using an LLM for qualitative interpretation. For *LLMs as Subjects*, researchers MUST discuss the fundamental inability of LLMs to truly simulate human behavior, the risk of stereotype amplification, and the limited ecological validity of simulated responses. For *Studying LLM Usage*, researchers MUST discuss generalizability constraints across different tools and user populations, and acknowledge how observed usage patterns may not transfer to other contexts. For *LLMs for Tools*, researchers MUST discuss replicability constraints arising from dependencies on commercial models, the impact of model updates on tool behavior, and limitations of the evaluation setup. For *Benchmarking LLMs*, researchers MUST discuss potential data contamination, benchmark scope limitations, and the extent to which benchmark performance generalizes to real-world tasks.

5.8.4 Advice for Reviewers:

Reviewers should verify that the limitation section is comprehensive and appropriate for the specific study type, checking that: (1) limitations address the specific threats relevant to the study type (e.g., label reliability for annotation studies, simulation fidelity for studies using LLMs as subjects); (2) mitigations are concrete and correspond to identified limitations rather than being generic statements; (3) the impact of LLM non-determinism on findings is discussed; (4) generalizability constraints—across models, configurations, time periods, and populations—are acknowledged. When important limitations are missing, reviewers should request they be added. The absence of a limitation section, or one that is formulaic or insufficiently specific, is a more serious concern than any individual missing limitation and may warrant a major revision.

Table 4 Rationale and key recommendations for each guideline. ● = MUST, ● = SHOULD.

ID	Rationale	Core Recommendations
G1	Transparency enables informed assessment of scope and limitations.	● Declare which LLM, how it was used, and where in the research process.
G2	Reproducibility requires precise identification of the system used in a study.	● Report exact version, date, configuration, and fine-tuning details.
G3	Complex systems require holistic documentation beyond the model itself.	● Describe full architecture, agent roles, and infrastructure.
G4	Prompts shape LLM behavior and must be verifiable.	● Publish all prompts, context files, and development strategy. ● Share interaction logs where feasible.
G5	Automated metrics alone cannot ensure validity of subjective constructs.	● Validate against human judgment with inter-rater reliability, where applicable.
G6	Reproducibility depends on access to the model under study.	● Include open LLM baseline; provide replication package.
G7	Meaningful evaluation requires reasoned valid measurement.	● Justify metric and benchmark choices; discuss their validity. ● Repeat experiments and report descriptive statistics.
G8	Honest acknowledgment of threats strengthens a study.	● Report threats by validity type. ● Employ mitigations where possible.

5.8.5 See Also:

- Section [Version and Configuration](#): model configuration, cost accounting, and sensitivity analysis.
- Section [Architecture](#): architecture-level triangulation and sensitivity analysis.
- Section [Prompts and Logs](#): prompt reporting, audit trails, and longitudinal re-runs.
- Section [Benchmarks and Metrics](#): metric selection and benchmark limitations.
- Section [Human Validation](#): human validation of subjective constructs.

6 Conclusion

LLM non-determinism, opaque training data, and rapidly evolving models threaten the reproducibility and replicability of empirical SE studies. This paper presents a taxonomy of seven study types that organizes how LLMs are used in SE research, together with eight guidelines for designing and reporting such studies. Each guideline distinguishes requirements (MUST) from recommended practices (SHOULD) and is contextualized by the study types it

applies to. Our guidelines recommend that researchers: (1) declare LLM usage and role; (2) report model versions, configurations, and customizations; (3) document the tool architecture beyond the model; (4) disclose prompts, their development, and, where feasible, interaction logs; (5) validate LLM outputs with humans; (6) include an open LLM as a baseline; (7) use suitable baselines, benchmarks, and metrics; and (8) articulate limitations and mitigations, including costs, potential biases, and environmental impact. An applicability matrix maps guidelines to study types, and a reporting checklist (Section A) supports authors and reviewers. The study types and guidelines were derived collaboratively by 22 co-authors, following a process similar to that of guidelines in other disciplines (see, e.g., [17]).

Following the quality dimensions of the *SIGSOFT Empirical Standards* [96], we consider our guidelines (1) *comprehensive*, as they cover seven study types that we identified and were derived from 22 co-authors' collective expertise; (2) *useful*, as they provide actionable recommendations with an applicability matrix and a concrete reporting checklist (Section A); (3) *well-argued*, as each guideline is motivated by a dedicated rationale and supported by references to the literature; and (4) *integrated with previous work*, as they build on and complement established empirical SE standards while addressing LLM-specific challenges. Adopting these practices will not eliminate LLM non-determinism, but it should make claims auditable and results easier to replicate and compare across studies. Because models and tools evolve, we maintain the guidelines as a living resource and invite the community to refine and extend them at llm-guidelines.org.

Acknowledgments

Our study types and guidelines are based on discussion sessions with researchers at the 2024 *International Software Engineering Research Network* (ISERN) meeting and at the 2nd *Copenhagen Symposium on Human-Centered Software Engineering AI* (supported by the Carlsberg Foundation with grant CF24-0693 and the Alfred P. Sloan Foundation with grant G-2024-22586 to Daniel Russo). LLM-based tools, including ChatGPT (GPT-5) and Claude Code (Opus 4.6), were used interactively for language editing and structural revision suggestions across all sections of this paper. All LLM-generated suggestions were reviewed and revised by the authors.

Appendix A Reporting Checklist

The following checklist, inspired by CONSORT [114], summarizes actionable items from the guidelines based on the *tl;dr* sections. The checklist is organized along typical paper sections. Items marked ● are requirements (MUST); items marked ● are recommendations (SHOULD). Each item references its source

guideline (G1–G8). Items annotated with *paper* or *supplementary material* indicate where information should be reported; unmarked items may be reported in either.

Introduction

- ☐ ● Disclose any use of LLMs in the empirical study, specifying which LLM, how, and where it was used (G1).
- ☐ ● Report the purpose of using LLMs, automated tasks, and expected benefits in the *paper* (G1).

Research Design and Methods

Model Selection and Configuration

- ☐ ● Report the exact LLM model or tool version, configuration, and experiment date in the *paper* (G2).
- ☐ ● For fine-tuned models, describe the fine-tuning goal, dataset, and procedure in the *paper* (G2).
- ☐ ● Report default parameters and explain model and version choices (G2).
- ☐ ● For quantized models, report the quantization level (e.g., 4-bit, 8-bit) and method (e.g., GPTQ or AWQ) (G2).
- ☐ ● Compare base and fine-tuned models using suitable metrics and benchmarks; share fine-tuning data and weights as *supplementary material* (or justify in the *paper* why they cannot be shared) (G2).

Architecture

- ☐ ● Describe the full architecture of LLM-based tools in the *paper*, including the role of the LLM, interactions with other components, and overall system behavior (G3).
- ☐ ● For ensemble architectures, explain the coordination logic between models in the *paper* (G3).
- ☐ ● If autonomous agents are used, specify agent roles, reasoning frameworks, and communication flows (G3).
- ☐ ● Report hosting, hardware setup, and latency implications (G3).
- ☐ ● For tools using retrieval or augmentation methods, describe data sources, integration mechanisms, and update and versioning strategies (G3).
- ☐ ● Include architectural diagrams and justify design decisions (G3).

Prompts and Interactions

- ☐ ● Describe prompt development strategies (e.g., zero-shot, few-shot), rationale, and selection process in the *paper* (G4).
- ☐ ● Publish all prompts or, when using templates, prompt templates with representative instances, including their structure, content, formatting, and dynamic components, as *supplementary material* (G4).
- ☐ ● Document input handling and token optimization strategies when prompts are long or complex (G4).

- ☐ ● Report generation and collection processes for dynamically generated or user-authored prompts (G4).
- ☐ ● Specify prompt reuse across models and configurations (G4).
- ☐ ● If full prompt disclosure is not feasible, provide summaries or examples (G4).
- ☐ ● Report prompt revisions and pilot testing insights (G4).
- ☐ ● Include full interaction logs (prompts and responses) as *supplementary material* if privacy and confidentiality can be ensured (G4).
- ☐ ● For agentic systems, report all context files (e.g., `AGENTS.md`) used to configure AI agents as *supplementary material* (G4).
- ☐ ● For agentic systems, report developed plans and generalize interaction logs to include human-agent interaction traces as *supplementary material* (G4).

Human Validation

- ☐ ● If using human validation, define the measured construct (e.g., usability, maintainability) and describe the measurement instrument in the *paper* (G5).
- ☐ ● When developing or adapting measurement instruments, share them (G5).
- ☐ ● Consider human validation early in the study design and build on established reference models for human-LLM comparison (G5).
- ☐ ● Validate LLM judgments against human judgment, report aggregation methods, and assess human-LLM agreement (G5).
- ☐ ● Control for confounding factors and conduct power analysis to ensure statistical robustness (G5).

Benchmarks and Metrics

- ☐ ● Justify all benchmark and metric choices in the *paper* (G7).
- ☐ ● Explain in the *paper* why the selected metrics are suitable for the specific study (G7).
- ☐ ● Include an open LLM as a baseline when using commercial models and report inter-model agreement (G6).
- ☐ ● Summarize benchmark structure, task types, and limitations (G7).
- ☐ ● Justify the number of experiment repetitions, for example through a power analysis or by monitoring convergence of descriptive statistics (G7).

Reproducibility, Ethics, and Resources

- ☐ ● Provide raw LLM outputs for reproducibility and discuss sensitive data handling (G8).
- ☐ ● Justify LLM usage in light of its resource demands (G8).
- ☐ ● If full data sharing is not possible, include a subset of the validation data for partial replication (G8).
- ☐ ● Provide a full replication package with step-by-step instructions as *supplementary material* (G6).

Results

- ☐ ● If comparing models or tools, use appropriate inferential statistics (e.g., hypothesis tests, effect sizes) rather than relying solely on summary statistics (G7).
- ☐ ● Repeat experiments due to the inherent non-determinism of LLMs and report the result distribution using descriptive statistics (G7).
- ☐ ● Use traditional (non-LLM) baselines for comparison where possible (G7).
- ☐ ● Report established metrics to make study results comparable; additional metrics may be reported where appropriate (G7).

Limitations and Threats to Validity

- ☐ ● Describe measurement constructs and methods; disclose any data leakage risks and avoid leaking evaluation data into LLM improvement pipelines (G8).
- ☐ ● Acknowledge non-disclosed confidential or proprietary components as reproducibility limitations (G3).
- ☐ ● Transparently report study limitations, including the impact of non-determinism and generalizability constraints (G8).
- ☐ ● Specify whether generalization across LLMs or across time was assessed, and discuss model and version differences (G8).
- ☐ ● Employ and report strategies to mitigate identified validity and reproducibility threats, such as replication packages, human validation, longitudinal re-runs, triangulation, and sensitivity analysis (G8).

7 Statements and Declarations

Competing Interests: The authors declare no competing interests. Multiple co-authors are members of the EMSE Editorial Board, which is disclosed here for transparency.

Authorship: Each author contributed or reviewed content and, most importantly, read the complete guidelines and confirmed that they stand behind all recommendations, not only their own contributions.

Funding: No funding was received for writing this article. However, the *2nd Copenhagen Symposium on Human-Centered Software Engineering AI*, which played a central role in forming the team of co-authors, was supported by the Carlsberg Foundation with grant CF24-0693 and the Alfred P. Sloan Foundation with grant G-2024-22586 to Daniel Russo.

Ethical Approval: No ethics approval was required for writing this article.

Data Availability: This article has no associated data. The guidelines are, however, [hosted online](#) and [maintained on GitHub](#).

References

1. ACM Publications Board (2020) Artifact review and badging – current. <https://www.acm.org/publications/policies/artifact-review-and-badging-current>, accessed 2025-01-13
2. Agarwal R, Schwarzer M, Castro PS, Courville AC, Bellemare MG (2021) Deep reinforcement learning at the edge of the statistical precipice. In: Ranzato M, Beygelzimer A, Dauphin YN, Liang P, Vaughan JW (eds) Advances in Neural Information Processing Systems 34: Annual Conference on Neural Information Processing Systems 2021, NeurIPS 2021, December 6-14, 2021, virtual, pp 29304–29320, URL <https://proceedings.neurips.cc/paper/2021/hash/f514cec81cb148559cf475e7426eed5e-Abstract.html>
3. Ahmed N, Wahed M, Thompson NC (2023) The growing influence of industry in AI research. *Science* 379(6635):884–886, DOI 10.1126/science.ade2420
4. Ahmed T, Devanbu PT, Treude C, Pradel M (2025) Can llms replace manual annotation of software engineering artifacts? In: 22nd IEEE/ACM International Conference on Mining Software Repositories, MSR@ICSE 2025, Ottawa, ON, Canada, April 28-29, 2025, IEEE, pp 526–538, DOI 10.1109/MSR66628.2025.00086
5. Ahuja S, Gumma V, Sitaram S (2024) Contamination report for multilingual benchmarks. *CoRR* abs/2410.16186, DOI 10.48550/ARXIV.2410.16186, URL <https://doi.org/10.48550/arXiv.2410.16186>, 2410-16186
6. Anandayavaraj D, Campbell M, Tewari A, Davis JC (2024) FAIL: analyzing software failures from the news using llms. In: Filkov V, Ray B, Zhou M (eds) Proceedings of the 39th IEEE/ACM International Conference on Automated Software Engineering, ASE 2024, Sacramento, CA, USA, October 27 - November 1, 2024, ACM, pp 506–518, DOI 10.1145/3691620.3695022
7. Angermeir F, Amougou M, Kreitz M, Bauer A, Linhuber M, Fucci D, Constante FM, Méndez D, Gorschek T (2025) Reflections on the reproducibility of commercial LLM performance in empirical software engineering studies. *CoRR* abs/2510.25506, DOI 10.48550/ARXIV.2510.25506, URL <https://doi.org/10.48550/arXiv.2510.25506>, 2510-25506
8. Argyle LP, Busby EC, Fulda N, Gubler J, Rytting CM, Wingate D (2022) Out of one, many: Using language models to simulate human samples. *CoRR* abs/2209.06899, DOI 10.48550/ARXIV.2209.06899, URL <https://doi.org/10.48550/arXiv.2209.06899>, 2209-06899
9. Association for Computing Machinery (2023) ACM Policy on Authorship. <https://www.acm.org/publications/policies/new-acm-policy-on-authorship>, accessed 2025-01-13
10. Austin J, Odena A, Nye MI, Bosma M, Michalewski H, Dohan D, Jiang E, Cai CJ, Terry M, Le QV, Sutton C (2021) Program synthe-

- sis with large language models. CoRR abs/2108.07732, URL <https://arxiv.org/abs/2108.07732>, [2108.07732](https://arxiv.org/abs/2108.07732)
11. Azanza M, Pereira J, Irastorza A, Galdos A (2024) Can llms facilitate onboarding software developers? an ongoing industrial case study. In: 36th International Conference on Software Engineering Education and Training, CSEE&T 2024, IEEE, pp 1–6, DOI 10.1109/CSEET62301.2024.10662989
 12. Bano M, Zowghi D, Whittle J (2023) Exploring qualitative research using llms. CoRR abs/2306.13298, DOI 10.48550/ARXIV.2306.13298, URL <https://doi.org/10.48550/arXiv.2306.13298>, [2306.13298](https://doi.org/10.48550/arXiv.2306.13298)
 13. Bano M, Hoda R, Zowghi D, Treude C (2024) Large language models for qualitative research in software engineering: exploring opportunities and challenges. *Autom Softw Eng* 31(1):8, DOI 10.1007/S10515-023-00407-8, URL <https://doi.org/10.1007/s10515-023-00407-8>
 14. Bano M, Gunatilake H, Hoda R (2025) What does a software engineer look like? exploring societal stereotypes in llms. In: 47th IEEE/ACM International Conference on Software Engineering: Software Engineering in Society, ICSE-SEIS 2025, Ottawa, ON, Canada, April 27 - May 3, 2025, IEEE, pp 173–184, DOI 10.1109/ICSE-SEIS66351.2025.00023
 15. Barros CF, Azevedo BB, Neto VVG, Kassab M, Kalinowski M, do Nascimento HAD, Bandeira MCGSP (2025) Large language model for qualitative research: A systematic mapping study. In: IEEE/ACM International Workshop on Methodological Issues with Empirical Studies in Software Engineering, WSESE@ICSE 2025, Ottawa, ON, Canada, May 3, 2025, IEEE, pp 48–55, DOI 10.1109/WSESE66602.2025.00015
 16. Bavaresco A, Bernardi R, Bertolazzi L, Elliott D, Fernández R, Gatt A, Ghaleb E, Giulianelli M, Hanna M, Koller A, Martins AFT, Mondorf P, Neplenbroek V, Pezzelle S, Plank B, Schlangen D, Suglia A, Surikuchi AK, Takmaz E, Testoni A (2024) Llms instead of human judges? A large scale empirical study across 20 NLP evaluation tasks. CoRR abs/2406.18403, DOI 10.48550/ARXIV.2406.18403, URL <https://doi.org/10.48550/arXiv.2406.18403>, [2406.18403](https://doi.org/10.48550/arXiv.2406.18403)
 17. Begg C, Cho M, Eastwood S, Horton R, Moher D, Olkin I, Pitkin R, Rennie D, Schulz KF, Simel D, Stroup DF (1996) Improving the quality of reporting of randomized controlled trials. The CONSORT statement. *JAMA* 276(8):637–639, DOI 10.1001/jama.276.8.637
 18. Bjarnason BH, Silva A, Monperrus M (2026) On randomness in agentic evals. URL <https://arxiv.org/abs/2602.07150>, [2602.07150](https://arxiv.org/abs/2602.07150)
 19. Blackwell RE, Barry J, Cohn AG (2024) Towards reproducible llm evaluation: Quantifying uncertainty in llm benchmark scores. URL <https://arxiv.org/abs/2410.03492>, [2410.03492](https://arxiv.org/abs/2410.03492)
 20. Borsboom D (2005) *Measuring the Mind: Conceptual Issues in Contemporary Psychometrics*. Cambridge University Press, DOI 10.1017/CBO9780511490026
 21. Cai Y, Yadavally A, Mishra A, Montejo G, Nguyen TN (2024) Programming assistant for exception handling with codebert. In: *Proceedings of*

- the 46th IEEE/ACM International Conference on Software Engineering, ICSE 2024, Lisbon, Portugal, April 14-20, 2024, ACM, pp 94:1–94:13, DOI 10.1145/3597503.3639188
22. Cao J, Chan YK, Ling Z, Wang W, Li S, Liu M, Qiao R, Han Y, Wang C, Yu B, He P, Wang S, Zheng Z, Lyu MR, Cheung SC (2025) How should we build a benchmark? revisiting 274 code-related benchmarks for LLMs. CoRR abs/2501.10711, [2501.10711](https://arxiv.org/abs/2501.10711)
23. Chandra S (2025) Benchmarks for AI in Software Engineering (BLOG@CACM). <https://cacm.acm.org/blogcacm/benchmarks-for-ai-in-software-engineering/>, accessed 2025-08-24
24. Chann S (2023) Non-determinism in GPT-4 is caused by Sparse MoE. <https://152334h.github.io/blog/non-determinism-in-gpt-4/>, accessed 2025-01-13
25. Chen L, Zaharia M, Zou J (2024) How is chatgpt’s behavior changing over time? Harvard Data Science Review 6(2), DOI 10.1162/99608f92.5317da47
26. Chen M, Tworek J, Jun H, Yuan Q, de Oliveira Pinto HP, Kaplan J, Edwards H, Burda Y, Joseph N, Brockman G, Ray A, Puri R, Krueger G, Petrov M, Khlaaf H, Sastry G, Mishkin P, Chan B, Gray S, Ryder N, Pavlov M, Power A, Kaiser L, Bavarian M, Winter C, Tillet P, Such FP, Cummings D, Plappert M, Chantzis F, Barnes E, Herbert-Voss A, Guss WH, Nichol A, Paino A, Tezak N, Tang J, Babuschkin I, Balaji S, Jain S, Saunders W, Hesse C, Carr AN, Leike J, Achiam J, Misra V, Morikawa E, Radford A, Knight M, Brundage M, Murati M, Mayer K, Welinder P, McGrew B, Amodei D, McCandlish S, Sutskever I, Zaremba W (2021) Evaluating large language models trained on code. CoRR abs/2107.03374, URL <https://arxiv.org/abs/2107.03374>, [2107.03374](https://arxiv.org/abs/2107.03374)
27. Chen Y, Ding Z, Alowain L, Chen X, Wagner DA (2023) Diversevul: A new vulnerable source code dataset for deep learning based vulnerability detection. In: Proceedings of the 26th International Symposium on Research in Attacks, Intrusions and Defenses, RAID 2023, Hong Kong, China, October 16-18, 2023, ACM, pp 654–668, DOI 10.1145/3607199.3607242
28. Choi HK, Khanov M, Wei H, Li Y (2025) How contaminated is your benchmark? quantifying dataset leakage in large language models with kernel divergence. CoRR abs/2502.00678, DOI 10.48550/ARXIV.2502.00678, URL <https://doi.org/10.48550/arXiv.2502.00678>, [2502.00678](https://arxiv.org/abs/2502.00678)
29. Cohen J (1992) A power primer. Psychological Bulletin 112(1):155–159, DOI 10.1037/0033-2909.112.1.155
30. Dhar R, Vaidhyathan K, Varma V (2024) Can llms generate architectural design decisions? - an exploratory empirical study. In: 21st IEEE International Conference on Software Architecture, ICSA 2024, Hyderabad, India, June 4-8, 2024, IEEE, pp 79–89, DOI 10.1109/ICSA59870.2024.00016

31. Du X, Liu M, Wang K, Wang H, Liu J, Chen Y, Feng J, Sha C, Peng X, Lou Y (2023) Classeval: A manually-crafted benchmark for evaluating llms on class-level code generation. CoRR abs/2308.01861, DOI 10.48550/ARXIV.2308.01861, URL <https://doi.org/10.48550/arXiv.2308.01861>, 2308.01861
32. Dybå T, Kampenes VB, Sjøberg DIK (2006) A systematic review of statistical power in software engineering experiments. Inf Softw Technol 48(8):745–755, DOI 10.1016/J.INFSOF.2005.08.009, URL <https://doi.org/10.1016/j.infsof.2005.08.009>
33. Eghbali A, Pradel M (2022) Crystalbleu: Precisely and efficiently measuring the similarity of code. In: 37th IEEE/ACM International Conference on Automated Software Engineering, ASE 2022, ACM, pp 28:1–28:12, DOI 10.1145/3551349.3556903
34. El-Hajjaji A, Salinesi C (2025) How good are synthetic requirements? evaluating llm-generated datasets for AI4RE. CoRR abs/2506.21138, DOI 10.48550/ARXIV.2506.21138, URL <https://doi.org/10.48550/arXiv.2506.21138>, 2506.21138
35. Gallegos IO, Rossi RA, Barrow J, Tanjim MM, Kim S, Dernoncourt F, Yu T, Zhang R, Ahmed N (2024) Bias and fairness in large language models: A survey. Computational Linguistics 50:1097–1179, DOI 10.1162/coli_a_00524
36. Gallifant J, Afshar M, Ameen S, Aphinyanaphongs Y, Chen S, Cacciamani G, Demner-Fushman D, Dligach D, Daneshjou R, Fernandes C, Hansen LH, Landman A, Lehmann L, McCoy LG, Miller T, Moreno A, Munch N, Restrepo D, Savova G, Umeton R, Gichoya JW, Collins GS, Moons KGM, Celi LA, Bitterman DS (2025) The tripod-llm reporting guideline for studies using large language models. Nature Medicine 31(1):60–69, DOI 10.1038/s41591-024-03425-5
37. Gerosa MA, Trinkenreich B, Steinmacher I, Sarma A (2024) Can AI serve as a substitute for human subjects in software engineering research? Autom Softw Eng 31(1):13, DOI 10.1007/S10515-023-00409-6, URL <https://doi.org/10.1007/s10515-023-00409-6>
38. Gibney E (2024) Not all ‘open source’ AI models are actually open. Nature News DOI 10.1038/d41586-024-02012-5
39. Gilardi F, Alizadeh M, Kubli M (2023) ChatGPT outperforms crowd workers for text-annotation tasks. Proceedings of the National Academy of Sciences 120(30):e2305016120, DOI 10.1073/pnas.2305016120
40. Gonçalves J, Silva M, Cabral B, Dias T, Maia E, Praça I, Severino R, Ferreira LL (2025) Evaluating llama 3.2 for software vulnerability detection. In: Praça I, Bernardi S, Inácio PRM (eds) Cybersecurity - 9th European Interdisciplinary Cybersecurity Conference, EICC 2025, Rennes, France, June 18-19, 2025, Proceedings, Springer, Communications in Computer and Information Science, vol 2500, pp 38–51, DOI 10.1007/978-3-031-94855-8_3
41. Graziotin D (2024) acmsigsoft/open-science-policies: v1.0.0. DOI 10.5281/zenodo.10796477

42. Guba EG (1981) Criteria for assessing the trustworthiness of naturalistic inquiries. *ECTJ* 29(2):75–91, DOI 10.1007/BF02766777
43. Guo D, Zhu Q, Yang D, Xie Z, Dong K, Zhang W, Chen G, Bi X, Wu Y, Li YK, Luo F, Xiong Y, Liang W (2024) Deepseek-coder: When the large language model meets programming - the rise of code intelligence. *CoRR* abs/2401.14196, DOI 10.48550/ARXIV.2401.14196, URL <https://doi.org/10.48550/arXiv.2401.14196>, [2401.14196](https://arxiv.org/abs/2401.14196)
44. Harding J, D'Alessandro W, Laskowski NG, Long R (2024) AI language models cannot replace human research participants. *AI Soc* 39(5):2603–2605, DOI 10.1007/S00146-023-01725-X, URL <https://doi.org/10.1007/s00146-023-01725-x>
45. He J, Shi J, Zhuo TY, Treude C, Sun J, Xing Z, Du X, Lo D (2025) Llm-as-a-judge for software engineering: Literature review, vision, and the road ahead. *CoRR* abs/2510.24367, DOI 10.48550/ARXIV.2510.24367, URL <https://doi.org/10.48550/arXiv.2510.24367>, [2510.24367](https://arxiv.org/abs/2510.24367)
46. He Z, Huang C, Ding CC, Rohatgi S, Huang TK (2024) If in a crowd-sourced data annotation pipeline, a GPT-4. In: Mueller FF, Kyburz P, Williamson JR, Sas C, Wilson ML, Dugas POT, Shklovski I (eds) *Proceedings of the CHI Conference on Human Factors in Computing Systems*, CHI 2024, ACM, pp 1040:1–1040:25, DOI 10.1145/3613904.3642834
47. Holtzman A, Buys J, Du L, Forbes M, Choi Y (2020) The curious case of neural text degeneration. In: 8th International Conference on Learning Representations, ICLR 2020, Addis Ababa, Ethiopia, April 26–30, 2020, OpenReview.net, URL <https://openreview.net/forum?id=rygGQyrFvH>
48. Hou X, Zhao Y, Liu Y, Yang Z, Wang K, Li L, Luo X, Lo D, Grundy J, Wang H (2024) Large language models for software engineering: A systematic literature review. *ACM Trans Softw Eng Methodol* 33(8):1–79, DOI 10.1145/3695988
49. Hu X, Niu F, Chen J, Zhou X, Zhang J, He J, Xia X, Lo D (2025) Assessing and advancing benchmarks for evaluating large language models in software engineering tasks. *CoRR* abs/2505.08903, DOI 10.48550/ARXIV.2505.08903, URL <https://doi.org/10.48550/arXiv.2505.08903>, [2505.08903](https://arxiv.org/abs/2505.08903)
50. Huang F, Kwak H, An J (2023) Is chatgpt better than human annotators? potential and limitations of chatgpt in explaining implicit hate speech. In: Ding Y, Tang J, Sequeda JF, Aroyo L, Castillo C, Houben G (eds) *Companion Proceedings of the ACM Web Conference 2023*, WWW 2023, ACM, pp 294–297, DOI 10.1145/3543873.3587368
51. Huang J, Yi B, Sun W, Wan B, Xu Y, Feng Y, Ye W, Qin Q (2024) Enhancing review classification via llm-based data annotation and multi-perspective feature representation learning. *SSRN Electronic Journal* pp 1–15, DOI 10.2139/ssrn.5002351, URL <https://ssrn.com/abstract=5002351>

52. Hui B, Yang J, Cui Z, Yang J, Liu D, Zhang L, Liu T, Zhang J, Yu B, Dang K, Yang A, Men R, Huang F, Ren X, Ren X, Zhou J, Lin J (2024) Qwen2.5-coder technical report. CoRR abs/2409.12186, DOI 10.48550/ARXIV.2409.12186, URL <https://doi.org/10.48550/arXiv.2409.12186>, [arXiv:2409.12186](https://arxiv.org/abs/2409.12186)
53. Husain H, Wu H, Gazit T, Allamanis M, Brockschmidt M (2019) Codesearchnet challenge: Evaluating the state of semantic code search. CoRR abs/1909.09436, URL <http://arxiv.org/abs/1909.09436>, [arXiv:1909.09436](https://arxiv.org/abs/1909.09436)
54. Hymel C, Johnson H (2025) Analysis of llms vs human experts in requirements engineering. CoRR abs/2501.19297, DOI 10.48550/ARXIV.2501.19297, URL <https://doi.org/10.48550/arXiv.2501.19297>, [arXiv:2501.19297](https://arxiv.org/abs/2501.19297)
55. Jahic J, Sami A (2024) State of practice: Llms in software engineering and software architecture. In: 21st IEEE International Conference on Software Architecture, ICSA 2024 - Companion, Hyderabad, India, June 4-8, 2024, IEEE, pp 311–318, DOI 10.1109/ICSA-C63560.2024.00059
56. Jain N, Han K, Gu A, Li W, Yan F, Zhang T, Wang S, Solar-Lezama A, Sen K, Stoica I (2024) Livecodebench: Holistic and contamination free evaluation of large language models for code. CoRR abs/2403.07974, DOI 10.48550/ARXIV.2403.07974, URL <https://doi.org/10.48550/arXiv.2403.07974>, [arXiv:2403.07974](https://arxiv.org/abs/2403.07974)
57. Jimenez CE, Yang J, Wettig A, Yao S, Pei K, Press O, Narasimhan KR (2024) Swe-bench: Can language models resolve real-world github issues? In: The Twelfth International Conference on Learning Representations, ICLR 2024, Vienna, Austria, May 7-11, 2024, OpenReview.net, URL <https://openreview.net/forum?id=VTF8yNQM66>
58. Jowsey T, Braun V, Clarke V, Clarke V, Lupton D, Fine M (2025) We reject the use of generative artificial intelligence for reflexive qualitative research. SSRN DOI 10.2139/ssrn.5676462, URL <https://ssrn.com/abstract=5676462>
59. Kang S, Yoon J, Yoo S (2023) Large language models are few-shot testers: Exploring llm-based general bug reproduction. In: 45th IEEE/ACM International Conference on Software Engineering, ICSE 2023, IEEE, pp 2312–2323, DOI 10.1109/ICSE48619.2023.00194
60. Khojah R, Mohamad M, Leitner P, de Oliveira Neto FG (2024) Beyond code generation: An observational study of chatgpt usage in software engineering practice. Proc ACM Softw Eng 1(FSE):1819–1840, DOI 10.1145/3660788
61. Khraisha Q, Put S, Kappenberg J, Warraitch A, Hadfield K (2024) Can large language models replace humans in systematic reviews? evaluating gpt-4's efficacy in screening and extracting data from peer-reviewed and grey literature in multiple languages. Research Synthesis Methods 15(4):616–626, DOI 10.1002/jrsm.1715, URL <https://onlinelibrary.wiley.com/doi/abs/10.1002/jrsm.1715>, <https://onlinelibrary.wiley.com/doi/pdf/10.1002/jrsm.1715>

62. v Kistowski J, Arnold JA, Huppler K, Lange KD, Henning JL, Cao P (2015) How to build a benchmark. In: Proceedings of the 6th ACM/SPEC International Conference on Performance Engineering, Association for Computing Machinery, New York, NY, USA, ICPE '15, p 333336, DOI 10.1145/2668930.2688819
63. Kong A, Zhao S, Chen H, Li Q, Qin Y, Sun R, Zhou X (2023) Better zero-shot reasoning with role-play prompting. CoRR abs/2308.07702, DOI 10.48550/ARXIV.2308.07702, URL <https://doi.org/10.48550/arXiv.2308.07702>, [2308.07702](https://arxiv.org/abs/2308.07702)
64. Krippendorff K (2018) Content Analysis: An Introduction to Its Methodology, 4th edn. SAGE Publications
65. Kübler JM, Budhathoki K, Kleindessner M, Zhou X, Yin J, Khetan A, Karypis G (2026) When llms get significantly worse: A statistical approach to detect model degradations. In: Proc. of the International Conference on Learning Representations (ICLR)
66. Kulal S, Pasupat P, Chandra K, Lee M, Padon O, Aiken A, Liang P (2019) Spoc: Search-based pseudocode to code. CoRR abs/1906.04908, URL <http://arxiv.org/abs/1906.04908>, [1906.04908](https://arxiv.org/abs/1906.04908)
67. Lachaux M, Rozière B, Chansussat L, Lample G (2020) Unsupervised translation of programming languages. CoRR abs/2006.03511, URL <https://arxiv.org/abs/2006.03511>, [2006.03511](https://arxiv.org/abs/2006.03511)
68. Li D, Gupta K, Bhaduri M, Sathiadoss P, Bhatnagar S, Chong J (2024) Comparing gpt-3.5 and gpt-4 accuracy and drift in radiology diagnosis please cases. Radiology 310(1):e232411, DOI 10.1148/radiol.232411, <https://doi.org/10.1148/radiol.232411>
69. Li R, Allal LB, Zi Y, Muennighoff N, Kocetkov D, Mou C, Marone M, Akiki C, Li J, Chim J, Liu Q, Zheltonozhskii E, Zhuo TY, Wang T, Dehaene O, Davaadorj M, Lamy-Poirier J, Monteiro J, Shliazhko O, Gontier N, Meade N, Zebaze A, Yee M, Umapathi LK, Zhu J, Lipkin B, Oblokulov M, Wang Z, V RM, Stillerman J, Patel SS, Abulkhanov D, Zocca M, Dey M, Zhang Z, Moustafa-Fahmy N, Bhattacharyya U, Yu W, Singh S, Lucioni S, Villegas P, Kunakov M, Zhdanov F, Romero M, Lee T, Timor N, Ding J, Schlesinger C, Schoelkopf H, Ebert J, Dao T, Mishra M, Gu A, Robinson J, Anderson CJ, Dolan-Gavitt B, Contractor D, Reddy S, Fried D, Bahdanau D, Jernite Y, Ferrandis CM, Hughes S, Wolf T, Guha A, von Werra L, de Vries H (2023) Starcoder: may the source be with you! CoRR abs/2305.06161, DOI 10.48550/ARXIV.2305.06161, URL <https://doi.org/10.48550/arXiv.2305.06161>, [2305.06161](https://arxiv.org/abs/2305.06161)
70. Liang JT, Badea C, Bird C, DeLine R, Ford D, Forsgren N, Zimmermann T (2024) Can gpt-4 replicate empirical software engineering research? Proceedings of the ACM on Software Engineering 1(FSE):1330–1353
71. Lu Q, Zhu L, Xu X, Xing Z, Whittle J (2024) Toward responsible AI in the era of generative AI: A reference architecture for designing foundation model-based systems. IEEE Softw 41(6):91–100, DOI 10.1109/MS.2024.3406333

72. Lubos S, Felfernig A, Tran TNT, Garber D, Mansi ME, Erdeniz SP, Le V (2024) Leveraging llms for the quality assurance of software requirements. In: Liebel G, Hadar I, Spoletini P (eds) 32nd IEEE International Requirements Engineering Conference, RE 2024, Reykjavik, Iceland, June 24-28, 2024, IEEE, pp 389–397, DOI 10.1109/RE59067.2024.00046
73. Madampe K, Grundy J, Hoda R, Obie HO (2024) The struggle is real! the agony of recruiting participants for empirical software engineering studies. In: 2024 IEEE Symposium on Visual Languages and Human-Centric Computing (VL/HCC), Liverpool, UK, September 2-6, 2024, IEEE, pp 417–422, DOI 10.1109/VL/HCC60511.2024.00065
74. McDonald N, Schoenebeck S, Forte A (2019) Reliability and inter-rater reliability in qualitative research: Norms and guidelines for CSCW and HCI practice. *Proc ACM Hum Comput Interact* 3(CSCW):72:1–72:23, DOI 10.1145/3359174
75. Menzies T (2025) The case for compact AI. *Commun ACM* 68(9):6–7, DOI 10.1145/3746057
76. Microsoft (2023) Codebert on github. <https://github.com/microsoft/CodeBERT>, accessed 2025-08-15
77. Microsoft (2025) Azure OpenAI Service models. <https://learn.microsoft.com/en-us/azure/ai-services/openai/concepts/models>, accessed 2025-01-13
78. Mitu NE, Mitu GT (2024) The hidden cost of AI: Carbon footprint and mitigation strategies. *Revista de Stiinte Politice Revue des Sciences Politiques* 84:9–16, DOI 10.2139/ssrn.5036344
79. Mohamed S, Parvin A, Parra E (2024) Chatting with AI: deciphering developer conversations with chatgpt. In: Spinellis D, Bacchelli A, Constantinou E (eds) 21st IEEE/ACM International Conference on Mining Software Repositories, MSR 2024, Lisbon, Portugal, April 15-16, 2024, ACM, pp 187–191, DOI 10.1145/3643991.3645078
80. Mohsenimofidi S, Galster M, Treude C, Baltes S (2026) Context engineering for AI agents in open-source software. In: Proceedings of the 23rd IEEE/ACM International Conference on Mining Software Repositories (MSR 2026), URL <https://arxiv.org/abs/2510.21413>, 2510.21413
81. Montes CM, Feldt R, Martos CM, Ouhbi S, Premanandan S, Graziotin D (2025) Large language models in thematic analysis: Prompt engineering, evaluation, and guidelines for qualitative software engineering research. *CoRR* abs/2510.18456, DOI 10.48550/ARXIV.2510.18456, URL <https://doi.org/10.48550/arXiv.2510.18456>, 2510.18456
82. de Moraes Leça M, Valença L, Santos R, de Souza Santos R (2025) Applications and implications of large language models in qualitative analysis: A new frontier for empirical software engineering. In: IEEE/ACM International Workshop on Methodological Issues with Empirical Studies in Software Engineering, WSESE@ICSE 2025, Ottawa, ON, Canada, May 3, 2025, IEEE, pp 36–43, DOI 10.1109/WSESE66602.2025.00013
83. Moumoula MB, Kaboré AK, Klein J, Bissyandé TF (2024) Large language models for cross-language code clone detection. *CoRR*

- abs/2408.04430, DOI 10.48550/ARXIV.2408.04430, URL <https://doi.org/10.48550/arXiv.2408.04430>, 2408.04430
84. Open Source Initiative (OSI) (2025) Open Source AI Definition 1.0. <https://opensource.org/ai/open-source-ai-definition>, accessed 2025-08-26
85. OpenAI (2023) How to make your completions outputs consistent with the new seed parameter. https://cookbook.openai.com/examples/reproducible_outputs_with_the_seed_parameter, accessed 2025-01-13
86. OpenAI (2025) OpenAI API Introduction. <https://platform.openai.com/docs/api-reference/chat/streaming>, accessed 2025-01-13
87. OpenCode Contributors (2025) OpenCode: The open source AI coding agent. <https://opencode.ai/>, accessed 2026-02-27
88. Ornelas T, Araújo AA, Araújo J, Araújo M, Trinkenreich B, Kalinowski M (2025) Llm-assisted thematic analysis: Opportunities, limitations, and recommendations. CoRR abs/2511.14528, DOI 10.48550/ARXIV.2511.14528, URL <https://doi.org/10.48550/arXiv.2511.14528>, 2511.14528
89. Pan Q, Ashktorab Z, Desmond M, Cooper MS, Johnson J, Nair R, Daly E, Geyer W (2024) Human-Centered Design Recommendations for LLM-as-a-Judge. In: ACL 2024 Workshop HuCLLM, arXiv, DOI 10.48550/arXiv.2407.03479, 2407.03479
90. Pangakis N, Wolken S, Fasching N (2023) Automated annotation with generative AI requires validation. CoRR abs/2306.00176, DOI 10.48550/ARXIV.2306.00176, URL <https://doi.org/10.48550/arXiv.2306.00176>, 2306.00176
91. Panickssery A, Bowman S, Feng S (2024) Llm evaluators recognize and favor their own generations. In: Globerson A, Mackey L, Belgrave D, Fan A, Paquet U, Tomczak J, Zhang C (eds) Advances in Neural Information Processing Systems, Curran Associates, Inc., vol 37, pp 68772–68802, URL https://proceedings.neurips.cc/paper_files/paper/2024/file/7f1f0218e45f5414c79c0679633e47bc-Paper-Conference.pdf
92. Papineni K, Roukos S, Ward T, Zhu W (2002) Bleu: a method for automatic evaluation of machine translation. In: Proceedings of the 40th Annual Meeting of the Association for Computational Linguistics, ACL, pp 311–318, DOI 10.3115/1073083.1073135, URL <https://aclanthology.org/P02-1040/>
93. Pezeshkpour P, Hruschka E (2024) Large language models sensitivity to the order of options in multiple-choice questions. In: Duh K, Gómez-Adorno H, Bethard S (eds) Findings of the Association for Computational Linguistics: NAACL 2024, Mexico City, Mexico, June 16-21, 2024, Association for Computational Linguistics, Findings of ACL, vol NAACL 2024, pp 2006–2017, DOI 10.18653/V1/2024.FINDINGS-NAACL.130, URL <https://doi.org/10.18653/v1/2024.findings-naacl.130>
94. Rabbi MF, Champa AI, Zibran MF, Islam MR (2024) AI writes, we analyze: The chatgpt python code saga. In: Spinellis D, Bacchelli A, Con-

- stantinou E (eds) 21st IEEE/ACM International Conference on Mining Software Repositories, MSR 2024, Lisbon, Portugal, April 15-16, 2024, ACM, pp 177–181, DOI 10.1145/3643991.3645076
95. Ralph P, Tempero ED (2018) Construct validity in software engineering research and software metrics. In: Rainer A, MacDonell SG, Keung JW (eds) Proceedings of the 22nd International Conference on Evaluation and Assessment in Software Engineering, EASE2018, ACM, pp 13–23, DOI 10.1145/3210459.3210461
96. Ralph P, bin Ali N, Baltes S, Bianculli D, Diaz J, Dittrich Y, Ernst N, Felderer M, Feldt R, Filieri A, de França BBN, Furia CA, Gay G, Gold N, Graziotin D, He P, Hoda R, Juristo N, Kitchenham B, Lenarduzzi V, Martínez J, Melegati J, Mendez D, Menzies T, Molleri J, Pfahl D, Robbes R, Russo D, Saarimäki N, Sarro F, Taibi D, Siegmund J, Spinellis D, Staron M, Stol K, Storey MA, Taibi D, Tamburri D, Torchiano M, Treude C, Turhan B, Wang X, Vegas S (2021) Empirical standards for software engineering research. URL <https://arxiv.org/abs/2010.03525>, 2010.03525
97. Ralph P, Kuuttila M, Arif H, Ayoola B (2024) Teaching software metrology: The science of measurement for software engineering. In: Méndez D, Avgeriou P, Kalinowski M, Ali NB (eds) Handbook on Teaching Empirical Software Engineering, Springer Nature Switzerland, pp 101–154, DOI 10.1007/978-3-031-71769-7_5
98. Rasheed Z, Waseem M, Ahmad A, Kemell K, Wang X, Nguyen-Duc A, Abrahamsson P (2024) Can large language models serve as data analysts? A multi-agent assisted approach for qualitative data analysis. CoRR abs/2402.01386, DOI 10.48550/ARXIV.2402.01386, URL <https://doi.org/10.48550/arXiv.2402.01386>, 2402.01386
99. Reiss MV (2023) Testing the reliability of chatgpt for text annotation and classification: A cautionary remark. CoRR abs/2304.11085, DOI 10.48550/ARXIV.2304.11085, URL <https://doi.org/10.48550/arXiv.2304.11085>, 2304.11085
100. Ren S, Guo D, Lu S, Zhou L, Liu S, Tang D, Sundaresan N, Zhou M, Blanco A, Ma S (2020) Codebleu: a method for automatic evaluation of code synthesis. CoRR abs/2009.10297, URL <https://arxiv.org/abs/2009.10297>, 2009.10297
101. Renze M (2024) The effect of sampling temperature on problem solving in large language models. In: Findings of the Association for Computational Linguistics: EMNLP 2024, pp 7346–7356, DOI 10.18653/v1/2024.findings-emnlp.432
102. Richards J, Wessel M (2024) What you need is what you get: Theory of mind for an llm-based code understanding assistant. In: IEEE International Conference on Software Maintenance and Evolution, ICSME 2024, IEEE, pp 666–671, DOI 10.1109/ICSME58944.2024.00070
103. Ronanki K, Berger C, Horkoff J (2023) Investigating chatgpts potential to assist in requirements elicitation processes. In: 2023 49th Euromicro Conference on Software Engineering and Advanced Applications (SEAA),


- IEEE, pp 354–361
104. Rozière B, Gehring J, Gloeckle F, Sootla S, Gat I, Tan XE, Adi Y, Liu J, Remez T, Rapin J, Kozhevnikov A, Evtimov I, Bitton J, Bhatt M, Canton-Ferrer C, Grattafiori A, Xiong W, Défossez A, Copet J, Azhar F, Touvron H, Martin L, Usunier N, Scialom T, Synnaeve G (2023) Code llama: Open foundation models for code. CoRR abs/2308.12950, DOI 10.48550/ARXIV.2308.12950, URL <https://doi.org/10.48550/arXiv.2308.12950>, 2308.12950
105. Russo D (2024) Navigating the complexity of generative ai adoption in software engineering. ACM Transactions on Software Engineering and Methodology 33(5):1–50
106. Sallou J, Durieux T, Panichella A (2024) Breaking the silence: the threats of using llms in software engineering. In: Proceedings of the 2024 ACM/IEEE 44th International Conference on Software Engineering: New Ideas and Emerging Results, NIER@ICSE 2024, Lisbon, Portugal, April 14–20, 2024, ACM, pp 102–106, DOI 10.1145/3639476.3639764
107. Santos A, Vegas S, Oivo M, Juristo N (2021) A procedure and guidelines for analyzing groups of software engineering replications. IEEE Trans Software Eng 47(9):1742–1763, DOI 10.1109/TSE.2019.2935720
108. Schäfer M, Nadi S, Eghbali A, Tip F (2024) An empirical evaluation of using large language models for automated unit test generation. IEEE Trans Software Eng 50(1):85–105, DOI 10.1109/TSE.2023.3334955
109. Schilke O, Reimann M (2025) The transparency dilemma: How ai disclosure erodes trust. Organizational Behavior and Human Decision Processes 188:104405, DOI 10.1016/j.obhdp.2025.104405, URL <https://www.sciencedirect.com/science/article/pii/S0749597825000172>
110. Schneider K, Fotrousi F, Wohlrab R (2025) A reference model for empirically comparing llms with humans. In: 47th IEEE/ACM International Conference on Software Engineering: Software Engineering in Society, ICSE-SEIS 2025, Ottawa, ON, Canada, April 27 - May 3, 2025, IEEE, pp 130–134, DOI 10.1109/ICSE-SEIS66351.2025.00018
111. Schröder S, Morgenroth T, Kuhl U, Vaquet V, Paaßen B (2025) Large language models do not simulate human psychology. CoRR abs/2508.06950, DOI 10.48550/ARXIV.2508.06950, URL <https://doi.org/10.48550/arXiv.2508.06950>, 2508.06950
112. Schroeder K, Wood-Doughty Z (2024) Can you trust LLM judgments? reliability of llm-as-a-judge. CoRR abs/2412.12509, DOI 10.48550/ARXIV.2412.12509, URL <https://doi.org/10.48550/arXiv.2412.12509>, 2412.12509
113. Schulhoff S, Ilie M, Balepur N, Kahadze K, Liu A, Si C, Li Y, Gupta A, Han H, Schulhoff S, Dulepet PS, Vidyadhara S, Ki D, Agrawal S, Pham C, Kroiz GC, Li F, Tao H, Srivastava A, Costa HD, Gupta S, Rogers ML, Goncarenco I, Sarli G, Galynker I, Peskoff D, Carpuat M, White J, Anadkat S, Hoyle AM, Resnik P (2024) The prompt report: A systematic survey of prompting techniques. CoRR abs/2406.06608, DOI 10.48550/ARXIV.2406.06608, URL <https://doi.org/10.48550/>

- [arXiv:2406.06608](#), [2406.06608](#)
114. Schulz KF, Altman DG, Moher D (2010) CONSORT 2010 statement: updated guidelines for reporting parallel group randomised trials. *BMJ* 340:c332, DOI 10.1136/bmj.c332
 115. Schwartz R, Dodge J, Smith NA, Etzioni O (2020) Green AI. *Communications of the ACM* 63(12):54–63, DOI 10.1145/3381831
 116. Sclar M, Choi Y, Tsvetkov Y, Suhr A (2024) Quantifying language models’ sensitivity to spurious features in prompt design or: How I learned to start worrying about prompt formatting. In: *The Twelfth International Conference on Learning Representations, ICLR 2024, Vienna, Austria, May 7–11, 2024*, OpenReview.net, URL <https://openreview.net/forum?id=RIu5lyNXjT>
 117. Shull F, Singer J, Sjøberg DIK (eds) (2008) *Guide to Advanced Empirical Software Engineering*. Springer, DOI 10.1007/978-1-84800-044-5
 118. Silva A, Monperrus M (2025) Repairbench: Leaderboard of frontier models for program repair. In: *IEEE/ACM International Workshop on Large Language Models for Code, LLM4Code@ICSE 2025, Ottawa, ON, Canada, May 3, 2025*, IEEE, pp 9–16, DOI 10.1109/LLM4CODE66737.2025.00006, URL <https://doi.org/10.1109/LLM4Code66737.2025.00006>
 119. Sjøberg DIK, Bergersen GR (2023) Construct validity in software engineering. *IEEE Trans Software Eng* 49(3):1374–1396, DOI 10.1109/TSE.2022.3176725
 120. Song Y, Wang G, Li S, Lin BY (2025) The good, the bad, and the greedy: Evaluation of llms should not ignore non-determinism. In: Chiruzzo L, Ritter A, Wang L (eds) *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies, NAACL 2025 - Volume 1: Long Papers*, Association for Computational Linguistics, pp 4195–4206, DOI 10.18653/v1/2025.NAACL-LONG.211, URL <https://doi.org/10.18653/v1/2025.naacl-long.211>
 121. Sonnekalb T, Gruner B, Brust C, Mäder P (2022) Generalizability of code clone detection on codebert. In: *37th IEEE/ACM International Conference on Automated Software Engineering, ASE 2022, ACM, pp 143:1–143:3*, DOI 10.1145/3551349.3561165
 122. Stack Overflow (2025) 2025 developer survey. <https://survey.stackoverflow.co/2025/>, 49,009 respondents from 166 countries, May–June 2025
 123. Strubell E, Ganesh A, McCallum A (2019) Energy and policy considerations for deep learning in NLP. In: *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pp 3645–3650, DOI 10.18653/v1/P19-1355
 124. Sullivan GM, Sargeant J (2011) Qualities of qualitative research: part I. *J Grad Med Educ* 3(4):449–452
 125. Sumers TR, Yao S, Narasimhan K, Griffiths TL (2024) Cognitive architectures for language agents. *Trans Mach Learn Res* 2024, URL

- <https://openreview.net/forum?id=1i6ZCvflQJ>
126. Takerngsaksiri W, Pasuksmit J, Thongtanunam P, Tantithamthavorn C, Zhang R, Jiang F, Li J, Cook E, Chen K, Wu M (2025) Human-in-the-loop software development agents. In: 47th IEEE/ACM International Conference on Software Engineering: Software Engineering in Practice, SEIP@ICSE 2025, Ottawa, ON, Canada, April 27 - May 3, 2025, IEEE, pp 342–352, DOI 10.1109/ICSE-SEIP66354.2025.00036
 127. Tam TYC, Sivarajkumar S, Kapoor S, Stolyar AV, Polanska K, McCarthy KR, Osterhoudt H, Wu X, Visweswaran S, Fu S, et al (2024) A framework for human evaluation of large language models in healthcare derived from literature review. *NPJ digital medicine* 7(1):258
 128. Wagner S, Barón MM, Falessi D, Baltes S (2025) Towards evaluation guidelines for empirical studies involving llms. In: IEEE/ACM International Workshop on Methodological Issues with Empirical Studies in Software Engineering, WSESE@ICSE 2025, May 3, 2025, IEEE, pp 24–27, DOI 10.1109/WSESE66602.2025.00011
 129. Wan M, Safavi T, Jauhar SK, Kim Y, Counts S, Neville J, Suri S, Shah C, White RW, Yang L, Andersen R, Buscher G, Joshi D, Rangan N (2024) Tnt-llm: Text mining at scale with large language models. In: Baeza-Yates R, Bonchi F (eds) *Proceedings of the 30th ACM SIGKDD Conference on Knowledge Discovery and Data Mining, KDD 2024*, ACM, pp 5836–5847, DOI 10.1145/3637528.3671647
 130. Wang A, Morgenstern J, Dickerson JP (2025) Large language models that replace human participants can harmfully misportray and flatten identity groups. *Nature Machine Intelligence* 7:400–411, DOI 10.1038/s42256-025-00986-z
 131. Wang C, Huang K, Zhang J, Feng Y, Zhang L, Liu Y, Peng X (2024) How and why llms use deprecated apis in code completion? an empirical study. *CoRR abs/2406.09834*, DOI 10.48550/ARXIV.2406.09834, URL <https://doi.org/10.48550/arXiv.2406.09834>, [2406.09834](https://arxiv.org/abs/2406.09834)
 132. Wang F, Arora C, Liu Y, Huang K, Tantithamthavorn C, Aleti A, Sambathkumar D, Lo D (2025) Multi-modal requirements data-based acceptance criteria generation using llms. In: 40th IEEE/ACM International Conference on Automated Software Engineering, ASE 2025, Seoul, Korea, Republic of, November 16–20, 2025, IEEE, pp 3334–3345, DOI 10.1109/ASE63991.2025.00275
 133. Wang S, Liu Y, Xu Y, Zhu C, Zeng M (2021) Want to reduce labeling cost? GPT-3 can help. In: Moens M, Huang X, Specia L, Yih SW (eds) *Findings of the Association for Computational Linguistics: EMNLP 2021, Virtual Event / Punta Cana, Dominican Republic, 16–20 November, 2021, Association for Computational Linguistics, Findings of ACL*, vol EMNLP 2021, pp 4195–4205, DOI 10.18653/v1/2021.FINDINGS-EMNLP.354, URL <https://doi.org/10.18653/v1/2021.findings-emnlp.354>
 134. Wang X, Kim H, Rahman S, Mitra K, Miao Z (2024) Human-llm collaborative annotation through effective verification of LLM labels. In: Mueller FF, Kyburz P, Williamson JR, Sas C, Wilson ML, Dugas POT,

- Shklovski I (eds) Proceedings of the CHI Conference on Human Factors in Computing Systems, CHI 2024, ACM, pp 303:1–303:21, DOI 10.1145/3613904.3641960
135. Widder DG, Whittaker M, West SM (2024) Why openai systems are actually closed, and why this matters. *Nature* 635(8040):827–833
136. Wiesinger J, Marlow P, Vuskovic V (2025) Agents. Google DeepMind URL <https://gemini.google.com>, contributors: Evan Huang, Emily Xue, Olcan Sercinoglu, Sebastian Riedel, Satinder Baveja, Antonio Gulli, Anant Nawalgaria
137. Wohlin C, Runeson P, Höst M, Ohlsson MC, Regnell B, Wesslén A (2024) Experimentation in Software Engineering, Second Edition. Springer, DOI 10.1007/978-3-662-69306-3
138. Xia Y, Shao H, Deng X (2024) Vulcobert: A codebert-based system for source code vulnerability detection. In: 2024 International Conference on Generative Artificial Intelligence and Information Security, GAIIS 2024, Kuala Lumpur, Malaysia, May 10-12, 2024, ACM, DOI 10.1145/3665348.3665391
139. Xiao T, Treude C, Hata H, Matsumoto K (2024) Devgpt: Studying developer-chatgpt conversations. In: Spinellis D, Bacchelli A, Constantinou E (eds) 21st IEEE/ACM International Conference on Mining Software Repositories, MSR 2024, Lisbon, Portugal, April 15-16, 2024, ACM, pp 227–230, DOI 10.1145/3643991.3648400
140. Xu C, Guan S, Greene D, Kechadi MT (2024) Benchmark data contamination of large language models: A survey. *CoRR* abs/2406.04244, DOI 10.48550/ARXIV.2406.04244, URL <https://doi.org/10.48550/arxiv.2406.04244>, 2406.04244
141. Xu R, Sun Y, Ren M, Guo S, Pan R, Lin H, Sun L, Han X (2024) AI for social science and social science of AI: A survey. *Inf Process Manag* 61(2):103665, DOI 10.1016/J.IPM.2024.103665, URL <https://doi.org/10.1016/j.ipm.2024.103665>
142. Yan L, Hwang A, Wu Z, Head A (2024) Ivie: Lightweight anchored explanations of just-generated code. In: Mueller FF, Kyburz P, Williamson JR, Sas C, Wilson ML, Dugas POT, Shklovski I (eds) Proceedings of the CHI Conference on Human Factors in Computing Systems, CHI 2024, ACM, pp 140:1–140:15, DOI 10.1145/3613904.3642239
143. Yang G, Zhou Y, Chen X, Zhang X, Han T, Chen T (2023) Exploitgen: Template-augmented exploit code generation based on codebert. *J Syst Softw* 197:111577, DOI 10.1016/J.JSS.2022.111577, URL <https://doi.org/10.1016/j.jss.2022.111577>
144. Yuan J, Li H, Ding X, Xie W, Li YJ, Zhao W, Wan K, Shi J, Hu X, Liu Z (2025) Understanding and mitigating numerical sources of nondeterminism in LLM inference. In: Advances in Neural Information Processing Systems 38: Annual Conference on Neural Information Processing Systems 2025, NeurIPS 2025, URL <https://openreview.net/forum?id=Q3qAsZAEZw>, 2506.09501

145. Zhao C, Habule M, Zhang W (2025) Large language models (llms) as research subjects: Status, opportunities and challenges. *New Ideas in Psychology* 79:101167, DOI 10.1016/j.newideapsych.2025.101167, URL <https://www.sciencedirect.com/science/article/pii/S0732118X25000236>
146. Zhou R, Chen L, Yu K (2024) Is LLM a reliable reviewer? A comprehensive evaluation of LLM on automatic paper reviewing tasks. In: Calzolari N, Kan M, Hoste V, Lenci A, Sakti S, Xue N (eds) *Proceedings of the 2024 Joint International Conference on Computational Linguistics, Language Resources and Evaluation, LREC/COLING 2024*, 20-25 May, 2024, Torino, Italy, ELRA and ICCL, pp 9340–9351, URL <https://aclanthology.org/2024.lrec-main.816>
147. Zhou W, Jiang YE, Li L, Wu J, Wang T, Qiu S, Zhang J, Chen J, Wu R, Wang S, Zhu S, Chen J, Zhang W, Zhang N, Chen H, Cui P, Sachan M (2023) Agents: An open-source framework for autonomous language agents. *CoRR abs/2309.07870*, DOI 10.48550/ARXIV.2309.07870, URL <https://doi.org/10.48550/arXiv.2309.07870>, [2309.07870](https://arxiv.org/abs/2309.07870)
148. Zhu Y, Zhang P, ul Haq E, Hui P, Tyson G (2023) Can chatgpt reproduce human-generated labels? A study of social computing tasks. *CoRR abs/2304.10145*, DOI 10.48550/ARXIV.2304.10145, URL <https://doi.org/10.48550/arXiv.2304.10145>, [2304.10145](https://arxiv.org/abs/2304.10145)



[Click here to access/download](#)
Supplementary Material
diff.pdf

