

Empirical Software Engineering

Guidelines for Empirical Studies in Software Engineering involving Large Language Models

--Manuscript Draft--

Manuscript Number:	EMSE-D-25-00637R2
Full Title:	Guidelines for Empirical Studies in Software Engineering involving Large Language Models
Article Type:	Research Papers
Keywords:	software engineering; Large Language Models; empirical research; meta-science; guidelines
Corresponding Author:	Sebastian Baltes Heidelberg University Bayreuth, GERMANY
Corresponding Author Secondary Information:	
Corresponding Author's Institution:	Heidelberg University
Corresponding Author's Secondary Institution:	
First Author:	Sebastian Baltes
First Author Secondary Information:	
Order of Authors:	Sebastian Baltes
	Florian Angermeir
	Chetan Arora
	Marvin Muñoz Barón
	Chunyang Chen
	Lukas Böhme
	Fabio Calefato
	Neil Ernst
	Davide Falessi
	Brian Fitzgerald
	Davide Fucci
	Junda He
	Christoph Treude
	Marcos Kalinowski
	Stefano Lambiase
	Daniel Russo
	Mircea Lungu
	Cristina Martinez Montes
	Lutz Prechelt
	Paul Ralph
	Rijnard van Tonder
	Stefan Wagner

Order of Authors Secondary Information:		
Funding Information:	Carlsbergfondet (CF24-0693)	Prof. Dr. Daniel Russo
	Alfred P. Sloan Foundation (G-2024-22586)	Prof. Dr. Daniel Russo
Abstract:	Large Language Models (LLMs) are widely used in software engineering (SE) research and practice, yet their non-determinism, opaque training data, and rapidly evolving models threaten the reproducibility and replicability of empirical studies. We address this challenge through a collaborative effort of 22 researchers, presenting a taxonomy of seven study types that organizes how LLMs are used in SE research, together with eight guidelines for designing and reporting such studies. Each guideline distinguishes requirements (must) from recommended practices (should) and is contextualized by the study types it applies to. Our guidelines recommend that researchers: (1) declare LLM usage and role; (2) report model versions, configurations, and customizations; (3) document the tool architecture beyond the model; (4) disclose prompts, their development, and interaction logs; (5) validate LLM outputs with humans; (6) include an open LLM as a baseline; (7) use suitable baselines, benchmarks, and metrics; and (8) articulate limitations and mitigations. We complement the guidelines with an applicability matrix mapping guidelines to study types and a reporting checklist for authors and reviewers. We maintain the study types and guidelines online as a living resource for the community to use and shape (llm-guidelines.org).	

Response Letter (Minor Revision)

Guidelines for Empirical Studies in Software Engineering Involving Large Language Models

EMSE-D-25-00637R1

S. Baltes et al.

Dear editor and reviewers,

Thank you for the favorable evaluation of our revision and for the additional, focused feedback from Reviewer 1. We have addressed each of the remaining points; brief responses follow below.

Best regards,
Sebastian Baltes (on behalf of all co-authors)

RESPONSE TO EDITOR

Comment 0.1:

Thank you for the revision. Please address the remaining feedback by Reviewer 1.

Response:

We thank the editors and reviewers for the positive assessment. We addressed every point raised by Reviewer 1, as detailed below. We also thank Reviewers 2 and 3 for confirming that their concerns have been resolved.

Beyond Reviewer 1’s points, we took the opportunity to incorporate further input from discussions during ICSE 2026 (April 2026) and the 3rd Copenhagen Symposium on Human-Centered Software Engineering AI (May 2026), and to strengthen the construct-validity guidance for benchmarks based on two recent surveys (Cao et al.’s review of 572 benchmarks and Bean et al.’s review of 445 benchmarks). The substantive additions include new **should** recommendations for benchmarks (defining the phenomenon, justifying the sampling strategy, isolating confounders, error analysis, contamination prevention) and human validation (rater demographics for value-laden constructs); a new **should** on recording runtime traces in an OTLP-compatible format such as the OpenTelemetry GenAI semantic conventions or OpenInference; bracketed conditional tags (e.g., `[fine-tuning]`, `[agents]`) prefixing applicability-conditioned items in the reporting checklist; and a bibliography pass that upgraded seven arXiv preprints to their peer-reviewed venues. The companion website now hosts a changelog documenting all changes between guideline versions, available at <https://llm-guidelines.org/changelog/>.

RESPONSE TO REVIEWERS

Reviewer 1

Comment 1.1:

The authors are to be commended for essentially rewriting the entire paper following the suggestions of the first-round reviews. Most of my comments were sufficiently addressed or rebutted. The following minor ones can be easily accommodated in a minor revision without further review.

Response:

We thank the reviewer for the encouraging assessment and for the precise, actionable feedback. We have addressed all remaining points item by item below.

Comment 1.2:

In common with other reviewers, please refrain from using uppercase MUST/SHOULD in the text. You already agreed that this publication is not a standard, so keeping these in the text is simply distracting. Retaining the distinction and emphasis in Table 3 is more than enough, though even there, I’d prefer to use the terms “mandatory” and “desirable” as is done in the SIGSOFT empirical standards.

Response:

We agree that the uppercase rendering read as standards-style language and have dropped it. As a compromise we now set **must**, **must not**, **should**, and **should not** in lowercase bold throughout the text rather than plain text: the lowercase resolves the standards-style concern, while keeping the keywords visually distinct preserves their role as scan-anchors that let a reader skim a recommendation paragraph and pick out the strength of each clause without reading every sentence in full.

On the suggestion to use *mandatory* and *desirable* in Table 3, we hope the reviewer agrees with us keeping **must** and **should**, for the same reason we gave in the previous round: after dropping the RFC reference, we retained **must/should** as plain indicators of essential vs. desired practices, in the same spirit as the SIGSOFT Empirical Standards’ *mandatory/desirable* categories. The categorical distinction is also carried visually in Table 3 by the filled (●) and open (○) icons defined in the caption, so the table conveys the same essential/desired split that the SIGSOFT Empirical Standards’ terms encode. Using a single vocabulary across Table 3, the guideline recommendations, and the reporting checklist (Appendix C) lets readers map matrix cells directly to the corresponding sentences in the body, which we believe is more useful than aligning Table 3 with the SIGSOFT Empirical Standards’ terminology while the surrounding prose uses different words.

Comment 1.3:

In Table 3, try to name (in short form, e.g. Annotator, Judge, Synthesis, Subject, Usage, Tools, Benchmarks) the study types. The labels are already angled, so it should be easy to make these fit.

Response:

Done. The column headers in Table 3 now read *Annotator*, *Judge*, *Synthesis*, *Subject*, *Usage*, *Tools*, and *Benchmarks*, exactly as suggested. Each label is hyperlinked to the corresponding study-type subsection. As part of this change, we also dropped the S1–S7 and G1–G8 identifiers from the body text, since the short-form names are sufficient throughout.

Comment 1.4:

Rather than the cute but possibly unfamiliar “tl;dr”, simply name the corresponding boxes as “Summary”.

Response:

Done. The framed box at the start of each guideline is now labeled *Summary* instead of *tl;dr*, in the paper as well as on the companion website.

Comment 1.5:

Rather than naming sections as “Example(s)” name them as Example / Examples according to their contents.

Response:

Done. Every subsection contains multiple examples, so all *Example(s)* subheadings now read *Examples* across the guidelines and study-type sections.

Comment 1.6:

The names of the “See Also” elements are now confusing because they don’t match the section titles. Please use the same section titles. If the change was made to avoid the unsightliness of long hyperlinks, simply hyperlink only the section number.

Response:

Done, using the approach the reviewer suggested. All “See Also” entries now read “Section *n* (*Full Section Title*): *rationale*”, with only the section number hyperlinked.

Comment 1.7:

In Appendix A, eliminate the confusing and redundant empty boxes, in favor of using the shall/must symbols as the bullet figure.

Response:

Done. The empty checkbox markers (☐) have been removed from all itemize lists in Appendix A. Each item is now introduced directly by its **must** (●) or **should** (◦) icon, which serves as the bullet figure, matching the symbology of Table 3.

Comment 1.8:

Please fix the capitalization in the references; e.g. The chatgpt python code saga → The ChatGPT Python code saga

Response:

Done. We reviewed the bibliography for title-case errors and corrected the capitalization of proper nouns and acronyms (e.g., *ChatGPT*, *Python*, *AI*, *API*). The example flagged by the reviewer now renders as “AI Writes, We Analyze: The ChatGPT Python Code Saga”, and similar fixes were applied across the bibliography.

Reviewer 2

Comment 2.1:

Thanks for the revision. This paper will help the community, my own PhD student independently found the arxiv version and shared it with the group :)

Response:

Thank you for this very encouraging note. We are glad that the work is already proving useful in practice, and we will continue to evolve the guidelines through the companion website (<https://llm-guidelines.org>) as the community engages with them.

Reviewer 3

Comment 3.1:

Thank you for the thorough and thoughtful major revision of your manuscript. I appreciate the care you took in addressing the comments and concerns raised in the previous review. The revisions have significantly improved the clarity and quality of the work, and I am satisfied that all of my concerns have been adequately addressed.

Response:

Thank you for the careful first-round review and for confirming that the revision addressed your concerns.

Guidelines for Empirical Studies in Software Engineering involving Large Language Models

Sebastian Baltes, Florian Angermeir, Chetan Arora, Marvin Muñoz Barón,
Chunyang Chen, Lukas Böhme, Fabio Calefato, Neil Ernst, Davide Falessi,
Brian Fitzgerald, Davide Fucci, Junda He, Christoph Treude,
Marcos Kalinowski, Stefano Lambiase, Daniel Russo, Mircea Lungu,
Cristina Martinez Montes, Lutz Prechelt, Paul Ralph, Rijnard van Tonder,
and Stefan Wagner

Detailed author information can be found on the next page.

Guidelines for Empirical Studies in Software Engineering involving Large Language Models

Sebastian Baltes · Florian Angermeir ·
Chetan Arora · Marvin Muñoz Barón ·
Chunyang Chen · Lukas Böhme ·
Fabio Calefato · Neil Ernst ·
Davide Falessi · Brian Fitzgerald ·
Davide Fucci · Junda He ·
Christoph Treude ·
Marcos Kalinowski · Stefano Lambiase ·
Daniel Russo · Mircea Lungu ·
Cristina Martinez Montes ·
Lutz Prechelt · Paul Ralph ·
Rijnard van Tonder · Stefan Wagner

Submitted: 26/08/2025

Revised: 27/02/2026, 05/05/2026

Sebastian Baltes
Heidelberg University, Germany
E-mail: sebastian.baltes@uni-heidelberg.de
(corresponding author)

Florian Angermeir
fortiss, Germany
Blekinge Institute of Technology, Sweden
E-mail: angermeir@fortiss.org

Chetan Arora
Monash University, Australia
E-mail: chetan.arora@monash.edu

Marvin Muñoz Barón · Chunyang Chen
Technical University of Munich, Germany
E-mail: {marvin.munoz-baron, chun-yang.chen}@tum.de

Lukas Böhme
Hasso-Plattner-Institut, Germany
University of Potsdam, Germany
E-mail: lukas.boehme@hpi.de

Fabio Calefato
University of Bari, Italy
E-mail: fabio.calefato@uniba.it

Neil Ernst
University of Victoria, Canada
E-mail: nernst@uvic.ca

Abstract Large Language Models (LLMs) are widely used in software engineering (SE) research and practice, yet their non-determinism, opaque training data, and rapidly evolving models threaten the reproducibility and replicability of empirical studies. We address this challenge through a collaborative effort of 22 researchers, presenting a taxonomy of seven study types that organizes how LLMs are used in SE research, together with eight guidelines for designing and reporting such studies. Each guideline distinguishes requirements (must) from recommended practices (should) and is contextualized by the study types it applies to. Our guidelines recommend that researchers: (1) declare LLM usage and role; (2) report model versions, configurations, and cus-

Davide Falessi

University of Rome Tor Vergata, Italy

E-mail: falessi@ing.uniroma2.it

Brian Fitzgerald

Lero, Ireland

University of Limerick, Ireland

E-mail: brian.fitzgerald@ul.ie

Davide Fucci

Blekinge Institute of Technology, Sweden

E-mail: davide.fucci@bth.se

Junda He · Christoph Treude

Singapore Management University, Singapore

E-mail: {jundahe, ctreude}@smu.edu.sg

Marcos Kalinowski

PUC Rio de Janeiro, Brazil

E-mail: kalinowski@inf.puc-rio.br

Stefano Lambiase · Daniel Russo

Aalborg University in Copenhagen, Denmark

E-mail: {stla, daniel.russo}@cs.aau.dk

Mircea Lungu

IT University of Copenhagen, Denmark

E-mail: mlun@itu.dk

Cristina Martinez Montes

Chalmers University of Technology, Sweden

University of Gothenburg, Sweden

E-mail: montesc@chalmers.se

Lutz Prechelt

Freie Universität Berlin, Germany

E-mail: prechelt@inf.fu-berlin.de

Paul Ralph

Dalhousie University, Canada

E-mail: paulralph@dal.ca

Rijnard van Tonder

Independent Researcher, Antigua and Barbuda

E-mail: rvantonder@gmail.com

Stefan Wagner

Technical University of Munich, Germany

E-mail: stefan.wagner@tum.de

tomizations; (3) document the tool architecture beyond the model; (4) disclose prompts, their development, and interaction logs; (5) validate LLM outputs with humans; (6) include an open LLM as a baseline; (7) use suitable baselines, benchmarks, and metrics; and (8) articulate limitations and mitigations. We complement the guidelines with an applicability matrix mapping guidelines to study types and a reporting checklist for authors and reviewers. We maintain the study types and guidelines online as a living resource for the community to use and shape (llm-guidelines.org).

Keywords software engineering, large language models, empirical research, meta-science, guidelines

1 Introduction

Since the release of ChatGPT in November 2022, large language models (LLMs) have been adopted widely across software engineering (SE) research [57], yet the reproducibility and replicability of empirical studies involving LLMs remains uncertain. Recent findings indicate a low reproducibility of SE studies involving LLMs [6]. Three characteristics of LLMs make this particularly challenging: their inherent non-determinism causes variability across runs [22, 131], with slight changes leading to substantially different results [2, 123]; commercial models evolve beyond version identifiers, so reported performance can change over time [30]; and even for “open” models, training data and fine-tuning details often remain undisclosed [46]. Moreover, prompt formatting choices alone can shift accuracy by tens of percentage points [127], and configured parameters such as temperature affect output variability [111]. Hence, not reporting these settings directly affects reproducibility.

Although the SE research community has developed guidelines for conducting and reporting specific types of empirical studies such as controlled experiments (e.g., [128, 149]), their replications (e.g., [118]), or empirical studies in general (e.g., the *ACM SIGSOFT Empirical Standards* [106]), none of these address the LLM-specific aspects described above. Previously, a position paper highlighted these issues [139], but there was no comprehensive community-developed guidance for designing and reporting empirical studies involving LLMs in SE.

Therefore, we present community-developed guidelines for conducting and reporting studies involving LLMs in SE research, co-developed by 22 researchers. After outlining our *Scope*, we introduce a taxonomy of *Study Types*, then present eight *Guidelines*. We complement these with an applicability matrix mapping guidelines to study types and a reporting checklist for authors and reviewers. For each study type and guideline, we identify relevant examples, both within and outside of SE research. We maintain the guidelines online as a living resource for the community to use and shape (llm-guidelines.org).

2 Scope and Conventions

Software Engineering as our Target Discipline: We target SE research because existing cross-disciplinary guidelines do not address its needs. For instance, Gallifant et al’s [43] guidelines for LLM use in healthcare research include discipline-specific items irrelevant to most SE studies (e.g., clinical-care usability) while lacking guidance on tool architectures, benchmarking, and code-specific evaluation that SE research requires. Moreover, SE research employs a wide variety of empirical methods [106, 149]. Our work builds on Sallou et al’s vision paper on mitigating validity threats in LLM-based SE research [117] and our own position paper [139], providing more detailed advice for a wider variety of research methods organized into a taxonomy of *Study Types* (Section 4).

Focus on Text-Based Use Cases: While multi-modal foundation models that use or generate images, audio, or video may also support SE research and practice, we focus on textual use cases of LLMs (e.g., in natural language or programming languages). Many of our guidelines, particularly those concerning model reporting, prompt documentation, and reproducibility, are likely applicable to multi-modal settings as well, but we leave their explicit validation to future work.

Focus on Direct Development or Research Support: While researchers may use LLMs for many peripheral tasks (e.g., proof-reading, spell-checking, translation), our guidelines focus on their direct role in empirical research and engineering practice. For engineers, we focus on the use of LLMs to automate SE tasks, that is, artificial intelligence (AI) for software engineering (AI4SE) (see *LLMs for SE*). This includes agentic systems that autonomously plan and execute multi-step tasks using LLMs (see *System and Prompt Design*). For researchers, we focus on the use of LLMs to automate empirical research tasks such as data collection, processing, or analysis (see *LLMs for Research*).

Researchers as our Target Audience: Our guidelines are intended to help SE researchers design, plan, conduct, and report empirical studies involving LLMs, and to support scholarly peer review of such studies. Each guideline includes an *Advice for Reviewers* subsection with targeted assessment suggestions. Our guidelines focus on what to report and how; they complement but do not replace methodological guidance for designing specific types of empirical studies.

Reporting Locations: We use *paper* to refer to the manuscript PDF, including any appendices bound with it. We use *supplementary material* to refer to any artifact, replication package, dataset, or other resource external to the manuscript PDF (e.g., hosted on Zenodo or Figshare). What belongs in the main body versus an appendix is a presentation choice for authors and venues; what belongs inside the manuscript versus outside it is the reporting-location

decision the guidelines make. When a recommendation does not specify a location, either *paper* or *supplementary material* is acceptable.

Navigating These Guidelines: These guidelines are structured to support different reading strategies depending on the reader’s goal. *Researchers planning a new study* may start with the taxonomy of study types (see *Study Types*) to identify which types apply to their planned work, then consult Table 3 to determine which guidelines are requirements (**must**) and which are recommendations (**should**) for those study types. Each guideline section opens with a brief summary in a shaded box, allowing readers to quickly assess relevance before reading the full text. *Researchers writing up results* may prefer to start with the checklist in Appendix C, which organizes actionable items by typical paper sections (Introduction, Research Design and Methods, Results, etc.). Table 2 provides a quick reference to all eight guidelines, and Table 4 maps each guideline’s rationale to its recommendations. *Reviewers* can use the *Advice for Reviewers* subsection at the end of each guideline for targeted guidance on assessing manuscripts. Table 3 helps reviewers identify which guidelines apply to the study type under review.

3 Methodology

This project was initiated at the 2024 meeting of the *International Software Engineering Research Network* (ISERN) in Barcelona, Spain, where the first and last author organized a session on guidelines for empirical studies involving LLMs. These discussions led to a position paper at *WSESE 2025* [139], which was presented at the *2nd Copenhagen Symposium on Human-Centered Software Engineering AI* (CHASEAI 2024), forming a workstream to collaboratively refine and extend the preliminary guidelines. Our process follows an established tradition of developing reporting guidelines through expert collaboration, as exemplified by CONSORT for clinical trials [21] and, more recently, TRIPOD-LLM for LLM-based prediction studies in medicine [43].

We held bi-weekly meetings, in which we decided on the structure of the study type taxonomy and the guidelines. During these discussions, we added a new study type (*LLMs for Synthesis*), refined and extended the guideline sections (in particular *Version and Configuration*, *System and Prompt Design*, *Session Traces*) and covered *Benchmarks and Metrics* as well as *Limitations and Mitigations*.

Then, at least two co-authors were assigned to each study type and guideline to refine and extend the content, followed by a peer review and refinement phase. Afterwards, the first author reviewed all sections and discussed potential changes with the topic leads. A requirement for co-authorship was that each author contributed or reviewed content and, most importantly, read the complete guidelines and confirmed that they stand behind all recommendations, not only their own contributions.

Table 1 Overview of study types.

Section	Title	Short Name
4.1	LLMs as Tools for Software Engineering Researchers	
4.1.1	LLMs as Annotators	Annotator
4.1.2	LLMs as Judges	Judge
4.1.3	LLMs for Synthesis	Synthesis
4.1.4	LLMs as Subjects	Subject
4.2	LLMs as Tools for Software Engineers	
4.2.1	Studying LLM Usage in Software Engineering	Usage
4.2.2	LLMs for New Software Engineering Tools	Tools
4.2.3	Benchmarking LLMs for Software Engineering Tasks	Benchmarks

To select illustrative examples for each study type and guideline, co-author pairs independently searched for relevant papers using academic databases. This initial pool was extended based on suggestions from other co-authors. Identified papers were assessed by one author and subsequently reviewed by a second author; papers that did not add new insights beyond those already covered were not included. Multiple review rounds were then carried out on the guidelines as a whole, during which selected papers were cross-checked by additional authors and, where necessary, further examples were added or replaced. The examples are intended to be illustrative, not the result of a systematic review; the involvement of multiple authors at different stages helped mitigate individual selection bias.

After completing the internal review, we invited external experts in empirical SE to review the study types and guidelines. During the review process of the *Empirical Software Engineering* journal, the structure and content were improved based on reviewers' suggestions. We self-assessed our guidelines against the quality dimensions of the *SIGSOFT Empirical Standards* [106]; this assessment is reported in the conclusion.

4 Study Types

Empirical guidelines are needed for SE studies involving LLMs to ensure the validity and reproducibility of their results. However, such guidelines must be tailored to different study types that each pose unique challenges. Therefore, we created a taxonomy of study types to contextualize our recommendations. We use the term *study type* to refer to categories of LLM involvement in empirical research, rather than to research methodologies (e.g., experiments, case studies). A single empirical study may involve multiple such study types. Each study type section starts with a **description**, followed by **examples** from the SE research community and beyond. The **advantages** and **challenges** of using LLMs are discussed in a summary subsection at the end of this section, synthesizing cross-cutting themes across all study types. See Table 1 for an overview of study types.

4.1 LLMs as Tools for Software Engineering Researchers

LLMs can serve as tools to help researchers conduct empirical studies. They can automate various tasks such as data collection, pre-processing, and analysis. For example, LLMs can apply predefined coding guides to qualitative datasets (*LLMs as Annotators*), assess the quality of software artifacts (*LLMs as Judges*), generate summaries of research papers (*LLMs for Synthesis*), or simulate human behavior in empirical studies (*LLMs as Subjects*). If LLMs could complete these tasks well enough, they would reduce the time and effort required to conduct a study. However, these applications each pose challenges to validity and reproducibility.

4.1.1 LLMs as Annotators

Description: In qualitative data analysis, manually annotating (“coding”) natural language text (e.g., requirements, interview transcripts, open-ended survey responses) is a time-consuming manual process [16]. LLMs can augment human coding, suggest new codes, and label artifacts based on a predefined coding guide much faster than humans can [55]. The extent to which, or under what conditions, LLMs can perform these tasks *effectively* remains an open research question [4]. Indeed, measuring their effectiveness is practically and philosophically challenging. From an interpretivist philosophical perspective, one cannot measure the quality of analysis by comparing one (human or machine) analyst’s work to another. From a realist perspective, triangulating across multiple human judges, LLM judges, and other data sources (e.g., whether a pull request is marked as having resolved an issue) improves confidence in the findings but does not prove that any one judge is valid.

Examples: Huang et al [60] used multiple LLMs for joint annotation of mobile application reviews. They used three models of comparable size with an absolute majority voting rule (i.e., a label is only accepted if it receives more than half of the total votes from the models). This approach slightly outperformed the best individual model tested. Meanwhile, Ahmed et al [4] examined LLMs as annotators in SE research across five datasets, six LLMs, and ten annotation tasks. They found that model-model agreement strongly correlates with human-model agreement; models performed poorly in tasks where humans also frequently disagreed. They proposed to use model confidence scores to identify specific samples that could be safely delegated to LLMs, potentially reducing human annotation effort without compromising inter-rater agreement.

4.1.2 LLMs as Judges

Description: LLMs can rate properties of software artifacts (e.g. assess code readability, adherence to coding standards, or comment quality) or sort multiple solutions by some attribute. These kinds of judgments are distinct from annotating unstructured text (see *LLMs as Annotators*).

Examples: Lubos et al [81] used Llama-2 to evaluate the quality of software requirements statements. They prompted the LLM with the text below, where the words in braces reflect the study parameters:

```
Your task is to evaluate the quality of a software requirement.
Evaluate whether the following requirement is
    {quality_characteristic}.
{quality_characteristic} means:
    {quality_characteristic_explanation}
The evaluation result must be: 'yes' or 'no'.
Request: Based on the following description of the project:
    {project_description}
Evaluate the quality of the following requirement:
    {requirement}.
Explain your decision and suggest an improved version.
```

They evaluated LLM output against expert human judges and found moderate agreement for simple requirements and poor agreement for more complex requirements. In contrast, Wang et al [143] used an LLM to generate acceptance criteria for user stories, and provided a rubric to an LLM to judge the generated acceptance criteria on interpretable scales (0 to 4).

4.1.3 LLMs for Synthesis

Description: Unlike annotation (see *LLMs as Annotators*), which focuses on categorizing or labeling individual data points, synthesis refers to the process of integrating and interpreting information from multiple sources to generate higher-level insights, identify patterns across datasets, and develop conceptual frameworks or theories. LLMs may be able to support synthesis tasks in SE research by processing and distilling information from qualitative data sources. Although synthesis in the preceding notion refers to abstraction and interpretation across multiple data sources, the term is sometimes also used to refer to generating synthetic content (e.g., source code, bug-fix pairs, or requirements) that is then used in downstream tasks to train, fine-tune, or evaluate existing models or tools. In this case, the synthesis is done primarily using the LLM and its training data; the input is limited to basic instructions and examples.

Examples: Published examples of applying LLMs for synthesis in SE remain scarce; however, some recent work in other domains is instructive [16]. Barros et al [18] conducted a systematic mapping study on using LLMs for qualitative research and found that most studies focused on the healthcare and social sciences. In these studies, LLMs supported qualitative methods, such

as grounded theory and thematic analysis, by aiding in pattern identification. In SE, de Moraes Leça et al [91] explored how LLMs have been applied for qualitative data analysis (QDA) and proposed general strategies and guidelines for their application. Ornelas et al [98] complemented this perspective by studying the opportunities and limitations of introducing LLM-based support into QDA, and by formulating recommendations for embedding humanAI collaboration across the thematic analysis phases. Building on these, subsequent work has proposed hybrid frameworks combining LLM support with human-led QDA. Rasheed et al [108] designed an LLM-driven multi-agent system that integrates AI with human decision-making to automate qualitative data analysis methods. Their system generated initial codes, developed themes, and summarized text. Similarly, Montes et al [90] compared the performance of humans and LLMs in coding, theme development, definition, and refinement, creating guidelines for a hybrid-LLM framework. Finally, El-Hajjami and Salinesi's work is an example of using LLMs to create synthetic datasets. They present an approach to generate synthetic requirements, showing that they “*can match or surpass human-authored requirements for specific classification tasks*” [41].

4.1.4 LLMs as Subjects

Description: In empirical studies, data is collected from participants through methods such as surveys, interviews, or controlled experiments. LLMs can serve as virtual *subjects* by simulating human behavior and interactions. If LLMs can generate responses that approximate those of human participants, they could be valuable for research involving user interactions, collaborative coding environments, and software usability assessments [158]. To achieve this, prompt engineering techniques are widely employed; for instance, the *Personas Pattern* [72] involves tailoring LLM responses to align with predefined profiles or roles that emulate specific user archetypes. To serve as virtual subjects, generated responses should be indistinguishable from human-produced texts, consistent with the attitudes and sociodemographic information of the conditioning context (e.g., junior vs. senior developers), naturally aligned with the form, tone, and content of the simulated scenario, and reflect patterns in relationships between ideas, demographics, and behavior observed in comparable human data [9].

Examples: Xu et al [153] compiled a list of ways LLMs can support social science research, some of which transfer to empirical SE research. For example, LLMs can emulate human responses and behaviors in simulated interviews and focus groups [45]. Similarly, Bano et al [17], investigated biases in LLM-generated candidate profiles in SE recruitment processes. They found biases favoring male candidates, lighter skin tones, and slim physiques, particularly for senior roles. LLMs may be able to simulate end-user feedback and behavior in usability studies, identify usability issues, and offer suggestions for improvement based on predefined user personas.

4.2 LLMs as Tools for Software Engineers

LLM-based assistants have become a widely adopted tool for software engineers [133], supporting them in various tasks such as code generation and debugging. Researchers have studied how software engineers use LLMs (*Studying LLM Usage*), developed new tools that integrate LLMs (*LLMs for Tools*), and benchmarked LLMs for software engineering tasks (*Benchmarking LLMs*).

4.2.1 Studying LLM Usage in Software Engineering

Description: Studying how software engineers use LLMs and LLM-based tools is important for understanding the current state of practice in SE. Researchers can observe software engineers' usage of LLM-based tools in the field, or study if and how they adopt such tools, their usage patterns, as well as perceived benefits and challenges. Surveys, interviews, observational studies, or analysis of usage logs can provide insights into how LLMs are integrated into development processes, how they influence decision making, and what factors affect their acceptance and effectiveness. Such studies can inform improvements for existing LLM-based tools, motivate the design of novel tools, or derive best practices for LLM-assisted software engineering. They can also uncover risks or deficiencies of existing tools.

Examples: Based on a convergent mixed-methods study, Russo has found that early adoption of generative AI by software engineers is primarily driven by compatibility with existing workflows [116]. Khojah et al investigated the use of ChatGPT (GPT-3.5) by professional software engineers in a week-long observational study [69]. They found that most developers do not use the code generated by ChatGPT directly but instead use the output as a guide to implement their own solutions. Azanza et al's case study found that LLMs could “*enhance personalized, instant onboarding support; however, relying on proprietary external LLMs poses significant data privacy risks*” [13]. Jahic and Sami found that most participants from 15 software companies they surveyed had already adopted AI (especially ChatGPT) for SE tasks; but cited copyright and privacy issues, as well as inconsistent or low-quality outputs, as barriers to adoption [64]. Retrospective studies that analyze data generated while developers use LLMs can provide additional insights into human-LLM interactions. For example, researchers can employ data mining methods to build large-scale conversation datasets, such as the DevGPT dataset introduced by Xiao et al [151]. Conversations can then be analyzed using quantitative [104] and qualitative [88] analysis methods.

4.2.2 LLMs for New Software Engineering Tools

Description: LLMs are being integrated into new tools that support software engineers in their daily tasks (e.g., to assist in code comprehension [154] and test case generation [119]). One way of integrating LLM-based tools into software engineers' workflows is using GenAI agents. Unlike traditional LLM-based tools, these agents are capable of acting autonomously and proactively, are often tailored to meet specific user needs, and can interact with external environments [147, 156]. From an architectural perspective, GenAI agents can be implemented in various ways [147]. However, they generally share three key components: (1) a reasoning mechanism that guides the LLM (often enabled by advanced prompt engineering), (2) a set of tools to interact with external systems (e.g., APIs or databases), and (3) a user communication interface that extends beyond traditional chat-based interactions [136, 160]. Researchers can also test and compare different tool architectures to increase artifact quality and developer satisfaction.

Examples: Yan et al proposed *IVIE*, a tool integrated into the VS Code graphical interface that generates and explains code using LLMs [154]. The authors focused more on the presentation, providing a user-friendly interface to interact with the LLM. Schäfer et al [119] presented a large-scale empirical evaluation on the effectiveness of LLMs for automated unit test generation. They presented *TestPilot*, a tool that implements an approach in which the LLM is provided with prompts that include the signature and implementation of a function under test, along with usage examples extracted from the documentation. Richards and Wessel introduced a preliminary GenAI agent designed to assist developers in understanding source code by incorporating a reasoning component grounded in the theory of mind [112]. Takerngsaksiri et al [137] present *HULA*, a multi-agent system deployed in Atlassian JIRA that lets engineers refine LLM-generated coding plans and source code, and report acceptance and modification rates from real users.

4.2.3 Benchmarking LLMs for Software Engineering Tasks

Description: Benchmarking is the process of evaluating an LLM's performance on standardized tasks, using standardized metrics, on a standardized dataset. The LLM output is compared to a supposed ground truth from the benchmark dataset. Typical tasks include code generation, code summarization, code completion, and code repair [156], but also natural language processing tasks, such as anaphora resolution (i.e., the task of identifying the referring expression of a word or phrase occurring earlier in the text). Metrics may include general metrics for text generation, such as *ROUGE*, *BLEU*, or *METEOR* [57], or task-specific metrics, such as *CodeBLEU* for code generation. Benchmarking requires high-quality reference datasets.

Examples: In SE, benchmarking may include evaluating an LLM’s ability to produce accurate and reliable outputs for a given input (usually a task description, possibly accompanied by data obtained from curated real-world projects or from synthetic SE-specific datasets). *RepairBench* [129], for example, contains 574 buggy Java methods and their corresponding fixed versions, which can be used to evaluate the performance of LLMs in code repair tasks. It uses the *Plausible@1* metric (i.e., the probability that the first generated patch passes all test cases) and the *AST Match@1* metric (i.e., the probability that the abstract syntax tree of the first generated patch matches the ground truth patch). *SWE-Bench* [66] is a more generic benchmark that contains 2,294 SE Python tasks extracted from GitHub pull requests. To score the LLM’s task performance, the benchmark validates whether the generated patch successfully compiles and calculates the percentage of passed test cases. Meanwhile, *HumanEval* [31] is often used to assess code generation.

4.3 Advantages and Challenges

Using LLMs in empirical SE research, whether as tools for researchers or for software engineers, offers several potential advantages but also raises fundamental challenges that cut across the study types described above.

Advantages: The primary advantage of using LLMs for research tasks is *speed, cost reduction, and scalability*. LLMs can annotate, judge, synthesize, and simulate faster and at lower cost than human researchers, with studies showing cost reductions of 50–96% on various natural language tasks [144] (e.g., He et al found that GPT-4 annotation required only two days and 122.08 USD compared to several weeks and 4,508 USD for a comparable MTurk pipeline [55]). Similarly, augmenting or replacing human participants with LLM-generated virtual participants would reduce recruitment effort [82]. This efficiency unlocks scalability: qualitative research traditionally does not scale well to large samples, but LLMs allow researchers to analyze larger datasets [16, 18, 91] and support larger judgment datasets.

LLMs can also *automate* tasks such as coding qualitative data, assessing artifact quality, and generating summaries, reducing cognitive demands and resources required for qualitative research. Some studies suggest that LLMs may improve *consistency*: ChatGPT’s accuracy exceeded crowd workers by approximately 25%, and LLMs can achieve higher inter-rater agreement than crowd workers and trained annotators [47]. However, no compelling evidence supports claims that LLMs are less biased than human annotators [47].

LLMs could potentially provide *access to otherwise-inaccessible research contexts*. If virtual participants’ behavior is sufficiently similar to human participants, LLMs could access underrepresented and hard-to-reach populations, strengthen generalizability and inclusiveness, impute missing data (see *LLMs for Synthesis*), and enable research that is ethically problematic with real hu-

mans (e.g., questions that would force a human to relive past trauma do not harm an LLM).

Studying real-world LLM-based tools allows researchers to *understand the state of practice*, uncovering usage patterns, adoption rates, and contextual factors, and *generate hypotheses* about how LLMs affect developer productivity, collaboration, and decision-making. From an engineering perspective, developing LLM-based tools *requires less task-specific engineering* than traditional SE approaches such as static analysis or symbolic execution, because a single model can handle diverse inputs without language-specific parsers or hand-crafted rules. Good benchmarks provide *standardized evaluation and model comparison*, reducing research effort and supporting open science by providing common ground for sharing data and results. Benchmarks built for specific SE tasks can help identify LLM weaknesses and support optimization and fine-tuning.

Challenges: The stochastic nature of LLM responses, where identical prompts may yield different outputs, *undermines replicability* across all study types, complicating interpretation of experimental results and test-retest reliability. More broadly, *reliability* is a persistent concern: conceptualizing an LLM-as-judge or LLM-as-annotator as a measurement instrument, researchers should expect validity, test-retest reliability, inter-rater reliability, minimal error, and measurement invariance [107]. However, LLMs show substantial variability depending on the dataset and task [100], are sensitive to prompt variations [109] and option order [103], can behave differently when reviewing their own outputs [101], and can be unreliable for high-stakes labeling [145]. Crucially, *reliability does not imply validity*: a reliable LLM might be reliably inaccurate [159], and for tasks with no single correct answer, the statistical framework pushes outputs toward the most likely answer, which may not be the best one [19]. See *Version and Configuration* and *Limitations and Mitigations* for reporting guidance on replicability and reliability.

The *rapid evolution* of LLMs and LLM-based tools, combined with professionals rapidly learning how to employ them, complicates longitudinal comparisons and may quickly render study findings obsolete (e.g., GPT-4’s accuracy on a simple mathematical task dropped from 84% to 51% within three months [30]). Findings tied to a specific model version may not transfer to later releases even within the same model family, complicating reproducibility and longitudinal comparisons.

The prevalence of *proprietary tools and opaque training data* limits researchers’ ability to assess and mitigate biases, and enables benchmark contamination [5]. LLMs exhibit *bias* in multiple forms: tendencies to overestimate certain labels [161], well-documented fairness issues [42], and the potential to reinforce prejudices. When simulating human participants, LLMs are “*likely to both misportray and flatten the representations of demographic groups*” [141]. See *Open LLMs* and *Limitations and Mitigations* for mitigation strategies.

Evaluation difficulty is inherent in studying LLM-based tools and benchmarks. Benchmarks may lack construct validity [107], usually do not capture

the full complexity of software engineering work [28], and may lead to *overconfidence* and overfitting [140]. LLMs are also susceptible to adversarial manipulation where semantics-preserving changes may deceive them into accepting flawed artifacts [54]. See *Benchmarks and Metrics* for detailed guidance on benchmark design, including Cao et al’s [27] guidelines for coding task benchmarks.

A fundamental challenge is *philosophical and methodological incongruence* with qualitative research. In a recent open letter, 419 experienced qualitative researchers argued “*that analytic approaches such as reflexive thematic analysis are human research practices requiring a subjective, positioned, and reflexive researcher and therefore the use of GenAI in such approaches is not methodologically congruent*” [67]. LLMs lack the capacity for genuinely reflexive qualitative analysis because they operate on statistical prediction without understanding the meaning of the data being analyzed [15, 90]. Effective use requires structured prompts and careful human oversight [18]; see *Human Validation* and *Limitations and Mitigations* for guidance.

There is *insufficient evidence* for the effectiveness of LLMs in most research roles. No compelling evidence exists that LLMs can accurately simulate human participants [122] or reliably judge most relevant properties of SE artifacts; gathering such evidence for each specific usage may be quite difficult [52]. LLMs may be better suited for augmenting rather than replacing human researchers [55, 144], but even then, limited evidence exists that augmentation increases *effectiveness* as well as *efficiency* [122]. When LLM outputs are incorrect, they can negatively *affect human judgment* [59, 99]; see *Human Validation* for mitigation strategies.

Majority voting across multiple outputs improves reliability but increases *cost and environmental impact* [109, 144]. While open models are available, the most capable ones require substantial hardware; relying on *cloud-based APIs* introduces concerns related to data privacy, security, and replicability. See *Limitations and Mitigations* for sustainability considerations.

Field study findings face *generalizability* challenges because outcomes may be highly sensitive to individual differences, usage patterns, goals, and contexts. Field studies must be “dependable” [135] beyond traditional validity criteria, which complicates methodology; see *Limitations and Mitigations* for a detailed threat taxonomy.

5 Guidelines

A primary goal of our guidelines is to *enable reproducibility and replicability* of empirical SE studies involving LLMs. As repeating LLM-focused research to verify results lies somewhere between the ACM’s definitions of reproducibility (different team, same research artifacts) and repeatability (different team, different research artifacts) [1] due to potential model changes and imperfect research artifacts, we follow Angermeir et al [6] and use the terms interchangeably. While previous guidelines regarding open science and

empirical studies still apply, LLM-specific characteristics (e.g. inherent non-determinism [131, 157], opaque and proprietary models) present additional replicability challenges, which, in turn, demand new guidance.

Each guideline below begins with a brief **summary**, followed by its **rationale**, **recommendations**, **examples**, **benefits**, **challenges**, links to the relevant **study types**. The *rationale* articulates the underlying principle, i.e., why the guideline matters, while the *recommendations* provide concrete, actionable practices. Table 4 in the appendix provides a compact overview of this mapping.

The guidelines further contain **advice for reviewers**. Broadly, reviewers should use our guidelines to help them interrogate the extent to which a manuscript’s authors have done what was practically possible to improve reproducibility. We must neither accept research absent reasonable efforts to improve reproducibility, nor reject research for failing to obtain an impossible goal. We borrow this principle of “reasonable efforts” from the SIGSOFT Empirical Standards [106], where it applies to methodological rigor more generally. Of course, papers should acknowledge their limitations, but to determine whether these limitations are reasonable, reviewers should ask “have the authors done what they could to minimize limitations?” Our guidelines attempt to capture what “reasonable efforts” practically means for LLM-based studies.

To distinguish essential criteria from recommendations, our guidelines use two tiers. A **must** criterion is a *requirement*. Studies that intend to follow these guidelines are expected to meet all **must** criteria. A **should** criterion represents a desired practice that strengthens a study’s rigor or transparency. However, there may be valid reasons to deviate in particular circumstances (e.g., resource constraints, inapplicability to a specific study context or type). Nonetheless, studies that deviate from a **should** criterion should briefly justify the deviation and discuss its potential impact (e.g., on validity or reproducibility). Table 2 lists our eight guidelines. Table 3 in the appendix provides an overview of how each guideline applies to each study type. Cells marked ● indicate that the guideline is a **must** requirement for the study type and cells marked ● indicate that the guideline **should** be followed. Cells marked – indicate that the guideline is typically not applicable or not feasible for the study type.

The following sections indicate which information we expect researchers to report, and whether it should be in the *paper* or *supplementary material*. Where a publication venue’s page limits hinder reporting all expected elements in the *paper*, it is better to report essential information in the *supplementary material* than not at all. The *supplementary material* should be published according to the *ACM SIGSOFT Open Science Policies* [49].

Table 2 Overview of guidelines.

Section	Title	Short Name
5.1	Declare LLM Usage and Role	Declare Usage
5.2	Report Model Version, Configuration, and Customizations	Model Version
5.3	Report System and Prompt Design	Design
5.4	Report Session Traces	Traces
5.5	Use Suitable Baselines, Benchmarks, and Metrics	Benchmarks
5.6	Use an Open LLM as a Baseline	Open LLM
5.7	Use Human Validation for LLM Outputs	Human Validation
5.8	Report Limitations and Mitigations	Limitations

5.1 Declare LLM Usage and Role

Summary: Researchers **must** disclose any use of LLMs to support empirical studies in their *paper*, specifying which LLM was used, how it was used, and where in the research process it was employed. They **should** report the exact purpose, the tasks that were automated, and the expected benefits in the *paper*. When the LLM is central to the study, the declaration **should** be prominent and detailed in the methodology section; for tangential uses, a brief statement in the methodology or acknowledgments suffices.

5.1.1 *Rationale:*

Transparency about LLM involvement is a prerequisite for informed assessment of a study’s scope, limitations, and potential biases. Without explicit disclosure, readers cannot evaluate how the LLM’s characteristics may have influenced the research process or its outcomes.

5.1.2 *Recommendations:*

When conducting any kind of empirical study involving LLMs, researchers **must** clearly declare that an LLM was used (see *Scope* for what we consider relevant research support). This **should** be done in a suitable section of the *paper*, for example, in the introduction or research methods section; Cheng et al [33] argue specifically for the methods section, since acknowledgments are easily missed at the end of a paper. For authoring scientific articles, this transparency is, for example, required by the ACM Policy on Authorship: “*The use of generative AI tools and technologies to create content is permitted but must be fully disclosed in the Work*” [11].

Beyond generic declarations, researchers **should** report the exact purpose of using an LLM in a study, the tasks it was used to automate, and the expected benefits in the *paper*. A sufficient declaration specifies not only *that* an LLM was used, but also *which* LLM (name and version), *how* it was used (e.g., as

an annotator, code generator, or judge), and *where* in the research process it was employed (e.g., data collection, analysis, or synthesis).

When the LLM is central to the study (e.g., as the main tool being evaluated or as a core component of the research method), the declaration **should** be prominent and detailed, appearing in the methodology section with cross-references to the specific guidelines that apply (e.g., Sections *Version and Configuration*, *System and Prompt Design*, and *Session Traces*). When the LLM’s role is more tangential (e.g., used for a single preprocessing step), a brief but explicit statement in the methodology or acknowledgments section is sufficient. In either case, the disclosure **must** be specific enough for readers to assess how the LLM’s involvement may affect the study’s validity and reproducibility.

5.1.3 *Examples:*

The *ACM Policy on Authorship* [11] suggests disclosing GenAI usage in the acknowledgments section of the *paper*, advising to “err on the side of caution, and include a disclosure in the acknowledgments section of the Work” when uncertain about the need [11]. For double-blind review, researchers can add a temporary “AI Disclosure” section where the acknowledgments would appear. An example of an LLM disclosure beyond writing support can be found in a recent paper by Lubos et al [81], in which they write in the methodology section:

“We conducted an LLM-based evaluation of requirements utilizing the Llama 2 language model with 70 billion parameters, fine-tuned to complete chat responses...”

A more contemporary declaration could similarly state:

“We used Claude Opus 4.7 via the Anthropic API to synthesize themes from interview transcripts, with all prompts and conversation logs published as supplementary material.”

5.1.4 *Benefits:*

Transparency in the use of LLMs helps other researchers understand the context and scope of the study, supporting better interpretation and comparison of the results. Beyond this declaration, we recommend researchers to be explicit about the LLM version they used (see *Version and Configuration*) and the LLM’s exact role.

5.1.5 Challenges:

Declaring LLM usage requires only a brief statement and no additional experiments, making compliance straightforward. One challenge might be authors' reluctance to disclose LLM usage for valid use cases, because they fear that AI-generated content makes reviewers think that the authors' work is less original. In fact, there is evidence suggesting that AI disclosure can negatively affect trust in authors [120]. However, the *ACM Policy on Authorship* is very clear in that any use of GenAI tools to create content **must** be disclosed. Our guidelines focus on research support beyond proof-reading and writing support (see *Scope*), but the threshold of what must be declared continues to evolve as organizations such as the ACM update their authorship policies.

5.1.6 Study Types:

Researchers **must** follow this guideline for all study types. The specific focus of the declaration varies by study type. For *LLMs as Annotators*, *LLMs as Judges*, *LLMs for Synthesis*, and *LLMs as Subjects*, researchers **must** declare the specific role assigned to the LLM (e.g., annotator, judge, synthesizer, or simulated participant). For *Studying LLM Usage*, researchers **must** clarify which LLM(s) the observed participants used and under which conditions. For *LLMs for Tools*, researchers **must** declare the LLM's role within the tool architecture and its contribution to the tool's functionality. For *Benchmarking LLMs*, researchers **must** declare which LLMs were benchmarked and for which tasks.

5.1.7 Advice for Reviewers:

The most common problem with disclosure is incompleteness or vagueness about how the LLM was used. If the paper says "we used LLM X to help with task Y" without specifying how, reviewers should request clarification. Such requests are typically minor revisions unless the missing details may reveal methodological problems.

Using an LLM as a copyeditor becomes problematic when authors do not or cannot take responsibility for the resulting text. If a reviewer finds clear evidence that LLM-generated text was not carefully reviewed (e.g., text like "As a Large Language Model, I..."), this may warrant rejection. If a reviewer suspects an author *cannot* take responsibility for AI-generated text, they should raise the concern with their editor or program chair. Due process requires that authors not be accused of misconduct without clear evidence.

If undisclosed LLM use is suspected, the reviewer should similarly consult their editor or program chair. When the evidence is conclusive, the key question is the degree to which undisclosed use affects the study's contribution, ranging from negligible (e.g., word choice in a single sentence) to severe (e.g., generating ostensibly empirical data or statistical analyses).

5.1.8 *See Also:*

- Section 5.2 (Report Model Version, Configuration, and Customizations): Disclosing that an LLM was used is incomplete without naming which model and version.
- Section 5.3 (Report System and Prompt Design): When the LLM lives inside a tool or agent, authors must also describe that tool’s architecture and prompts.
- Section 5.4 (Report Session Traces): Session traces show what the LLM did during the study.

5.2 Report Model Version, Configuration, and Customizations

Summary: Researchers **must** report the exact LLM model or tool version, configuration, and experiment date in the *paper*. When using quantized models, researchers **should** report the quantization level and method. For fine-tuned models, they **must** describe the fine-tuning goal, dataset, and procedure in the *paper*. Researchers **should** include default parameters, explain model choices, compare base- and fine-tuned model using suitable metrics and benchmarks, and share fine-tuning data and weights as *supplementary material* (or alternatively justify in the *paper* why they cannot share them).

5.2.1 *Rationale:*

LLMs and LLM-based tools are frequently updated, and configuration parameters such as temperature or seed values affect content generation. This guideline focuses on documenting the *model-specific* aspects of empirical studies involving LLMs, concentrating on the models themselves, their version, configuration parameters, and customizations (e.g., fine-tuning). While the *System and Prompt Design* section addresses system-level integration and the authored artifacts (including prompts) that the model uses on each call, the information outlined here is essential for reproducibility whenever an LLM is involved.

5.2.2 Recommendations:

Researchers **must** document in the *paper* which model or tool version they used in their study, along with the date when the experiments were carried out and the configured parameters that affect output generation. Since default values might change over time, researchers **should** report all configuration values, even if they used the defaults. Checksums and fingerprints **should** be reported since they identify specific versions and configurations. Depending on the study context, other properties such as the context window size (number of tokens) **should** be reported. When using quantized models, researchers **should** report the quantization level (e.g., 4-bit, 8-bit) and method (e.g., GPTQ or AWQ), as different quantization approaches produce different outputs, affecting both output quality and reproducibility. Researchers **should** motivate in the *paper* why they selected certain models, versions, and configurations. Reasons may be monetary, technical, or methodological (e.g., planned comparison to previous work). Depending on the specific study context, additional information regarding the experiment or tool architecture **should** be reported.

A common customization approach for existing LLMs is fine-tuning. If a model was fine-tuned, researchers **must** describe the fine-tuning goal (e.g., improving the performance for a specific task), the fine-tuning procedure (e.g., full fine-tuning versus Low-Rank Adaptation (LoRA), selected hyperparameters, loss function, learning rate, and batch size), and the fine-tuning dataset (e.g., data sources, the preprocessing pipeline, dataset size) in the *paper*. Researchers **should** either share the fine-tuning dataset as part of the *supplementary material* or explain in the *paper* why the data cannot be shared (e.g., because it contains confidential or personal data that could not be anonymized). The same applies to the fine-tuned model weights. Suitable benchmarks and metrics **should** be used to compare the base model with the fine-tuned model.

In summary, researchers **must** report in the *paper* at minimum (1) the exact model or tool name and version, (2) all configured parameters that affect output generation, (3) the experiment date, and, for fine-tuned models, (4) the fine-tuning goal, dataset characterization, and parameters and procedure. Beyond these requirements, researchers **should** additionally report default parameter values, checksums or fingerprints, model properties relevant to the study (e.g., context-window size), and quantization level and method where applicable; for fine-tuned models, they **should** share the dataset and model weights as *supplementary material* (unless legal or privacy constraints prevent disclosure) and report validation metrics and benchmarks.

Commercial models (e.g., GPT-5) or LLM-based tools (e.g., ChatGPT) might not give researchers access to all required information. For these tools, researchers **should** report what is available and openly acknowledge limitations that hinder reproducibility.

5.2.3 *Examples:*

Based on the documentation that OpenAI and Azure provide [86, 95], researchers might, for example, report:

“We integrated a `gpt-4` model in version 0125-Preview via the Azure OpenAI Service, and configured it with a temperature of 0.7, `top_p` set to 0.8, a maximum token length of 512, and the seed value 23487. We ran our experiment on 10th January 2025. The system fingerprint was `fp_6b68a8204b`.

Kang et al provide a similar statement in their paper on exploring LLM-based bug reproduction [68]:

“We access OpenAI Codex via its closed beta API, using the `code-davinci-002` model. For Codex, we set the temperature to 0.7, and the maximum number of tokens to 256.”

Our guidelines additionally recommend reporting a checksum/fingerprint and exact dates; otherwise, this example is close to our recommendations.

Dhar et al [37] assessed whether LLMs can generate architectural design decisions, detailing the system architecture and the LLM’s role within it. They provide information on the fine-tuning approach and datasets, including the source of architectural decision records, preprocessing methods, and data selection criteria.

For self-hosted models, the *supplementary material* can become a true replication package. For example, for models provisioned using ollama, one can report the specific tag and checksum, e.g., “`llama3.3`, tag `70b-instruct-q8_0`, checksum `d5b5e1b84868`.” Given suitable hardware, running the model is then as easy as executing the following command:
`ollama run llama3.3:70b-instruct-q8_0`

5.2.4 *Benefits:*

Reporting this information is a prerequisite for the verification, reproduction, and replication of LLM-based studies. While LLMs are inherently non-deterministic, this cannot excuse dismissing reproducibility. Although exact reproducibility is hard to achieve, reporting the information outlined in this guideline helps researchers come as close as possible to that standard.

5.2.5 Challenges:

Different model providers and modes of operating the models allow for varying degrees of information. For example, OpenAI provides a model version and a system fingerprint describing the backend configuration, which can also influence the output. However, the fingerprint is intended only to detect changes in the model or its configuration; one cannot go back to a certain fingerprint. As a beta feature, OpenAI lets users set a seed parameter to receive “(mostly) consistent output” [94]. However, the seed value does not allow for full reproducibility and the fingerprint changes frequently. Although, as motivated above, open models substantially simplify re-running experiments, they also come with challenges in terms of reproducibility, as generated outputs can be inconsistent despite setting the temperature to 0 and using a seed value (see GitHub issue for Llama3). Setting the temperature to 0 configures greedy decoding (always selecting the most probable next token), which minimizes output variability but can degrade quality by producing repetitive text and missing higher-quality responses [56].

Even with a temperature of 0, full determinism is rarely guaranteed: floating-point arithmetic on GPUs causes slight numerical differences that cascade into divergent token selections [157], Sparse Mixture-of-Experts routing amplifies this effect [29], silent backend changes in commercial APIs produce different outputs over time [30], and even self-hosted open models with identical settings do not always yield consistent outputs [6, 131]. Researchers **should not** treat a temperature of 0 as a guarantee of reproducibility, but as one measure among several, including fixed seed values [94], system fingerprints, and archiving of raw outputs. When a temperature of 0 is chosen primarily for reproducibility, this motivation **should** be stated explicitly, along with an acknowledgment of its potential impact on output quality.

5.2.6 Study Types:

This guideline **must** be followed for all study types for which the researcher has access to (parts of) the model’s configuration. They **must** always report the configuration that is visible to them, acknowledging the reproducibility challenges of commercial tools and models that are offered as-a-service. Depending on the specific study type, researchers **should** provide additional information on the system and prompt design (see *System and Prompt Design*), session traces (see *Session Traces*), and specific limitations and mitigations (see *Limitations and Mitigations*).

For example, when *Studying LLM Usage* by focusing on commercial tools such as ChatGPT or GitHub Copilot, researchers **must** be as specific as possible in describing their study setup. The configured model name, version, and the date when the experiment was conducted **must** always be reported. See *System and Prompt Design* for prompt reporting and *Session Traces* for interaction logs.

For *LLMs as Annotators*, *LLMs as Judges*, and *LLMs for Synthesis*, researchers **must** report the model configuration used for the respective annotation, judging, or synthesis tasks, including temperature and other sampling parameters that affect output variability. For *LLMs as Subjects*, researchers **must** report any persona-related configuration settings and parameters that configure the simulated behavior. For *LLMs for Tools*, researchers **must** report the configuration for each model integrated in the tool’s architecture, including any model-specific parameter choices. For *Benchmarking LLMs*, researchers **must** report the configuration for all benchmarked models to enable fair cross-model comparisons.

5.2.7 Advice for Reviewers:

Missing version, configuration, or parameter information is typically a minor revision request. Before concluding that information is absent, reviewers should check appendices and supplementary materials, as details are sometimes reported there rather than in the main text. Rejection over missing details is rarely warranted unless the omissions obscure deeper methodological problems.

5.2.8 See Also:

- Section 5.3 (Report System and Prompt Design): Beyond the model itself, authors must also document the architecture and prompts that use it.
- Section 5.4 (Report Session Traces): Session traces show what the reported model and configuration produced at runtime.
- Section 5.5 (Use Suitable Baselines, Benchmarks, and Metrics): Benchmark comparisons require identifying the exact model version under test.
- Section 5.6 (Use an Open LLM as a Baseline): An open model gives full version visibility, which some commercial products do not.
- Section 5.8 (Report Limitations and Mitigations): Hidden or shifting commercial versions become reproducibility threats in their own right.

5.3 Report System and Prompt Design

Summary: Researchers **must** describe in the *paper* the full architecture of LLM-based tools, from standalone uses to agentic systems, including the LLM’s role and its interactions with other components. Hosting, infrastructure, and latency **must** be reported. Researchers **must** publish all prompts as *supplementary material*, including prompt templates with representative instances and the dynamic generation process where applicable; sensitive content **must** be anonymized. If full prompt disclosure is not feasible, summaries or examples **should** be provided. The prompting strategy and prompt reuse across models and configurations **must** be specified, and

how the prompts were developed **should** be described. Researchers **must** describe context-file mechanisms used and summarize in the *paper* which tools and skills were exposed; full file contents and the complete tool catalog (schemas, definitions, Model Context Protocol (MCP) servers) **should** be published as *supplementary material*. For agentic systems, researchers **must** specify the agents’ roles, reasoning frameworks, communication flows, and agent behavior traceability across LLM output, tool calls, and user interactions. For retrieval-augmented generation (RAG) and similar methods, researchers **must** describe how external data was retrieved, stored, and integrated. Where legally possible, the implementation **should** be open-sourced; non-disclosed proprietary components **must** be acknowledged as reproducibility limitations.

5.3.1 *Rationale:*

LLM-based studies rest on artifacts that researchers design, author, or configure before any model is invoked: *software layers* that pre-process data, prepare prompts, filter user requests, or post-process responses, and the *context* those layers feed into each call [44]. Context includes prompt templates, context files, tool and skill schemas, and retrieval mechanisms that bring external data into the call. For example, ChatGPT and GitHub Copilot use the same underlying models, but their outputs differ substantially because Copilot automatically adds project context. Researchers can also build tools using models directly via APIs.

Prompts are central to any LLM-based study [127]. A *prompt* is a concrete input to an LLM that guides its output [124]. Depending on the task, a prompt may include instructions (e.g., “*classify the following bug report*”), background information, input data, and output format specifications (e.g., “*respond as JSON with fields ‘category’ and ‘justification’*”), with outputs ranging from unstructured text to structured formats such as JSON [36]. A *prompt template* is a parameterized structure containing static elements (e.g., instructions, output format specifications) and placeholders for variable content (e.g., source code under analysis) that are filled in at runtime to construct concrete prompts [124]. In automated studies, researchers typically design prompt templates from which individual prompts are then instantiated. Prompts substantially influence a model’s output, so how they are formatted and integrated into an LLM-based study is essential for transparency, verifiability, and reproducibility. Sclar et al [127] question the methodological validity of comparing models with “*an arbitrarily chosen, fixed prompt format*”, because their research has shown that the performance of different prompt formats only weakly correlates between different models.

This guideline addresses everything researchers put in place before runs begin, complementing *Version and Configuration* (model-specific details) and *Session Traces* (runtime behavior). It does not apply when LLMs are used

solely for language polishing, paraphrasing, translation, tone or style adaptation, or layout edits (see *Scope*).

5.3.2 Recommendations:

Researchers **must** clearly describe the tool architecture and what exactly the LLM (or ensemble of LLMs) contributes to the tool or method presented in a research paper, including any dependencies on proprietary tools that affect reproducibility. Researchers **should** explain design decisions and which retrieval mechanisms were implemented (e.g., keyword search, semantic similarity matching, rule-based extraction). Researchers **should** describe how the models were hosted and accessed; for *time-sensitive measurements*, the hosting choice (e.g., self-hosting on local hardware, an aggregator such as OpenRouter, or a vendor API such as the OpenAI API) can substantially impact results, so researchers **must** clarify whether local infrastructure or cloud services were used, including detailed infrastructure specifications and latency considerations. Where legally possible (e.g., not restricted by industry-partner agreements), researchers **should** release the source code of their implementation under an open-source license.

Standalone LLMs require limited architectural documentation, while study setups that integrate LLMs with other components, configure agents, or depend on particular context files or tool schemas require more detailed reporting. In the topic-specific paragraphs that follow, the *paper must* contain a high-level description of any reported component, with full details provided as *supplementary material*.

Prompt Reporting: Researchers **must** report all prompts used in an empirical study, including instructions, background information, input data, and output indicators. The complete set **must** be made publicly available as *supplementary material*; when the paper cannot include all prompts in full, it **should** provide summaries and representative examples. The only exception to this disclosure requirement is when privacy, anonymity, or confidentiality concerns (e.g., when working with industry partners) prevent complete disclosure; in such cases, researchers **should** provide summaries and representative examples instead. For prompts containing sensitive information, researchers **must** anonymize personal identifiers, replace proprietary code with placeholders, and clearly highlight modified sections. When prompt templates are used, researchers **must** report them alongside representative instances, specifying which parts are static and which are dynamically filled. When prompts are generated dynamically (e.g., through preprocessing or retrieval-augmented generation (RAG)), the generation process **must** be thoroughly documented, including any automated algorithms or rules. Researchers **should** specify the exact formatting of prompts, including how code snippets were enclosed (e.g., triple backticks), whether whitespace was preserved, and how other artifacts such as error messages and stack traces were presented. For studies involv-

ing human participants who create or modify prompts, researchers **should** describe how these prompts were collected and analyzed.

Prompt Development: Researchers **should** explain in the *paper* how they developed the prompts and why they decided to follow certain prompting strategies. If prompts from the early phases of a research project are unavailable, researchers **should** at least summarize the prompt evolution.

Prompt development is often iterative, involving collaboration between human researchers and AI tools. Researchers **should** report any instances in which LLMs were used to suggest prompt refinements and how these suggestions were incorporated. A prompt changelog can track prompt evolution, including revisions and reasons for changes (e.g., v1.0: initial prompt; v2.0: incorporated examples of ideal responses). Because prompt effectiveness varies between models and model versions, researchers **must** make clear which prompts were used for which models in which versions and with which configuration.

Prompting Strategy and Input Handling: Researchers **must** specify whether zero-shot, one-shot, or few-shot prompting was used. For few-shot prompts, the examples provided to the model **should** be clearly outlined, along with the rationale for selecting them. If multiple versions of a prompt were tested, researchers **should** describe how these variations were evaluated and how the final design was chosen.

When dealing with extensive or complex prompt context, researchers **should** describe the strategies they used to handle input length constraints (e.g., truncating, summarizing, or splitting prompts into multiple parts). Token optimization measures, such as simplifying code formatting or removing unnecessary comments, **should** also be documented if applied.

Context Files and Agent Configuration: Agentic coding tools expose a range of configuration and extension mechanisms including *context files*, skills, sub-agents, hooks, settings, rules, and Model Context Protocol (MCP) servers [44]. Context files such as `AGENTS.md` or `CLAUDE.md` are version-controlled Markdown files containing persistent, project-specific instructions automatically included in the context on each call [89]. Researchers **must** describe which context-file mechanisms were used and **should** publish the full contents as *supplementary material*, because they steer agent behavior similarly to system prompts. Because context files are version-controlled, their evolution across a study is recoverable from the project’s git history and from the runtime traces reported under *Session Traces*.

Tool Catalog and Skill Definitions: Agentic systems expose a set of callable tools, skills, and sub-agents to the model on each call, along with their schemas: parameter names, types, and natural-language descriptions as the model sees them. Tools vary in granularity, from one operation per tool to generic execution interfaces. Claude Code, for example, does not register individual

command-line utilities but defines a single **Bash** tool (**PowerShell** on Windows) through which the agent runs arbitrary shell commands such as **grep** or **ls** [8]. The catalog, the wording of descriptions, and the order in which tools are presented all influence which tool the model selects, so the exact serialized form matters for reproducibility. Researchers **must** summarize in the *paper* which tools and skills were exposed to the model, and **should** include the complete list (names with one-line purposes), tool schemas, skill definitions, sub-agent definitions, and the identities of any connected Model Context Protocol (MCP) servers as *supplementary material*. Where sub-agents are used, researchers **should** also describe the delegation pattern: which sub-agent handles which kind of task, and how control returns to the orchestrator.

Pipelines and Agentic Systems: If an LLM is used as a *standalone system*, for example by sending prompts directly to a GPT-5 model via the OpenAI API without pre-processing the prompts or post-processing the responses, a brief explanation is usually sufficient. For *complex systems* with pre-processing, retrieval mechanisms, or autonomous agents, aspects to consider include how the LLM interacts with other components such as databases, external APIs, and frameworks. If the LLM is part of an *agent-based system* that autonomously plans or executes tasks, researchers **must** describe its exact architecture, including the agents' roles (e.g., planner, executor, coordinator), whether it is a single-agent or multi-agent system, how it interacts with external tools and users, and the reasoning framework used (e.g., chain-of-thought, self-reflection, multi-turn dialogue). For agent-based systems that use external tools (e.g., Claude Code and its subagents), researchers **must** describe *agent behavior traceability* by delineating: (1) *LLM Input/Output*, the internal deliberation, planning, and interpretation performed by the model; (2) *External tool calls*, the agent's explicit invocations of APIs, databases, file systems, or other external services (e.g., via the Model Context Protocol (MCP)); (3) *User or system interactions*, human-in-the-loop feedback, environment responses, or multi-agent communication. This separation lets readers attribute task success to the LLM, the external tools, or their interaction; runtime trace-reporting for these components is covered in *Session Traces*.

RAG and Ensembles: If retrieval-augmented generation (RAG) or related methods were used (e.g., rule-based retrieval, structured query generation, or hybrid approaches), researchers **must** describe how external data was retrieved, stored, and integrated into the LLM's responses, including the type of storage (e.g., vector databases, relational databases, knowledge graphs) and how retrieved information was selected. Stored data used for context augmentation **should** be reported, including data preprocessing, versioning, and update frequency; if not confidential, an anonymized snapshot **should** be made available as *supplementary material*. For *ensemble models*, in addition to following the *Version and Configuration* guideline for each model, researchers **should** describe the architecture connecting them: the routing logic that determines which model handles which input, model interactions, and the

output combination strategy (e.g., majority voting, weighted averaging, sequential processing).

5.3.3 *Examples:*

Schäfer et al [119] evaluated LLMs for automated unit test generation, providing a detailed description of the system architecture including code parsing, prompt formulation, LLM interaction, and test suite integration. They also detail the datasets used, including sources, selection criteria, and preprocessing steps.

A second example is Yan et al’s *IVIE* tool [154], which integrates LLMs into the VS Code interface. The authors document the tool architecture, detailing the IDE integration, context extraction from code editors, and the formatting pipeline for LLM-generated explanations.

Liang et al’s paper [78] is a good example of comprehensive prompt reporting. The authors make the exact prompts available in their *supplementary material* on Figshare, including details such as code blocks enclosed in triple backticks. The *paper* explains the rationale behind the prompt design and the data output format, and it includes an overview figure and two concrete examples, keeping the main text concise while remaining reproducible.

5.3.4 *Benefits:*

Documenting the architecture and supplemental data of LLM-based systems strengthens reproducibility and transparency [80], enabling experiment replication, result validation, and cross-study comparison. Detailed prompt documentation strengthens verifiability, reproducibility, and comparability of LLM-based studies, letting other researchers replicate studies, refine prompts, and evaluate how different content types and formatting choices influence LLM behavior. Reporting the tool catalog and context files enables other researchers to reconstruct the exact context the model saw, which is often the dominant factor in agent behavior.

5.3.5 *Challenges:*

Documenting LLM-based architectures involves challenges such as proprietary APIs and dependencies that restrict disclosure, managing large-scale retrieval databases, and ensuring efficient query execution. Researchers must also balance transparency with data privacy concerns, adapt to the evolving nature of LLM integrations, and handle the complexity of multi-agent interactions and decision-making logic, all of which can impact reproducibility and system clarity. Prompts themselves are challenging to document because they often combine multiple components such as code, error messages, and explanatory text, and privacy or confidentiality concerns can hinder sharing.

Not all systems allow reporting of complete (system) prompts, context files, and tool schemas. For commercial tools, researchers **must** report all available information and acknowledge unknown aspects as limitations. Disclosure practices vary across vendors: some publish their system prompts, others keep them proprietary. Where prompts are published, they also change between releases [148], so researchers **should** record the tool version and date of use even when the prompt content itself is unavailable. Understanding suggestions of commercial tools such as *GitHub Copilot* might require recreating the exact state of the codebase at the time the suggestion was made, which is a challenging context to report. One solution is to use version control to capture the exact state of the codebase when a recommendation was made, keeping track of the files that were automatically added as context. We also recommend exploring open-source tools such as *OpenCode* [96], which expose more of the configuration that controls agent behavior.

5.3.6 Study Types:

This guideline **must** be followed for all studies that involve tools with system-level components beyond bare LLMs, from lightweight wrappers that pre-process user input or post-process model outputs, to systems employing retrieval-augmented methods or complex agent-based architectures. It also **must** be followed by all studies that use concrete prompts or prompt templates.

For *LLMs for Tools*, this guideline is of primary importance: researchers **must** describe the tool’s full architecture, explain how prompts were generated and structured within the tool, report the context files, tool catalog, and skill definitions used, and document how the role of each LLM fits into the overall system behavior. For *Benchmarking LLMs*, researchers **should** describe the evaluation harness and infrastructure if it goes beyond bare model API calls (e.g., custom sandboxing, orchestration layers, or post-processing pipelines); see [7] for a practitioner account of evaluation harness design for agent-based systems. Researchers using pre-defined prompts (e.g., *HumanEval*, *SWE-Bench*) **must** specify the benchmark version and any modifications made to the prompts or evaluation setup. If prompt tuning, retrieval-augmented generation (RAG), or other methods were used to adapt prompts, researchers **must** disclose and justify those changes, and **should** make the relevant code publicly available. For *Studying LLM Usage*, researchers **should** describe the tool architecture of the studied tool to the extent that it is accessible, as architectural details may influence observed usage patterns. For controlled experiments under *Studying LLM Usage*, exact prompts **must** be reported for all conditions. For *LLMs as Annotators*, researchers **must** document any pre-defined coding guides or instructions included in prompts, as these influence how the model labels artifacts. For *LLMs as Judges*, researchers **must** report the evaluation criteria, scales, and examples embedded in the prompt to ensure consistent interpretation. For *LLMs for Synthesis* tasks (e.g., summarization, aggregation), researchers **must** document the sequence of prompts used to

generate and refine outputs, including follow-ups and clarification queries. For *LLMs as Subjects* (e.g., simulating human participants), researchers **must** report any role-playing instructions, constraints, or personas used to guide LLM behavior. For *LLMs as Annotators*, *LLMs as Judges*, *LLMs for Synthesis*, and *LLMs as Subjects*, if the research setup involves a custom pipeline (e.g., retrieval-augmented generation for annotation or chained prompts for synthesis), the architecture **should** also be reported.

5.3.7 Advice for Reviewers:

As with other guidelines, missing architectural or prompt information is typically a minor revision request unless so much is missing that methodological rigor cannot be assessed. Reviewers may ask authors to move key details from supplements into the paper body to ensure the main text is self-contained.

Regarding proprietary or confidential components, three principles apply: (1) empirically evaluating an opaque artifact is a valid scientific contribution; (2) the less available and inspectable the artifact, the weaker the contribution; (3) *scientific instruments* including tools, metrics, scales, and experimental materials must be fully disclosed, even when the object under study cannot be.

A challenging aspect of prompt reporting is the description of prompt development, which is inherently iterative and creative. Reviewers should *not* expect justification of each word choice or post-hoc rationalized accounts of prompt generation. Instead, reviewers should focus on whether (1) a new research team could *use* (not reproduce) exactly the same prompts in exactly the same way, and (2) potential biases or validity issues are transparent.

5.3.8 See Also:

- Section 5.2 (Report Model Version, Configuration, and Customizations): Every architecture runs on at least one specific model; name its version.
- Section 5.4 (Report Session Traces): Architecture and prompts are static; session traces show how they behave at runtime.
- Section 5.7 (Use Human Validation for LLM Outputs): Human evaluation often complements automated metrics for tool outputs.
- Section 5.6 (Use an Open LLM as a Baseline): Open models let other researchers run the reported prompts and schemas on the exact same model weights.
- Section 5.8 (Report Limitations and Mitigations): When a tool hides internal prompts or schemas, authors must report the gap as a limitation.

5.4 Report Session Traces

Summary: To address model non-determinism and ensure reproducibility, especially when targeting SaaS-based commercial tools, researchers **should** include full interaction logs (prompts and responses) as *supplementary material* if privacy and confidentiality can be ensured. For agentic systems, interaction logs **should** be extended to cover exchanges between humans and the agent and between external tools and the agent, including human-in-the-loop feedback and approval decisions. Researchers **should** report the complete runtime trace as *supplementary material*, covering both external tool invocations (tool name, arguments, result, ordering) and which configured artifacts (skills, context files, sub-agents reported under *System and Prompt Design*) were activated, so that readers can attribute task outcomes to the model, the tools, or their interaction pattern. Developed plans **should** be reported as *supplementary material* if available. When full trace disclosure is not feasible, representative or anonymized examples **should** be provided, and unobservable aspects of commercial-tool runs **must** be acknowledged as reproducibility limitations.

5.4.1 *Rationale:*

A *session* is any bounded stretch of activity during which an LLM is invoked. Sessions cover one-shot prompts, batch runs, multi-turn conversations, and agentic runs that span many tool calls. A *session trace* records everything that crosses the LLM boundary or is produced around the model during that stretch: prompts received, responses returned, tool calls the model made together with their arguments and results, plans the model developed, and which of the statically configured artifacts (e.g., skills, context files, sub-agents, tools) were actually picked up at runtime.

Two kinds of session trace are important to capture. An *interaction log* captures what humans or other software tools sent to the LLM and what the LLM returned. A *runtime trace* captures what the LLM did at runtime: calls to external tools (e.g., APIs, file systems, databases, MCP servers, sub-agents) and activations of configured artifacts (e.g., context files, skills). Reporting both kinds is needed for an agentic run, because neither tells the full story on its own.

Even with the exact same prompts, decoding strategies, and parameters, LLMs can behave non-deterministically. Non-determinism can arise from probabilistic sampling and, even with greedy decoding (temperature = 0), from batching, input preprocessing, and floating-point arithmetic on GPUs [29]. For other researchers to verify the conclusions drawn from LLM interactions, the traces of what happened at runtime are often as important as the design of the system itself. This matters particularly for studies targeting commercial software-as-a-service (SaaS) solutions such as ChatGPT, and for agentic runs

where behavior depends on how the model chose among many possible tool calls.

The rationale is similar to reporting interview transcripts in qualitative research. Just as a human participant might give different answers to the same question asked two months apart, the responses from tools such as ChatGPT can also vary over time, and the trace of an agent’s decisions on any given run is often not reproducible at a later date. This guideline addresses runtime reporting and complements *System and Prompt Design*, which covers the static artifacts that determine what the model sees.

5.4.2 Recommendations:

Interaction Logs: Researchers **should** report the full interaction logs (prompts sent to the LLM and responses returned) as part of their *supplementary material*. For agentic systems, interaction logs generalize to cover exchanges between humans and the agent as well as between external tools and the agent, including human-in-the-loop feedback, approval or rejection decisions, and iterative refinements. These **should** also be reported as *supplementary material* so that readers can reconstruct the sequence of exchanges and assess human oversight decisions. For traces containing sensitive information, researchers **must** anonymize personal identifiers, replace proprietary code with placeholders, and clearly highlight modified sections.

Runtime Traces: When an LLM calls out to external tools (e.g., APIs, file systems, databases, MCP servers, sub-agents) or activates configured artifacts (e.g., context files, skills, sub-agents reported under *System and Prompt Design*), this runtime activity forms a *runtime trace* distinct from the interaction log. Researchers **should** report the complete runtime trace as *supplementary material*, including for each entry the tool or artifact name, arguments (if any), result, and ordering relative to surrounding interaction-log entries. This lets readers attribute task success to the model, the external tools, or their interaction pattern, and distinguish artifacts that were configured from those that actually influenced a given run. Researchers **should** record runtime traces in an open format compatible with the OpenTelemetry Protocol (OTLP), such as the OpenTelemetry GenAI semantic conventions [97] or OpenInference [10], and report the version used, rather than a study-specific schema.

Agentic Plans: For agentic systems that autonomously plan and execute tasks, the developed plans **should** be reported as *supplementary material* if available. In Claude Code, for example, a plan is a short Markdown document the user can open and edit during a session, listing the proposed steps and the files or commands the agent intends to touch. Other frameworks such as LangGraph keep plans inside the agent’s internal execution state.

All reported traces **must** be made publicly available as *supplementary material*, subject to privacy and confidentiality constraints. When full trace log-

ging is not feasible, researchers **should** provide representative examples or anonymized traces.

5.4.3 *Examples:*

An example of reporting full interaction logs is the study by Ronanki et al [113], for which the authors reported the full answers of ChatGPT and uploaded them to Zenodo. For agentic systems, Bouzenia and Pradel [25] unified the runtime trajectories of three SE agents (*RepairAgent*, *AutoCodeRover*, *Open-Hands*) into a common format and released the resulting 120 trajectories with 2,822 LLM interactions as a public dataset.

5.4.4 *Benefits:*

Unlike human participant conversations, which often cannot be reported because of confidentiality, LLM interaction logs can be shared. This enables reproduction studies, tracking of response changes over time or across model versions, and secondary research on LLM consistency for specific SE tasks.

For agentic systems, reporting runtime traces alongside interaction logs lets readers attribute task outcomes to the right component: the model, the tool, or the orchestration pattern. Usage traces complement the static configuration reported under *System and Prompt Design* by showing which of the configured artifacts actually influenced a given run.

5.4.5 *Challenges:*

Not all systems allow reporting of complete interaction logs with ease, and this hinders transparency and verifiability. For commercial tools, researchers **must** report all available information and acknowledge unknown aspects as limitations. Tool-call traces for commercial SaaS agents are often opaque: the user sees the final response but not the sequence of internal tool calls. When this is the case, authors should document what was and was not observable. For agents running locally or via open-source tools, these traces are usually more accessible and **should** be reported whenever available. Agent frameworks differ in whether and how they log agent-to-agent communication, so reporting practices vary across studies.

5.4.6 Study Types:

Runtime trace reporting requirements depend heavily on the study type and on the accessibility of the underlying system. The general expectation is that interaction logs are reported as *supplementary material* whenever feasible; runtime traces are additionally expected when agentic execution or evaluation harnesses go beyond direct API calls.

For *Studying LLM Usage*, especially observational studies targeting commercial tools, researchers **must** report the full interaction logs except when transcripts might identify anonymous participants or reveal personal or confidential information. If complete interaction logs cannot be shared (e.g., because they contain confidential information), the prompts and responses **must** at least be summarized and described in the *paper*. For *LLMs for Tools* that use agentic execution, researchers **should** report runtime traces for the runs used to evaluate the tool. For *Benchmarking LLMs* that use agent-based harnesses, researchers **should** report runtime traces for representative runs, letting readers understand how task success depends on the agent’s decision sequence rather than the model’s raw output alone. For *LLMs as Annotators*, *LLMs as Judges*, *LLMs for Synthesis*, and *LLMs as Subjects*, when the research setup involves multi-turn interaction or agentic orchestration, researchers **should** report the corresponding interaction logs and, where applicable, runtime traces.

5.4.7 Advice for Reviewers:

As with other guidelines, missing trace information is typically a minor revision request unless so much is missing that methodological rigor cannot be assessed. Reviewers should recognize that complete trace reporting is easier for local and open-source setups than for commercial SaaS agents. When reviewing commercial-tool studies, reviewers should focus on whether authors have reported everything the system exposed and have been explicit about what remains unobservable.

5.4.8 See Also:

- Section 5.3 (Report System and Prompt Design): Session traces show runtime behavior; system and prompt design covers the static artifacts that produce it.
- Section 5.2 (Report Model Version, Configuration, and Customizations): A trace cannot be reproduced or compared across studies without the model identity that produced it.
- Section 5.7 (Use Human Validation for LLM Outputs): Agentic traces produce outputs that often warrant human evaluation.
- Section 5.6 (Use an Open LLM as a Baseline): Open models let other researchers reproduce a reported trace on the exact same model weights.
- Section 5.8 (Report Limitations and Mitigations): When a tool hides parts of the trace, authors must report the gap as a limitation.

5.5 Use Suitable Baselines, Benchmarks, and Metrics

Summary: Researchers **must** justify all benchmark and metric choices in the *paper* and **should** summarize benchmark structure, task types, and limitations. They **should** operationally define the phenomenon being measured, justify the sampling strategy used to select problems for inclusion in the benchmark, isolate the target capability from confounders where possible, and report an error analysis of the observed failures. Where possible, traditional (non-LLM) baselines **should** be used for comparison. Researchers **must** explain in the *paper* why the selected metrics are suitable for the specific study; prior adoption in related work alone does not constitute sufficient justification. They **should** report established metrics to make study results comparable, but can report additional metrics that they consider appropriate. Due to the inherent non-determinism of LLMs, experiments **should** be repeated; the result distribution **should** then be reported using descriptive statistics. If comparing models or tools, researchers **must** use appropriate inferential statistics rather than relying solely on summary statistics. Researchers **should** justify the number of experiment repetitions, for example through a power analysis or by monitoring convergence of descriptive statistics.

5.5.1 *Rationale:*

Meaningful evaluation requires well-understood, valid measurement instruments. Without justified benchmarks and metrics, claims about LLM performance lack the rigor needed for scientific comparison and cumulative progress. A *benchmark* is a “standard tool for the competitive evaluation and comparison of competing systems or components according to specific characteristics, such as performance, dependability, or security” [71]. “A *metric* is a method, algorithm, or procedure for assigning one or more numbers to a phenomenon” [107]. A baseline is a reference point, enabling comparison of LLMs against traditional algorithms with lower computational costs.

5.5.2 Recommendations:

Benchmark and metric selection requires understanding the benchmark tasks, what exactly is being measured, and how it relates to the (often latent) variables researchers actually care about. When one or more metrics or benchmarks are used, researchers **must** briefly justify in the *paper* why each selected benchmark and metric is suitable for the given task or study, and **must** discuss the reliability and validity (especially construct validity) of those choices. Beyond these requirements, researchers **should**:

- provide an operational definition of the phenomenon the benchmark is intended to measure (e.g., functional correctness, code maintainability, vulnerability detection), including its scope and any sub-components [20];
- summarize the structure and tasks of the selected benchmark(s), including the programming language(s) and descriptive statistics such as the number of contained tasks and test cases;
- describe and justify the sampling strategy used to select problems for inclusion in the benchmark (e.g., function-completion problems in *HumanEval*, GitHub issues in *SWE-bench*, vulnerable functions in *DiverseVul*); if non-probability sampling (e.g., convenience) is used, discuss its implications for the generalizability of conclusions [14, 20];
- discuss the limitations of the selected benchmark(s) (e.g., widely used benchmarks such as *HumanEval* [31] and *MBPP* [12] only test short Python functions, which is not representative of the full breadth of SE work [28]);
- include an example of a task and the corresponding test case(s) to illustrate the structure of the benchmark.

If multiple benchmarks exist for the same task, researchers **should** compare both performance and design choices (e.g., which tasks are included, how outputs are scored, what aspect of the phenomenon is covered) across benchmarks [20]. When selecting only a subset of all available benchmarks, researchers **should** use the most specific benchmarks given the context. When adapting an existing benchmark, researchers **should** document what was changed and why, and **should** report performance on both the original and the adapted version where feasible [20].

Benchmark scores often confound the target capability with unrelated capabilities the task happens to require. For code translation benchmarks, prompt formatting alone can shift performance by up to 40% [27, 53], conflating prompt format sensitivity with translation capability. More broadly, output format compliance (e.g., specific JSON schemas or unit-test conventions) and general instruction-following capability are routinely bundled into the score [20]. Researchers **should** identify which capabilities a benchmark conflates, isolate the target capability where possible (e.g., by reporting per-subtask breakdowns or by adopting benchmarks designed to test the target capability in isolation), and acknowledge remaining confounders as threats to construct validity.

After running a benchmark, researchers **should** perform an error analysis by categorizing the failures observed and reporting the relative frequency of each category. If failures concentrate on inputs that demand capabilities other than the target capability (e.g., reading across many files rather than fixing the bug the benchmark targets), this is a construct-validity threat and **should** be reported alongside the primary scores [20].

In addition to disclosing data sources and collection dates (see Challenges below), researchers creating or releasing a new benchmark **should** adopt concrete contamination-prevention mechanisms: maintaining a held-out subset of items for ongoing, uncontaminated evaluation; embedding canary strings, that is, unique markers that downstream tools can later search for in model outputs to detect inclusion in training data; and investigating whether the benchmark’s source materials may already appear in common LLM training corpora [20, 27]. Researchers using existing benchmarks **should** additionally discuss contamination as a study limitation (*Limitations and Mitigations*).

Furthermore, researchers **should** check whether a less resource-intensive approach (e.g., for static analysis tasks or program repair) can serve as a baseline. If so, the LLM or LLM-based tool **should** be compared with such baselines using suitable metrics. Even if LLM-based tools outperform baselines, researchers **should** discuss whether the resources consumed justify the (often marginal) improvements [84].

To compare traditional and LLM-based approaches or different LLM-based tools, researchers **should** report established metrics whenever possible, as this allows secondary research. They can report additional metrics that they consider appropriate. We briefly discuss common metrics in the Examples subsection below. As mentioned, researchers **must** motivate why they chose a certain metric or variant thereof for their particular study. Prior adoption alone does not constitute sufficient justification; researchers **must not** justify metric choices solely by citing their use in prior work.

Due to LLM non-determinism, researchers **should** repeat experiments and report descriptive statistics of model or tool performance (e.g., arithmetic mean, median, confidence intervals, standard deviations) [2, 22]. If comparing models or tools, researchers **must** use appropriate inferential statistics (e.g., hypothesis tests such as the Mann-Whitney U test, McNemar’s test for binary outcomes [74], or bootstrap-based comparisons, along with effect sizes) rather than relying solely on differences in means or other summary statistics. When scoring uses ratings that vary across raters or runs (e.g., human raters, LLM-as-judge), researchers **should** report the distribution of ratings per item rather than only aggregated point estimates or exact-match agreement, since aggregation can mask systematic disagreement [20].

The number of required repetitions depends on factors such as the study type, the variability of the task, and the desired precision of estimates. As with sample sizes for human validation (*Human Validation*), researchers **should** justify their chosen number of repetitions, for example, through a power analysis [22, 35] or by monitoring the convergence of descriptive statistics across

incremental runs [23]. A pilot study can help estimate the expected variability and inform this decision.

From a measurement perspective, researchers **should** reflect on the theories, values, and measurement models on which the benchmarks and metrics they have selected for their study are based. For example, labeling phenomena as “bugs” in a large open dataset reflects a certain theory of what constitutes a bug, as well as the values and perspectives of the people who labeled the dataset. Reflecting on the context in which these labels were assigned and discussing whether and how the labels generalize to a new study context is crucial.

Researchers building or releasing new SE benchmarks **should** consult operational checklists. Cao et al [27] provide *HOW2BENCH*, a 55-item checklist covering benchmark design, construction, evaluation, analysis, and release. Bean et al [20] provide a complementary domain-agnostic checklist organized around the construct-validity recommendations summarized above.

5.5.3 Examples:

Common Metrics: Two common metrics used for generation tasks are *BLEU-N* and *pass@k*. *BLEU-N* [102] was originally developed for evaluating machine translation quality by measuring n-gram precision between a candidate and reference text, ranging from 0 (dissimilar) to 1 (similar). It has been widely adopted in SE for code generation tasks, though its validity in this context is debatable: n-gram overlap does not capture functional correctness, and syntactically different code can be semantically equivalent (see Challenges below). Code-specific variations such as *CodeBLEU* [110] and *CrystalBLEU* [40] attempt to address these limitations by introducing additional heuristics such as AST matching.

The metric *pass@k* reports the likelihood that a model correctly completes a code snippet at least once within k attempts. Originally introduced by Kulal et al [75] as the success rate within a budget of B program attempts (the role now played by k in *pass@k*) and popularized by Chen et al [31], it ranges from 0 (no correct solution in k tries) to 1 (at least one correct solution), with correctness typically defined by passing test cases.

The *pass@k* metric is defined as:

$$\text{pass@}k = 1 - \frac{\binom{n-c}{k}}{\binom{n}{k}}, \text{ where:}$$

- n is the total number of generated samples per prompt,
- c is the number of correct samples among n , and
- k is the number of samples drawn.

The choice of k depends on the downstream task: *pass@1* is critical for single-suggestion scenarios such as code completion, while higher k values (e.g., 2, 5, 10) assess multi-attempt capability and are commonly reported in technical reports of recent code LLMs (e.g., Code Llama [115], DeepSeek-Coder [51],

Qwen2.5-Coder [61], StarCoder [77]). However, $pass@k$ is not a universal metric suitable for all generation tasks. It requires a binary notion of correctness, making the metric appropriate for code synthesis evaluated via unit tests, but unsuitable for open-ended generation tasks such as comment generation, where multiple valid outputs exist.

A complementary perspective from industry practice is $pass^k$ (also written $pass^{\wedge k}$), which measures the probability that *all* k trials succeed rather than *at least one* [7]. While $pass@k$ increases with k , $pass^k$ decreases, making it useful for assessing reliability in deployment scenarios where consistent success matters (e.g., customer-facing agents).

If a study evaluates an LLM-based tool for supporting humans, a relevant metric is the acceptance rate, meaning the ratio of all accepted artifacts (e.g., test cases, code snippets) in relation to all artifacts that were generated and presented to the user. Another way of evaluating LLM-based tools is calculating inter-model agreement, which reveals how much a tool’s performance depends on specific models and versions. Metrics used to measure inter-model agreements include general agreement (percentage), *Cohen’s κ* , and *Krippendorff’s α* (see *Human Validation* for recommended thresholds and best practices for measuring agreement).

Common problem types for LLM-based studies are classification, recommendation, and generation, each requiring different metrics [57]. Hu et al [58] categorized 191 LLM benchmarks by SE task, providing a valuable reference. For an overview of code foundation models, agents, and their evaluation, see Yang et al [156]. From a practitioner perspective, Anthropic [7] categorize evaluation approaches for LLM-based agents into code-based graders (e.g., test suite execution, static analysis), model-based graders (e.g., rubric-scored LLM judgments), and human graders. This taxonomy may help researchers systematically design evaluation strategies for agent-based tools. Common metrics include *BLEU*, $pass@k$, *Accuracy@k*, and *Exact Match* for generation; *Mean Reciprocal Rank* for recommendation; and *Precision*, *Recall*, *F1-score*, and *Accuracy* for classification.

Benchmark Examples: Benchmarks used for code generation include *HumanEval* (available on GitHub) [31], *MBPP* (available on Hugging Face) [12], *ClassEval* (available on GitHub) [38], *LiveCodeBench* (available on GitHub) [65], and *SWE-bench* (available on GitHub) [66]. An example of a code translation benchmark is *TransCoder* [114] (available on GitHub).

Most benchmarks focus on generation tasks, but benchmarks for classification and recommendation also exist. For classification, *DiverseVul* (available on GitHub) [32] provides vulnerable and non-vulnerable functions evaluated using standard classification metrics. For recommendation, *CodeSearchNet* (available on GitHub) [62] contains code-documentation pairs evaluated using *Mean Reciprocal Rank*.

5.5.4 *Benefits:*

Benchmarks and metrics improve reproducibility and comparability across studies, leading to faster improvements in the cumulative body of knowledge. It is simply easier to assess a new system against one or a few benchmarks than hundreds of competing systems, and when most systems are assessed against the same benchmarks and metrics, we can see which approaches work best (at least, on those benchmarks). This comparison enables progress tracking; for example, when researchers iteratively improve a new LLM-based tool and test it against benchmarks after substantial changes. For practitioners, leaderboards (published benchmark results of models) support selecting models for downstream tasks. Meanwhile, baselines help us understand just how much LLMs improve performance over non-LLM-enabled alternatives.

5.5.5 *Challenges:*

Computer scientists commonly use simple metrics (e.g., lines of code, CPU time) as proxies for complex, multidimensional latent variables (e.g., system size, environmental sustainability), without empirically validating that the metrics capture the intended construct [105]. Unlike fields such as psychology where measurement theory has a longer tradition [24], many AI metrics and benchmarks lack both theoretical underpinnings and empirical validation of their construct, measurement, and ecological validity.

These validity issues are widespread, not isolated. In a survey of 572 code benchmarks released between 2014 and 2025, Cao et al [27] found that 84.2% did not consider test-suite coverage when constructing test cases, 64.0% reported single-pass evaluations without controlling for randomness, and 82.5% did not address data contamination. Bean et al [20] reported comparable patterns in a parallel review of 445 LLM benchmarks from ML and NLP venues. As a result, model performance on these benchmarks can reflect benchmark artifacts rather than the capability the benchmark claims to test.

For example, *pass@k* is intended to measure a model’s *functional correctness*, that is, its capability to generate code that produces the expected output for a given specification. However, functional correctness is only one dimension of code quality. *pass@k* does not capture maintainability, readability, security, or efficiency of the generated code, all of which are critical for downstream use. Furthermore, “correctness” is defined entirely by test suites, whose coverage is itself unvalidated: a solution that passes all provided tests may still be incorrect for untested inputs. Whether a test suite adequately operationalizes correctness for a given task is a subjective judgment that is rarely examined.

Similarly, *HumanEval* is intended to measure a model’s capability to *synthesize short Python functions from docstring specifications*. This operationalizes a narrow slice of “code generation capability”: it covers neither multi-file tasks, nor debugging, nor the use of existing codebases. These tasks dominate real-world software engineering [28, 156]. Notably, neither *pass@k* nor

HumanEval has undergone rigorous empirical validation of its construct validity; their widespread adoption rests on face validity and convenience rather than on evidence that they reliably measure what researchers intend them to measure [27]. *HumanEval* also contains implementation, documentation, and test-case bugs in its original release, which directly affect score interpretation [79]; comparable issues have been documented in other widely-used code benchmarks [27].

Benchmarks and metrics also have generalizability problems. Prominent LLM benchmarks such as *HumanEval* and *MBPP* use Python, so researchers can optimize for Python’s idiosyncrasies, producing gains that do not generalize to other languages. Similarly, the metric *BLEU-N* is a syntactic metric, so code can score highly without being executable. The metric *Exact Match*, meanwhile, does not account for functional equivalence of syntactically different code. Both *BLEU-N* and *Exact Match* are influenced by code formatting, which confounds their intended use.

Execution-based metrics such as *pass@k* directly evaluate correctness by running test cases, but they require an execution environment. When metric values look surprising, examine the specific test cases driving them: common causes include outliers, bugs in test suites or scoring code, and items that exercise capabilities other than the one being measured.

Finally, benchmark data contamination, where the benchmark itself is part of the training data, may lead to artificially high performance if the model remembers the solution from the training data rather than deriving it from the input [152]. Therefore, for proprietary LLMs that do not release their training data, researchers **should** consider using human validation, curating new data, or refactoring existing data [152].

To mitigate contamination, researchers can create new benchmark datasets by collecting data after a specified cutoff date; researchers **must** disclose data sources and collection dates for each release. However, this temporal approach requires continuous updates as new models may include the benchmark in future training data. Alternatively, keeping the benchmark private prevents inclusion in training sets, but requires trust in the benchmark creator and a system to execute it without leaking data. Researchers can also evaluate how contaminated their benchmark is [34].

5.5.6 Study Types:

This guideline **must** be followed for all study types that automatically evaluate the performance of LLMs or LLM-based tools. The design of a benchmark and the selection of appropriate metrics are highly dependent on the specific study type and research goal. Recommending specific metrics for specific study types is beyond the scope of these guidelines, but Hu et al provide a good overview of existing metrics for evaluating LLMs [58].

For *Benchmarking LLMs*, this guideline is of primary importance: researchers **must** use established benchmarks or rigorously justify the creation of new ones, **must** report standard metrics to enable cross-study comparison,

and **should** report an error analysis alongside the primary scores so that readers can assess whether reported gains reflect the target capability or other capabilities the task happens to require. For *LLMs as Annotators*, the research goal might be to assess which model comes close to a ground truth dataset created by human annotators. Especially for open annotation tasks, selecting suitable metrics to compare LLM-generated and human-generated labels is important. In general, annotation tasks can vary widely. Are multiple labels allowed for the same sequence? Are the available labels predefined, or should the LLM generate a set of labels independently? Due to this task dependence, researchers **must** justify their metric choice, explaining what aspects of the task it captures together with known limitations. For *LLMs for Tools*, researchers **should** benchmark the tool against suitable baselines using appropriate metrics that capture the tool’s intended contribution. For *LLMs as Judges*, researchers **should** report inter-rater agreement metrics and validity measures to demonstrate the reliability and quality of LLM judgments. For *LLMs for Synthesis*, researchers **should** specify metrics for comparing synthesized outputs (e.g., coverage, faithfulness) and justify their appropriateness for the synthesis task. For *Studying LLM Usage*, researchers **should** justify the measurement instruments and metrics used for studying LLM usage patterns, including any survey scales or behavioral measures. For *LLMs as Subjects*, researchers **should** compare simulated and real human responses using appropriate metrics to assess simulation fidelity. If researchers assess a well-established task such as code generation, they **should** report standard metrics such as *pass@k* and compare the performance between models. If non-standard metrics are used, researchers **must** state their reasoning.

5.5.7 Advice for Reviewers:

Reviewers should expect manuscripts to: (1) clearly identify the constructs or variables the study aims to measure (e.g., LLM performance, quality of generated code), including independent, dependent, and control variables; (2) present their measurement model, i.e., which metrics, benchmarks, or baselines are used and how they relate to the target constructs; (3) justify the selection of metrics, benchmarks, and baselines; (4) discuss *in detail* the assumptions, reliability, and validity (*especially construct validity*) of each benchmark and metric; (5) articulate any limitations regarding construct and measurement validity.

As with other guidelines, missing information about baselines or metrics is typically a revision request. However, vague descriptions that conflate broad concerns (e.g., effectiveness, quality) with specific counting methods should be questioned. Ubiquity of a benchmark or metric does not imply validity or appropriateness for a given context. Manuscripts should convey a solid understanding of construct and measurement validity by explaining and justifying their measurement models.

5.5.8 *See Also:*

- Section 5.6 (Use an Open LLM as a Baseline): Using an open LLM as a baseline supports reproducible benchmark comparisons across studies.
- Section 5.7 (Use Human Validation for LLM Outputs): Human evaluation captures aspects of LLM outputs that automated metrics cannot measure.
- Section 5.8 (Report Limitations and Mitigations): Every benchmark and metric choice carries construct-validity threats that authors must discuss.

5.6 Use an Open LLM as a Baseline

Summary: Researchers **should** include an open LLM as a baseline when using commercial models and report inter-model agreement. By open LLM, we mean a model that everyone with the required hardware can deploy and operate. Such models are usually “open weight.” A full replication package with step-by-step instructions **should** be provided as part of the *supplementary material*.

5.6.1 *Rationale:*

Reproducibility depends on access to the model under study. When research relies exclusively on proprietary models, other researchers cannot independently verify or build upon the findings. Including an open LLM as a baseline ensures that at least part of the study can be fully replicated.

5.6.2 *Recommendations:*

Empirical studies using LLMs in SE, especially those that target commercial tools or models, **should** incorporate an open LLM as a baseline and report established metrics for inter-model agreement. We acknowledge that including an open LLM baseline might not always be possible, for example, if the study involves human participants, and letting them work on the tasks using two different models might not be feasible. Using an open model as a baseline is also not necessary if the use of the LLM is tangential to the study goal.

Open models allow other researchers to verify research results and build upon them, even without access to commercial models. A comparison of commercial and open models also allows researchers to contextualize model performance. Researchers **should** provide a complete replication package as part of their *supplementary material*, including clear step-by-step instructions on how to verify and reproduce the results reported in the paper.

Open LLMs are available on platforms such as *Hugging Face* and can be hosted locally using frameworks such as *Ollama* or *LM Studio*, or on cloud services such as *Together AI*, AWS, Azure, and Google Cloud. Researchers

can also explore open-source tools such as *Continue*, *Cline*, and *opencode* [96], which publish their full agent code, system prompts, and tool catalog. Vendor-hosted services such as *GitHub Copilot* and *Claude Code* make parts of their tooling public (e.g., editor extensions, hook examples), but their deployed agent loops, system prompts, and tool catalogs typically remain proprietary.

The term “open” can have different meanings in the context of LLMs. Widder et al discuss three types of openness (i.e., transparency, reusability, and extensibility) and what openness can and cannot provide [146]. The *Open Source Initiative* (OSI) [93] defines open-source AI as having access to everything needed to understand, modify, share, retrain, and recreate the model.

5.6.3 Examples:

An increasing number of studies have adopted open LLMs as baseline models. For example, Wang et al evaluated seven advanced LLMs, six of which were open-source, testing 145 API mappings drawn from eight popular Python libraries across 28,125 completion prompts aimed at detecting deprecated API usage in code completion [142]. Moumoula et al compared four LLMs on a cross-language code clone detection task [92]. Three evaluated models were open-source. Gonçalves et al fine-tuned the open LLM *LLaMA* 3.2 on a refined version of the *DiverseVul* dataset to benchmark vulnerability detection performance [48]. *CodeBERT*, a bimodal transformer pre-trained by Microsoft Research, is published with model weights, source code, and data-processing scripts on GitHub [85]. It has been used as an open baseline across diverse SE tasks, including exploit code generation [155], vulnerability detection [150], code clone detection [132], and programming assistance for exception handling [26].

5.6.4 Benefits:

A true open LLM baseline improves reproducibility by exposing model architectures, parameter settings, and ideally training data, enabling independent verification of results. Moreover, by adopting an open-source baseline, researchers can directly compare novel methods against established performance metrics without the variability introduced by proprietary systems. These models allow detailed inspection of data processing pipelines and decision-making routines, helping identify biases and model limitations. Furthermore, unlike closed-source alternatives, which can be withdrawn or altered without notice, open LLMs ensure long-term accessibility and stability, preserving critical resources for future studies. Finally, the licensing requirements associated with open-source implementations lower financial barriers, making advanced language models attainable for research groups operating under constrained budgets.

5.6.5 Challenges:

Open-source LLMs face several notable challenges. First, they often lag behind the most advanced proprietary “frontier” models in common benchmarks, making it difficult for researchers to demonstrate clear improvements when evaluating new methods using open LLMs alone. Additionally, deploying and experimenting with these models typically requires substantial hardware resources, in particular high-performance GPUs, which may be beyond reach for many academic groups. The notion of “openness” remains in flux: many models release only trained weights without training data or methodological details (“open weight” openness) [46], which is why we reference the OSI definition in our recommendations [93]. Finally, unlike APIs provided by proprietary vendors (e.g., the OpenAI API), installing, configuring, and fine-tuning open-source models can be technically demanding. Documentation is often sparse or fragmented, placing a high barrier to entry for researchers without specialized engineering support.

5.6.6 Study Types:

This guideline applies primarily to study types in which the researcher controls which LLM is used. In formal benchmarking studies (see *Benchmarking LLMs*), an open LLM **must** be one of the models under evaluation; without one, the reported scores cannot be re-run by others, defeating the cumulative comparability that distinguishes a benchmark from a one-off measurement. For controlled experiments, an open LLM **should** be one of the models under test; if the experimental design requires capabilities that only a specific commercial model exhibits, researchers **should** acknowledge this as a limitation. When evaluating *LLMs for Tools*, researchers **should** use an open LLM as a baseline whenever it is technically feasible; if integration proves too complex, they **should** report the initial benchmarking results of open models. For *LLMs as Annotators* and *LLMs as Judges*, researchers **should** compare annotation or judgment quality from open vs. commercial models to assess the extent to which results depend on a specific proprietary model. For *LLMs for Synthesis*, researchers **should** compare synthesis results from open and commercial models to evaluate robustness of the findings. For *Studying LLM Usage*, using an open LLM as a baseline is often not feasible when the study observes participants using specific commercial tools; in such cases, investigators **should** explicitly acknowledge its absence and discuss how this limitation might affect their conclusions. For *LLMs as Subjects*, using an open LLM as a baseline may similarly be impractical if the study design requires the capabilities of a specific model; researchers **should** acknowledge this limitation when applicable.

5.6.7 *Advice for Reviewers:*

Reviewers should distinguish between LLM use that is central to the research (e.g., building LLM-driven SE tools, using LLMs for synthesis) and use that is tangential (e.g., generating recruiting materials). An open LLM baseline is expected only when LLM use is central; otherwise, its absence need not be justified. When an open baseline is expected, reviewers should look for either the use of an open LLM or a convincing argument for why it is impractical. If authors claim openness, some justification for that characterization is appropriate. Where practical, reviewers should examine whether the replication package contains sufficiently detailed instructions. Reviewers should *not* penalize studies for performance differences between open and proprietary models, as this is beyond the authors' control.

5.6.8 *See Also:*

- Section 5.2 (Report Model Version, Configuration, and Customizations): Authors must report version, configuration, and parameters for open models, not just commercial ones.
- Section 5.3 (Report System and Prompt Design): Self-hosting an open model adds hardware, hosting, and infrastructure to document.
- Section 5.5 (Use Suitable Baselines, Benchmarks, and Metrics): Inter-model agreement metrics make open-vs-commercial comparisons interpretable.
- Section 5.8 (Report Limitations and Mitigations): When using an open LLM as a baseline is impractical, authors must report this as a limitation.

5.7 Use Human Validation for LLM Outputs

Summary: If assessing the quality of generated artifacts is important and no reference datasets or suitable comparison metrics exist, researchers **should** use human validation for LLM outputs. If they do, they **must** define the measured construct (e.g., usability, maintainability) and describe the measurement instrument in the *paper*. Researchers **should** consider human validation early in the study design (not as an afterthought) and build on established reference models for human-LLM comparison. When developing or adapting measurement instruments, researchers **must** share them. When aggregating LLM judgments, methods and rationale **should** be reported and inter-rater agreement **should** be assessed. Confounding factors **should** be controlled for, and power analysis **should** be conducted to ensure statistical robustness. For value-laden or culturally contingent constructs, researchers **should** describe rater demographics beyond expertise and discuss potential demographic biases.

5.7.1 *Rationale:*

Even with well-justified benchmarks and metrics (*Benchmarks and Metrics*), automated measurement captures only the constructs the benchmark was designed to measure. Subjective constructs and constructs that no current benchmark covers require human assessment.

5.7.2 *Recommendations:*

When LLMs automate research or software development tasks previously performed by humans, the LLM's performance needs to be assessed. Where no benchmark adequately operationalizes the target construct, researchers **should** validate LLM outputs against human judgment.

Study Design Considerations: Studies that include human participants need additional considerations, including a recruitment strategy, annotation guidelines, training sessions, or ethical approvals. Therefore, researchers **should** consider human validation early in the study design, not as an afterthought. Authors **must** clearly define in the *paper* the constructs that the human and LLM annotators evaluate [105]. When designing custom instruments to assess LLM output (e.g., questionnaires, scales), researchers **must** share their instruments.

Subjective Judgment and Agreement: When the judgment is subjective (i.e., depends on the judge's values or theories), the LLM output **should** be evaluated against judgments aggregated from multiple human judges, and researchers **should** clearly describe their aggregation method and reasoning. One systematic approach is to randomly order the objects requiring judgment and put them in groups small enough to judge in one to three hours. Two or three human experts review the first group together, simultaneously judging, discussing, and creating a set of decision rules documenting their reasoning. Then, the experts iterate among rounds of:

1. independently rating one group of objects,
2. calculating inter-rater agreement (IRA) or reliability (IRR),
3. meeting to discuss disagreements and reach consensus,
4. updating the decision rules so the disagreement will not recur.

The goal is to reach a target IRA or IRR (e.g., Krippendorff's $\alpha > 0.8$), indicating sufficient decision rules. Once this target is reached and sustained for two to three groups, it is permissible to continue with a single rater. Researchers **should** report measures of IRA or IRR, preferably broken down by round.

Confounding factors **should** be discussed and, where feasible, controlled for (e.g., by categorizing participants according to their level of experience or expertise). For value-laden or culturally contingent constructs (e.g., judging code-style appropriateness or comment helpfulness), researchers **should**

describe rater demographics beyond expertise (e.g., geographic, linguistic, or professional background) and discuss potential demographic biases in rater recruitment and instructions [20]. Where applicable, researchers may perform a power analysis [35, 39] to estimate the required sample size, ensuring sufficient statistical power in their experimental design. However, we also note that to our knowledge, there is limited established guidance specific to determining sample sizes for LLM-human comparisons. In other fields, 100 comparisons are often made without further explanation [138].

Researchers **should** use established reference models to compare humans with LLMs. For example, Schneider et al [121] outline design considerations for studies comparing LLMs with humans. If studies aim to automate annotating software artifacts using an LLM, researchers **should** follow systematic approaches to decide whether and how human annotators can be replaced (e.g., showing that a jury of three LLMs exhibit model-to-model agreement comparable to trained human annotators, such as Krippendorff’s $\alpha > 0.8$). Model-to-model agreement requires the same threshold as human-to-human agreement, since low thresholds may simply reflect shared model biases rather than annotation quality.

Agentic Tools: Besides generating and modifying content, agentic software development tools such as *Claude Code* can autonomously call command-line tools or pull in additional information from MCP servers. For such tools, it makes sense to separate textual output (e.g., file changes) from output related to tool execution. When evaluating agentic tools, researchers **should** assess the feedback that users provided for proposed changes, report statistics on how frequently they accepted the content, and how they modified it. Agentic human-in-the-loop interaction is a built-in human validation, with larger degrees of freedom than traditional experiments that use LLMs directly.

5.7.3 Examples:

Ahmed et al [4] proposed a systematic method for deciding whether LLMs can replace human annotators on a given task, using model-to-model agreement as an initial screening criterion (with a threshold of $\alpha > 0.5$) and model confidence for sample-level decisions. However, their threshold is well below the levels generally considered acceptable for inter-rater agreement. Krippendorff recommends discarding data with $\alpha < 0.667$, considers $0.667 \leq \alpha < 0.8$ sufficient only for tentative conclusions, and requires $\alpha \geq 0.8$ for reliable data [73]. Researchers adopting similar screening approaches should apply thresholds consistent with these established standards. Notably, while three LLMs exhibited higher inter-model agreement than human annotators on some tasks, human-model agreement remained low on others. This illustrates that high model-model reliability does not guarantee alignment with human judgment, reinforcing the need for human validation before assuming that LLM annotations are valid. Moreover, a high model-model agreement could merely indicate

that the models share systematic biases and hence reliably agree on the wrong answer.

Hymel and Johnson [63] evaluated ChatGPT’s capability to generate requirements documents by comparing an LLM-generated and a human-generated document based on the same business use case. Domain experts reviewed both documents and attempted to distinguish their origin. For agentic tools, Takerngsaksiri et al [137] report acceptance and modification rates from real users for *HULA*, an agentic plan-and-code system deployed in Atlassian JIRA.

5.7.4 *Benefits:*

Validation against human judgments builds confidence in the LLM’s accuracy and validity, increasing the trustworthiness of findings, especially when IRA or IRR metrics are reported [70]. Furthermore, incorporating feedback from human judges may help to improve the LLM or LLM-based tool based on the reported experiences.

5.7.5 *Challenges:*

Assessment of a construct such as LLM performance is challenging because ensuring that a construct is defined well and operationalized using an appropriate measurement model requires a deep understanding of (1) the construct, (2) construct validity in general, and (3) instrumentation [107, 130]. Comparing an LLM to human judges is typically slower and more expensive than machine-generated measures. More fundamentally, neither human judgment nor machine-generated measures provides an objective ground truth against which LLM accuracy can be firmly determined. Human validation is worth its added cost only when automated measures alone fail to capture the constructs the study evaluates.

Human judgments exhibit variability due to differences in experience, expertise, interpretations, and personal biases [83]. When diverse humans rate items reliably given clear decision rules, we assume that reliability implies validity, but it does not. Measuring reliability is *much* easier than measuring validity, and often the best researchers can do is argue conceptually for why their judges, decision rules, and constructs *should* produce valid ratings.

5.7.6 Study Types:

This guideline applies to all study types, although the need for human validation varies. For *LLMs as Annotators*, researchers **should** validate LLM-generated annotations against human annotators to assess labeling quality and identify systematic biases. When using *LLMs as Judges*, researchers **should** co-create initial rating criteria with humans and validate a sample of LLM judgments against human expert assessments. For *LLMs for Synthesis*, researchers **should** employ human oversight to verify that qualitative interpretations and synthesized outputs faithfully represent the underlying data. For *LLMs as Subjects*, researchers **should** validate simulated responses against real human data to assess the fidelity of the simulation. For *Studying LLM Usage*, researchers **should** carefully reflect on the validity of their evaluation criteria and validate subjective assessments with human experts. For *LLMs for Tools* whose output is supposed to match human expectations, researchers **should** validate the LLM output against human judgment. For *Benchmarking LLMs*, there is less need for human validation when using extensively validated and widely-used benchmarks, but researchers **should** employ human validation when creating or adapting new benchmarks.

5.7.7 Advice for Reviewers:

Human validation may be the most challenging of our guidelines to assess because it often requires evaluating conceptual arguments. If LLM output is validated only by comparison with other LLMs, reviewers should look for *quantitative empirical evidence* that such comparison is reliable and valid. High inter-model agreement alone is insufficient, as reliability does not imply validity. Similarly, reviewers should expect evidence that any employed benchmarks are reliable and valid. Absent such evidence, human validation is warranted. A single human judge is appropriate only when judgments depend on widely accepted theories and involve limited value conflict (e.g., tagging method names containing abbreviations). For multiple judges, reviewers should expect IRA/IRR improvement techniques as described in the recommendations above (experienced raters, organized rounds, consensus meetings, updated decision rules). Low IRA or IRR (e.g., Krippendorff's $\alpha < 0.8$) without these techniques is a concern. Conversely, if authors have followed best practices and still obtained mediocre results (e.g., $0.66 < \alpha < 0.8$), this should be noted as a limitation. Beyond reliability, reviewers should expect authors to explain conceptually why their human judgments should be valid, considering construct definitions, decision rules, and judge expertise. As with other guidelines, missing information is typically a revision request. Absent judge instructions, instruments, decision rules, or construct definitions may prevent assessment of rigor and validity, while missing recruitment details are less critical. Clarification requests about construct definitions are routine and should not alone warrant rejection.

5.7.8 *See Also:*

- Section 5.5 (Use Suitable Baselines, Benchmarks, and Metrics): Human evaluation complements automated metrics where benchmarks cannot sufficiently capture the target construct.
- Section 5.2 (Report Model Version, Configuration, and Customizations): Human-validation findings characterize only the specific model and configuration that produced the outputs.
- Section 5.3 (Report System and Prompt Design): Prompt and architecture choices produce the outputs that humans then validate.
- Section 5.4 (Report Session Traces): Stored interaction logs and runtime traces are the artifacts human reviewers examine.
- Section 5.8 (Report Limitations and Mitigations): Rater demographics, threshold choices, and the validity of human judgments are limitations to discuss.

5.8 Report Limitations and Mitigations

Summary: Researchers **must** transparently report study limitations, including the impact of non-determinism and generalizability constraints. The *paper* **must** specify whether generalization across LLMs or across time was assessed, and discuss model and version differences. Authors **must** describe measurement constructs and methods, disclose any data leakage risks, and avoid leaking evaluation data into LLM improvement pipelines. For studies involving sensitive data, they **must** discuss data governance mechanisms. They **should** justify LLM usage in light of its resource demands. Mitigation strategies such as replication packages, human validation, longitudinal re-runs, triangulation, and sensitivity analysis **should** be employed and reported where applicable. Where full data sharing is not possible, a subset of the validation data **should** be included to enable partial replication.

5.8.1 *Rationale:*

When using LLMs for empirical studies in SE, researchers face unique challenges and potential limitations that can influence the validity, reliability, and reproducibility of their findings [117]. Researchers must openly discuss these limitations and explain how their impact was mitigated. These limitations are relative to current LLM capabilities and tool architectures; speculating about future improvements is beyond the scope of a paper's limitation section. Nevertheless, risk management and threat mitigation **should** be planned during study design, not as an afterthought.

5.8.2 Recommendations:

Researchers **must** clearly present the limitations of their work without defensiveness or obfuscation. These limitations may concern diverse topics including generalizability, internal validity (e.g. data leakage), reliability (i.e. non-determinism), and reproducibility (e.g. resource requirements). When deterministic reproducibility is structurally unattainable (e.g., SaaS-based models with opaque versioning), researchers **should** adopt trustworthiness criteria from qualitative research to substantiate dependability and confirmability of findings. LLM-based studies may involve many kinds of generalization including the following:

1. From tested LLMs to others; i.e., whether the performance characteristics of the LLM(s) studied transfer to LLM(s) not included in the study.
2. From tested configurations to others; i.e., how sensitive results are to the specific configuration(s) of the LLM under test.
3. From research to practice; i.e., whether developers using the same tools obtain results comparable to those reported under researcher-supervised conditions.
4. From a sampled population to a larger one; i.e., whether findings from a specific group of human validators or participants apply to a broader population.
5. From one time period to another; i.e., whether the same LLM, or other versions of it, produce similar results at a different point in time.

The following concerns present an overview that has to be tailored to the individual study context. This section does not repeat requirements from other recommendations.

Internal Validity: The primary threats to internal validity are:

- *Data leakage and contamination.* Inter-dataset duplication can produce training-evaluation overlap, yielding overly optimistic results.
- *Evaluation data entering model-improvement pipelines.* Evaluation samples can unintentionally feed retraining or fine-tuning, especially in longitudinal studies involving LLMs.
- *Incomplete architecture, prompt, or pipeline reporting.* Undisclosed components introduce hidden confounders.

Inter-dataset duplication is prevalent in SE, particularly for code-related benchmarks. As transparency on training data is limited for LLMs, researchers **must** discuss potential data leakage effects and their impact on results.

Construct Validity: The primary threats to construct validity are:

- *Metric-construct mismatch.* Traditional metrics such as BLEU or ROUGE may miss SE-specific aspects such as functional correctness or behavioral equivalence.

- *Over-reliance on benchmark-specific metrics.* Optimizing for a single benchmark may produce dataset-specific shortcuts that pass the benchmark without exhibiting the capability it was designed to test, overstating real-world utility.
- *Benchmark scope limitations.* Benchmarks commonly ignore runtime behaviors, security implications, readability, testability, and maintainability, yielding results that may not transfer to realistic development settings.

If constructs are based on subjective interpretations, purely automated metrics are insufficient. Researchers **must** discuss how they ensured quality of subjective results, similarly to qualitative research.

External Validity: The primary threats to external validity are:

- *Cross-model transfer limitations.* Results obtained with one LLM or family of LLMs may not generalize to others due to differences in training data, architecture, and post-training procedures (e.g., fine-tuning and RLHF).
- *Tool-architecture specificity.* Tools built around vendor-specific APIs or features (e.g., function calling, structured output, or context-window size) may not transfer to other models without substantial re-engineering.
- *Limited domain coverage.* Studies often focus on a narrow set of programming languages, task types, or application domains, limiting generalizability to other SE contexts.
- *Limited participant diversity.* Study participants such as human validators or developers may not represent the broader population in terms of expertise, geographic location, or cultural background.
- *Cross-time instability.* Performance of proprietary models can change over time, leading to non-generalizable study outcomes [30, 76].

Generalizability is particularly critical for proprietary and non-deterministic systems whose behavior is subject to drift (i.e., silent changes in model output over time) [30]. Researchers **must** discuss the limitations and mitigations of external validity.

Reliability & Reproducibility: The primary threats to reliability and reproducibility are:

- *Non-deterministic outputs.* Identical prompts and configurations can yield different outputs across runs due to factors such as floating-point arithmetic, batching, and stochastic decoding strategies.
- *Infrastructure dependence.* Results may vary depending on the hardware, software stack, and hosting environment used; vendor-imposed quotas, throttling, or pricing changes can further prevent re-running experiments at the original scale, making exact replication challenging across different infrastructure setups.
- *Resource inequality.* LLM research is resource-intensive and remains predominantly in the domain of private companies or well-funded research institutions [3, 126], excluding researchers from under-resourced institutions.

Researchers **must** discuss the measures taken to increase reliability and reproducibility. However, non-deterministic reproducibility is not inherently disqualifying. Qualitative traditions such as ethnography, grounded theory, and action research ensure trustworthiness through *credibility*, *transferability*, *dependability*, and *confirmability* [50]. The same applies to SaaS-based LLM research, where providers frequently deprecate model versions without guaranteeing stable behavior.

Ethical & Regulatory Boundaries: The primary concerns for ethical and regulatory matters are:

- *Use of sensitive or proprietary data.* Studies involving proprietary code, confidential business data, or personally identifiable information may face restrictions on data sharing that limit reproducibility.
- *Jurisdictional obligations.* Data protection regulations such as GDPR or CCPA and institutional policies may impose constraints on data collection, processing, and sharing in LLM-based studies.
- *Implicit model bias.* Especially for qualitative research, LLMs might “reinforce dominant paradigms and biases” and “identify, replicate and reinforce dominant language and patterns” [67].

Studies involving sensitive data **must** discuss data governance mechanisms tailored towards LLM environments, compliant with applicable juristical obligations. If applicable, researchers **should** discuss how model biases potentially impact the study outcomes and how those biases were evaluated.

Environmental & Sustainability Constraints:

- *Energy consumption.* With growing model size, the environmental impact of experiments with LLMs increases, and the substantial energy costs of LLM experiments warrant consideration in study design [134].
- *Trade-off between repetition and sustainability.* Repeating experiments increases reliability but also energy consumption, requiring trade-offs during study design.

Researchers **should** justify the LLM’s resource consumption against the benefits over traditional approaches.

Mitigation Strategies: The following mitigation strategies can help address the threats described above, depending on the study scope and feasibility.

- *Replication Packages.* Cover prompt and architecture specifications, model outputs, and representative examples for partial replicability, accompanied by an implementation using an open model for long-term stability.
- *Human Validation.* Validate subjective constructs following quality criteria known from qualitative research.
- *Longitudinal Re-Runs.* Repeat experiments with LLMs over time, complemented by statistical analyses.

- *Methodological Trustworthiness Measures*. Researchers **should** consider triangulation, reflexivity, audit trails, and peer debriefing as complementary measures when deterministic reproduction is structurally impossible.
- *Triangulation*. Use multiple models (e.g., proprietary and open), multiple independent datasets, and multiple complementary metrics.
- *Cost Accounting*. Report input and output tokens, token and service costs, or hardware specifications.
- *Energy Preservation*. Select smaller or newer less resource-intensive models and apply techniques such as input/output token reduction, model pruning, quantization, or knowledge distillation [87] where feasible. Carbon footprint estimation is desirable, but still difficult.
- *Ethical and Regulatory Considerations*. Apply data governance mechanisms, ethical reviews, and bias assessment procedures.
- *Sensitivity Analysis*. Vary LLM configurations, prompts, architecture decisions, datasets, and where applicable human participant backgrounds.

5.8.3 Examples:

Sallou et al [117] catalog three categories of LLM-specific threats to validity (i.e., closed-source models, implicit data leakage, and reproducibility) and pair each with concrete mitigation strategies (e.g., versioned model archives, metamorphic test data, multiple replication runs with variability metrics, and detailed execution metadata). Du et al [38] pair each threat in their *ClassEval* evaluation with a concrete mitigation: manually constructing the benchmark with multiple annotators to limit data leakage, piloting prompts on held-out tasks to control for prompt sensitivity, and reporting greedy-decoding results to control for non-determinism.

5.8.4 Benefits:

Transparent reporting of limitations and mitigations helps readers calibrate confidence in the findings, makes explicit which threats were addressed and which remain open, and documents design decisions that other authors can borrow or refine. Transparent threat discussion also keeps a paper’s claims proportionate to its evidence.

5.8.5 Challenges:

Identifying limitations one is not already aware of is the hardest part of writing a threats section, particularly for methodological threats outside the team’s primary expertise. Publication and reviewing norms can pressure authors to downplay weaknesses, while page limits make exhaustive treatment impractical. Threat lists that recite generic LLM-research issues (e.g., model bias, non-determinism, or contamination) without showing how each one applies to specific design choices in this study leave reviewers unable to tell which risks actually applied.

5.8.6 Study Types:

Researchers **must** follow this guideline for all study types. Transparently reporting limitations and mitigations is a universal requirement, but specific concerns vary by study type. For *LLMs as Annotators*, researchers **must** discuss potential biases in label assignment, label reliability limitations, and sensitivity of annotations to prompt wording and model choice. For *LLMs as Judges*, researchers **must** address measurement validity concerns, known biases such as position bias or verbosity bias, and the extent to which LLM judgments align with human expert assessments. For *LLMs for Synthesis*, researchers **must** discuss the risk of contextual misinterpretation, potential loss of nuance in summarized or aggregated outputs, and reflexivity limitations inherent in using an LLM for qualitative interpretation. For *LLMs as Subjects*, researchers **must** discuss the fundamental inability of LLMs to truly simulate human behavior, the risk of stereotype amplification, and the limited ecological validity of simulated responses. For *Studying LLM Usage*, researchers **must** discuss generalizability constraints across different tools and user populations, and acknowledge how observed usage patterns may not transfer to other contexts. For *LLMs for Tools*, researchers **must** discuss replicability constraints arising from dependencies on commercial models, the impact of model updates on tool behavior, and limitations of the evaluation setup. For *Benchmarking LLMs*, researchers **must** discuss potential data contamination, benchmark scope limitations, and the extent to which benchmark performance generalizes to real-world tasks.

5.8.7 Advice for Reviewers:

Reviewers should verify that the limitation section is comprehensive and appropriate for the specific study type, checking that: (1) limitations address the specific threats relevant to the study type (e.g., label reliability for annotation studies, simulation fidelity for studies using LLMs as subjects); (2) mitigations are concrete and correspond to identified limitations rather than being generic statements; (3) the impact of LLM non-determinism on findings is discussed; (4) generalizability constraints, across models, configurations, time periods, and populations, are acknowledged. When important limitations are missing, reviewers should request they be added. The absence of a limitation section, or one that is formulaic or insufficiently specific, is a more serious concern than any individual missing limitation and may warrant a major revision.

5.8.8 See Also:

- Section 5.2 (Report Model Version, Configuration, and Customizations): Authors can re-run with different models or configurations to check whether the results depend on those specific choices.
- Section 5.3 (Report System and Prompt Design): Triangulation across architectures and prompts is one mitigation strategy.

- Section 5.4 (Report Session Traces): Stored session traces serve as a baseline against which authors can monitor LLM behavior drift over time.
- Section 5.5 (Use Suitable Baselines, Benchmarks, and Metrics): Benchmark and metric choices are one source of construct-validity threats authors must discuss.
- Section 5.7 (Use Human Validation for LLM Outputs): When automated metrics cannot validly measure a construct, human validation is an alternative.

6 Conclusion

LLM non-determinism, opaque training data, and rapidly evolving models threaten the reproducibility and replicability of empirical SE studies. This paper presents a taxonomy of seven study types that organizes how LLMs are used in SE research, together with eight guidelines for designing and reporting such studies. Each guideline distinguishes requirements (**must**) from recommended practices (**should**) and is contextualized by the study types it applies to. Our guidelines recommend that researchers: (1) declare LLM usage and role; (2) report model versions, configurations, and customizations; (3) document the tool architecture beyond the model; (4) disclose prompts, their development, and, where feasible, interaction logs; (5) validate LLM outputs with humans; (6) include an open LLM as a baseline; (7) use suitable baselines, benchmarks, and metrics; and (8) articulate limitations and mitigations, including costs, potential biases, and environmental impact. An applicability matrix maps guidelines to study types, and a reporting checklist (Section C) supports authors and reviewers. The study types and guidelines were derived collaboratively by 22 co-authors, following a process similar to that of guidelines in other disciplines (see, e.g., [21]).

Following the quality dimensions of the *SIGSOFT Empirical Standards* [106], we consider our guidelines (1) *comprehensive*, as they cover seven study types that we identified and were derived from 22 co-authors' collective expertise; (2) *useful*, as they provide actionable recommendations with an applicability matrix and a concrete reporting checklist (Section C); (3) *well-argued*, as each guideline is motivated by a dedicated rationale and supported by references to the literature; and (4) *integrated with previous work*, as they build on and complement established empirical SE standards while addressing LLM-specific challenges. Adopting these practices will not eliminate LLM non-determinism, but it should make claims auditable and results easier to replicate and compare across studies. Because models and tools evolve, we maintain the guidelines as a living resource and invite the community to refine and extend them at llm-guidelines.org.

Acknowledgments

Our study types and guidelines are based on discussion sessions with researchers at the *2024 International Software Engineering Research Network (ISERN)* meeting and at the *2nd Copenhagen Symposium on Human-Centered Software Engineering AI* (supported by the Carlsberg Foundation with grant CF24-0693 and the Alfred P. Sloan Foundation with grant G-2024-22586 to Daniel Russo). LLM-based tools, including ChatGPT (GPT-5) and Claude Code (Opus 4.6 and 4.7), were used interactively for language editing and structural revision suggestions across all sections of this paper. All LLM-generated suggestions were reviewed and revised by the authors.

Appendix A Applicability Matrix

Table 3 Applicability of guidelines to study types. ● = the guideline’s core recommendations **must** be followed for this study type, ● = **should** be followed, – = are not directly applicable. Each guideline’s study-type-specific guidance is detailed in the corresponding subsection.

	Annotator	Judge	Synthesis	Subject	Usage	Tools	Benchmarks
Declare Usage	●	●	●	●	●	●	●
Model Version	●	●	●	●	●	●	●
Design	●	●	●	●	●	●	●
Traces	●	●	●	●	●	●	●
Benchmarks	●	●	●	●	●	●	●
Open LLM	●	●	●	–	–	●	●
Human Validation	●	●	●	●	●	●	●
Limitations	●	●	●	●	●	●	●

Appendix B Rationale and Recommendations

Table 4 Rationale and key recommendations for each guideline. ● = must, ● = should.

Guideline	Rationale	Core Recommendations
Declare Usage	Transparency enables informed assessment of scope and limitations.	● Declare which LLM, how it was used, and where in the research process.
Model Version	Reproducibility requires precise identification of the system used in a study.	● Report exact version, date, configuration, and fine-tuning details. ● Report defaults, checksums, and quantization; motivate model choice; acknowledge commercial-model reproducibility limits.
Design	Static artifacts determine what the model sees on every call and must be documented in full.	● Describe architecture, infrastructure, prompts and templates (with development strategy), context files, tool catalog, agent architecture, and retrieval mechanisms. ● For LLM usage, describe the tool architecture to the extent accessible.
Traces	Runtime traces make LLM and agent behavior verifiable despite non-determinism and tool opacity.	● For studies of LLM usage, share full interaction logs subject to privacy constraints. ● Otherwise, share interaction logs, runtime traces, and plans where feasible.
Benchmarks	Meaningful evaluation requires reasoned valid measurement.	● Justify metric and benchmark choices; discuss their validity. ● Define the phenomenon and sampling strategy; isolate confounders; analyze errors; repeat experiments and report distributions.
Open LLM	Reproducibility depends on access to the model under study.	● For benchmarking studies, include an open LLM. ● Otherwise, include an open LLM as a baseline; provide a replication package.
Human Validation	Automated metrics alone cannot ensure validity of subjective constructs.	● Define the measured construct and share any custom measurement instruments. ● Validate against human judgment with inter-rater reliability; describe rater demographics for value-laden constructs.
Limitations	Honest acknowledgment of threats strengthens a study.	● Discuss threats to internal validity (data leakage), reliability (non-determinism), construct and external validity. ● Employ mitigations where possible.

Appendix C Reporting Checklist

The following checklist, inspired by CONSORT [125], summarizes actionable items from the guidelines based on the **summary** sections. The checklist is organized along typical paper sections. Items marked ● are requirements (**must**), and items marked ● are recommendations (**should**). Each item references its source guideline by short name. Items annotated with *paper* or *supplementary material* indicate where we expect the information to be reported. Unmarked items may be reported either in the paper or as supplementary material. Items prefixed with a bracketed tag apply only to studies with that characteristic (e.g., [fine-tuning], [agents]). Beyond these characteristic-tagged items, each guideline’s *Study Types* subsection lists study-type-specific recommendations where applicable.

Introduction

- Disclose any use of LLMs in the empirical study, specifying which LLM, how, and where it was used (Declare Usage).
- Report in the *paper* the purpose of using LLMs, the tasks they automate, and the expected benefits (Declare Usage).

Research Design and Methods

Model Selection and Configuration

- Report in the *paper* the exact LLM model or tool version, the configuration, and the experiment date (Model Version).
- [fine-tuning] For fine-tuned models, describe in the *paper* the fine-tuning goal, the dataset, and the procedure (Model Version).
- Report default parameters and explain model and version choices (Model Version).
- Report checksums and additional model properties where available; for commercial tools, openly acknowledge their reproducibility limits (Model Version).
- [quantization] For quantized models, report the quantization level (e.g., 4-bit, 8-bit) and method (e.g., GPTQ or AWQ) (Model Version).
- [fine-tuning] Compare base and fine-tuned models using suitable metrics and benchmarks; share fine-tuning data and weights as *supplementary material* (or justify in the *paper* why they cannot be shared) (Model Version).
- [commercial-models] Include an open LLM as a baseline when using commercial models and report inter-model agreement (Open LLM).

System and Prompt Design

- Describe in the *paper* the full architecture of LLM-based tools, including the role of the LLM, interactions with other components, and overall system behavior (Design).
- Specify in the *paper* whether zero-shot, one-shot, or few-shot prompting was used (Design).

- Specify prompt reuse across models and configurations (Design).
- Publish all prompts or, when using templates, prompt templates with representative instances, including their structure, content, formatting, and variable components, as *supplementary material* (Design).
- Report hosting, hardware setup, and latency implications (Design).
- [dynamic-prompts] For dynamically generated prompts, document the generation process thoroughly (Design).
- [context-files] Describe in the *paper* any context-file mechanisms used (e.g., AGENTS.md, CLAUDE.md) (Design).
- [tool-use] Summarize in the *paper* which tools and skills were exposed to the model (Design).
- [agents] If autonomous agents are used, specify agent roles, reasoning frameworks, and communication flows (Design).
- [context-augmentation] For retrieval-augmented generation (RAG) or related methods, describe how external data was retrieved, stored, and integrated (Design).
- Justify design decisions (Design).
- Describe prompt development rationale and selection process (Design).
- Report prompt evolution and any LLM-suggested refinements (Design).
- Where legally possible, release the source code of the implementation under an open-source license (Design).
- [participant-prompts] For user-authored prompts, describe how they were collected and analyzed (Design).
- [long-prompts] Document input handling and token optimization strategies when prompts are long or complex (Design).
- [restricted-sharing] If full prompt disclosure is not feasible, provide summaries or examples (Design).
- [ensemble] For ensemble architectures, explain in the *paper* the coordination logic between models (Design).
- [context-augmentation] Report data preprocessing, versioning, and update frequency for stored data used for context augmentation (Design).
- [context-files] Include context-file contents as *supplementary material* (Design).
- [tool-use] Include the complete list of tools and skills (names with one-line purposes), tool schemas, skill definitions, sub-agent definitions, and connected MCP servers as *supplementary material* (Design).

Session Traces

- Include full interaction logs (prompts and responses) as *supplementary material* if privacy and confidentiality can be ensured (Traces).
- [agents] For agentic systems, include interaction logs covering exchanges between humans and the agent and between external tools and the agent (human-in-the-loop feedback, approvals, refinements) as *supplementary material* (Traces).
- [agents] For agentic systems, report the complete runtime trace as *supplementary material*, including for each entry the tool or artifact name, ar-

guments, result, and ordering, and which configured artifacts (skills, context files, sub-agents) were activated (Traces).

● [agents] For agentic systems, report developed plans as *supplementary material* if available (Traces).

Benchmarks and Metrics

- Justify in the *paper* all benchmark and metric choices (Benchmarks).
- Explain in the *paper* why the selected metrics are suitable for the specific study (Benchmarks).
- Provide an operational definition of the phenomenon the benchmark is intended to measure, including its scope and any sub-components (Benchmarks).
- Summarize benchmark structure, task types, and limitations (Benchmarks).
- Identify the capabilities a benchmark conflates with the target phenomenon, isolate the target where possible, and acknowledge remaining confounders as construct-validity threats (Benchmarks).
- Perform an error analysis: categorize the failures observed and report their relative frequency; report failures that cluster on confounding capabilities as construct-validity threats (Benchmarks).
- Describe and justify the sampling strategy used to select problems for inclusion in the benchmark (Benchmarks).
- Justify the number of experiment repetitions, for example through a power analysis or by monitoring convergence of descriptive statistics (Benchmarks).
- [non-probability-sampling] For non-probability sampling (e.g., convenience), discuss the implications for the generalizability of conclusions (Benchmarks).
- [new-benchmark] For new or released benchmarks, adopt contamination-prevention mechanisms: held-out subset, canary strings, and pre-exposure investigation against common training corpora (Benchmarks).
- [multi-rater-scoring] For ratings that vary across raters or runs (human raters, LLM-as-judge), report the distribution of ratings per item rather than only aggregated point estimates (Benchmarks).

Human Validation

- [human-validation] If using human validation, define in the *paper* the measured construct (e.g., usability, maintainability) and describe the measurement instrument (Human Validation).
- [human-validation] When developing or adapting measurement instruments, share them (Human Validation).
- Consider human validation early in the study design and build on established reference models for human-LLM comparison (Human Validation).
- [human-validation] Validate LLM judgments against human judgment, report aggregation methods, and assess human-LLM agreement (Human Validation).
- [human-validation] Discuss and, where feasible, control for confounding factors (Human Validation).

- [human-validation] [subjective-constructs] For value-laden or culturally contingent constructs, describe rater demographics beyond expertise and discuss potential demographic biases (Human Validation).

Reproducibility, Ethics, and Resources

- [restricted-sharing] For studies involving sensitive data, discuss data governance mechanisms compliant with applicable jurisdictional obligations (Limitations).
- Justify LLM usage in light of its resource demands (Limitations).
- Provide a full replication package as *supplementary material*, including step-by-step instructions for verifying and reproducing the results (Open LLM).
- [restricted-sharing] Where full sharing of prompts, traces, or datasets is not feasible, share representative examples for partial replicability (Limitations).

Results

- [comparing-models] If comparing models or tools, use appropriate inferential statistics (e.g., hypothesis tests, effect sizes) rather than relying solely on summary statistics (Benchmarks).
- Repeat experiments due to the inherent non-determinism of LLMs and report the result distribution using descriptive statistics (Benchmarks).
- Use traditional (non-LLM) baselines for comparison where possible (Benchmarks).
- Report established metrics to make study results comparable; additional metrics may be reported where appropriate (Benchmarks).

Limitations and Threats to Validity

- Describe measurement constructs and methods; disclose any data leakage risks and avoid leaking evaluation data into LLM improvement pipelines (Limitations).
- Transparently report study limitations, including the impact of non-determinism and generalizability constraints (Limitations).
- Specify whether generalization across LLMs or across time was assessed, and discuss model and version differences (Limitations).
- [restricted-sharing] Acknowledge non-disclosed confidential or proprietary components as reproducibility limitations (Design).
- Employ and report strategies to mitigate identified validity and reproducibility threats, such as replication packages, human validation, longitudinal re-runs, triangulation, and sensitivity analysis (Limitations).

7 Statements and Declarations

Competing Interests: The authors declare no competing interests. Multiple co-authors are members of the EMSE Editorial Board, which is disclosed here for transparency.

Authorship: Each author contributed or reviewed content and, most importantly, read the complete guidelines and confirmed that they stand behind all recommendations, not only their own contributions.

Funding: No funding was received for writing this article. However, the *2nd Copenhagen Symposium on Human-Centered Software Engineering AI*, which played a central role in forming the team of co-authors, was supported by the Carlsberg Foundation with grant CF24-0693 and the Alfred P. Sloan Foundation with grant G-2024-22586 to Daniel Russo.

Ethical Approval: No ethics approval was required for writing this article.

Data Availability: This article has no associated data. The guidelines are hosted online and maintained on GitHub.

References

1. ACM Publications Board (2020) Artifact review and badging – current. <https://www.acm.org/publications/policies/artifact-review-and-badging-current>, accessed 2025-01-13
2. Agarwal R, Schwarzer M, Castro PS, Courville AC, Bellemare MG (2021) Deep reinforcement learning at the edge of the statistical precipice. In: Ranzato M, Beygelzimer A, Dauphin YN, Liang P, Vaughan JW (eds) *Advances in Neural Information Processing Systems 34: Annual Conference on Neural Information Processing Systems 2021, NeurIPS 2021, December 6-14, 2021, virtual*, pp 29304–29320, URL <https://proceedings.neurips.cc/paper/2021/hash/f514cec81cb148559cf475e7426eed5e-Abstract.html>
3. Ahmed N, Wahed M, Thompson NC (2023) The growing influence of industry in AI research. *Science* 379(6635):884–886, DOI 10.1126/science.ade2420
4. Ahmed T, Devanbu PT, Treude C, Pradel M (2025) Can LLMs replace manual annotation of software engineering artifacts? In: *22nd IEEE/ACM International Conference on Mining Software Repositories, MSR@ICSE 2025, Ottawa, ON, Canada, April 28-29, 2025, IEEE*, pp 526–538, DOI 10.1109/MSR66628.2025.00086
5. Ahuja S, Gumma V, Sitaram S (2024) Contamination report for multilingual benchmarks. CoRR abs/2410.16186, DOI 10.48550/ARXIV.2410.16186, URL <https://doi.org/10.48550/arXiv.2410.16186>, 2410.16186
6. Angermeier F, Amougou M, Kreitz M, Bauer A, Linhuber M, Fucci D, Constante FM, Méndez D, Gorschek T (2025) Reflections on the reproducibility of commercial LLM performance in empirical software engineer-

- ing studies. CoRR abs/2510.25506, DOI 10.48550/ARXIV.2510.25506, URL <https://doi.org/10.48550/arXiv.2510.25506>, 2510.25506
7. Anthropic (2025) Demystifying Evals for AI Agents. <https://www.anthropic.com/engineering/demystifying-evals-for-ai-agents>, accessed 2026-03-19
 8. Anthropic (2026) Claude Code Tools Reference. <https://code.claude.com/docs/en/tools-reference>, accessed 2026-04-28
 9. Argyle LP, Busby EC, Fulda N, Gubler J, Rytting CM, Wingate D (2022) Out of one, many: Using language models to simulate human samples. CoRR abs/2209.06899, DOI 10.48550/ARXIV.2209.06899, URL <https://doi.org/10.48550/arXiv.2209.06899>, 2209.06899
 10. Arize AI (2026) OpenInference: AI observability and evaluation specification. <https://github.com/Arize-ai/openinference>, accessed 2026-05-07
 11. Association for Computing Machinery (2023) ACM Policy on Authorship. <https://www.acm.org/publications/policies/new-acm-policy-on-authorship>, accessed 2025-01-13
 12. Austin J, Odena A, Nye MI, Bosma M, Michalewski H, Dohan D, Jiang E, Cai CJ, Terry M, Le QV, Sutton C (2021) Program synthesis with large language models. CoRR abs/2108.07732, URL <https://arxiv.org/abs/2108.07732>, 2108.07732
 13. Azanza M, Pereira J, Irastorza A, Galdos A (2024) Can LLMs facilitate onboarding software developers? an ongoing industrial case study. In: 36th International Conference on Software Engineering Education and Training, CSEET 2024, IEEE, pp 1–6, DOI 10.1109/CSEET62301.2024.10662989
 14. Baltes S, Ralph P (2022) Sampling in software engineering research: a critical review and guidelines. *Empir Softw Eng* 27(4):94, DOI 10.1007/s10664-021-10072-8, URL <https://doi.org/10.1007/s10664-021-10072-8>
 15. Bano M, Zowghi D, Whittle J (2023) Exploring qualitative research using llms. CoRR abs/2306.13298, DOI 10.48550/ARXIV.2306.13298, URL <https://doi.org/10.48550/arXiv.2306.13298>, 2306.13298
 16. Bano M, Hoda R, Zowghi D, Treude C (2024) Large language models for qualitative research in software engineering: exploring opportunities and challenges. *Autom Softw Eng* 31(1):8, DOI 10.1007/S10515-023-00407-8, URL <https://doi.org/10.1007/s10515-023-00407-8>
 17. Bano M, Gunatilake H, Hoda R (2025) What does a software engineer look like? exploring societal stereotypes in llms. In: 47th IEEE/ACM International Conference on Software Engineering: Software Engineering in Society, ICSE-SEIS 2025, Ottawa, ON, Canada, April 27 - May 3, 2025, IEEE, pp 173–184, DOI 10.1109/ICSE-SEIS66351.2025.00023, URL <https://doi.org/10.1109/ICSE-SEIS66351.2025.00023>
 18. Barros CF, Azevedo BB, Neto VVG, Kassab M, Kalinowski M, do Nascimento HAD, Bandeira MCGSP (2025) Large language model for qualitative research: A systematic mapping study. In: IEEE/ACM Interna-

- tional Workshop on Methodological Issues with Empirical Studies in Software Engineering, WSESE@ICSE 2025, Ottawa, ON, Canada, May 3, 2025, IEEE, pp 48–55, DOI 10.1109/WSESE66602.2025.00015, URL <https://doi.org/10.1109/WSESE66602.2025.00015>
19. Bavaresco A, Bernardi R, Bertolazzi L, Elliott D, Fernández R, Gatt A, Ghaleb E, Giulianelli M, Hanna M, Koller A, Martins AFT, Mondorf P, Neplenbroek V, Pezzelle S, Plank B, Schlangen D, Suglia A, Surikuchi AK, Takmaz E, Testoni A (2024) LLMs instead of human judges? A large scale empirical study across 20 NLP evaluation tasks. CoRR abs/2406.18403, DOI 10.48550/ARXIV.2406.18403, URL <https://doi.org/10.48550/arXiv.2406.18403>, 2406.18403
 20. Bean AM, Kearns RO, Romanou A, Hafner FS, Mayne H, Batzner J, Foroutan N, Schmitz C, Korgul K, Batra H, Deb O, Beharry E, Emde C, Foster T, Gausen A, Grandury M, Han S, Hofmann V, Ibrahim L, Kim H, Kirk HR, Lin F, Liu GKM, Luettgau L, Magomere J, Rysstrøm J, Sotnikova A, Yang Y, Zhao Y, Bibi A, Bosselut A, Clark R, Cohan A, Foerster J, Gal Y, Hale SA, Raji ID, Summerfield C, Torr PHS, Ududec C, Rocher L, Mahdi A (2025) Measuring what matters: Construct validity in large language model benchmarks. In: Advances in Neural Information Processing Systems 38: Annual Conference on Neural Information Processing Systems 2025, NeurIPS 2025, Datasets and Benchmarks Track, URL <https://arxiv.org/abs/2511.04703>
 21. Begg C, Cho M, Eastwood S, Horton R, Moher D, Olkin I, Pitkin R, Rennie D, Schulz KF, Simel D, Stroup DF (1996) Improving the quality of reporting of randomized controlled trials. The CONSORT statement. JAMA 276(8):637–639, DOI 10.1001/jama.276.8.637
 22. Bjarnason BH, Silva A, Monperrus M (2026) On randomness in agentic evals. CoRR abs/2602.07150, DOI 10.48550/ARXIV.2602.07150, URL <https://doi.org/10.48550/arXiv.2602.07150>, 2602.07150
 23. Blackwell RE, Barry J, Cohn AG (2024) Towards reproducible LLM evaluation: Quantifying uncertainty in LLM benchmark scores. CoRR abs/2410.03492, DOI 10.48550/ARXIV.2410.03492, URL <https://doi.org/10.48550/arXiv.2410.03492>, 2410.03492
 24. Borsboom D (2005) Measuring the Mind: Conceptual Issues in Contemporary Psychometrics. Cambridge University Press, DOI 10.1017/CBO9780511490026
 25. Bouzenia I, Pradel M (2025) Understanding software engineering agents: A study of thought-action-result trajectories. In: 40th IEEE/ACM International Conference on Automated Software Engineering, ASE 2025, Seoul, Korea, Republic of, November 16–20, 2025, IEEE, pp 2846–2857, DOI 10.1109/ASE63991.2025.00234, URL <https://doi.org/10.1109/ASE63991.2025.00234>
 26. Cai Y, Yadavally A, Mishra A, Montejo G, Nguyen TN (2024) Programming assistant for exception handling with CodeBERT. In: Proceedings of the 46th IEEE/ACM International Conference on Software Engineering, ICSE 2024, Lisbon, Portugal, April 14–20, 2024, ACM, pp 94:1–94:13,

- DOI 10.1145/3597503.3639188
27. Cao J, Chan YK, Ling Z, Wang W, Li S, Liu M, Qiao R, Han Y, Wang C, Yu B, He P, Wang S, Zheng Z, Lyu MR, Cheung SC (2026) Rigor, reliability, and reproducibility matter: A decade-scale survey of 572 code benchmarks. CoRR abs/2501.10711, URL <https://arxiv.org/abs/2501.10711>, 2501.10711
 28. Chandra S (2025) Benchmarks for AI in Software Engineering (BLOG@CACM). <https://cacm.acm.org/blogcacm/benchmarks-for-ai-in-software-engineering/>, accessed 2025-08-24
 29. Chann S (2023) Non-determinism in GPT-4 is caused by Sparse MoE. <https://152334h.github.io/blog/non-determinism-in-gpt-4/>, accessed 2025-01-13
 30. Chen L, Zaharia M, Zou J (2024) How is ChatGPT's behavior changing over time? Harvard Data Science Review 6(2), DOI 10.1162/99608f92.5317da47
 31. Chen M, Tworek J, Jun H, Yuan Q, de Oliveira Pinto HP, Kaplan J, Edwards H, Burda Y, Joseph N, Brockman G, Ray A, Puri R, Krueger G, Petrov M, Khlaaf H, Sastry G, Mishkin P, Chan B, Gray S, Ryder N, Pavlov M, Power A, Kaiser L, Bavarian M, Winter C, Tillet P, Such FP, Cummings D, Plappert M, Chantzis F, Barnes E, Herbert-Voss A, Guss WH, Nichol A, Paino A, Tezak N, Tang J, Babuschkin I, Balaji S, Jain S, Saunders W, Hesse C, Carr AN, Leike J, Achiam J, Misra V, Morikawa E, Radford A, Knight M, Brundage M, Murati M, Mayer K, Welinder P, McGrew B, Amodei D, McCandlish S, Sutskever I, Zaremba W (2021) Evaluating large language models trained on code. CoRR abs/2107.03374, URL <https://arxiv.org/abs/2107.03374>, 2107.03374
 32. Chen Y, Ding Z, Alowain L, Chen X, Wagner DA (2023) DiverseVul: A new vulnerable source code dataset for deep learning based vulnerability detection. In: Proceedings of the 26th International Symposium on Research in Attacks, Intrusions and Defenses, RAID 2023, Hong Kong, China, October 16-18, 2023, ACM, pp 654–668, DOI 10.1145/3607199.3607242
 33. Cheng A, Calhoun A, Reedy G (2025) Artificial intelligence-assisted academic writing: recommendations for ethical use. Advances in Simulation 10(1):26, DOI 10.1186/s41077-025-00350-6, URL <https://doi.org/10.1186/s41077-025-00350-6>
 34. Choi HK, Khanov M, Wei H, Li Y (2025) How contaminated is your benchmark? quantifying dataset leakage in large language models with kernel divergence. CoRR abs/2502.00678, DOI 10.48550/ARXIV.2502.00678, URL <https://doi.org/10.48550/arXiv.2502.00678>, 2502.00678
 35. Cohen J (1992) A power primer. Psychological Bulletin 112(1):155–159, DOI 10.1037/0033-2909.112.1.155
 36. DAIRAI (2024) Elements of a Prompt. <https://www.promptingguide.ai/introduction/elements>, accessed 2026-04-28

37. Dhar R, Vaidhyathan K, Varma V (2024) Can LLMs generate architectural design decisions? - an exploratory empirical study. In: 21st IEEE International Conference on Software Architecture, ICSA 2024, Hyderabad, India, June 4-8, 2024, IEEE, pp 79–89, DOI 10.1109/ICSA59870.2024.00016
38. Du X, Liu M, Wang K, Wang H, Liu J, Chen Y, Feng J, Sha C, Peng X, Lou Y (2024) Evaluating large language models in class-level code generation. In: Proceedings of the 46th IEEE/ACM International Conference on Software Engineering, ICSE 2024, Lisbon, Portugal, April 14-20, 2024, ACM, pp 81:1–81:13, DOI 10.1145/3597503.3639219, URL <https://doi.org/10.1145/3597503.3639219>
39. Dybå T, Kampenes VB, Sjøberg DIK (2006) A systematic review of statistical power in software engineering experiments. *Inf Softw Technol* 48(8):745–755, DOI 10.1016/J.INFSOF.2005.08.009, URL <https://doi.org/10.1016/j.infsof.2005.08.009>
40. Eghbali A, Pradel M (2022) CrystalBLEU: Precisely and efficiently measuring the similarity of code. In: 37th IEEE/ACM International Conference on Automated Software Engineering, ASE 2022, ACM, pp 28:1–28:12, DOI 10.1145/3551349.3556903
41. El-Hajjami A, Salinesi C (2025) How good are synthetic requirements? evaluating LLM-generated datasets for AI4RE. *CoRR* abs/2506.21138, DOI 10.48550/ARXIV.2506.21138, URL <https://doi.org/10.48550/arXiv.2506.21138>, 2506.21138
42. Gallegos IO, Rossi RA, Barrow J, Tanjim MM, Kim S, Dernoncourt F, Yu T, Zhang R, Ahmed NK (2024) Bias and fairness in large language models: A survey. *Comput Linguistics* 50(3):1097–1179, DOI 10.1162/COLI_A_00524, URL https://doi.org/10.1162/coli_a_00524
43. Gallifant J, Afshar M, Ameen S, Aphinyanaphongs Y, Chen S, Cacciamani G, Demner-Fushman D, Dligach D, Daneshjou R, Fernandes C, Hansen LH, Landman A, Lehmann L, McCoy LG, Miller T, Moreno A, Munch N, Restrepo D, Savova G, Umeton R, Gichoya JW, Collins GS, Moons KGM, Celi LA, Bitterman DS (2025) The TRIPOD-LLM reporting guideline for studies using large language models. *Nature Medicine* 31(1):60–69, DOI 10.1038/s41591-024-03425-5
44. Galster M, Mohsenimofidi S, Lulla JL, Abubakar MA, Treude C, Baltes S (2026) Configuring agentic AI coding tools: An exploratory study. In: Proceedings of the 3rd ACM International Conference on AI-powered Software (AIware 2026)
45. Gerosa MA, Trinkenreich B, Steinmacher I, Sarma A (2024) Can AI serve as a substitute for human subjects in software engineering research? *Autom Softw Eng* 31(1):13, DOI 10.1007/S10515-023-00409-6, URL <https://doi.org/10.1007/s10515-023-00409-6>
46. Gibney E (2024) Not all ‘open source’ AI models are actually open. *Nature News* DOI 10.1038/d41586-024-02012-5
47. Gilardi F, Alizadeh M, Kubli M (2023) ChatGPT outperforms crowd workers for text-annotation tasks. *Proceedings of the National Academy*

- of Sciences 120(30):e2305016120, DOI 10.1073/pnas.2305016120
48. Gonçalves J, Silva M, Cabral B, Dias T, Maia E, Praça I, Severino R, Ferreira LL (2025) Evaluating llama 3.2 for software vulnerability detection. In: Praça I, Bernardi S, Inácio PRM (eds) *Cybersecurity - 9th European Interdisciplinary Cybersecurity Conference, EICC 2025, Rennes, France, June 18-19, 2025, Proceedings*, Springer, Communications in Computer and Information Science, pp 38–51, DOI 10.1007/978-3-031-94855-8_3, URL https://doi.org/10.1007/978-3-031-94855-8_3
 49. Graziotin D (2024) acmsigsoft/open-science-policies: v1.0.0. DOI 10.5281/zenodo.10796477
 50. Guba EG (1981) Criteria for assessing the trustworthiness of naturalistic inquiries. *ECTJ* 29(2):75–91, DOI 10.1007/BF02766777
 51. Guo D, Zhu Q, Yang D, Xie Z, Dong K, Zhang W, Chen G, Bi X, Wu Y, Li YK, Luo F, Xiong Y, Liang W (2024) Deepseek-coder: When the large language model meets programming - the rise of code intelligence. *CoRR* abs/2401.14196, DOI 10.48550/ARXIV.2401.14196, URL <https://doi.org/10.48550/arXiv.2401.14196>, 2401.14196
 52. Harding J, D'Alessandro W, Laskowski NG, Long R (2024) AI language models cannot replace human research participants. *AI Soc* 39(5):2603–2605, DOI 10.1007/S00146-023-01725-X, URL <https://doi.org/10.1007/s00146-023-01725-x>
 53. He J, Rungta M, Koleczek D, Sekhon A, Wang FX, Hasan S (2024) Does prompt formatting have any impact on LLM performance? *CoRR* abs/2411.10541, DOI 10.48550/ARXIV.2411.10541, URL <https://doi.org/10.48550/arXiv.2411.10541>, 2411.10541
 54. He J, Shi J, Zhuo TY, Treude C, Sun J, Xing Z, Du X, Lo D (2026) LLM-as-a-judge for software engineering: Literature review, vision, and the road ahead. *ACM Transactions on Software Engineering and Methodology* DOI 10.1145/3797276, URL <https://doi.org/10.1145/3797276>
 55. He Z, Huang C, Ding CC, Rohatgi S, Huang TK (2024) If in a crowd-sourced data annotation pipeline, a GPT-4. In: Mueller FF, Kyburz P, Williamson JR, Sas C, Wilson ML, Dugas POT, Shklovski I (eds) *Proceedings of the CHI Conference on Human Factors in Computing Systems, CHI 2024, ACM*, pp 1040:1–1040:25, DOI 10.1145/3613904.3642834
 56. Holtzman A, Buys J, Du L, Forbes M, Choi Y (2020) The curious case of neural text degeneration. In: *8th International Conference on Learning Representations, ICLR 2020, Addis Ababa, Ethiopia, April 26-30, 2020*, OpenReview.net, URL <https://openreview.net/forum?id=rygGQyrFvH>
 57. Hou X, Zhao Y, Liu Y, Yang Z, Wang K, Li L, Luo X, Lo D, Grundy J, Wang H (2024) Large language models for software engineering: A systematic literature review. *ACM Trans Softw Eng Methodol* 33(8):220:1–220:79, DOI 10.1145/3695988, URL <https://doi.org/10.1145/3695988>
 58. Hu X, Niu F, Chen J, Zhou X, Zhang J, He J, Xia X, Lo D (2025) Assessing and advancing benchmarks for evaluating large lan-

- guage models in software engineering tasks. CoRR abs/2505.08903, DOI 10.48550/ARXIV.2505.08903, URL <https://doi.org/10.48550/arXiv.2505.08903>, 2505.08903
59. Huang F, Kwak H, An J (2023) Is ChatGPT better than human annotators? potential and limitations of ChatGPT in explaining implicit hate speech. In: Ding Y, Tang J, Sequeda JF, Aroyo L, Castillo C, Houben G (eds) Companion Proceedings of the ACM Web Conference 2023, WWW 2023, ACM, pp 294–297, DOI 10.1145/3543873.3587368
60. Huang J, Yi B, Sun W, Wan B, Xu Y, Feng Y, Ye W, Qin Q (2024) Enhancing review classification via LLM-based data annotation and multi-perspective feature representation learning. SSRN Electronic Journal pp 1–15, DOI 10.2139/ssrn.5002351, URL <https://ssrn.com/abstract=5002351>
61. Hui B, Yang J, Cui Z, Yang J, Liu D, Zhang L, Liu T, Zhang J, Yu B, Dang K, Yang A, Men R, Huang F, Ren X, Ren X, Zhou J, Lin J (2024) Qwen2.5-coder technical report. CoRR abs/2409.12186, DOI 10.48550/ARXIV.2409.12186, URL <https://doi.org/10.48550/arXiv.2409.12186>, 2409.12186
62. Husain H, Wu H, Gazit T, Allamanis M, Brockschmidt M (2019) CodeSearchNet challenge: Evaluating the state of semantic code search. CoRR abs/1909.09436, URL <http://arxiv.org/abs/1909.09436>, 1909.09436
63. Hymel C, Johnson H (2025) Analysis of llms vs human experts in requirements engineering. CoRR abs/2501.19297, DOI 10.48550/ARXIV.2501.19297, URL <https://doi.org/10.48550/arXiv.2501.19297>, 2501.19297
64. Jahic J, Sami A (2024) State of practice: LLMs in software engineering and software architecture. In: 21st IEEE International Conference on Software Architecture, ICSA 2024 - Companion, Hyderabad, India, June 4-8, 2024, IEEE, pp 311–318, DOI 10.1109/ICSA-C63560.2024.00059
65. Jain N, Han K, Gu A, Li W, Yan F, Zhang T, Wang S, Solar-Lezama A, Sen K, Stoica I (2024) LiveCodeBench: Holistic and contamination free evaluation of large language models for code. CoRR abs/2403.07974, DOI 10.48550/ARXIV.2403.07974, URL <https://doi.org/10.48550/arXiv.2403.07974>, 2403.07974
66. Jimenez CE, Yang J, Wettig A, Yao S, Pei K, Press O, Narasimhan KR (2024) SWE-bench: Can language models resolve real-world github issues? In: The Twelfth International Conference on Learning Representations, ICLR 2024, Vienna, Austria, May 7-11, 2024, OpenReview.net, URL <https://openreview.net/forum?id=VTF8yNQM66>
67. Jowsey T, Braun V, Clarke V, Clarke V, Lupton D, Fine M (2025) We reject the use of generative artificial intelligence for reflexive qualitative research. SSRN DOI 10.2139/ssrn.5676462, URL <https://ssrn.com/abstract=5676462>
68. Kang S, Yoon J, Yoo S (2023) Large language models are few-shot testers: Exploring LLM-based general bug reproduction. In: 45th IEEE/ACM

- International Conference on Software Engineering, ICSE 2023, IEEE, pp 2312–2323, DOI 10.1109/ICSE48619.2023.00194
69. Khojah R, Mohamad M, Leitner P, de Oliveira Neto FG (2024) Beyond code generation: An observational study of ChatGPT usage in software engineering practice. *Proc ACM Softw Eng* 1(FSE):1819–1840, DOI 10.1145/3660788
70. Khraisha Q, Put S, Kappenberg J, Warraitch A, Hadfield K (2024) Can large language models replace humans in systematic reviews? evaluating GPT-4’s efficacy in screening and extracting data from peer-reviewed and grey literature in multiple languages. *Research Synthesis Methods* 15(4):616–626, DOI 10.1002/jrsm.1715, URL <https://onlinelibrary.wiley.com/doi/abs/10.1002/jrsm.1715>
71. von Kistowski J, Arnold JA, Huppler K, Lange K, Henning JL, Cao P (2015) How to build a benchmark. In: John LK, Smith CU, Sachs K, Lladó CM (eds) *Proceedings of the 6th ACM/SPEC International Conference on Performance Engineering*, Austin, TX, USA, January 31 - February 4, 2015, ACM, pp 333–336, DOI 10.1145/2668930.2688819, URL <https://doi.org/10.1145/2668930.2688819>
72. Kong A, Zhao S, Chen H, Li Q, Qin Y, Sun R, Zhou X, Wang E, Dong X (2024) Better zero-shot reasoning with role-play prompting. In: Duh K, Gómez-Adorno H, Bethard S (eds) *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, NAACL 2024, Mexico City, Mexico, June 16-21, 2024, Association for Computational Linguistics, pp 4099–4113, DOI 10.18653/V1/2024.NAACL-LONG.228, URL <https://doi.org/10.18653/v1/2024.naacl-long.228>
73. Krippendorff K (2018) *Content Analysis: An Introduction to Its Methodology*, 4th edn. SAGE Publications
74. Kübler JM, Budhathoki K, Kleindessner M, Zhou X, Yin J, Khetan A, Karypis G (2026) When LLMs get significantly worse: A statistical approach to detect model degradations. In: *Proc. of the International Conference on Learning Representations (ICLR)*
75. Kulal S, Pasupat P, Chandra K, Lee M, Padon O, Aiken A, Liang P (2019) Spoc: Search-based pseudocode to code. In: Wallach HM, Larochelle H, Beygelzimer A, d’Alché-Buc F, Fox EB, Garnett R (eds) *Advances in Neural Information Processing Systems 32: Annual Conference on Neural Information Processing Systems 2019, NeurIPS 2019, December 8-14, 2019, Vancouver, BC, Canada*, pp 11883–11894, URL <https://proceedings.neurips.cc/paper/2019/hash/7298332f04ac004a0ca44cc69ecf6f6b-Abstract.html>
76. Li D, Gupta K, Bhaduri M, Sathiadoss P, Bhatnagar S, Chong J (2024) Comparing GPT-3.5 and GPT-4 accuracy and drift in radiology diagnosis please cases. *Radiology* 310(1):e232411, DOI 10.1148/radiol.232411
77. Li R, Allal LB, Zi Y, Muennighoff N, Kocetkov D, Mou C, Marone M, Akiki C, Li J, Chim J, Liu Q, Zheltonozhskii E, Zhuo TY, Wang T, De-

- haene O, Davaadorj M, Lamy-Poirier J, Monteiro J, Shliazhko O, Gontier N, Meade N, Zebaze A, Yee M, Umapathi LK, Zhu J, Lipkin B, Oblokulov M, Wang Z, V RM, Stillerman JT, Patel SS, Abulkhanov D, Zocca M, Dey M, Zhang Z, Fahmy N, Bhattacharyya U, Yu W, Singh S, Luccioni S, Villegas P, Kunakov M, Zhdanov F, Romero M, Lee T, Timor N, Ding J, Schlesinger C, Schoelkopf H, Ebert J, Dao T, Mishra M, Gu A, Robinson J, Anderson CJ, Dolan-Gavitt B, Contractor D, Reddy S, Fried D, Bahdanau D, Jernite Y, Ferrandis CM, Hughes S, Wolf T, Guha A, von Werra L, de Vries H (2023) Starcoder: may the source be with you! *Trans Mach Learn Res* 2023, URL <https://openreview.net/forum?id=KoF0g41haE>
78. Liang JT, Badea C, Bird C, DeLine R, Ford D, Forsgren N, Zimmermann T (2024) Can GPT-4 replicate empirical software engineering research? *Proc ACM Softw Eng* 1(FSE):1330–1353, DOI 10.1145/3660767, URL <https://doi.org/10.1145/3660767>
 79. Liu J, Xia CS, Wang Y, Zhang L (2023) Is your code generated by ChatGPT really correct? rigorous evaluation of large language models for code generation. In: Oh A, Naumann T, Globerson A, Saenko K, Hardt M, Levine S (eds) *Advances in Neural Information Processing Systems 36: Annual Conference on Neural Information Processing Systems 2023, NeurIPS 2023*, URL http://papers.nips.cc/paper_files/paper/2023/hash/43e9d647ccd3e4b7b5baab53f0368686-Abstract-Conference.html
 80. Lu Q, Zhu L, Xu X, Xing Z, Whittle J (2024) Toward responsible AI in the era of generative AI: A reference architecture for designing foundation model-based systems. *IEEE Softw* 41(6):91–100, DOI 10.1109/MS.2024.3406333
 81. Lubos S, Felfernig A, Tran TNT, Garber D, Mansi ME, Erdeniz SP, Le V (2024) Leveraging LLMs for the quality assurance of software requirements. In: Liebel G, Hadar I, Spoletini P (eds) *32nd IEEE International Requirements Engineering Conference, RE 2024, Reykjavik, Iceland, June 24-28, 2024, IEEE*, pp 389–397, DOI 10.1109/RE59067.2024.00046
 82. Madampe K, Grundy J, Hoda R, Obie HO (2024) The struggle is real! the agony of recruiting participants for empirical software engineering studies. In: *2024 IEEE Symposium on Visual Languages and Human-Centric Computing (VL/HCC)*, Liverpool, UK, September 2-6, 2024, IEEE, pp 417–422, DOI 10.1109/VL/HCC60511.2024.00065
 83. McDonald N, Schoenebeck S, Forte A (2019) Reliability and inter-rater reliability in qualitative research: Norms and guidelines for CSCW and HCI practice. *Proc ACM Hum Comput Interact* 3(CSCW):72:1–72:23, DOI 10.1145/3359174
 84. Menzies T (2025) The case for compact AI. *Commun ACM* 68(9):6–7, DOI 10.1145/3746057
 85. Microsoft (2023) CodeBERT on GitHub. <https://github.com/microsoft/CodeBERT>, accessed 2025-08-15
 86. Microsoft (2025) Azure OpenAI Service models. <https://learn.microsoft.com/en-us/azure/ai-services/openai/concepts/>

- models, accessed 2025-01-13
87. Mitu NE, Mitu GT (2024) The hidden cost of AI: Carbon footprint and mitigation strategies. *Revista de Stiinte Politice Revue des Sciences Politiques* 84:9–16, DOI 10.2139/ssrn.5036344
 88. Mohamed S, Parvin A, Parra E (2024) Chatting with AI: deciphering developer conversations with ChatGPT. In: Spinellis D, Bacchelli A, Constantinou E (eds) 21st IEEE/ACM International Conference on Mining Software Repositories, MSR 2024, Lisbon, Portugal, April 15-16, 2024, ACM, pp 187–191, DOI 10.1145/3643991.3645078
 89. Mohsenimofidi S, Galster M, Treude C, Baltes S (2026) Context engineering for AI agents in open-source software. In: Proceedings of the 23rd IEEE/ACM International Conference on Mining Software Repositories (MSR 2026), URL <https://arxiv.org/abs/2510.21413>
 90. Montes CM, Feldt R, Martos CM, Ouhbi S, Premanandan S, Graziotin D (2025) Large language models in thematic analysis: Prompt engineering, evaluation, and guidelines for qualitative software engineering research. CoRR abs/2510.18456, DOI 10.48550/ARXIV.2510.18456, URL <https://doi.org/10.48550/arXiv.2510.18456>, 2510.18456
 91. de Moraes Leça M, Valença L, Santos R, de Souza Santos R (2025) Applications and implications of large language models in qualitative analysis: A new frontier for empirical software engineering. In: IEEE/ACM International Workshop on Methodological Issues with Empirical Studies in Software Engineering, WSESE@ICSE 2025, Ottawa, ON, Canada, May 3, 2025, IEEE, pp 36–43, DOI 10.1109/WSESE66602.2025.00013, URL <https://doi.org/10.1109/WSESE66602.2025.00013>
 92. Moumoula MB, Kaboré AK, Klein J, Bissyandé TF (2024) Large language models for cross-language code clone detection. CoRR abs/2408.04430, DOI 10.48550/ARXIV.2408.04430, URL <https://doi.org/10.48550/arXiv.2408.04430>, 2408.04430
 93. Open Source Initiative (OSI) (2025) Open Source AI Definition 1.0. <https://opensource.org/ai/open-source-ai-definition>, accessed 2025-08-26
 94. OpenAI (2023) How to make your completions outputs consistent with the new seed parameter. https://cookbook.openai.com/examples/reproducible_outputs_with_the_seed_parameter, accessed 2025-01-13
 95. OpenAI (2025) OpenAI API Introduction. <https://platform.openai.com/docs/api-reference/chat/streaming>, accessed 2025-01-13
 96. OpenCode Contributors (2025) OpenCode: The open source AI coding agent. <https://opencode.ai/>, accessed 2026-02-27
 97. OpenTelemetry Authors (2026) OpenTelemetry semantic conventions for generative AI. <https://opentelemetry.io/docs/specs/semconv/gen-ai/>, accessed 2026-05-07
 98. Ornelas T, Araújo AA, Araújo J, Araújo M, Trinkenreich B, Kalinowski M (2025) Llm-assisted thematic analysis: Opportunities, limitations, and recommendations. CoRR abs/2511.14528, DOI 10.48550/ARXIV.2511.

- 14528, URL <https://doi.org/10.48550/arXiv.2511.14528>, 2511.14528
99. Pan Q, Ashktorab Z, Desmond M, Cooper MS, Johnson J, Nair R, Daly E, Geyer W (2024) Human-Centered Design Recommendations for LLM-as-a-Judge. In: ACL 2024 Workshop HuCLLM, arXiv, DOI 10.48550/arXiv.2407.03479
100. Pangakis N, Wolken S, Fasching N (2023) Automated annotation with generative AI requires validation. CoRR abs/2306.00176, DOI 10.48550/ARXIV.2306.00176, URL <https://doi.org/10.48550/arXiv.2306.00176>, 2306.00176
101. Panickssery A, Bowman SR, Feng S (2024) LLM evaluators recognize and favor their own generations. In: Globersons A, Mackey L, Belgrave D, Fan A, Paquet U, Tomczak JM, Zhang C (eds) Advances in Neural Information Processing Systems 38: Annual Conference on Neural Information Processing Systems 2024, NeurIPS 2024, Vancouver, BC, Canada, December 10 - 15, 2024, URL http://papers.nips.cc/paper_files/paper/2024/hash/7f1f0218e45f5414c79c0679633e47bc-Abstract-Conference.html
102. Papineni K, Roukos S, Ward T, Zhu W (2002) BLEU: a method for automatic evaluation of machine translation. In: Proceedings of the 40th Annual Meeting of the Association for Computational Linguistics, ACL, pp 311–318, DOI 10.3115/1073083.1073135, URL <https://aclanthology.org/P02-1040/>
103. Pezeshkpour P, Hruschka E (2024) Large language models sensitivity to the order of options in multiple-choice questions. In: Duh K, Gómez-Adorno H, Bethard S (eds) Findings of the Association for Computational Linguistics: NAACL 2024, Mexico City, Mexico, June 16-21, 2024, Association for Computational Linguistics, Findings of ACL, vol NAACL 2024, pp 2006–2017, DOI 10.18653/V1/2024.FINDINGS-NAACL.130, URL <https://doi.org/10.18653/v1/2024.findings-naacl.130>
104. Rabbi MF, Champa AI, Zibran MF, Islam MR (2024) AI writes, we analyze: The ChatGPT Python code saga. In: Spinellis D, Bacchelli A, Constantinou E (eds) 21st IEEE/ACM International Conference on Mining Software Repositories, MSR 2024, Lisbon, Portugal, April 15-16, 2024, ACM, pp 177–181, DOI 10.1145/3643991.3645076
105. Ralph P, Tempero ED (2018) Construct validity in software engineering research and software metrics. In: Rainer A, MacDonell SG, Keung JW (eds) Proceedings of the 22nd International Conference on Evaluation and Assessment in Software Engineering, EASE2018, ACM, pp 13–23, DOI 10.1145/3210459.3210461
106. Ralph P, bin Ali N, Baltes S, Bianculli D, Diaz J, Dittrich Y, Ernst N, Felderer M, Feldt R, Filieri A, de França BBN, Furia CA, Gay G, Gold N, Graziotin D, He P, Hoda R, Juristo N, Kitchenham B, Lenarduzzi V, Martínez J, Melegati J, Mendez D, Menzies T, Molleri J, Pfahl D, Robbes R, Russo D, Saarimäki N, Sarro F, Taibi D, Siegmund J, Spinellis D, Staron M, Stol K, Storey MA, Taibi D, Tamburri D, Torchiano M, Treude

- C, Turhan B, Wang X, Vegas S (2021) Empirical standards for software engineering research. URL <https://arxiv.org/abs/2010.03525>, 2010.03525
107. Ralph P, Kuuttila M, Arif H, Ayoola B (2024) Teaching software metrology: The science of measurement for software engineering. In: Méndez D, Avgeriou P, Kalinowski M, Ali NB (eds) Handbook on Teaching Empirical Software Engineering, Springer Nature Switzerland, pp 101–154, DOI 10.1007/978-3-031-71769-7_5, URL https://doi.org/10.1007/978-3-031-71769-7_5
108. Rasheed Z, Waseem M, Ahmad A, Kemell K, Wang X, Nguyen-Duc A, Abrahamsson P (2024) Can large language models serve as data analysts? A multi-agent assisted approach for qualitative data analysis. CoRR abs/2402.01386, DOI 10.48550/ARXIV.2402.01386, URL <https://doi.org/10.48550/arXiv.2402.01386>, 2402.01386
109. Reiss MV (2023) Testing the reliability of ChatGPT for text annotation and classification: A cautionary remark. CoRR abs/2304.11085, DOI 10.48550/ARXIV.2304.11085, URL <https://doi.org/10.48550/arXiv.2304.11085>, 2304.11085
110. Ren S, Guo D, Lu S, Zhou L, Liu S, Tang D, Sundaresan N, Zhou M, Blanco A, Ma S (2020) CodeBLEU: a method for automatic evaluation of code synthesis. CoRR abs/2009.10297, URL <https://arxiv.org/abs/2009.10297>, 2009.10297
111. Renze M (2024) The effect of sampling temperature on problem solving in large language models. In: Findings of the Association for Computational Linguistics: EMNLP 2024, pp 7346–7356, DOI 10.18653/v1/2024.findings-emnlp.432
112. Richards J, Wessel M (2024) What you need is what you get: Theory of mind for an LLM-based code understanding assistant. In: IEEE International Conference on Software Maintenance and Evolution, ICSME 2024, IEEE, pp 666–671, DOI 10.1109/ICSME58944.2024.00070
113. Ronanki K, Berger C, Horkoff J (2023) Investigating ChatGPTs potential to assist in requirements elicitation processes. In: 2023 49th Euromicro Conference on Software Engineering and Advanced Applications (SEAA), IEEE, pp 354–361
114. Rozière B, Lachaux M, Chausson L, Lample G (2020) Unsupervised translation of programming languages. In: Larochelle H, Ranzato M, Hadsell R, Balcan M, Lin H (eds) Advances in Neural Information Processing Systems 33: Annual Conference on Neural Information Processing Systems 2020, NeurIPS 2020, December 6-12, 2020, virtual, URL <https://proceedings.neurips.cc/paper/2020/hash/ed23fbf18c2cd35f8c7f8de44f85c08d-Abstract.html>
115. Rozière B, Gehring J, Gloeckle F, Sootla S, Gat I, Tan XE, Adi Y, Liu J, Remez T, Rapin J, Kozhevnikov A, Evtimov I, Bitton J, Bhatt M, Canton-Ferrer C, Grattafiori A, Xiong W, Défossez A, Copet J, Azhar F, Touvron H, Martin L, Usunier N, Scialom T, Synnaeve G (2023) Code Llama: Open foundation models for code. CoRR abs/2308.12950,

- DOI 10.48550/ARXIV.2308.12950, URL <https://doi.org/10.48550/arXiv.2308.12950>, 2308.12950
116. Russo D (2024) Navigating the complexity of generative AI adoption in software engineering. *ACM Trans Softw Eng Methodol* 33(5):135:1–135:50, DOI 10.1145/3652154, URL <https://doi.org/10.1145/3652154>
 117. Sallou J, Durieux T, Panichella A (2024) Breaking the silence: the threats of using llms in software engineering. In: *Proceedings of the 2024 ACM/IEEE 44th International Conference on Software Engineering: New Ideas and Emerging Results, NIER@ICSE 2024, Lisbon, Portugal, April 14-20, 2024, ACM*, pp 102–106, DOI 10.1145/3639476.3639764, URL <https://doi.org/10.1145/3639476.3639764>
 118. Santos A, Vegas S, Oivo M, Juristo N (2021) A procedure and guidelines for analyzing groups of software engineering replications. *IEEE Trans Software Eng* 47(9):1742–1763, DOI 10.1109/TSE.2019.2935720
 119. Schäfer M, Nadi S, Eghbali A, Tip F (2024) An empirical evaluation of using large language models for automated unit test generation. *IEEE Trans Software Eng* 50(1):85–105, DOI 10.1109/TSE.2023.3334955
 120. Schilke O, Reimann M (2025) The transparency dilemma: How ai disclosure erodes trust. *Organizational Behavior and Human Decision Processes* 188:104405, DOI 10.1016/j.obhdp.2025.104405, URL <https://www.sciencedirect.com/science/article/pii/S0749597825000172>
 121. Schneider K, Fotrousi F, Wohlrab R (2025) A reference model for empirically comparing llms with humans. In: *47th IEEE/ACM International Conference on Software Engineering: Software Engineering in Society, ICSE-SEIS 2025, Ottawa, ON, Canada, April 27 - May 3, 2025, IEEE*, pp 130–134, DOI 10.1109/ICSE-SEIS66351.2025.00018, URL <https://doi.org/10.1109/ICSE-SEIS66351.2025.00018>
 122. Schröder S, Morgenroth T, Kuhl U, Vaquet V, Paaßen B (2025) Large language models do not simulate human psychology. *CoRR abs/2508.06950*, DOI 10.48550/ARXIV.2508.06950, URL <https://doi.org/10.48550/arXiv.2508.06950>, 2508.06950
 123. Schroeder K, Wood-Doughty Z (2024) Can you trust LLM judgments? reliability of LLM-as-a-judge. *CoRR abs/2412.12509*, DOI 10.48550/ARXIV.2412.12509, URL <https://doi.org/10.48550/arXiv.2412.12509>, 2412.12509
 124. Schulhoff S, Ilie M, Balepur N, Kahadze K, Liu A, Si C, Li Y, Gupta A, Han H, Schulhoff S, Dulepet PS, Vidyadhara S, Ki D, Agrawal S, Pham C, Kroiz GC, Li F, Tao H, Srivastava A, Costa HD, Gupta S, Rogers ML, Goncarenco I, Sarli G, Galynker I, Peskoff D, Carpuat M, White J, Anadkat S, Hoyle AM, Resnik P (2024) The prompt report: A systematic survey of prompting techniques. *CoRR abs/2406.06608*, DOI 10.48550/ARXIV.2406.06608, URL <https://doi.org/10.48550/arXiv.2406.06608>, 2406.06608
 125. Schulz KF, Altman DG, Moher D (2010) CONSORT 2010 statement: updated guidelines for reporting parallel group randomised trials. *BMJ*

- 340:c332, DOI 10.1136/bmj.c332
126. Schwartz R, Dodge J, Smith NA, Etzioni O (2020) Green AI. *Communications of the ACM* 63(12):54–63, DOI 10.1145/3381831
127. Sclar M, Choi Y, Tsvetkov Y, Suhr A (2024) Quantifying language models’ sensitivity to spurious features in prompt design or: How I learned to start worrying about prompt formatting. In: *The Twelfth International Conference on Learning Representations, ICLR 2024, Vienna, Austria, May 7-11, 2024*, OpenReview.net, URL <https://openreview.net/forum?id=RIu5lyNXjT>
128. Shull F, Singer J, Sjøberg DIK (eds) (2008) *Guide to Advanced Empirical Software Engineering*. Springer, DOI 10.1007/978-1-84800-044-5
129. Silva A, Monperrus M (2025) Repairbench: Leaderboard of frontier models for program repair. In: *IEEE/ACM International Workshop on Large Language Models for Code, LLM4Code@ICSE 2025, Ottawa, ON, Canada, May 3, 2025*, IEEE, pp 9–16, DOI 10.1109/LLM4CODE66737.2025.00006, URL <https://doi.org/10.1109/LLM4Code66737.2025.00006>
130. Sjøberg DIK, Bergersen GR (2023) Construct validity in software engineering. *IEEE Trans Software Eng* 49(3):1374–1396, DOI 10.1109/TSE.2022.3176725
131. Song Y, Wang G, Li S, Lin BY (2025) The good, the bad, and the greedy: Evaluation of LLMs should not ignore non-determinism. In: Chiruzzo L, Ritter A, Wang L (eds) *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies, NAACL 2025 - Volume 1: Long Papers*, Association for Computational Linguistics, pp 4195–4206, DOI 10.18653/V1/2025.NAACL-LONG.211, URL <https://doi.org/10.18653/v1/2025.naacl-long.211>
132. Sonnekalb T, Gruner B, Brust C, Mäder P (2022) Generalizability of code clone detection on CodeBERT. In: *37th IEEE/ACM International Conference on Automated Software Engineering, ASE 2022, ACM*, pp 143:1–143:3, DOI 10.1145/3551349.3561165
133. Stack Overflow (2025) 2025 developer survey. <https://survey.stackoverflow.co/2025/>, 49,009 respondents from 166 countries, May–June 2025
134. Strubell E, Ganesh A, McCallum A (2019) Energy and policy considerations for deep learning in NLP. In: Korhonen A, Traum DR, Màrquez L (eds) *Proceedings of the 57th Conference of the Association for Computational Linguistics, ACL 2019, Florence, Italy, July 28- August 2, 2019, Volume 1: Long Papers*, Association for Computational Linguistics, pp 3645–3650, DOI 10.18653/V1/P19-1355, URL <https://doi.org/10.18653/v1/p19-1355>
135. Sullivan GM, Sargeant J (2011) Qualities of qualitative research: part I. *J Grad Med Educ* 3(4):449–452
136. Sumers TR, Yao S, Narasimhan K, Griffiths TL (2024) Cognitive architectures for language agents. *Trans Mach Learn Res* 2024, URL

- <https://openreview.net/forum?id=1i6ZCvflQJ>
137. Takerngsaksiri W, Pasuksmit J, Thongtanunam P, Tantithamthavorn C, Zhang R, Jiang F, Li J, Cook E, Chen K, Wu M (2025) Human-in-the-loop software development agents. In: 47th IEEE/ACM International Conference on Software Engineering: Software Engineering in Practice, SEIP@ICSE 2025, Ottawa, ON, Canada, April 27 - May 3, 2025, IEEE, pp 342–352, DOI 10.1109/ICSE-SEIP66354.2025.00036, URL <https://doi.org/10.1109/ICSE-SEIP66354.2025.00036>
 138. Tam TYC, Sivarajkumar S, Kapoor S, Stolyar AV, Polanska K, McCarthy KR, Osterhoudt H, Wu X, Visweswaran S, Fu S, Mathur P, Cacciamani GE, Sun C, Peng Y, Wang Y (2024) A framework for human evaluation of large language models in healthcare derived from literature review. *npj Digit Medicine* 7(1), DOI 10.1038/S41746-024-01258-7, URL <https://doi.org/10.1038/s41746-024-01258-7>
 139. Wagner S, Barón MM, Falessi D, Baltes S (2025) Towards evaluation guidelines for empirical studies involving LLMs. In: IEEE/ACM International Workshop on Methodological Issues with Empirical Studies in Software Engineering, WSESE@ICSE 2025, May 3, 2025, IEEE, pp 24–27, DOI 10.1109/WSESE66602.2025.00011
 140. Wan M, Safavi T, Jauhar SK, Kim Y, Counts S, Neville J, Suri S, Shah C, White RW, Yang L, Andersen R, Buscher G, Joshi D, Rangan N (2024) TnT-LLM: Text mining at scale with large language models. In: Baeza-Yates R, Bonchi F (eds) *Proceedings of the 30th ACM SIGKDD Conference on Knowledge Discovery and Data Mining, KDD 2024*, ACM, pp 5836–5847, DOI 10.1145/3637528.3671647
 141. Wang A, Morgenstern J, Dickerson JP (2025) Large language models that replace human participants can harmfully misportray and flatten identity groups. *Nat Mac Intell* 7(3):400–411, DOI 10.1038/S42256-025-00986-Z, URL <https://doi.org/10.1038/s42256-025-00986-z>
 142. Wang C, Huang K, Zhang J, Feng Y, Zhang L, Liu Y, Peng X (2024) How and why llms use deprecated apis in code completion? an empirical study. *CoRR abs/2406.09834*, DOI 10.48550/ARXIV.2406.09834, URL <https://doi.org/10.48550/arXiv.2406.09834>, 2406.09834
 143. Wang F, Arora C, Liu Y, Huang K, Tantithamthavorn C, Aleti A, Sambathkumar D, Lo D (2025) Multi-modal requirements data-based acceptance criteria generation using llms. In: 40th IEEE/ACM International Conference on Automated Software Engineering, ASE 2025, Seoul, Korea, Republic of, November 16-20, 2025, IEEE, pp 3334–3345, DOI 10.1109/ASE63991.2025.00275, URL <https://doi.org/10.1109/ASE63991.2025.00275>
 144. Wang S, Liu Y, Xu Y, Zhu C, Zeng M (2021) Want to reduce labeling cost? GPT-3 can help. In: Moens M, Huang X, Specia L, Yih SW (eds) *Findings of the Association for Computational Linguistics: EMNLP 2021, Virtual Event / Punta Cana, Dominican Republic, 16-20 November, 2021*, Association for Computational Linguistics, Findings of ACL, vol EMNLP 2021, pp 4195–4205, DOI 10.18653/V1/2021.FINDINGS-EMNLP.354,

- URL <https://doi.org/10.18653/v1/2021.findings-emnlp.354>
145. Wang X, Kim H, Rahman S, Mitra K, Miao Z (2024) Human-LLM collaborative annotation through effective verification of LLM labels. In: Mueller FF, Kyburz P, Williamson JR, Sas C, Wilson ML, Dugas POT, Shklovski I (eds) Proceedings of the CHI Conference on Human Factors in Computing Systems, CHI 2024, ACM, pp 303:1–303:21, DOI 10.1145/3613904.3641960
 146. Widder DG, Whittaker M, West SM (2024) Why openai systems are actually closed, and why this matters. *Nature* 635(8040):827–833
 147. Wiesinger J, Marlow P, Vuskovic V (2025) Agents. Google DeepMind URL <https://gemini.google.com>, contributors: Evan Huang, Emily Xue, Olcan Sercinoglu, Sebastian Riedel, Satinder Baveja, Antonio Gulli, Anant Nawalgaria
 148. Willison S (2026) Opus system prompt. <https://simonwillison.net/2026/Apr/18/opus-system-prompt/>, accessed 2026-04-28
 149. Wohlin C, Runeson P, Höst M, Ohlsson MC, Regnell B, Wesslén A (2024) Experimentation in Software Engineering, Second Edition. Springer, DOI 10.1007/978-3-662-69306-3
 150. Xia Y, Shao H, Deng X (2024) Vulcobert: A CodeBERT-based system for source code vulnerability detection. In: 2024 International Conference on Generative Artificial Intelligence and Information Security, GAIIS 2024, Kuala Lumpur, Malaysia, May 10-12, 2024, ACM, DOI 10.1145/3665348.3665391
 151. Xiao T, Treude C, Hata H, Matsumoto K (2024) DevGPT: Studying developer-ChatGPT conversations. In: Spinellis D, Bacchelli A, Constantinou E (eds) 21st IEEE/ACM International Conference on Mining Software Repositories, MSR 2024, Lisbon, Portugal, April 15-16, 2024, ACM, pp 227–230, DOI 10.1145/3643991.3648400
 152. Xu C, Guan S, Greene D, Kechadi MT (2024) Benchmark data contamination of large language models: A survey. CoRR abs/2406.04244, DOI 10.48550/ARXIV.2406.04244, URL <https://doi.org/10.48550/arXiv.2406.04244>, 2406.04244
 153. Xu R, Sun Y, Ren M, Guo S, Pan R, Lin H, Sun L, Han X (2024) AI for social science and social science of AI: A survey. *Inf Process Manag* 61(2):103665, DOI 10.1016/J.IPM.2024.103665, URL <https://doi.org/10.1016/j.ipm.2024.103665>
 154. Yan L, Hwang A, Wu Z, Head A (2024) Ivie: Lightweight anchored explanations of just-generated code. In: Mueller FF, Kyburz P, Williamson JR, Sas C, Wilson ML, Dugas POT, Shklovski I (eds) Proceedings of the CHI Conference on Human Factors in Computing Systems, CHI 2024, ACM, pp 140:1–140:15, DOI 10.1145/3613904.3642239
 155. Yang G, Zhou Y, Chen X, Zhang X, Han T, Chen T (2023) Exploitgen: Template-augmented exploit code generation based on CodeBERT. *J Syst Softw* 197:111577, DOI 10.1016/J.JSS.2022.111577, URL <https://doi.org/10.1016/j.jss.2022.111577>

156. Yang Y, Shin J, Choi B, Park M, Ju D, Lee C, Chun S, Kang D, Kim J, Yoon S (2025) From code foundation models to agents and applications: A comprehensive survey and practical guide to code intelligence. arXiv preprint arXiv:250211827
157. Yuan J, Li H, Ding X, Xie W, Li YJ, Zhao W, Wan K, Shi J, Hu X, Liu Z (2025) Understanding and mitigating numerical sources of nondeterminism in LLM inference. In: Advances in Neural Information Processing Systems 38: Annual Conference on Neural Information Processing Systems 2025, NeurIPS 2025, URL <https://openreview.net/forum?id=Q3qAsZAEZw>
158. Zhao C, Habule M, Zhang W (2025) Large language models (LLMs) as research subjects: Status, opportunities and challenges. *New Ideas in Psychology* 79:101167, DOI 10.1016/j.newideapsych.2025.101167, URL <https://www.sciencedirect.com/science/article/pii/S0732118X25000236>
159. Zhou R, Chen L, Yu K (2024) Is LLM a reliable reviewer? A comprehensive evaluation of LLM on automatic paper reviewing tasks. In: Calzolari N, Kan M, Hoste V, Lenci A, Sakti S, Xue N (eds) Proceedings of the 2024 Joint International Conference on Computational Linguistics, Language Resources and Evaluation, LREC/COLING 2024, 20-25 May, 2024, Torino, Italy, ELRA and ICCL, pp 9340–9351, URL <https://aclanthology.org/2024.lrec-main.816>
160. Zhou W, Jiang YE, Li L, Wu J, Wang T, Qiu S, Zhang J, Chen J, Wu R, Wang S, Zhu S, Chen J, Zhang W, Zhang N, Chen H, Cui P, Sachan M (2023) Agents: An open-source framework for autonomous language agents. CoRR abs/2309.07870, DOI 10.48550/ARXIV.2309.07870, URL <https://doi.org/10.48550/arXiv.2309.07870>, 2309.07870
161. Zhu Y, Zhang P, ul Haq E, Hui P, Tyson G (2023) Can ChatGPT reproduce human-generated labels? A study of social computing tasks. CoRR abs/2304.10145, DOI 10.48550/ARXIV.2304.10145, URL <https://doi.org/10.48550/arXiv.2304.10145>, 2304.10145



Change of authorship request form (pre-acceptance)

Please read the important information on page 4 before you begin

This form should be used by authors to request any change in authorship including changes in corresponding authors. Please fully complete all sections. Use black ink and block capitals and provide each author's full name with the given name first followed by the family name.

Please note: In author collaborations where there is formal agreement for representing the collaboration, it is sufficient for the representative or legal guarantor (usually the corresponding author) to complete and sign the Authorship Change Form on behalf of all authors.

Section 1: Please provide the current title of manuscript

(For journals: Please provide the manuscript ID, title and/or DOI if available.)

(For books: Please provide the title, ISBN and/or DOI if available.)

Manuscript ID no. in case of unpublished manuscript:
EMSE-D-25-00637R1

DOI in case of published manuscript:

ISBN (for books):

Title: Guidelines for Empirical Studies in Software Engineering involving Large Language Models

Section 2: Please provide the previous authorship, in the order shown on the manuscript before the changes were introduced. Please indicate the corresponding author by adding (CA) behind the name.

	First name(s)	Family name	ORCID or SCOPUS id, if available
1 st author	Sebastian	Baltes	
2 nd author	Florian	Angermeir	
3 rd author	Chetan	Arora	
4 th author	Marvin	Muñoz Barón	

5 th author	Chunyang	Chen	
6 th author	Lukas	Böhme	
7 th author	Fabio	Calefato	
8	Neil	Ernst	
9	Davide	Falessi	
10	Brian	Fitzgerald	
11	Davide	Fucci	
12	Marcos	Kalinowski	
13	Stefano	Lambiase	
14	Daniel	Russo	
15	Mircea	Lungu	
16	Lutz	Prechelt	
17	Paul	Ralph	
18	Christoph	Treude	
19	Stefan	Wagner	

Please use an additional sheet if there are more than 7 authors.

Change of authorship request form (pre-acceptance)

Section 3: Please provide a justification for change. Please use this section to explain your reasons for changing the authorship of your manuscript, e.g. what necessitated the change in authorship? Please refer to the (journal) policy pages for more information about authorship. Please explain why omitted authors were not originally included and/or why authors were removed on the submitted manuscript.

See response letter:

As part of this revision, we welcomed new co-authors (Cristina Martinez Montes, Rijnard van Tonder, and Junda He) who contributed additional expertise to the major rewrite of the guidelines. This project is a community effort, and we remain open to new contributors who can strengthen the breadth and depth of our recommendations.

Section 4: Proposed new authorship. Please provide your new authorship list in the order you would like it to appear on the manuscript. Please indicate the corresponding author by adding (CA) behind the name. If the corresponding author has changed, please indicate the reason under section 3.

	First name(s)	Family name (this name will appear in full on the final publication and will be searchable in various abstract and indexing databases)
1 st author	Sebastian	Baltes
2 nd author	Florian	Angermeir
3 rd author	Chetan	Arora
4 th author	Marvin	Muñoz Barón
5 th author	Chunyang	Chen
6 th author	Lukas	Böhme
7 th author	Fabio	Calefato

8	Neil	Ernst
9	Davide	Falessi
10	Brian	Fitzgerald
11	Davide	Fucci
12	Junda	He
13	Christoph	Treude
14	Marcos	Kalinowski
15	Stefano	Lambiase
16	Daniel	Russo
17	Mircea	Lungu
18	Cristina	Martinez Montes
19	Lutz	Prechelt
20	Paul	Ralph
21	Rijnard	van Tonder
22	Stefan	Wagner

Please use an additional sheet if there are more than 7 authors.

Change of authorship request form (pre-acceptance)

Section 5: Author contribution, Acknowledgement and Disclosures. Please use this section to provide a new disclosure statement and, if appropriate, acknowledge any contributors who have been removed as authors and ensure you state what contribution any new authors made (if applicable per the journal or book (series) policy). **Please ensure these are updated in your manuscript - after approval of the change(s) - as our production department will not transfer the information in this form to your manuscript.**



n/a


State 'Not applicable' if there are no new authors.

Change of authorship request form (pre-acceptance)

Section 6: Declaration of agreement. All authors, unchanged, new and removed *must* sign this declaration.

(NB: Please print the form, sign and return a scanned copy. Please note that signatures that have been inserted as an image file are acceptable as long as it is handwritten. Typed names in the signature box are unacceptable.) * Please delete as appropriate. Delete all of the bold if you were on the original authorship list and are remaining as an author.


	First name	Family name		Signature	Affiliated institute	Date
1 st author	Sebastian	Baltes	I agree to the proposed new authorship shown in section 4 /and the addition/removal*of my name to the authorship list.		Heidelberg University	03. March 2026
2 nd author	Florian	Angermeir	I agree to the proposed new authorship shown in section 4 /and the addition/removal*of my name to the authorship list.		fortiss Blekinge Institute of Technology	04. March 2026
3 rd author	Chetan	Arora	I agree to the proposed new authorship shown in section 4 /and the addition/removal*of my name to the authorship list.		Monash University	05. March 2026
4 th author	Marvin	Muñoz Barón	I agree to the proposed new authorship shown in section 4 /and the addition/removal*of my name to the authorship list.		Technical University of Munich	09. March 2026
5 th author	Chunyang	Chen	I agree to the proposed new authorship shown in section 4 /and the addition/removal*of my name to the authorship list.		Technical University of Munich	March 4, 2026

6 th author	Lukas	Böhme	I agree to the proposed new authorship shown in section 4 /and the addition/removal*of my name to the authorship list.		Hasso Plattner Institute	04. March 2026
7 th author	Fabio	Calefato	I agree to the proposed new authorship shown in section 4 /and the addition/removal*of my name to the authorship list.		University of Bari	03. March 2026
8	Neil	Ernst	I agree to the proposed new authorship shown in section 4 /and the addition/removal*of my name to the authorship list.		University of Victoria	Mar 3, 2026
9	Davide	Falessi	I agree to the proposed new authorship shown in section 4 /and the addition/removal*of my name to the authorship list.		University of Rome Tor Vergata	3 March 2026
10	Brian	Fitzgerald	I agree to the proposed new authorship shown in section 4 /and the addition/removal*of my name to the authorship list.		Lero and University of Limerick	3 March 2026
11	Davide	Fucci	I agree to the proposed new authorship shown in section 4 /and the addition/removal*of my name to the authorship list.		Blekinge Tekniska Högskola	March 4th, 2026
12	Junda	He	I agree to the proposed new authorship shown in section 4 /and the addition/removal*of my name to the authorship list.		Singapore Management University	March 8, 2026

13	Christoph	Treude	I agree to the proposed new authorship shown in section 4 /and the addition/removal*of my name to the authorship list.		Singapore Management University	March 4, 2026
14	Marcos	Kalinowski	I agree to the proposed new authorship shown in section 4 /and the addition/removal*of my name to the authorship list.		Pontifical Catholic University of Rio de Janeiro	March 3, 2026
15	Stefano	Lambiase	I agree to the proposed new authorship shown in section 4 /and the addition/removal*of my name to the authorship list.		Aalborg University	March 4, 2026
16	Daniel	Russo	I agree to the proposed new authorship shown in section 4 /and the addition/removal*of my name to the authorship list.		Aalborg University	3 March 2026
17	Mircea	Lungu	I agree to the proposed new authorship shown in section 4 /and the addition/removal*of my name to the authorship list.		IT University Copenhagen	March 3, '26
18	Cristina	Martinez Montes	I agree to the proposed new authorship shown in section 4 /and the addition/removal*of my name to the authorship list.		Chalmers University of Technology and University of Gothenburg	March 05, 2026
19	Lutz	Prechelt	I agree to the proposed new authorship shown in section 4 /and the addition/removal*of my name to		Freie Universität Berlin	2026-03-06

			the authorship list.			
20	Paul	Ralph	I agree to the proposed new authorship shown in section 4 /and the addition/removal*of my name to the authorship list.		Dalhousie University	03. March 2026
21	Rijnard	van Tonder	I agree to the proposed new authorship shown in section 4 /and the addition/removal*of my name to the authorship list.		Independent	03. March 2026
22	Stefan	Wagner	I agree to the proposed new authorship shown in section 4 /and the addition/removal*of my name to the authorship list.		Technical University of Munich	3 March 2026

Please use an additional sheet if there are more than 7 authors.


 Springer Nature is one of the world's leading global research, educational and professional publishers, created in May 2015 through the combination of Nature Publishing Group, Palgrave Macmillan, Macmillan Education and Springer Science+Business Media.



Change of authorship request form (pre-acceptance)

Important information. Please read.

- Please return this form, fully completed, to Springer Nature. We will consider the information you have provided to decide whether to approve the proposed change in authorship. We may choose to contact your institution for more information or undertake a further investigation, if appropriate, before making a final decision.
- By signing this declaration, all authors guarantee that the order of the authors are in accordance with their scientific contribution, if applicable as different conventions apply per discipline, and that only authors have been added who made a meaningful contribution to the work.
- Please note, we cannot investigate or mediate any authorship disputes. If you are unable to obtain agreement from all authors (including those who you wish to be removed) you must refer the matter to your institution(s) for investigation. Please inform us if you need to do this.
- If you are not able to return a fully completed form within **30 days** of the date that it was sent to the author requesting the change, we may have to withdraw your manuscript. We cannot publish manuscripts where authorship has not been agreed by all authors (including those who have been removed).
- Incomplete forms will be rejected.

 Springer Nature is one of the world's leading global research, educational and professional publishers, created in May 2015 through the combination of Nature Publishing Group, Palgrave Macmillan, Macmillan Education and Springer Science+Business Media. ---